

# High Performance Computing for Engineers

David Thomas

[dt10@ic.ac.uk](mailto:dt10@ic.ac.uk) / <https://github.com/m8pple>

Room 903

<https://github.com/HPCE/hpce-2016>

# *High Performance* **Computing for Engineers**

- Research
  - Testing communication protocols
  - Evaluating signal-processing filters
  - Simulating analogue and digital designs

# *High Performance* **Computing for Engineers**

- Research
  - Testing communication protocols
  - Evaluating signal-processing filters
  - Simulating analogue and digital designs
- Tools
  - CAD tools: synthesis, place-and-route, verification
  - Libraries/toolboxes: filter design, compressive sensing

# *High Performance* **Computing for Engineers**

- Research
  - Testing communication protocols
  - Evaluating signal-processing filters
  - Simulating analogue and digital designs
- Tools
  - CAD tools: synthesis, place-and-route, verification
  - Libraries/toolboxes: filter design, compressive sensing
- Products
  - Oil exploration and discovery
  - Mobile-phone apps
  - Financial computing

# High Performance *Computing for Engineers*

- Types of performance metrics

# High Performance *Computing for Engineers*

- Types of performance metrics
  - Throughput
  - Latency
  - Power
  - Design-time
  - Capital and running costs

# High Performance *Computing for Engineers*

- Types of performance metrics
  - Throughput
  - Latency
  - Power
  - Design-time
  - Capital and running costs
- Required versus desired performance
  - *Subject to a throughput of  $X$ , minimise average power*
  - *Subject to a budget of  $Y$ , maximise energy efficiency*
  - *Subject to  $Z$  development days, maximise throughput*

# What is available to you


- Types of compute device
  - Multi-core CPUs
  - GPUs (Graphics Processing Units)
  - MPPAs (Massively Parallel Processor Arrays)
  - FPGAs (Field Programmable Gate Arrays)



# What is available to you

- Types of compute device
  - Multi-core CPUs
  - GPUs (Graphics Processing Units)
  - MPPAs (Massively Parallel Processor Arrays)
  - FPGAs (Field Programmable Gate Arrays)
- Types of compute system
  - Embedded Systems
  - Mobile Phones
  - Tablets
  - Laptops
  - Grid computing
  - Cloud computing


# HTC Droid DNA



## Snapdragon S4 Pro

- CPU : Quad-core Krait (ARM derivative)
- GPU : Adreno 320 GPU (OpenCL compatible)

# Lenovo Thinkpad Edge E525



AMD Fusion A8-3500M

- CPU : Quad-Core 2.4GHz Phenom-II
- GPU : HD 6620G 400MHz (320 cores)

# Imperial HPC Cluster

- cx1 – cluster of networked machines
  - 1395 nodes (boxes) -> 13558 CPU cores
- cx2 - SGI Altix ICE 8200 EX
  - 456 nodes -> 5272 CPU cores
  - Optimised for MPI (message processing) tasks
- ax4 – one machine: 15TB of RAM + 1280 cores



- Grid-management system
  - Run program on 1000 PCs with one command
  - Available to researchers and undergrads (if they ask nicely)

# Cloud Computing

- There are now **big** commodity cloud providers
  - Amazon Web Services (EC2) : 10x bigger than anyone else
  - Microsoft Azure : 2x bigger than all the rest
  - Google Cloud Platform : public facing cloud is fairly small (?)

# Cloud Computing

- There are now **big** commodity cloud providers
  - Amazon Web Services (EC2) : 10x bigger than anyone else
  - Microsoft Azure : 2x bigger than all the rest
  - Google Cloud Platform : public facing cloud is fairly small (?)
- Multiple instance types
  - t2.micro: 1 CPU, 1 GB
  - g2.2xlarge: 8 CPUs, 15 GB + GPU
  - c3.8xlarge: 32 CPUs, 108 GB

# Cloud Computing

- There are now **big** commodity cloud providers
  - Amazon Web Services (EC2) : 10x bigger than anyone else
  - Microsoft Azure : 2x bigger than all the rest
  - Google Cloud Platform : public facing cloud is fairly small (?)
- Multiple instance types
  - t2.micro: 1 CPU, 1 GB
  - g2.2xlarge: 8 CPUs, 15 GB + GPU
  - c3.8xlarge: 32 CPUs, 108 GB
- Multiple pricing options : on-demand vs. spot-price
  - On-demand : fixed price for as long as you want

# Cloud Computing


- There are now **big** commodity cloud providers
  - Amazon Web Services (EC2) : 10x bigger than anyone else
  - Microsoft Azure : 2x bigger than all the rest
  - Google Cloud Platform : public facing cloud is fairly small (?)
- Multiple instance types
  - t2.micro: 1 CPU, 1 GB \$0.013 / hour
  - g2.2xlarge: 8 CPUs, 15 GB + GPU \$0.650 / hour
  - c3.8xlarge: 32 CPUs, 108 GB \$1.680 / hour
- Multiple pricing options : on-demand vs. spot-price
  - On-demand : fixed price for as long as you want




# Cloud Computing

- There are now **big** commodity cloud providers
  - Amazon Web Services (EC2) : 10x bigger than anyone else
  - Microsoft Azure : 2x bigger than all the rest
  - Google Cloud Platform : public facing cloud is fairly small (?)
- Multiple instance types
  - t2.micro: 1 CPU, 1 GB \$0.013 / hour
  - g2.2xlarge: 8 CPUs, 15 GB + GPU \$0.650 / hour \$0.103 / hour
  - c3.8xlarge: 32 CPUs, 108 GB \$1.680 / hour \$0.311 / hour
- Multiple pricing options : on-demand vs. spot-price
  - On-demand : fixed price for as long as you want
  - Spot-price : price fluctuates according to demand

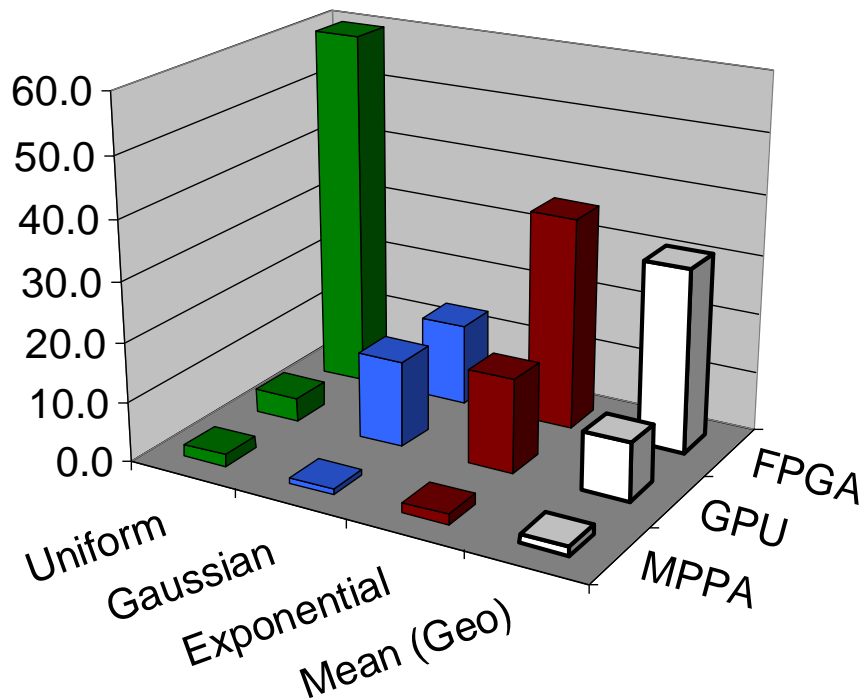
# Pricing can be volatile (and make no sense)




# Pricing can be volatile (and make no sense)



# Performance and Efficiency Relative to CPU



Performance




Power Efficiency


# Design tradeoffs




# Design tradeoffs



# Design tradeoffs



# Design tradeoffs





# Design tradeoffs

- **Task-based** parallelism vs threads
  - Easy to program (less time coding)
  - Easy to get right (less time testing)
- Many implementations and APIs
  - Intel Threaded Building Blocks (TBB)
  - Microsoft .NET Task Parallel Library
  - OpenCL

Performance

1


Design-time

ential SW


based SW

-based SW

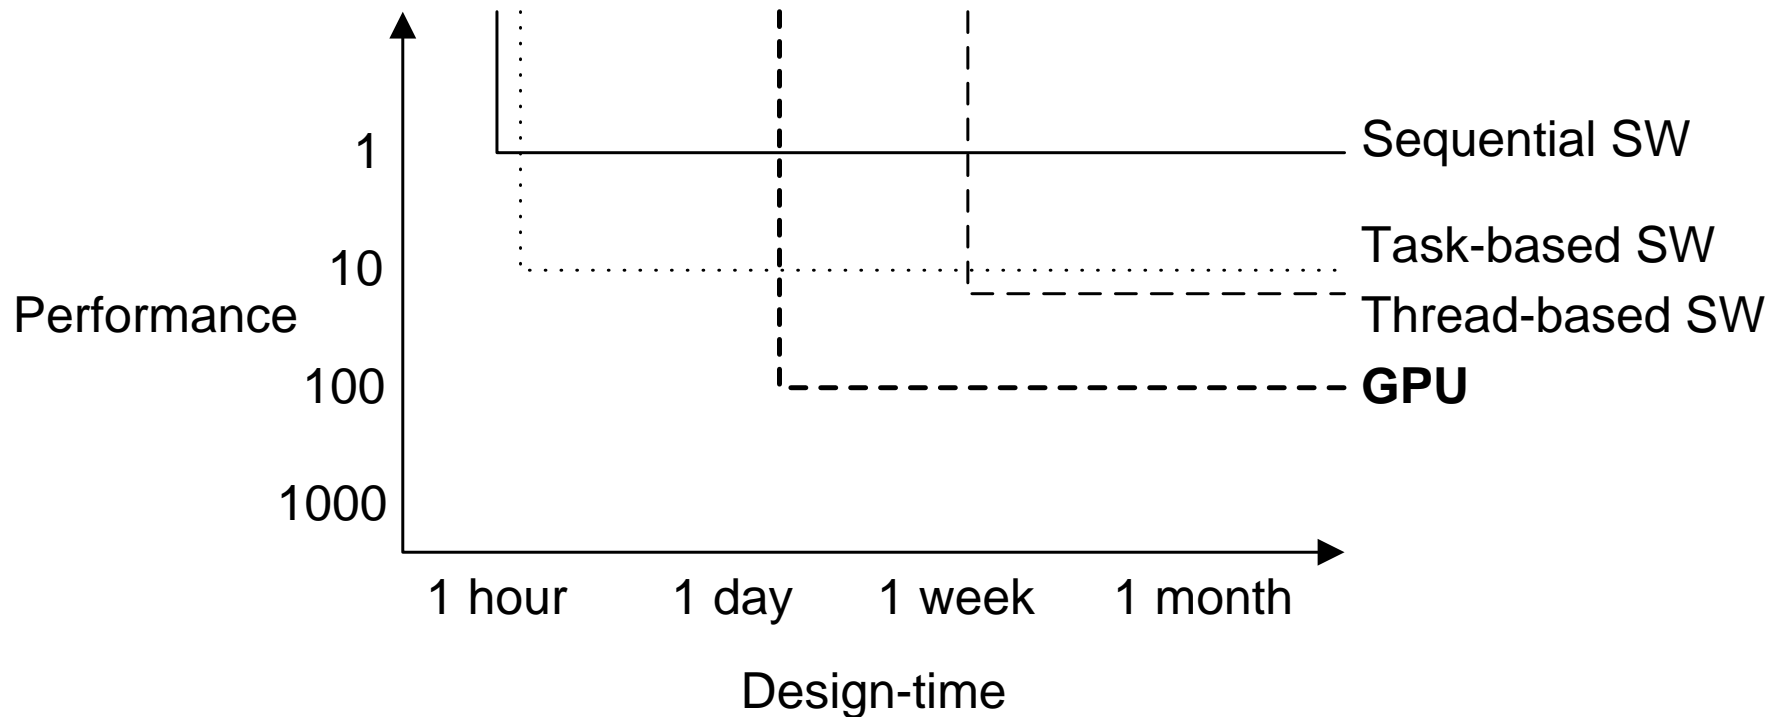
# Design tradeoffs



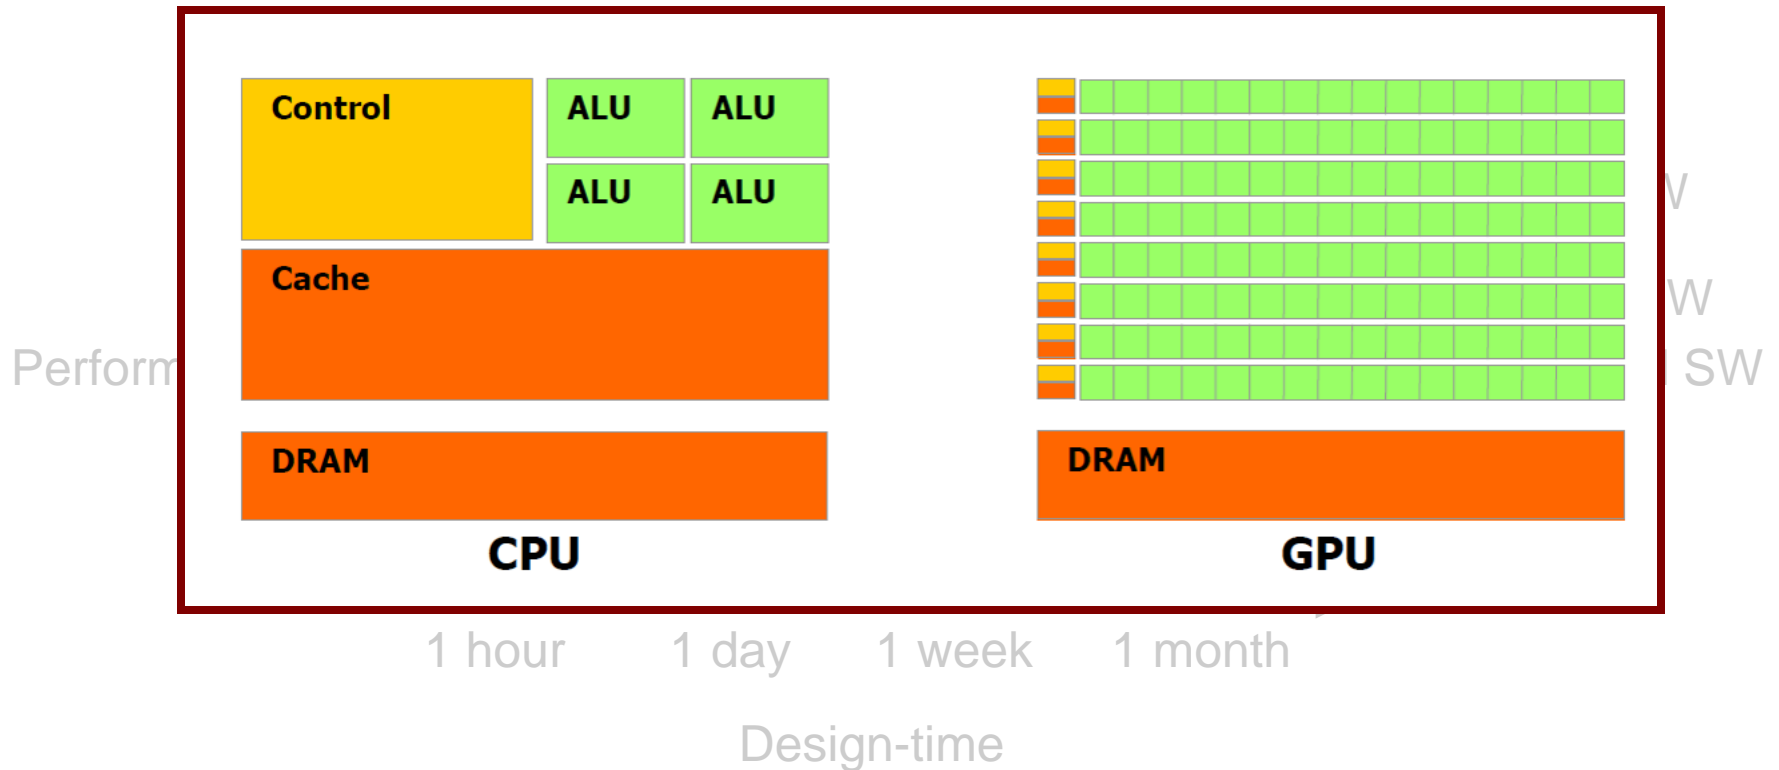
# Design tradeoffs




# Design tradeoffs



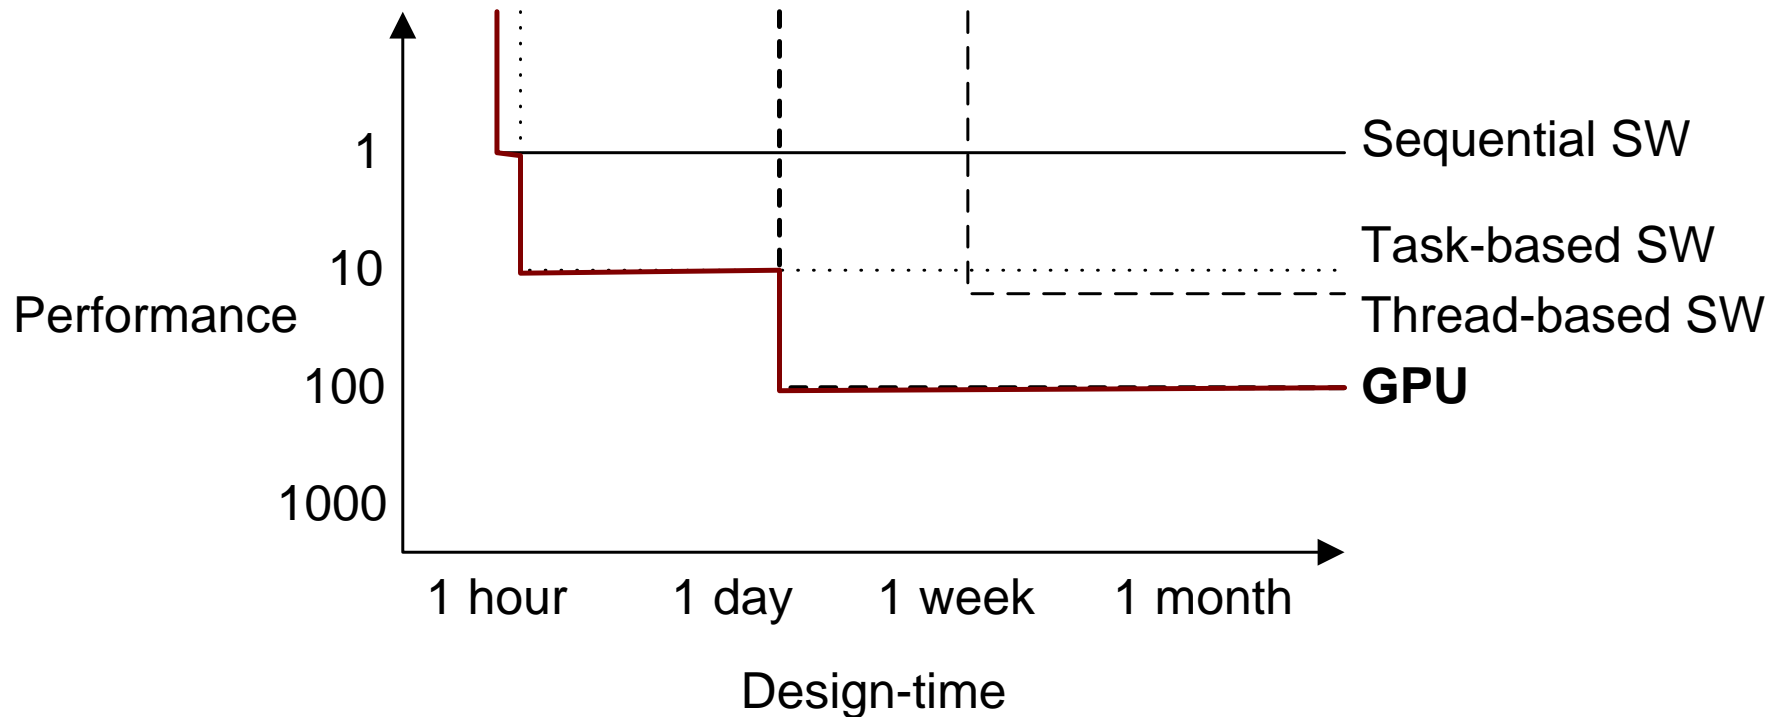
# Design tradeoffs




# Design tradeoffs



# Design tradeoffs

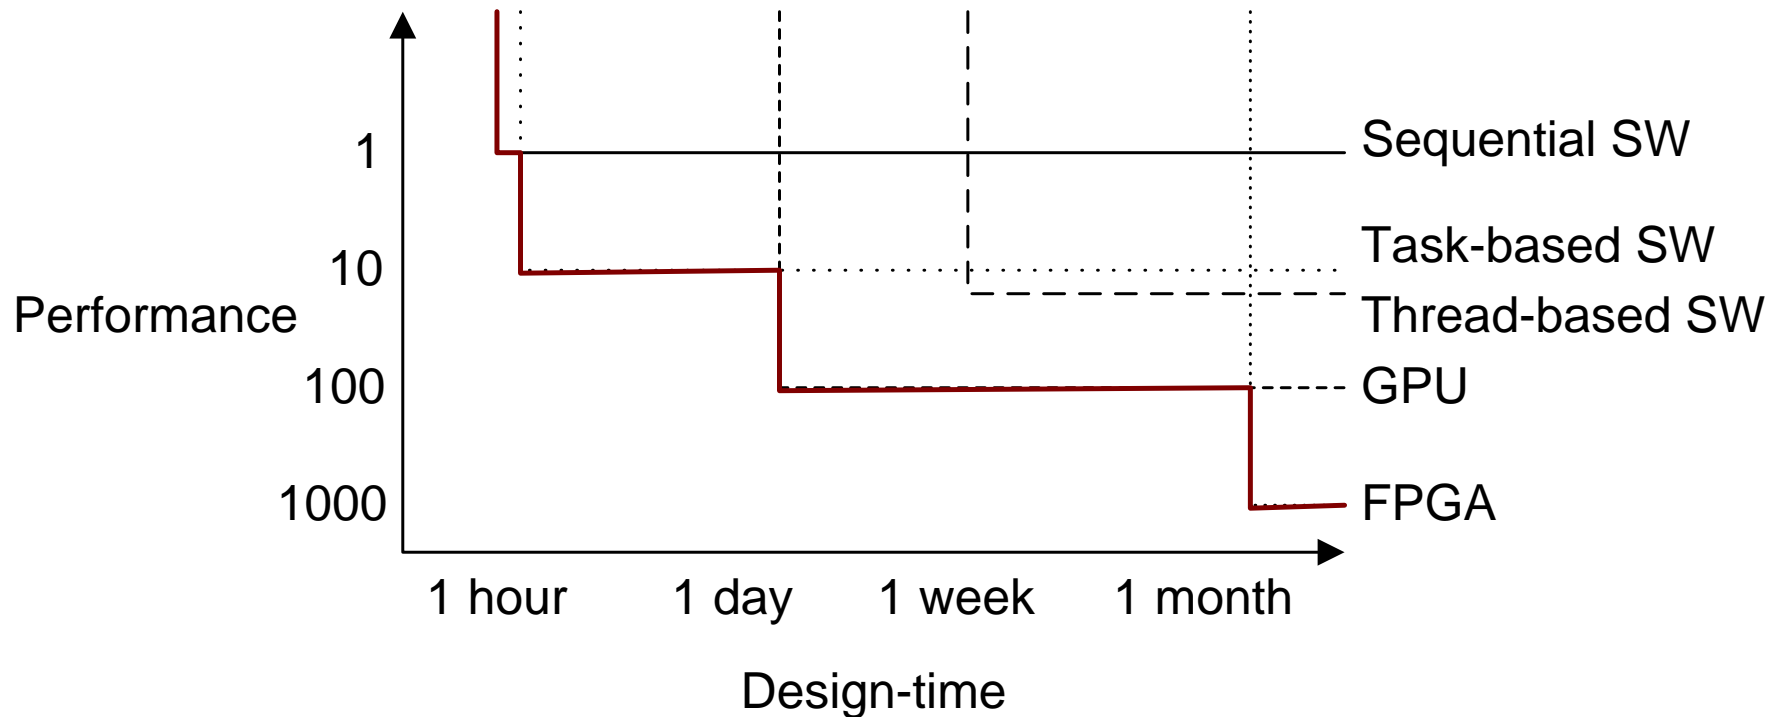


# Design tradeoffs






# Design tradeoffs



# What you will learn



- **Systems:** what high-performance systems are available
- **Methods:** how these systems can be programmed
- **Practise:** concrete experience with multi-core and GPUs
- **Analysis:** knowing what to use and when
- **Tools:** making better use of your time

"I SPEND A LOT OF TIME ON THIS TASK.  
I SHOULD WRITE A PROGRAM AUTOMATING IT!"



Developer productivity  
is also part of performance

"I SPEND A LOT OF TIME ON THIS TASK.  
I SHOULD WRITE A PROGRAM AUTOMATING IT!"



# *Re: XKCD* - My Professional Context

- Undergraduate degree and PhD from Computing
  - If pushed, I self-identify as a “programmer”
- Research focuses on hardware acceleration
  - Both academic and industrial applications
- My motivation for this course
  - Supervising final year project students
  - Working with PhD students
  - Talking to industry people

# Why are you here?

# Course Assessment

- 40% : Four short exercises to build skills
  - Get familiar with environments and how to do common tasks
  - Structured and quite linear – should not be taxing
  - Force people to do work earlier in term

# Course Assessment

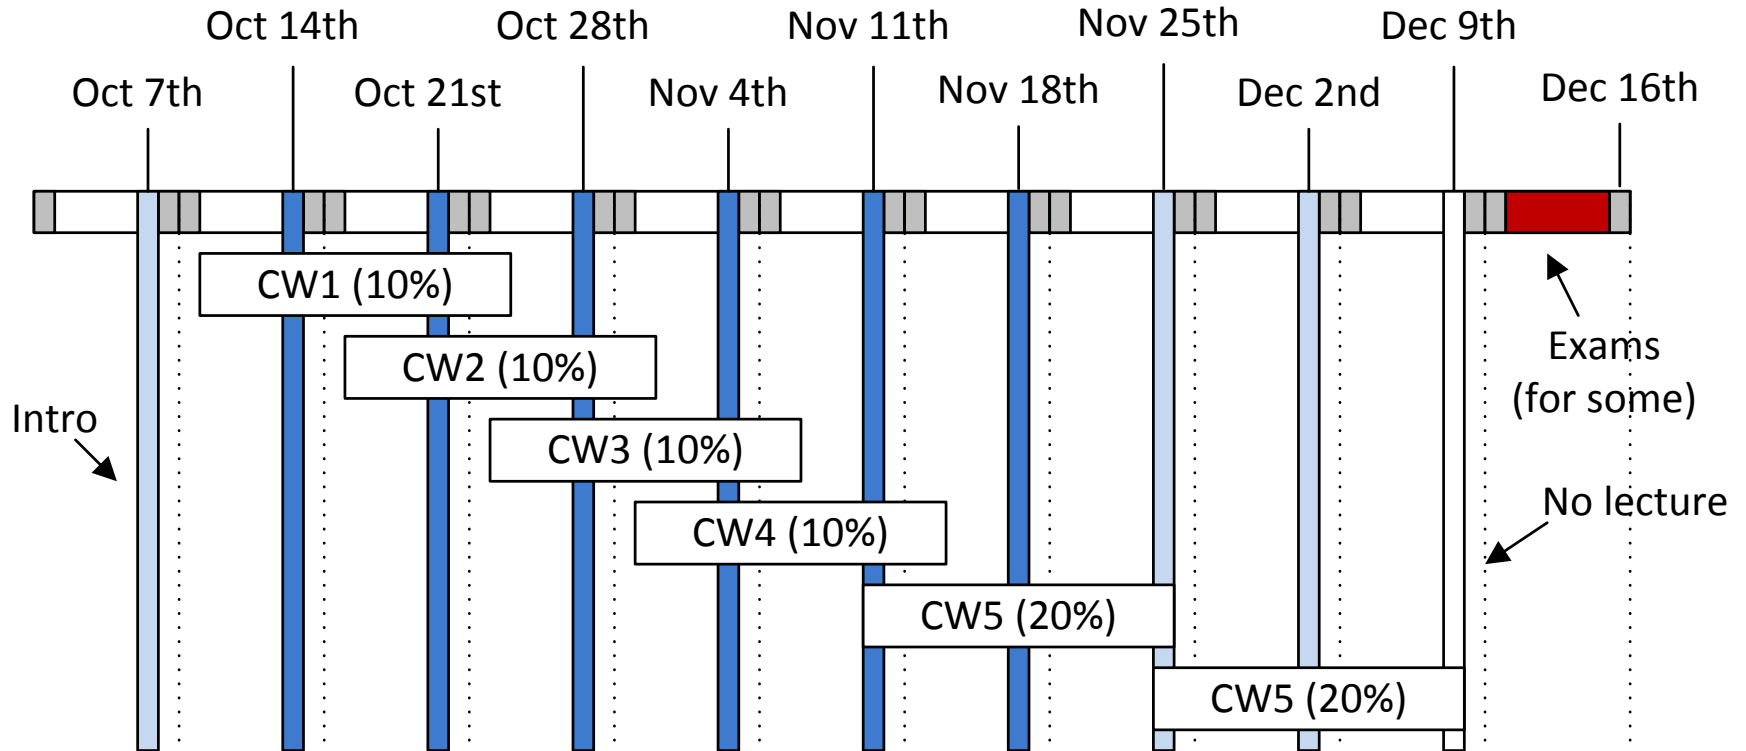
- 40% : Four short exercises to build skills
  - Get familiar with environments and how to do common tasks
  - Structured and quite linear – should not be taxing
  - Force people to do work earlier in term
- 40% : Two larger tasks to apply skills to real problems
  - Allow demonstration of knowledge and skills
  - Unstructured; open-ended; competitive; hard



# Course Assessment

- 40% : Four short exercises to build skills
  - Get familiar with environments and how to do common tasks
  - Structured and quite linear – should not be taxing
  - Force people to do work earlier in term
- 40% : Two larger tasks to apply skills to real problems
  - Allow demonstration of knowledge and skills
  - Unstructured; open-ended; competitive; hard
- 20% : Oral assessment; individual
  - 30 minutes each, will happen at the start of spring term
  - Test ability to communicate about your code and solutions
  - (Check that you did the work)

# Timetable



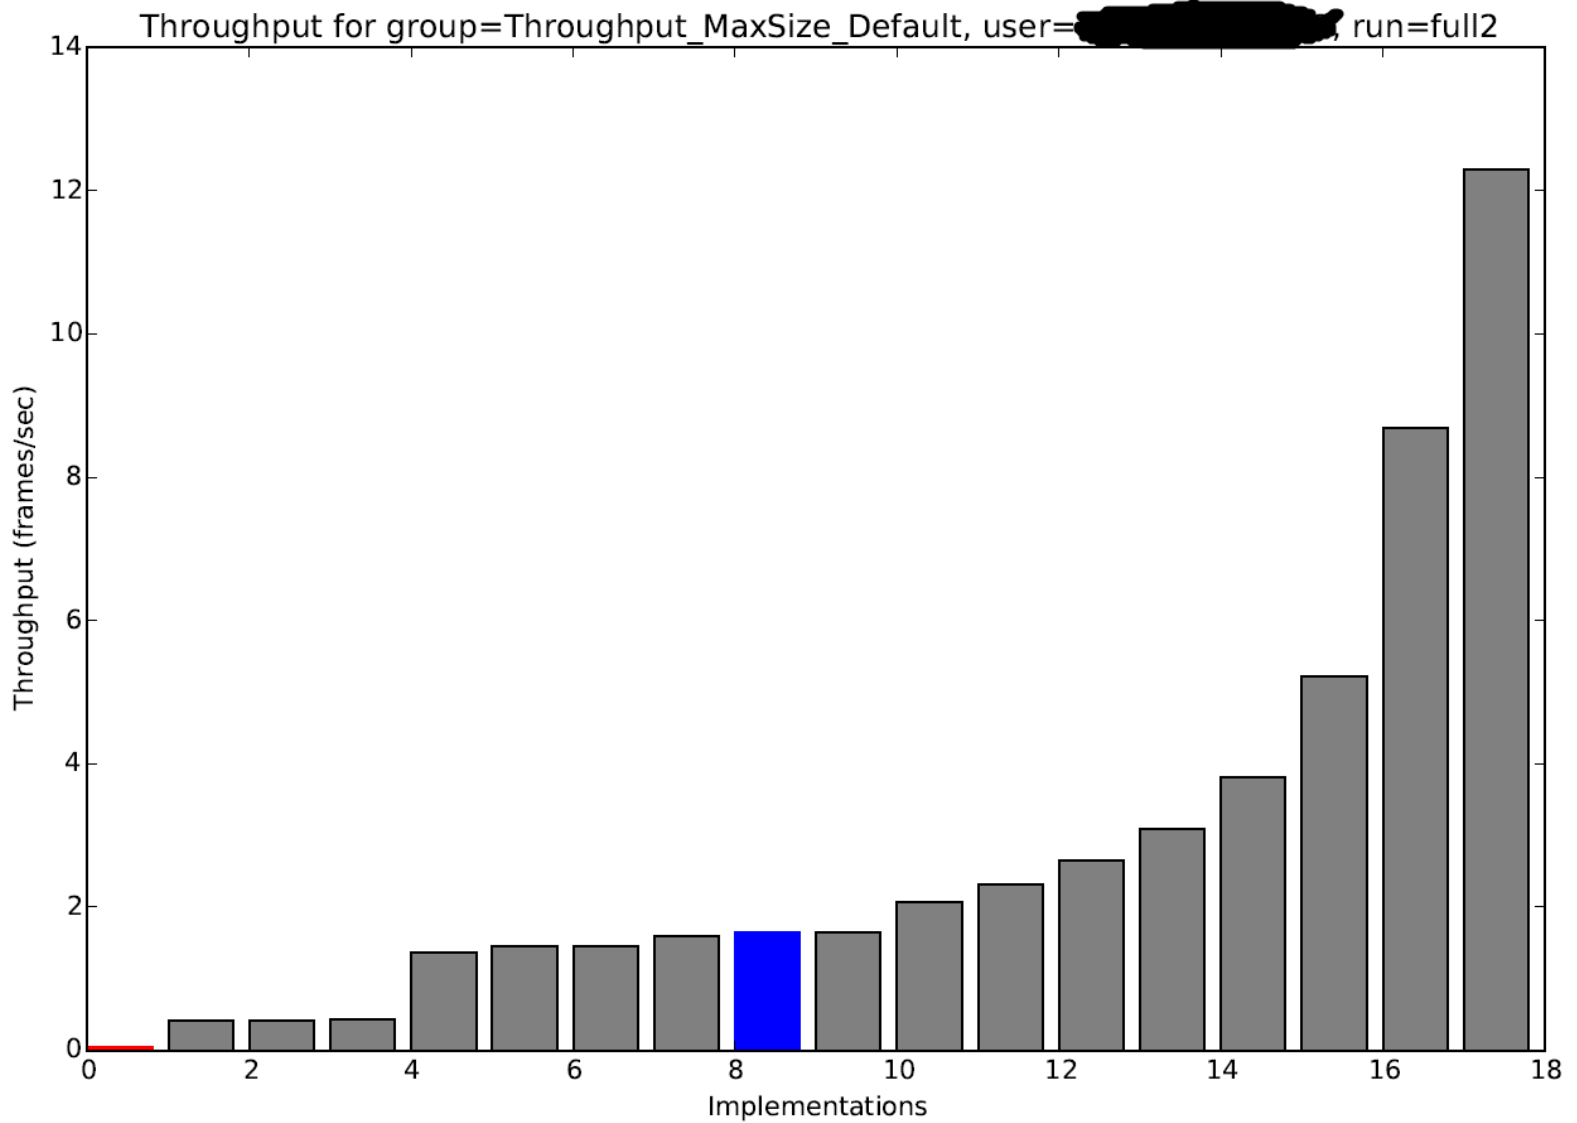
- Moved to 2-hour blocks (Bleh. Makes the timetable easier)
- Try to place lectures *within* first four courseworks
- Avoid exams at the end of term

# Feedback

- Feedback != grades
  - Feedback is formative : what worked, what is going well, ...
- 100% coursework isn't intended to give instant marks
  - But... it is supposed to have fast feedback
  - Should be fast enough to be useful during learning process

# Feedback

- Feedback != grades
  - Feedback is formative : what worked, what is going well, ...
- 100% coursework isn't intended to give instant marks
  - But... it is supposed to have fast feedback
  - Should be fast enough to be useful during learning process
- Looking at previous years:
  - Success : discussions with me + students via github issues and PRs
  - Success : feedback during CW5 and CW6 via git
  - Success : orals are a good point for reflection (students say so!)
  - Failure : timing of CW1-CW4. Too variable, takes too long



# Approach to CW1-CW4 feedback

- CW1-CW4 are ***supposed*** to be easy
  - Everyone should be able to get 100%
  - (as a consequence, CW5+CW6+oral are marked on wide range)

# Approach to CW1-CW4 feedback

- CW1-CW4 are ***supposed*** to be easy
  - Everyone should be able to get 100%
  - (as a consequence, CW5+CW6+oral are marked on wide range)
- Problem: assessment is mechanical but breaks
  - Student's code tends to fail in weird ways

# Approach to CW1-CW4 feedback

- CW1-CW4 are ***supposed*** to be easy
  - Everyone should be able to get 100%
  - (as a consequence, CW5+CW6+oral are marked on wide range)
- Problem: assessment is mechanical but breaks
  - Student's code tends to fail in weird ways
- Solution: enable self-assessment
  - Assessment scripts distributed in 2<sup>nd</sup> week of CW1-CW4
  - Students can run it locally and see how it fails
  - Can iterate on it till it works, get immediate result
  - If committed to git, it will get run remotely as well



# Approach to CW1-CW4 feedback

- CW1-CW4 are ***supposed*** to be easy
  - Everyone should be able to get 100%
  - (as a consequence, CW5+CW6+oral are marked on wide range)
- Problem: assessment is mechanical but breaks
  - Student's code tends to fail in weird ways
- Solution: enable self-assessment
  - Assessment scripts distributed in 2<sup>nd</sup> week of CW1-CW4
  - Students can run it locally and see how it fails
  - Can iterate on it till it works, get immediate result
  - If committed to git, it will get run remotely as well
- Formative feedback is on demand
  - Ask a question about submission on github
  - Ask a question in class

# Skills needed

- Basic programming
  - If you can't program in any language then worry

# Skills needed

- Basic programming
  - If you can't program in *\_any\_* language then worry
- Intel TBB uses C++ rather than C
  - Some weird C++ stuff, but not scary: *explained in lectures*
  - Setup and basics covered in coursework
- GPU programming uses OpenCL (C-like)
  - Let's you use whatever graphics card you happen to have
  - Working examples, explained in lectures
  - Language and compiler setup covered in coursework
- Not expected to become a guru, just make it faster

# Key Focus: Engineering

- How does this apply to you?
- Examples from Elec. Eng. problems
  - Mathematical analysis
  - Simulation of digital circuits
  - VLSI circuit layout
  - Communication channel evaluation
- Tools and languages used in EE
  - C / C++
  - MATLAB

# Course admin

- Slides on the course homepage
  - <https://github.com/HPCE/hpce-2016>
- Blackboard site is not used very much
  - (Why? Because I can automate git. No clicks)
- Other tools/sites we will be using
  - github : various forms of code distribution and submission
  - AWS (Amazon Web Services) for multi-core and GPUs later on
- Bring a device to lectures (laptop, tablet, charged phone)

# The almighty git

- Git (and github) is used extensively in this course
  - As a method of distributing information + coursework
  - As a means of communication and clarification (issues)
  - As a way to provide online feedback (pushes during CW)
  - To allow pair-working between students
  - As a way to submit code (for later courseworks)

# The almighty git

- Git (and github) is used extensively in this course
  - As a method of distributing information + coursework
  - As a means of communication and clarification (issues)
  - As a way to provide online feedback (pushes during CW)
  - To allow pair-working between students
  - As a way to submit code (for later courseworks)
- You don't need to know git already
  - It isn't that complicated anyway

# The almighty git

- Git (and github) is used extensively in this course
  - As a method of distributing information + coursework
  - As a means of communication and clarification (issues)
  - As a way to provide online feedback (pushes during CW)
  - To allow pair-working between students
  - As a way to submit code (for later courseworks)
- You don't need to know git already
  - It isn't that complicated anyway
- You **do** need a github account



# Platforms

- I don't care what platform/OS you use, as long as:
  - You have access to a bash-like command line
  - You have a fairly modern C++ compiler
  - There is more than one CPU
- Reasonable choices are:
  - Windows (*wsl* or *mingw* or vm+linux)
  - OS-X (using *brew* or *ports*)
  - Linux

# Platforms


- I don't care what platform/OS you use, as long as:
  - You have access to a bash-like command line
  - You have a fairly modern C++ compiler
  - There is more than one CPU
- Reasonable choices are:
  - Windows (*wsl* or *mingw* or *vm+linux*)
  - OS-X (using *brew* or *ports*)
  - Linux
- You are responsible for your platform
  - There is setup info in the coursework
  - I can help you, and you can help each other
  - Note: you can do dev on one platform, eval on another

# Platforms

- I don't care what platform/OS you use, as long as:
  - You have access to a bash-like command line
  - You have a fairly modern C++ compiler
  - There is more than one CPU
- Reasonable choices are:
  - Windows (*wsl* or *mingw* or *vm+linux*)
  - OS-X (using *brew* or *ports*)
  - Linux
- You are responsible for your platform
  - There is setup info in the coursework
  - I can help you, and you can help each other
  - Note: you can do dev on one platform, eval on another
- Eventually you will use AWS GPU instances
  - No GUI. Not even on this continent.

# Action

- If you want to take this course then:
  1. Get a github account
  2. Send me an email:
    - Subject: “[HPCE-github-request]”
    - Body: github id + your Imperial **login** (the short one)



# How do you do well in this course?

# Simple example : Totient function

- Eulers totient function:  $\text{totient}(n)$ 
  - Number of integers in range  $1..n$  which are relatively prime to  $n$
  - Integers  $i$  and  $j$  are relatively prime if  $\text{gcd}(i,j)=1$

# Version 0 : Simple loop

- Eulers totient function: totient(n)
  - Number of integers in range 1..n which are relatively prime to n
  - Integers i and j are relatively prime if  $\gcd(i,j)=1$

```
unsigned totient_v0(unsigned begin, unsigned end)
{
    unsigned count=0;


    for(unsigned i=begin; i<end; i++){
        count = count + gcd(i);
    }

    return count;
}
```

```
/vagrant/lec0
vagrant@debiancontrib-jessie /vagrant/lec0
$
vagrant@debiancontrib-jessie /vagrant/lec0
$ g++ -std=c++11 -o totient_v0 totient_v0.cpp
vagrant@debiancontrib-jessie /vagrant/lec0
$ ES=" 10 11 12 13 14 15 16 17 18";
vagrant@debiancontrib-jessie /vagrant/lec0
$ for e in $ES; do ./totient_v0 $e; done
e^10.000,      22026,      7340,  0.011545
e^11.000,      59874,     18752,  0.029131
e^12.000,     162754,     77076,  0.047173
e^13.000,     442413,    294936,  0.141477
e^14.000,    1202604,    369408,  0.384869
e^15.000,    3269017,    3264016,  1.241035
e^16.000,    8886110,    3367296,  3.366319
e^17.000,   24154952,   11599680,  9.706297
^C
vagrant@debiancontrib-jessie /vagrant/lec0
$ |
```



**v0**




# Turn on optimisation!

```
/vagrant/lec0
e^13.000,      442413,      294936,      0.141477
e^14.000,     1202604,     369408,      0.384869
e^15.000,     3269017,     3264016,      1.241035
e^16.000,     8886110,     3367296,      3.366319
e^17.000,    24154952,    11599680,      9.706297
^C

vagrant@debiancontrib-jessie /vagrant/lec0
$ g++ -O3 -std=c++11 -o totient_v0 totient_v0.cpp

vagrant@debiancontrib-jessie /vagrant/lec0
$ for e in $ES; do ./totient_v0 $e; done
e^10.000,      22026,       7340,      0.005334
e^11.000,      59874,      18752,      0.015950
e^12.000,     162754,      77076,      0.031679
e^13.000,      442413,      294936,      0.066495
e^14.000,     1202604,     369408,      0.180856
e^15.000,     3269017,     3264016,      0.574568
e^16.000,     8886110,     3367296,      1.570109
e^17.000,    24154952,    11599680,      4.524068
^C

vagrant@debiancontrib-jessie /vagrant/lec0
$
```



# Convert the for loop to parallel loop

```
#include "tbb/parallel_for.h"

uint64_t totient_v1(uint64_t n)
{
    uint64_t count=0;

    //for(uint64_t i=1; i<=n; i++){
    tbb::parallel_for(uint64_t(1), (n+1), [&](uint64_t i){
        if(gcd(i,n)==1){
            count = count + 1;
        }
    });

    return count;
}
```

# Compile with TBB


```
/vagrant/lec0
e^13.000,      442413,      294936,      0.066495
e^14.000,      1202604,      369408,      0.180856
e^15.000,      3269017,      3264016,      0.574568
e^16.000,      8886110,      3367296,      1.570109
e^17.000,      24154952,     11599680,      4.524068
^C

vagrant@debiancontrib-jessie /vagrant/lec0
$ g++ -O3 -std=c++11 -o totient_v1 totient_v1.cpp -ltbb

vagrant@debiancontrib-jessie /vagrant/lec0
$ for e in $ES; do ./totient_v1 $e; done
e^10.000,      22026,       6306,       0.002529
e^11.000,      59874,      16358,      0.005598
e^12.000,      162754,      62800,      0.008893
e^13.000,      442413,      227490,      0.023908
e^14.000,      1202604,      324572,      0.048735
e^15.000,      3269017,      2305111,      0.165825
e^16.000,      8886110,      3038523,      0.443429
e^17.000,      24154952,     10154214,      1.253405
e^18.000,      65659969,     48895328,      3.917649

vagrant@debiancontrib-jessie /vagrant/lec0
$ |
```

# Faster, but...



# Faster but wrong ☹

```
/vagrant/lec0
e^13.000, 442413, 294936, 0
e^14.000, 1202604, 369408, 0
e^15.000, 3269017, 3264016, 1
e^16.000, 8886110, 3367296, 3
e^17.000, 24154952, 11599680, 9
^C

vagrant@debiancontrib-jessie /vagrant
$ g++ -O3 -std=c++11 -o totient_v0 totient_v0.cpp -ltbb

vagrant@debiancontrib-jessie /vagrant
$ for e in $ES; do ./totient_v0 $e; done
e^10.000, 22026, 7340, 0
e^11.000, 59874, 18752, 0
e^12.000, 162754, 77076, 0
e^13.000, 442413, 294936, 0
e^14.000, 1202604, 369408, 0
e^15.000, 3269017, 3264016, 0
e^16.000, 8886110, 3367296, 1
e^17.000, 24154952, 11599680, 4
^C

vagrant@debiancontrib-jessie /vagrant
$

/vagrant/lec0
e^13.000, 442413, 294936, 0.066495
e^14.000, 1202604, 369408, 0.180856
e^15.000, 3269017, 3264016, 0.574568
e^16.000, 8886110, 3367296, 1.570109
e^17.000, 24154952, 11599680, 4.524068
^C

vagrant@debiancontrib-jessie /vagrant/lec0
$ g++ -O3 -std=c++11 -o totient_v1 totient_v1.cpp -ltbb

vagrant@debiancontrib-jessie /vagrant/lec0
$ for e in $ES; do ./totient_v1 $e; done
e^10.000, 22026, 6306, 0.002529
e^11.000, 59874, 16358, 0.005598
e^12.000, 162754, 62800, 0.008893
e^13.000, 442413, 227490, 0.023908
e^14.000, 1202604, 324572, 0.048735
e^15.000, 3269017, 2305111, 0.165825
e^16.000, 8886110, 3038523, 0.443429
e^17.000, 24154952, 10154214, 1.253405
e^18.000, 65659969, 48895328, 3.917649

vagrant@debiancontrib-jessie /vagrant/lec0
$
```

# Unsafe use of shared variable

```
#include "tbb/parallel_for.h"

uint64_t totient_v1(uint64_t n)
{
    uint64_t count=0;

    //for(uint64_t i=1; i<=n; i++){
    tbb::parallel_for(uint64_t(1), (n+1), [&](uint64_t i){
        if(gcd(i,n)==1){
            count = count + 1;
        }
    });

    return count;
}
```




# Make it atomic

```
uint64_t totient_v2(uint64_t n)
{
    std::atomic<uint64_t> count;


    //for(uint64_t i=1; i<=n; i++){
    tbb::parallel_for(uint64_t(1), (n+1), [&](uint64_t i){
        if(gcd(i,n)==1){
            count += 1;
        }
    });

    return count;
}
```

# Fast and correct



# Speedup (4 CPU machine)



# Initial Lessons

- Speeding up loops *can* be easy
- Need to watch out for shared variables
- The speedup in  $P$  cores is less than  $P$

# Reminder Action

- If you want to take this course then:
  1. Get a github account
  2. Send me an email:
    - Subject: “[HPCE-github-request]”
    - Body: github id + your Imperial **login** (the short one)

