

Course admin stuff

- Coursework 1 is under way
 - 80 people (!) signed up in github
- Where to find it
 - Spec is on github: <https://github.com/HPCE/hpce-2016-cw1>
 - Submission for this coursework is via blackboard
 - Submission is open

Expectations for coursework

- Coursework is not lab [1]
 - You have to manage when, where, and how long you spend on it
 - 100% coursework does not mean easy [1]
- Long hours are neither sufficient ***nor*** necessary for an A
 - Like anything else: some people are just good at it
 - But, a good correlation between organisation and marks
- You are expected to be reasonably independent
 - This is a masters level course

[1] – Though the earlier parts kind of are.

Working together

- The software community has a tradition of sharing
 - Many open-source projects, some of which you will rely on
 - Lots of forums for discussing problems: stack-overflow, ...
- Approach this work in the same way
 - You may encounter the same problems as other students
 - Discuss solutions with each other, help each other out
 - One-on-one discussions, github issues, whatever
 - <https://github.com/HPCE/hpce-2015-cw1/blob/master/background-bugs.md>
 - Give credit or thanks if appropriate: be excellent to each other
- But you have to balance co-operation and competition
 - The later courseworks require good ideas and strategies
 - Up to you to protect your IP.

Plagiarism


- All submitted material must be written by you
 - Do not share any code with each other (except within pairs)
- No plagiarism checking software: I just read the code
 - Some similarity of structure is expected, but there are limits
 - Students are amusingly bad at obfuscation
 - You need to be able to explain any code you submit in the oral
 - Suspected plagiarism will be passed to the plagiarism committee
- If necessary you **may** use code from third-party sources
 - e.g. open-source projects, samples, stack overflow, ...
 - Origin and extent must be very clearly shown
 - Need to be able to justify (orally) why it was used
 - Should be aware of potential licensing implications

More practical: image processing


- *Gamma correction*: adjust light/dark
 - $p_{[x,y]} = \text{pow}(p_{[x,y]}, \text{gamma})$



$\gamma = 2$



$\gamma = 1$



$\gamma = 0.5$

```

void process_frame(
    float gamma,
    unsigned width, unsigned height,
    const uint8_t *frameIn,
    uint8_t *frameOut
){
    for(unsigned x= 0u; x<width; x++){
        for(unsigned y=0; y<height; y++){
            double fIn = frameIn[y*width + x] * (1.0/256.0);

            double fOut = pow( fIn, gamma );

            frameOut[y*width + x] = (uint8_t)floor(fOut * 256.0);
        }
    }
}

```

```

void process_frame(
    float gamma,
    unsigned width, unsigned height,
    const uint8_t *frameIn,
    uint8_t *frameOut
){
    tbb::parallel_for(0u, width, [&](unsigned x){
        for(unsigned y=0; y<height; y++){
            double fIn = frameIn[y*width + x] * (1.0/256.0);

            double fOut = pow( fIn, gamma );

            frameOut[y*width + x] = (uint8_t)floor(fOut * 256.0);
        }
    });
}

```

f is a variable, but we let compiler decide on its type

Capture variables by reference
(can modify outer variables)

Lambda parameters,
just like function parameter.

```
void process_frame(  
    float gamma,  
    unsigned width, unsigned height,  
    const uint8_t *frameIn,  
    uint8_t *frameOut  
) {  
    auto f = [&](unsigned x) {  
        for(unsigned y=0; y<height; y++){  
            double fIn = frameIn[y*width + x] * (1.0/256.0);  
  
            double fOut = pow( fIn, gamma );  
  
            frameOut[y*width + x] = (uint8_t)floor(fOut * 256.0);  
        }  
    };  
  
    tbb::parallel_for(0u, width,  
        f  
    );  
}
```



```

void process_frame(
    float gamma,
    unsigned width, unsigned height,
    const uint8_t *frameIn,
    uint8_t *frameOut
){
    auto f = [&](unsigned x){
        for(unsigned y=0; y<height; y++){
            double fIn = frameIn[y*width + x] * (1.0/256.0);


            double fOut = pow( fIn, gamma );

            frameOut[y*width + x] = (uint8_t)floor(fOut * 256.0);
        }
    };

    for(unsigned i= 0u; i<width; i++){
        f(i);
    }
}

```

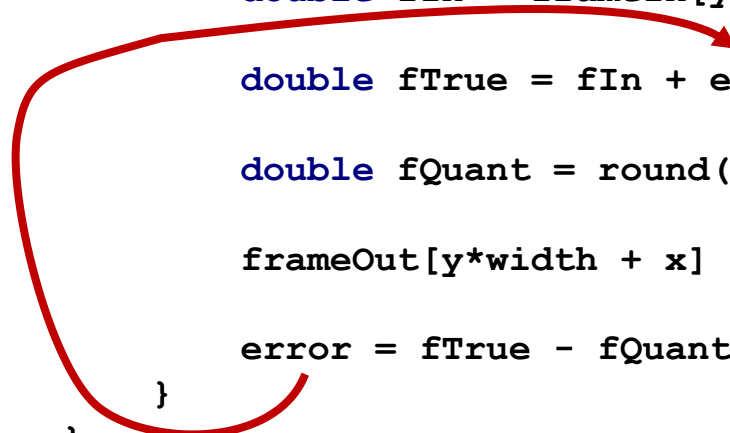
Quantisation via dithering



- Dithering: cumulative error due to quantisation is tracked

Quantisation via error dithering

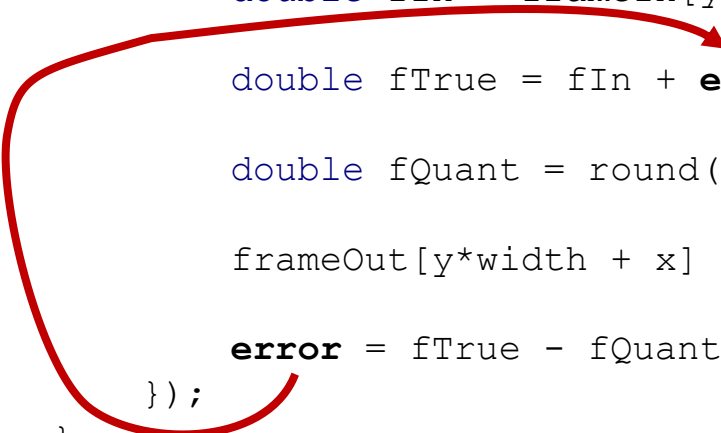
```
void process_frame(  
    unsigned levels,  
    unsigned width, unsigned height,  
    const uint8_t *frameIn,  
    uint8_t *frameOut  
)  
{  
    for(unsigned x=0; x<width; x++){  
        double error=0.0;  
        for(unsigned y=0; y<height; y++){  
            double fIn = frameIn[y*width + x] * (1.0/255.0);  
  
            double fTrue = fIn + error;  
  
            double fQuant = round( fTrue * levels ) / levels;  
  
            frameOut[y*width + x] = (uint8_t)floor(fQuant * 255.0);  
  
            error = fTrue - fQuant;  
        }  
    }  
}
```



Loop carried dependency through error

Parallelising the inner loop

```
void process_frame(  
    unsigned levels,  
    unsigned width, unsigned height,  
    const uint8_t *frameIn,  
    uint8_t *frameOut  
)  
{  
    for(unsigned x=0; x<width; x++){  
        double error=0.0;  
        tbb::parallel_for(0u, height, [&](unsigned y){  
            double fIn = frameIn[y*width + x] * (1.0/255.0);  
  
            double fTrue = fIn + error;  
  
            double fQuant = round( fTrue * levels ) / levels;  
  
            frameOut[y*width + x] = (uint8_t)floor(fQuant * 255.0);  
  
            error = fTrue - fQuant;  
        });  
    }  
}
```



Parallelising the outer loop

```
void process_frame(
    unsigned levels,
    unsigned width, unsigned height,
    const uint8_t *frameIn,
    uint8_t *frameOut
)
{
    tbb::parallel_for(0u, width, [&](unsigned x){
        double error=0.0;
        for(unsigned y=0; y<height; y++){
            double fIn = frameIn[y*width + x] * (1.0/255.0);

            double fTrue = fIn + error;

            double fQuant = round( fTrue * levels ) / levels;

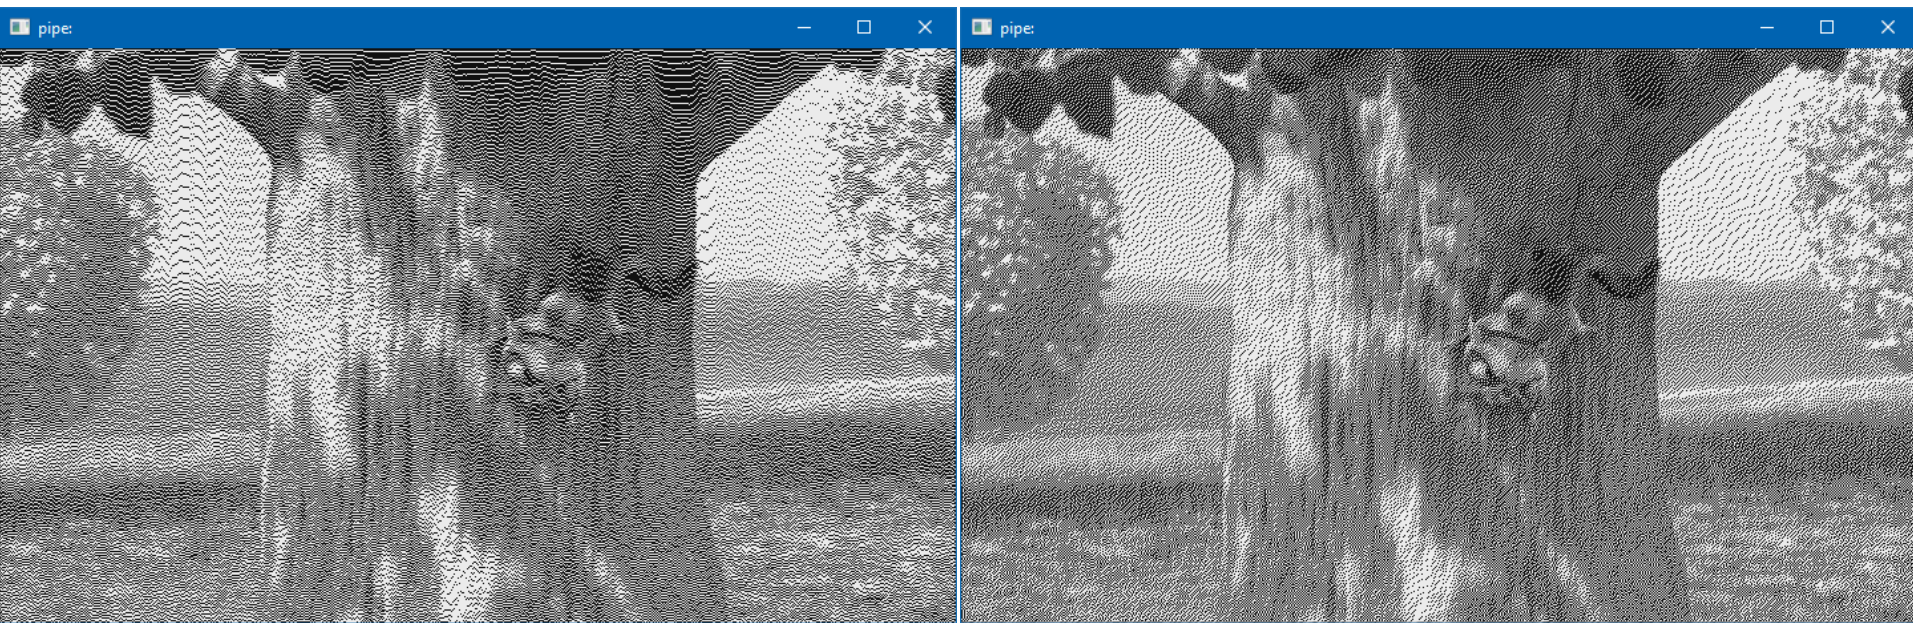
            frameOut[y*width + x] = (uint8_t)floor(fQuant * 255.0);

            error = fTrue - fQuant;
        }
    });
}
```

A solution for the inner loop?

```
void process_frame(  
    unsigned levels,  
    unsigned width, unsigned height,  
    const uint8_t *frameIn,  
    uint8_t *frameOut  
)  
{  
    std::vector<double> error(height, 0.0);  
  
    for(unsigned x=0; x<width; x++){  
        tbb::parallel_for(0u, height, [&](unsigned y){  
            double fIn = frameIn[y*width + x] * (1.0/255.0);  
            assert( (fIn>=0) && (fIn<=1.0));  
  
            double fTrue = fIn + error[y];  
  
            double fQuant = round( fTrue * levels ) / levels;  
  
            frameOut[y*width + x] = (uint8_t)floor(fQuant * 255.0);  
  
            error[y] = fTrue - fQuant;  
        });  
    }  
}
```


2D error diffusion



- Attempt to diffuse error both across and down image
- Reduce tendency towards banding effects

```


void process_frame(
    unsigned levels, unsigned width, unsigned height,
    double *frame
)
{
    for(unsigned x=0; x<width-1; x++){
        for(unsigned y=0; y<height-1; y++){
            double fIn = frame[y*width + x];

            double fQuant = round( fIn * levels ) / levels;
            frame[y*width + x] = fQuant;

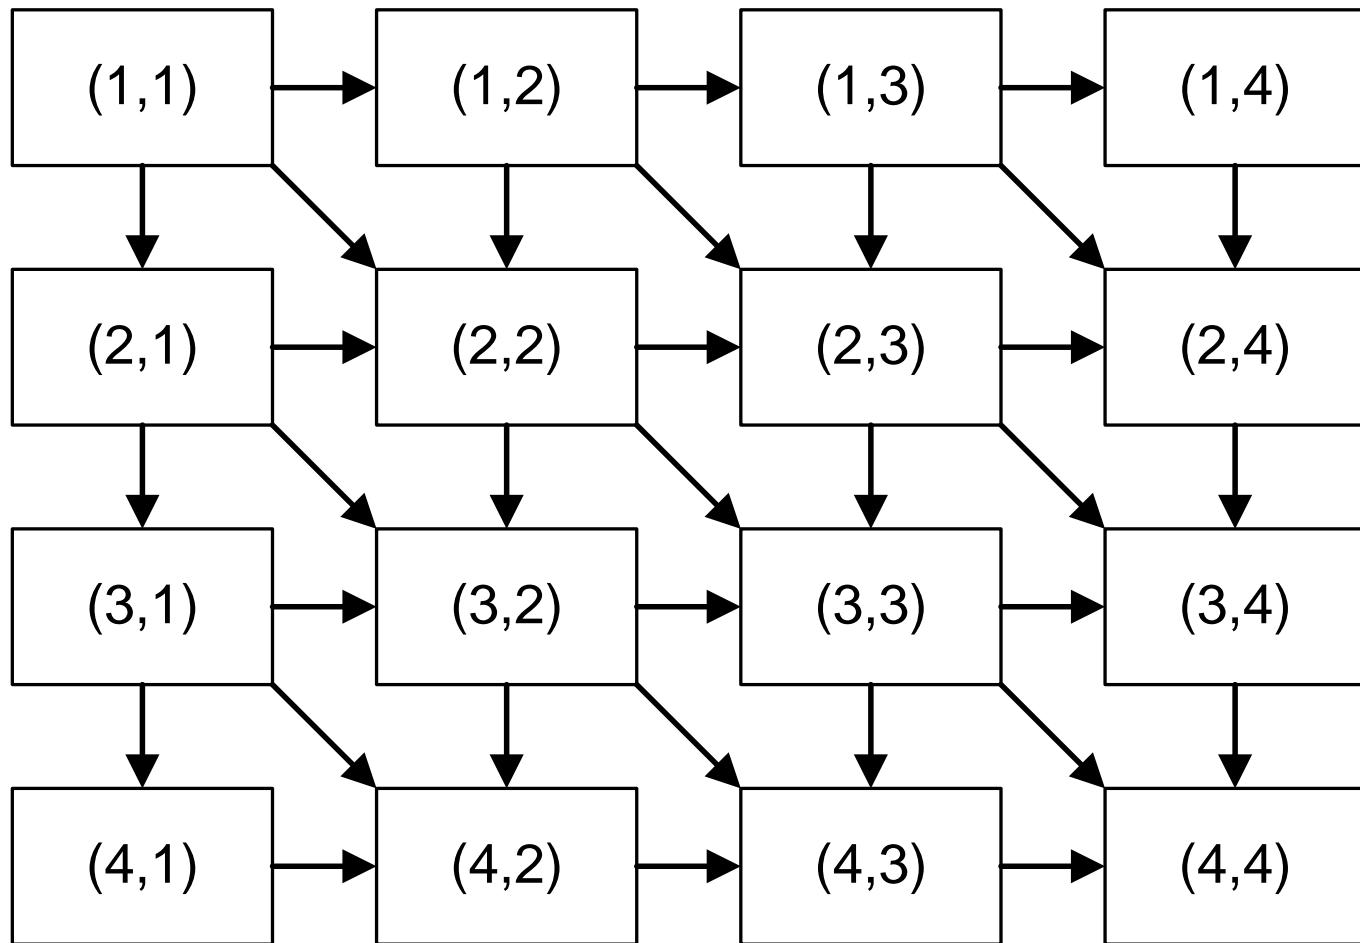
            double error = fIn - fQuant;
            frame[  y      * width + x+1 ] += error * 0.4;
            frame[ (y+1) * width + x  ] += error * 0.4;
            frame[ (y+1) * width + x+1 ] += error * 0.2;
        }
    }
}

```

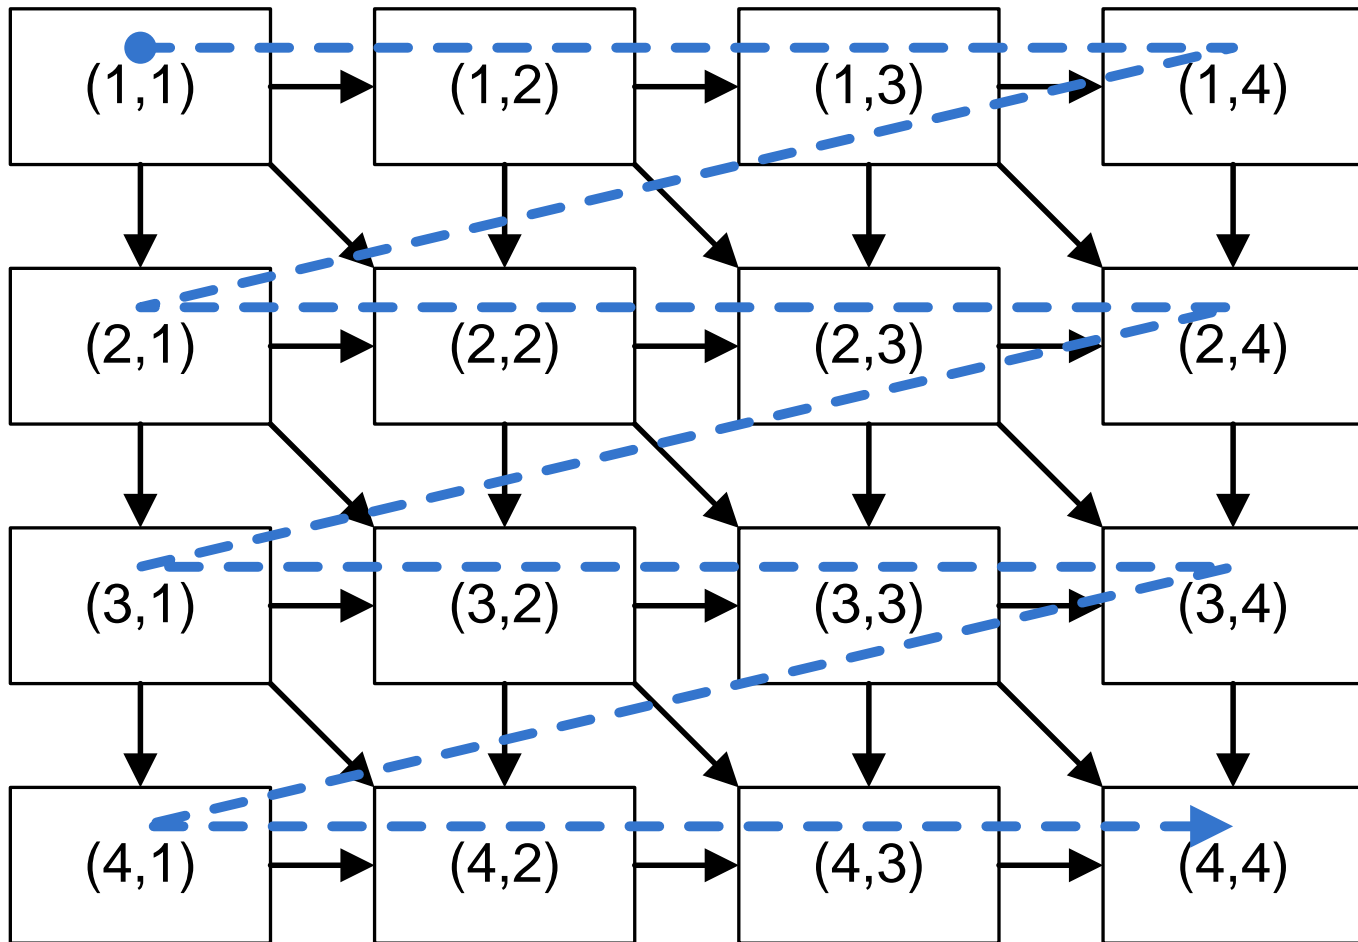
- More difficult loop carried dependency
- Have a write before read dependency
 - Current loop iteration reads from (x,y)
 - Writes (x,y+1), (x+1,y), and (x+1,y+1)
 - Three constraints per node




Generalise to the full grid




Serial execution: y then x




Serial execution: x then y




Parallelisation: can't do it along x




Parallelisation: can't do it along y



Skewing the loops




Or viewed another way




Iteration spaces

- We already have iteration as a primitive
 - for loops : bounded iteration over known range
 - while loops : possibly unbounded iteration (though maybe not)
- Iteration spaces assign unique labels to distinct iterations




Transforming iteration spaces

- We now want to map (x,y) to a new iteration space $[i,j]$
 - A different set of loop variables that expose parallel operations



Transforming iteration spaces

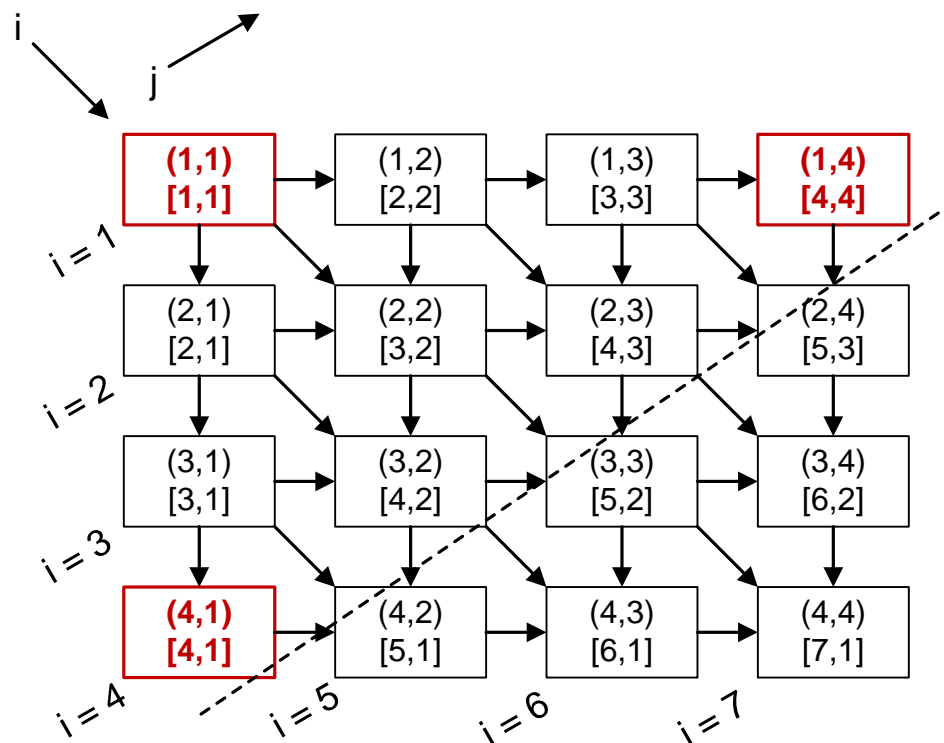
- Mapping must be distinct: all (x,y) go to a different $[i,j]$



Solving for the equations

- Mapping must be distinct: all (x,y) go to a different $[i,j]$

```
for(unsigned i=0; i<n-1; i++){  
    for(int j=i; j>=0; j--){  
        unsigned x=i;  
        unsigned y=i-j;
```



Back to the code!

```
void process_frame(unsigned levels, unsigned width, unsigned height, double *frame)
{
    unsigned n=std::min(width,height);

    for(unsigned i=0; i<n-1; i++){
        for(int j=i; j>=0; j--){
            unsigned x=i;
            unsigned y=i-j;

            double fIn = frame[y*width + x];

            double fQuant = round( fIn * levels ) / levels;

            frame[y*width + x] = fQuant;

            double error = fIn - fQuant;
            frame[ y      * width + x+1] += error * 0.4;
            frame[ (y+1) * width + x  ] += error * 0.4;
            frame[ (y+1) * width + x+1] += error * 0.2;
        }
    }
}
```

Note: this only handles *half* the iteration space

Parallel! Correct?

```
void process_frame(unsigned levels, unsigned width, unsigned height, double *frame) {
    unsigned n=std::min(width,height);

    for(unsigned i=0; i<n-1; i++){
        tbb::parallel_for(0u, i+1, [&](unsigned rev_j){
            int j=i-(int)rev_j;
            unsigned x=i;
            unsigned y=i-j;

            double fIn = frame[y*width + x];

            double fQuant = round( fIn * levels ) / levels;

            frame[y*width + x] = fQuant;

            double error = fIn - fQuant;
            frame[ y      * width + x+1] += error * 0.4;
            frame[ (y+1) * width + x  ] += error * 0.4;
            frame[ (y+1) * width + x+1] += error * 0.2;
        });
    }
}
```

Note: this only handles *half* the iteration space