

Recap: Overview of task groups

- A task group collects together a number of child tasks
 - The task creating the group is called the parent
 - One or more child tasks are created and `run()` by the parent
 - Child tasks **may** execute in parallel
 - Parent task must `wait()` for all child tasks before returning
- Some important differences between tasks and threads
 - Threads **must** execute in parallel
 - A thread may continue after its creator exits
 - Threads must be joined individually

parallel_for using tbb::task_group

```
#include "tbb/task_group.h"

template<class TI, class TF>
void parallel_for(const TI &begin, const TI &end, const TF &f)
{
    if(begin+1 == end){
        f(begin);
    }else{
        auto left=[&]() { parallel_for(begin, (begin+end)/2, f); }
        auto right=[&]() { parallel_for((begin+end)/2, end, f); }

        // Spawn the two tasks in a group
        tbb::task_group group;
        group.run(left);
        group.run(right);

        group.wait(); // Wait for both to finish
    }
}
```

More patterns: `tbb::parallel_invoke`

```
template<typename Func0, typename Func1>
void parallel_invoke(const Func0& f0, const Func1& f1);

template<typename Func0, typename Func1, typename Func2>
void parallel_invoke(const Func0& f0, const Func1& f1, const Func2& f2);
```

- Takes two or more functions and ***may*** run in parallel
 - Overloaded for different numbers of arguments
 - No overload for 1 argument for obvious reasons
- Interface is very clean, but also quite simple
 - Decision about number of tasks is completely static
 - You can't add more tasks once some starts
 - No choice about when to synchronise with tasks

parallel_invoke using task_group

- parallel_invoke can be implement using task_group
 - task_group supports a super-set of the functionality

```
template<typename Fc0, typename Fc1, typename Fc2>
void parallel_invoke(const Fc0& f0, const Fc1& f1, const Fc2& f2)
{
    tbb::task_group group;
    group.run(f0);
    group.run(f1);
    group.run(f2);
    group.wait();
}
```

Can't^[1] do task_group using parallel_invoke

- task_group is intrinsically dynamic
 - Decide how much work to add at run-time
 - Can add work even while tasks are running in the group

```
void my_function(int n, float *x)
{
    tbb::task_group group;
    for(unsigned i=0;i<n;i++){
        if(x[i]==0)
            group.run([=]() { f(i); });
        else
            group.run([=]() { g(x[i]); });
    }
    group.wait();
}
```

More complex loop nests

```
std::vector<uint8_t> process_frame(  
    int n, int w, int h,  
    std::vector<uint8_t> fIn // Receive frame (by value)  
) {  
    // Handle n==0 case  
    std::vector<uint8_t> fOut=fIn;  
  
    for(int i=1; i<n; i++){  
        fIn = fOut;  
  
        for(int y=1; y < h-1; y++){  
            for(int x=1; x < w-1; x++){  
                uint8_t nhood [5] = {  
                    fIn[(y-1)*w+x],  
                    fIn[y*w+(x-1)], fIn[ y      *w+x], fIn[y*w+(x+1)],  
                    fIn[(y+1)*w+x]  
                };  
                uint8_t value = min_of_array(5, nhood);  
                fOut[y*w+x] = value;  
            }  
        }  
    }  
    return fOut;  
}
```

How do you parallelise?

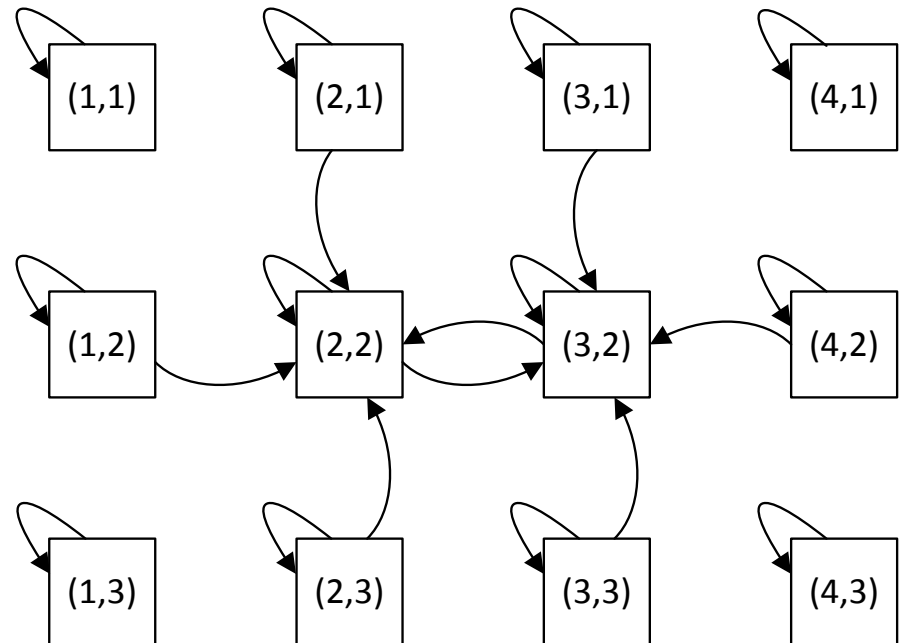
```


uint8_t nhoud[5] = {
    fIn[ (y-1)*w+x] ,
    fIn[ y*w+(x-1) ] , fIn[ (y+0)*w+x] , fIn[ y*w+(x+1) ] ,
    fIn[ (y+1)*w+x]
};


uint8_t value = min_of_array(5, nhoud);


fOut[y*w+x] = value;


```

















```

for(int i=1; i<n; i++){
    fIn = fOut;

    for(int y=1; y < h-1; y++){
        for(int x=1; x < w-1; x++){
            uint8_t nhoud [5] = {
                fIn[(y-1)*w+x] ,
                fIn[y*w+(x-1)] , fIn[ y      *w+x] , fIn[y*w+(x+1)] ,
                fIn[(y+1)*w+x]
            };
            uint8_t value = min_of_array(5, nhoud);
            fOut[y*w+x] = value;
        }
    }
}

```

- Can vectorise over both x and y
- Intermediate variable tmp saves us

Parallelising over a 2d space

```
for(int i=0; i<n; i++){
    fIn = fOut;

    tbb::blocked_range2d<int> r( 1,h-1, 1,w-1 );

    tbb::parallel_for( r, [&](const tbb::blocked_range2d<int> &xy) {
        for(int y=xy.rows().begin(); y < xy.rows().end(); y++){
            for(int x=xy.cols().begin(); x < xy.cols().end(); x++){
                uint8_t nhoud [5] = {
                    fIn[(y-1)*w+x],
                    fIn[y*w+(x-1)], fIn[ y      *w+x], fIn[y*w+(x+1)],
                    fIn[(y+1)*w+x]
                };

                uint8_t value = min_of_array(5, nhoud);

                fOut[y*w+x] = value;
            }
        }
    });
}
```

Outer loop?

```
for(int i=0; i<n; i++){
    fIn = fOut;

    tbb::blocked_range2d<int> r( 1,h-1, 1,w-1 );

    tbb::parallel_for( r, [&](const tbb::blocked_range2d<int> &xy) {
        for(int y=xy.rows().begin(); y < xy.rows().end(); y++){
            for(int x=xy.cols().begin(); x < xy.cols().end(); x++){
                uint8_t nhoud [5] = {
                    fIn[(y-1)*w+x],
                    fIn[y*w+(x-1)], fIn[ y      *w+x], fIn[y*w+(x+1)],
                    fIn[(y+1)*w+x]
                };

                uint8_t value = min_of_array(5, nhoud);

                fOut[y*w+x] = value;
            }
        }
    });
}
```

- Strict loop carried dependency
- Pure cyclic chain
 - Impossible to break
- No parallelism...?
- We'll come back to it

The underlying primitive: `tbb::task`

- TBB has a basic primitive called `tbb::task`
- This is the raw unit of scheduling understood by the lib.
 - Other high-level wrappers create tasks internally
 - The TBB run-time takes tasks and schedules them to a CPU
- Tasks are very flexible, with a lot of power
 - Can express complicated dependency graphs
 - Build non-local synchronisation and barriers
- With power comes responsibility
 - They allow you to make mistakes
 - Possible (though not likely) to mess up the TBB run-time
- Better to create wrappers on top that hide tasks
 - `parallel_for`, `parallel_invoke`, `task_group`, `parallel_reduce`, ...

```

class MyTask
: public tbb::task
{
    int start, end;

    MyTask(int _start, int _end)
    { start=_start; end=_end; }

    tbb::task * execute()
    {
        if(cond())
            return 0;
        set_ref_count(3);
        MyTask &t1=*new(allocate_child()) MyTask(start, (start+end)/2);
        spawn(t1);
        MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
        spawn(t2);
        DoSomethingFirst();
        wait_for_all();
        DoSomethingElse();
    }
};

```

```

void MyTask(int start, int end)
{
    if(cond())
        return 0;
    tbb::task_group group;
    group.run([=]() {MyTask(start, (start+end)/2); });
    group.run([=]() {MyTask((start+end)/2, end); });
    DoSomethingFirst();
    group.wait();
    DoSomethingElse();
    return 0;
}

```

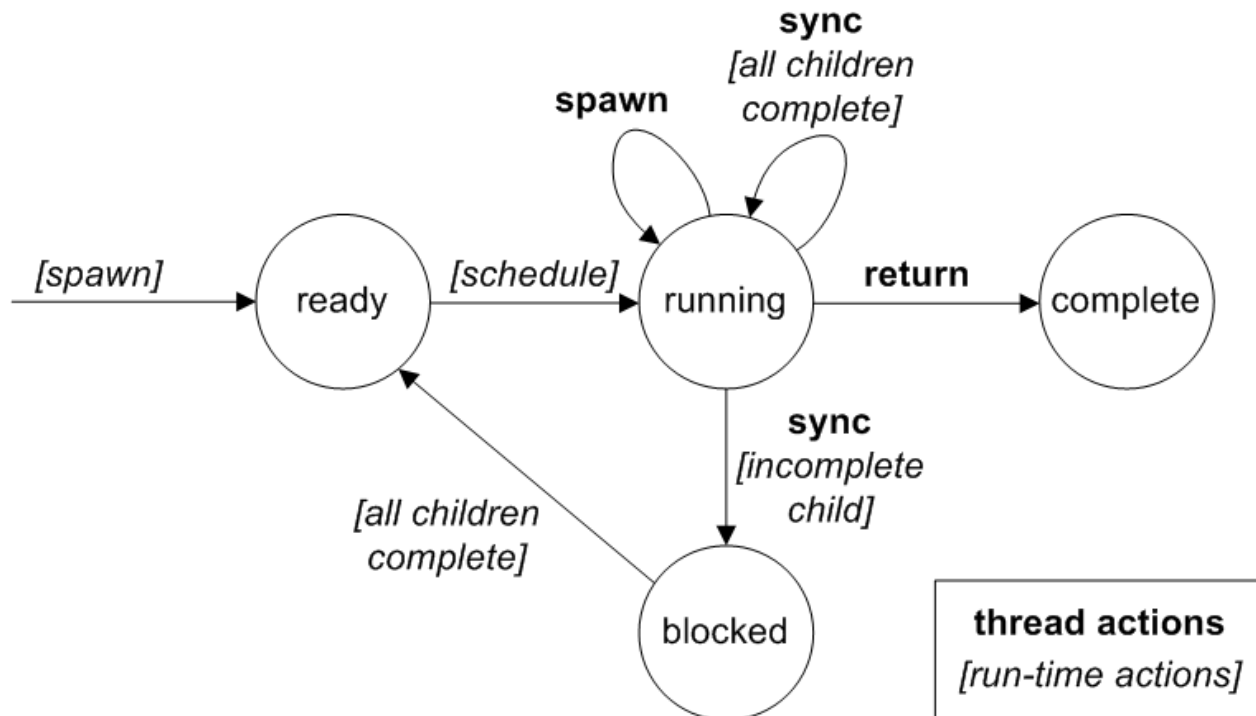
```

void CreateTasks(int start, int end)
{
    MyTask &root=*new(allocate_root()) MyTask(start, end);
    tbb::task::spawn_root_and_wait();
}

```

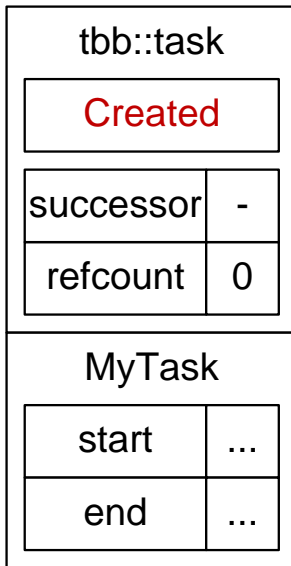
Life-cycle of a task

- Life-cycle of task due to interaction between task and run-time
 - Individual task calls **spawn**, **wait_for_all (sync)**, **return**
 - TBB run-time will keep track of a task's children (*dependencies*)

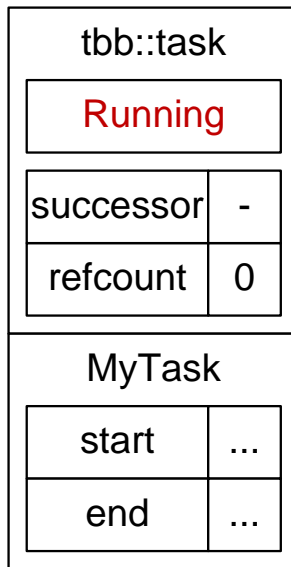


Scheduling through reference counts

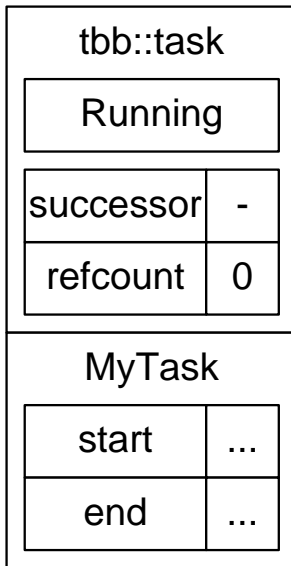
- Each task has a reference count and a successor task
- The reference count identifies whether a task is blocked
 - If the reference count is zero then the task could be run
 - But only if it has been given to the task scheduler
 - Legal to create a task and not give it to the scheduler
 - *Note the difference: “reference count” vs “C++ reference”*
- Successor task identifies the task blocked by this task
 - Generally the successor task is the creator, or parent
 - When a task completes it decrements the count of its successor




```
void CreateTasks(int start, int end)
{
    MyTask &root=*new(allocate_root()) MyTask(start,end);
    tbb::task::spawn_root_and_wait();
}
```



```
void CreateTasks(int start, int end)
{
    MyTask &root=*new(allocate_root()) MyTask(start,end);
    tbb::task::spawn_root_and_wait();
}
```



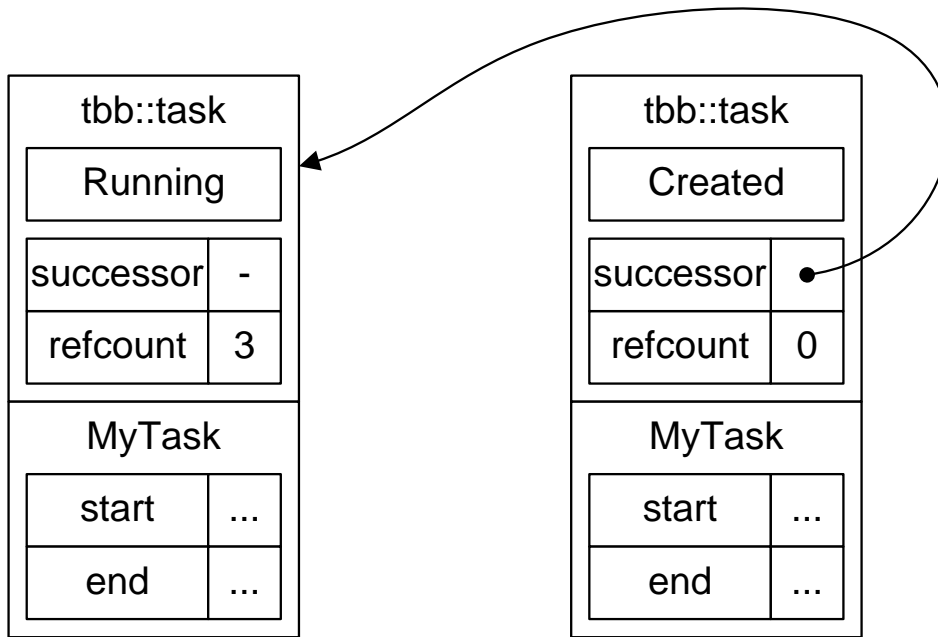
```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start, (start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```



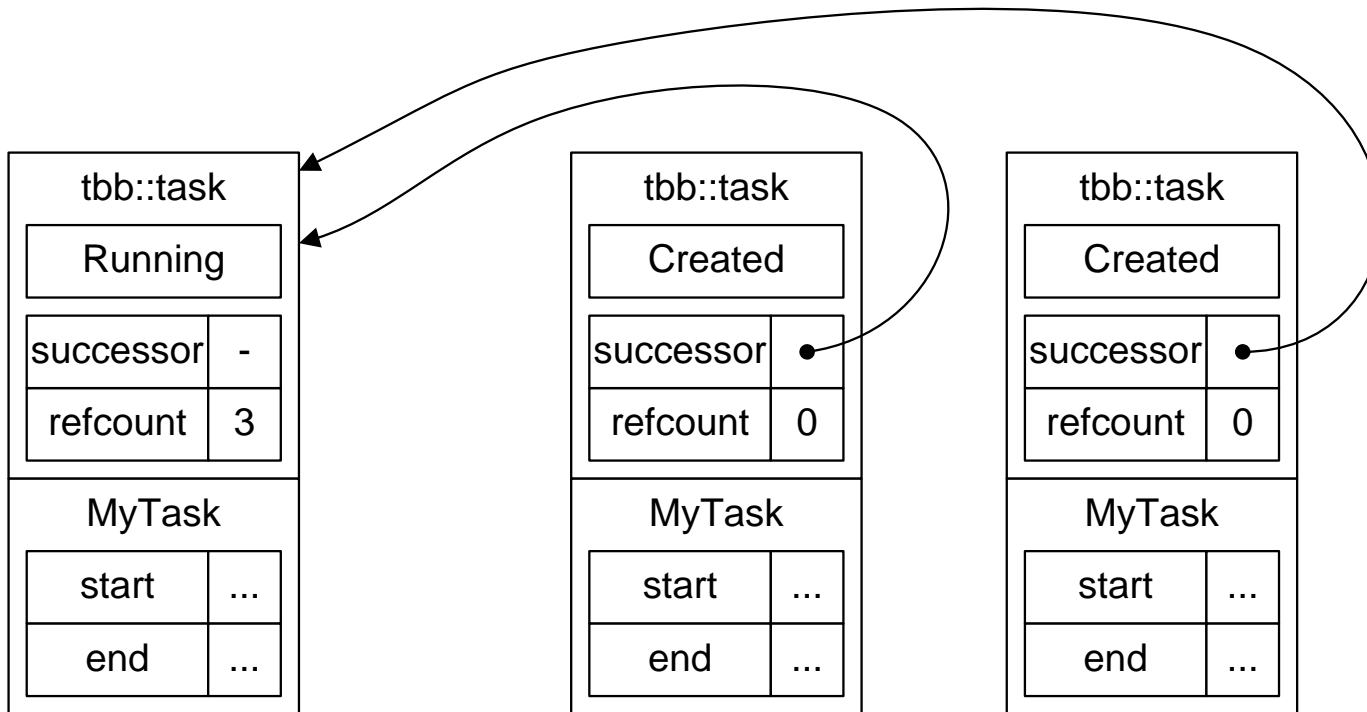
```

tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start,(start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}


```

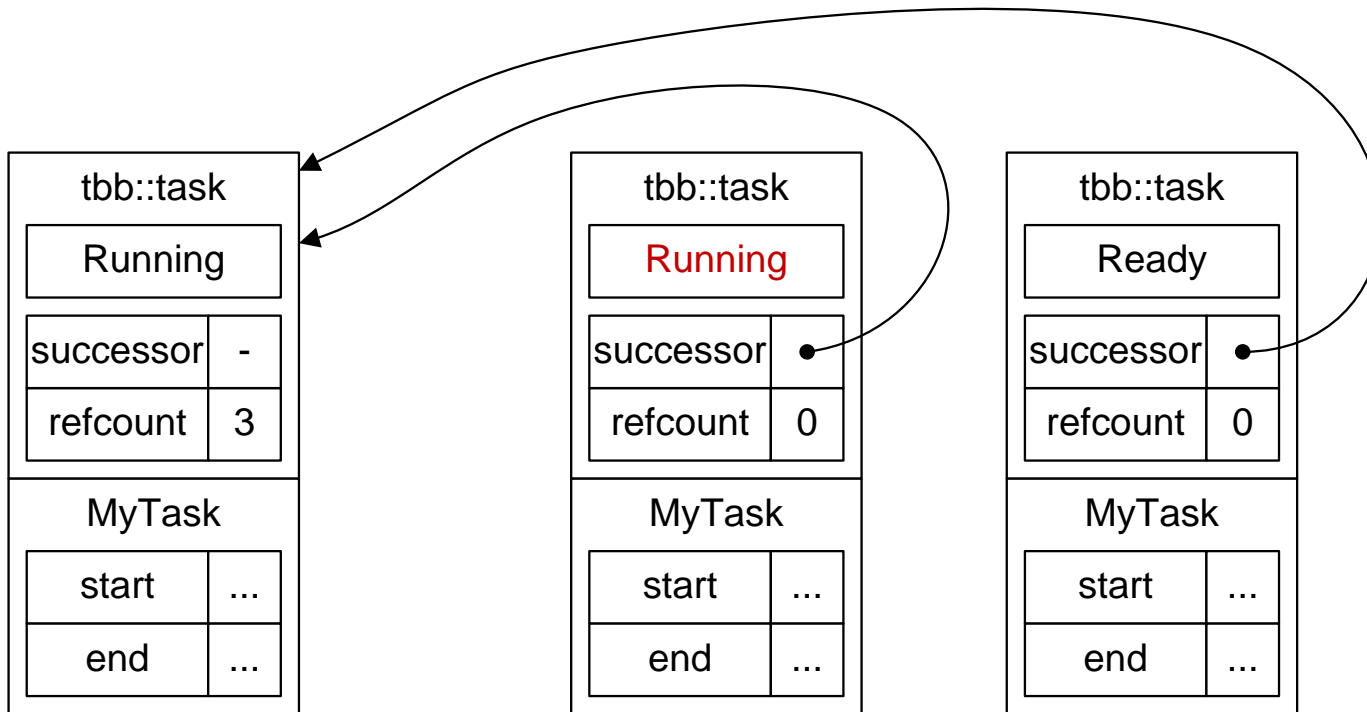
```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start, (start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```




```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start,(start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```




```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start,(start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```




```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start,(start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```




```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start,(start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```



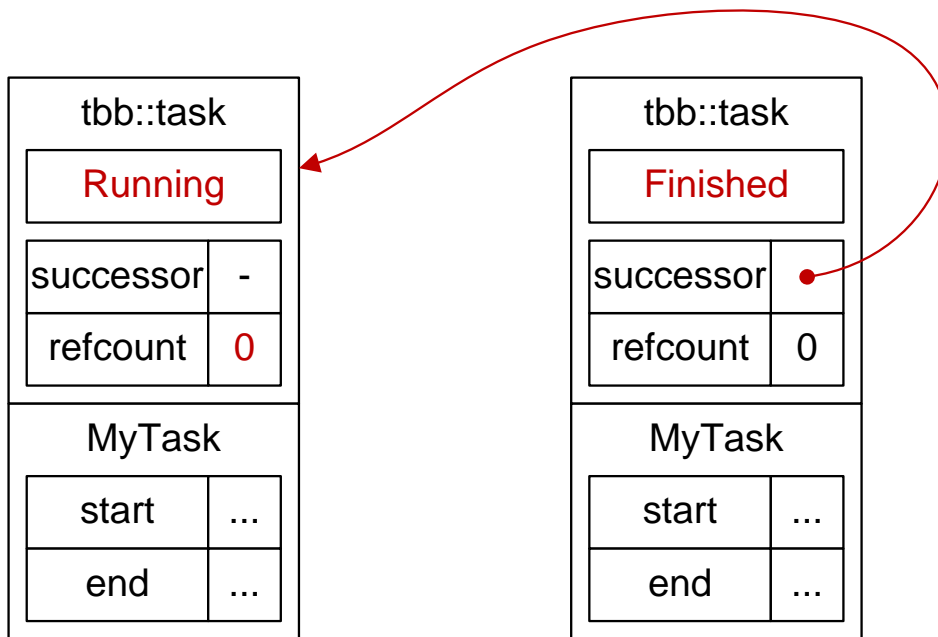
```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start,(start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```




```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start, (start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```




```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start,(start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```

```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start, (start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```



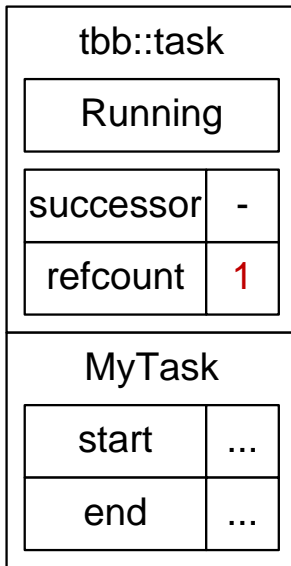
```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start,(start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```



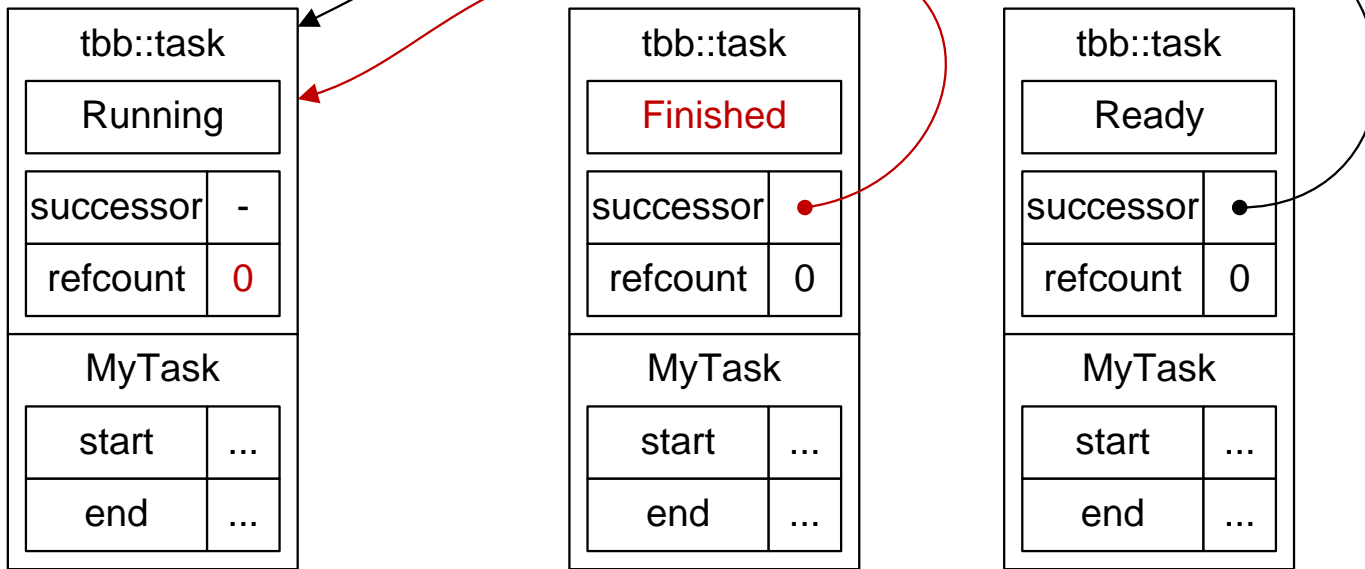
```
void CreateTasks(int start, int end)
{
    MyTask &root=*new(allocate_root()) MyTask(start,end);
    tbb::task::spawn_root_and_wait();
}
```

Managing reference counts

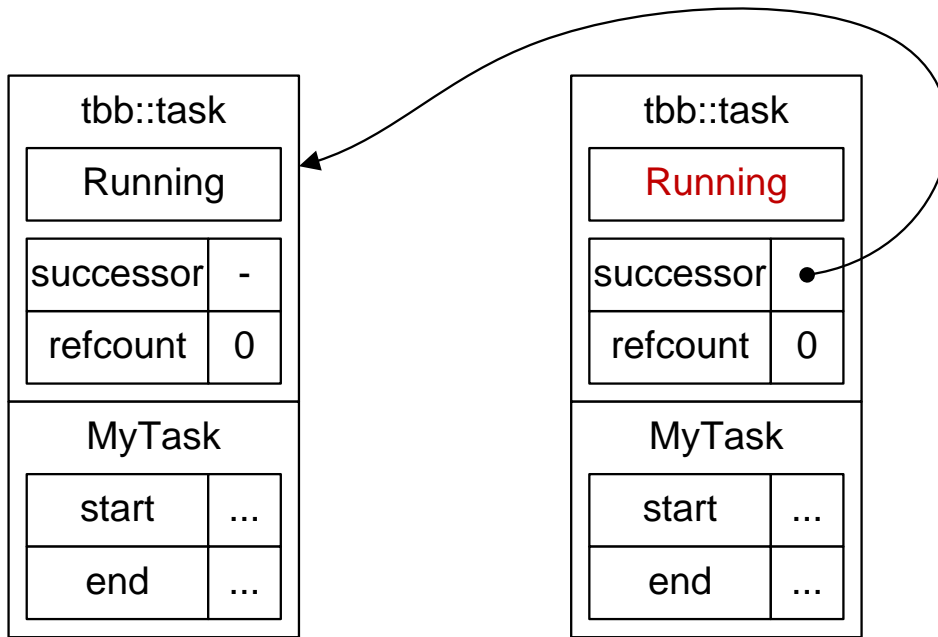
- What happens if we get the reference count wrong?
- Finishing task calls `decrement_ref_count` on successor
 - Automatically returns task to scheduler if count becomes zero



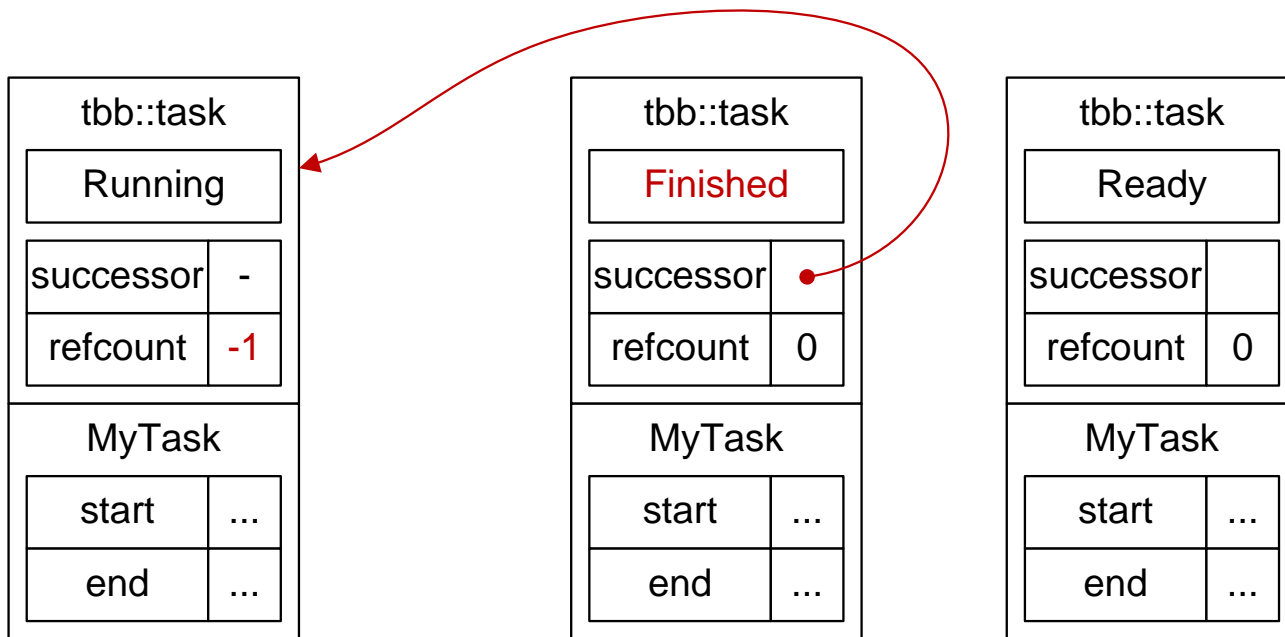
```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(1);
    MyTask &t1=*new(allocate_child()) MyTask(start, (start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```



```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(1);
    MyTask &t1=*new(allocate_child()) MyTask(start, (start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```



```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    MyTask &t1=*new(allocate_child()) MyTask(start, (start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    set_ref_count(3);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```



```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    MyTask &t1=*new(allocate_child()) MyTask(start, (start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    set_ref_count(3);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```


Many design patterns are built on tasks

- Iteration in various forms
 - `parallel_for`, `parallel_for_each`
- Reduction and accumulation
 - `parallel_reduce`
- Data-dependent looping and queue processing
 - `parallel_do`
- Support for heterogeneous tasks
 - `parallel_invoke`, `task_group`
- Heterogeneous tasks and token-based data-flow
 - `parallel_pipeline`
- Goal: turn design patterns in concrete functions