

CW5

- Timing:
 - Issued today : <https://github.com/HPCE/hpce-2016-cw5>
 - Due Fri 25th Oct (2 weeks)
- Pair work:
 - You'll work from two private repositories
 - Give each other push permission to your private repos
 - Keep the two repositories in sync
 - I'll infer the graph of pairs from access rights
- If you don't have a pair sorted out...
 - Hang around after the lecture

Suggestions

1. Read the code
 - What are the dependencies, what is the complexity?
2. Parallelism – Low-hanging fruit
 - Parallel for, recursive parallelism
3. Obvious inefficiencies
 - Excessive copying or re-initialisation (sloppy code)
 - Caching / eliminating computation (sloppy algorithms)
4. Parallelism – Larger-scale changes
 - Moving code to GPU
 - More aggressive re-structuring of code
5. Deep analysis
 - Algorithmic opportunities (change complexity)
 - Batching of computation
 - Maximise compute per data transfer (memory bandwidth)

Feasibility:

Learning to say no


Making sure metrics are meaningful

- Some things are quantifiable, but not very useful
 - CPU performance: MHz is not the same as performance
 - Cameras: Mega-Pixels is not the same as quality
- Consistent and quantifiable metrics provide open competition
 - Suppliers of systems always want to use the “best” metrics
 - Metrics should be defined by users, not suppliers
- People will optimise for metrics (it’s what they are for!)
 - Poor metrics lead to poor design and optimisation
 - Part of the specification problem


Feasibility studies

- People come up with demands
 - “*I want real-time spectral analysis of a 0hz-1GHz signal*”
 - “*We must process HD video within a latency of 1ms*”
 - “*This base-station must beam-form 32 channels*”
- Is it feasible to meet those demands?
 - Will it be easy?
 - Will it require optimisation?
 - Will it require a specialised platform?
 - Is it fundamentally impossible?
- You need some estimates before you start implementation
 - Execution time is the most basic check to be made

Circuit Placement





Logical components in circuit




Physical resources in device

- Take the graph of circuit, and find a valid placement onto physical resources
- Make sure that all logical **components** have a unique physical location
- Make sure that all logical **connections** map to a physical channel

Circuit Placement




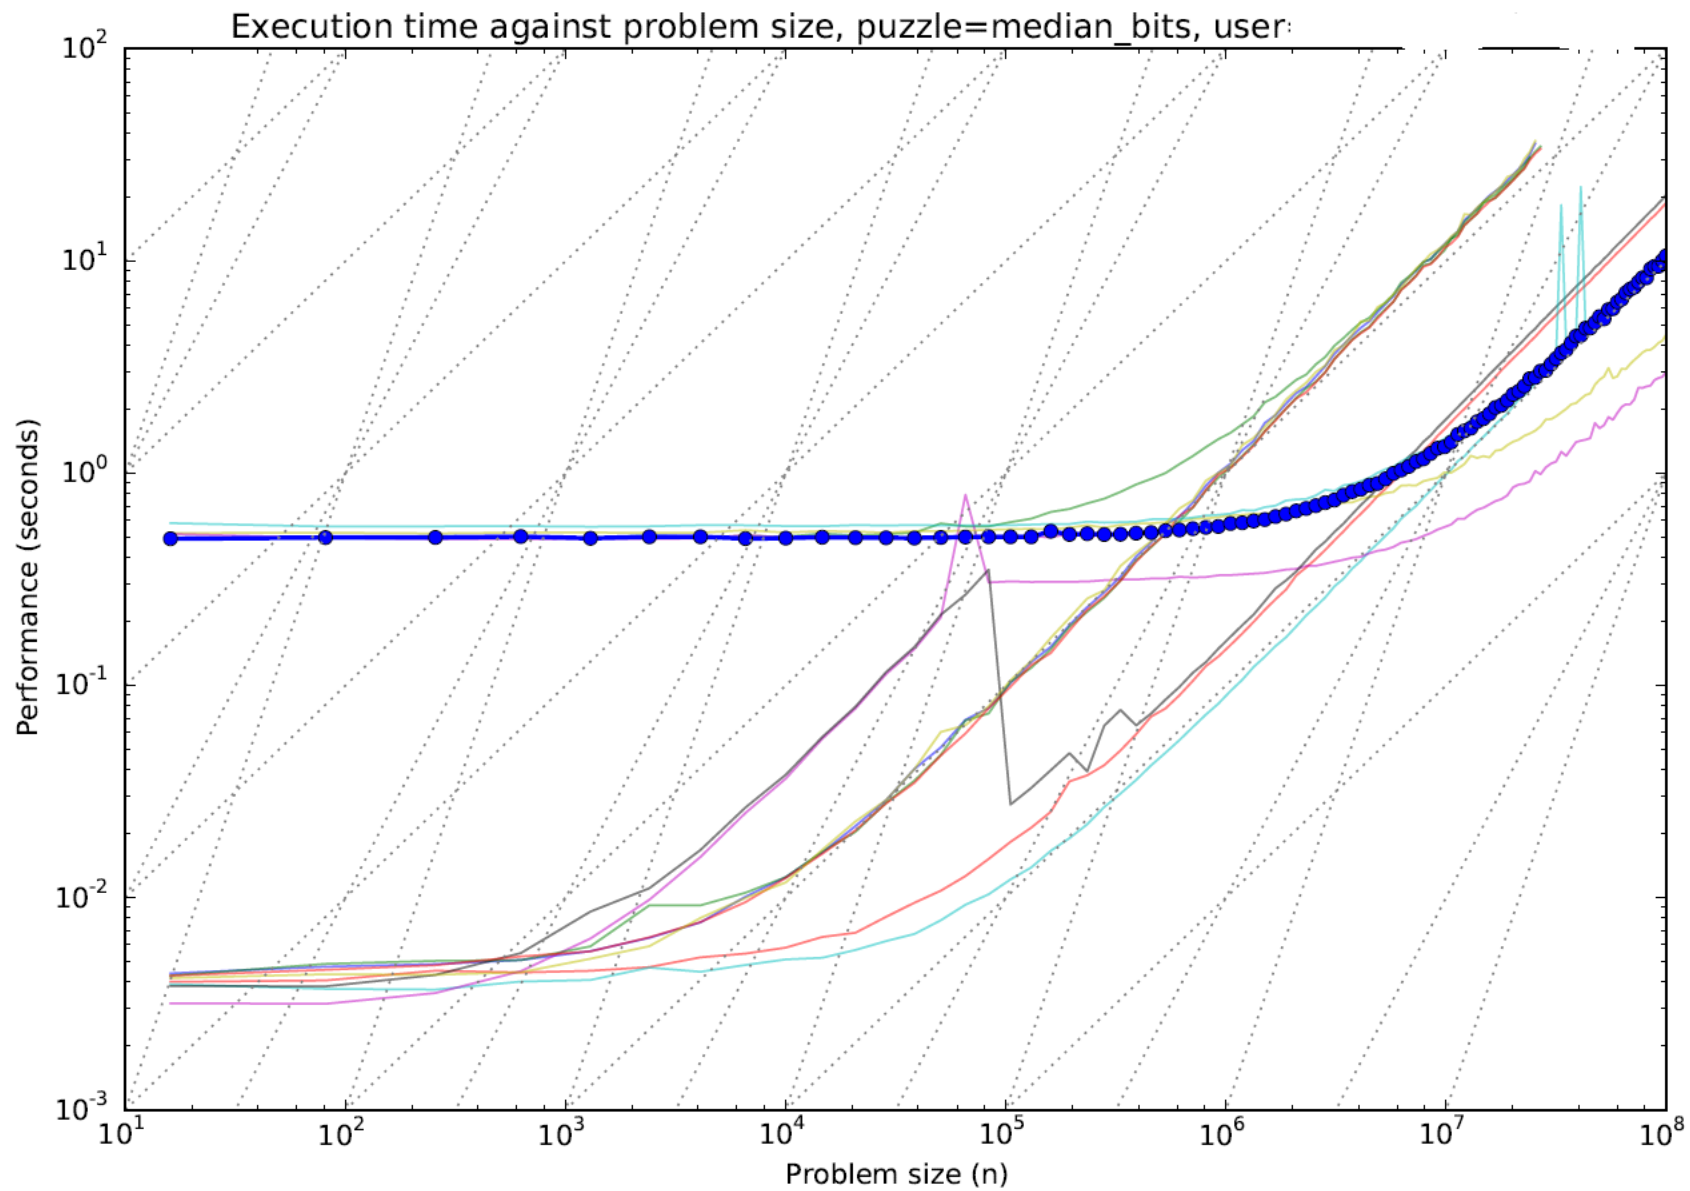
How difficult is a big problem?

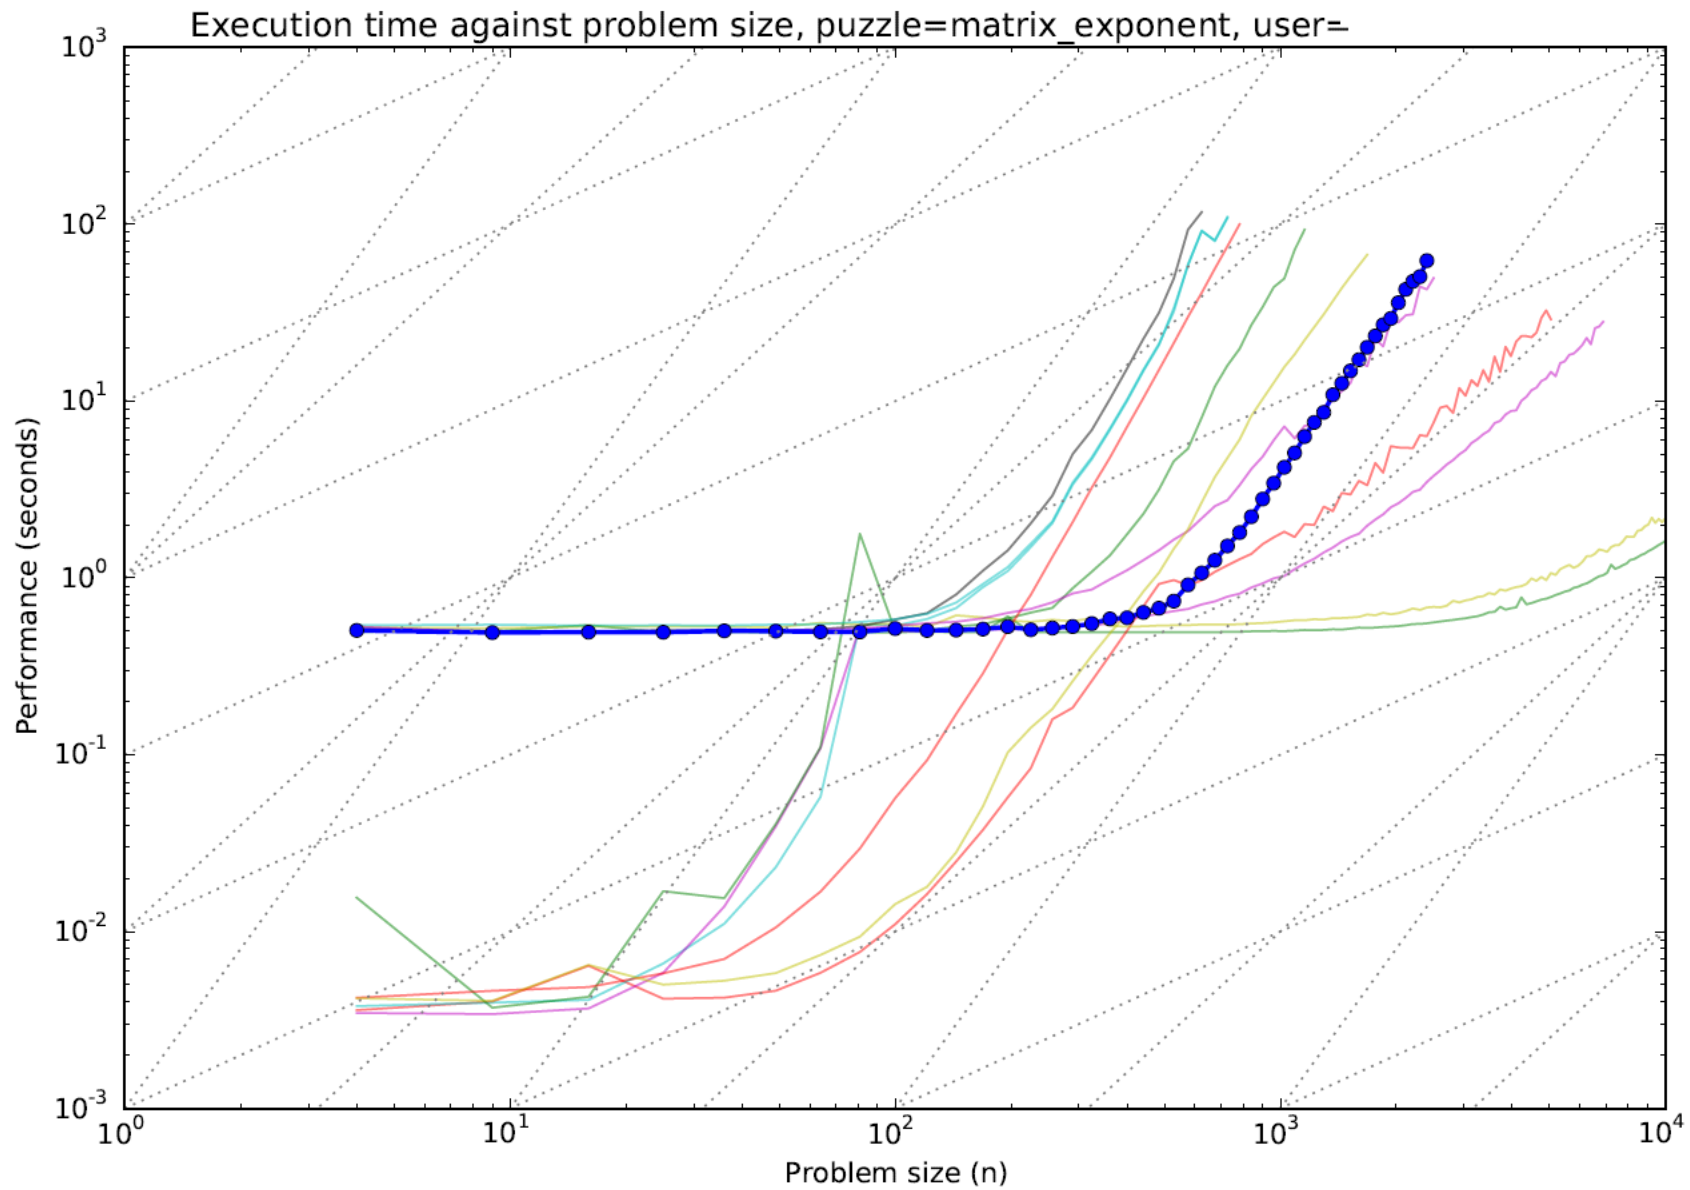


Trying to measure complexity

- Want to capture complexity of a task using equations
 - ***Time complexity***: how many “steps” does it require
 - ***Space complexity***: how much “storage” does it require
- We could derive exact equations for each
 - How many instructions does the task take in total?
 - How many bytes of memory are allocated during execution
- Means we have to worry about lots of details
 - Language: did you use C++, Fortran, VHDL?
 - Compiler: what optimisation flags were used?
 - Architecture: are integers 32-bit or 64-bit?
- Exact equations are sometimes possible, but often impractical







Complexity and Big-O (recap)

- Let's assume the existence of a function $g(n)$
 - $g(n)$ specifies *exactly* how many steps are taken
 - *Note: this function doesn't need to be stated explicitly!*
 - n is the “size” of the problem; e.g. an input vector of length n
- Goal: find a simple function $f(n)$, such that $g(n) \in O(f(n))$

$$g(n) \in O(f(n)) \quad \text{iff} \quad \exists n_c > 0, m > 0 : [\forall n > n_c : [g(n) < m \times f(n)]]$$

“There must exist a critical value n_c , and a positive constant m , such that for all $n > n_c$ the relation $g(n) \leq m \times f(n)$ holds”

“For increasing n , eventually you'll reach a point where $f(n)$ times a constant is always bigger than $g(n)$ ”

Complexity and Big-O

- $O(f(n))$ is a set of functions with the same or lower complexity
 - $n \in O(n)$ $n^2 \notin O(n)$
 - $n \in O(n^2)$ $n^2 \in O(n^2)$
 - $O(n) \subset O(n^2) \subset O(n^3)$
- It is sort of correct to claim that everything is $O(\infty)$
 - But it's really not very useful...
- Try to find the smallest complexity class containing a function
 - $n^2 + 2 \in O(n^2)$
 - $100 n^2 + 0.1 n^3 - 100 \in O(n^3)$
 - $2^n + n^4 \in O(2^n)$
- Find the fastest growing component, and choose that

Reduction Rules

```
function F(n:integer) : integer
```

```
begin
```

```
  for i = 0..n/2
```

← Executes $n/2$ times: $O(n)$

```
    acc = acc + init(i)
```

```
    for j = 64..n/3
```

← Executes $(n/3-64)*n/2$

```
      tmp = tmp + func(j, acc)
```

```
    next j
```

```
  next i
```

```
  return tmp
```

$= n^2 / 6 - 32n \in O(n^2)$

```
end
```

$O(a \times g(n)) \equiv O(g(n)),$ a is independent of n

$O(a + g(n)) \equiv O(g(n))$

$O((a \times n)^2) \equiv O(n^2)$

Common complexity classes

```
int SUM(int N)
{
    int acc=0;
    for(int i=0;i<N;i++)
        acc=acc+D[i];
    return acc;
}
```

```
function [R]=randMMM(N)
    A=rand(N,N);
    B=rand(N,N);
    R=A*B;
end
```

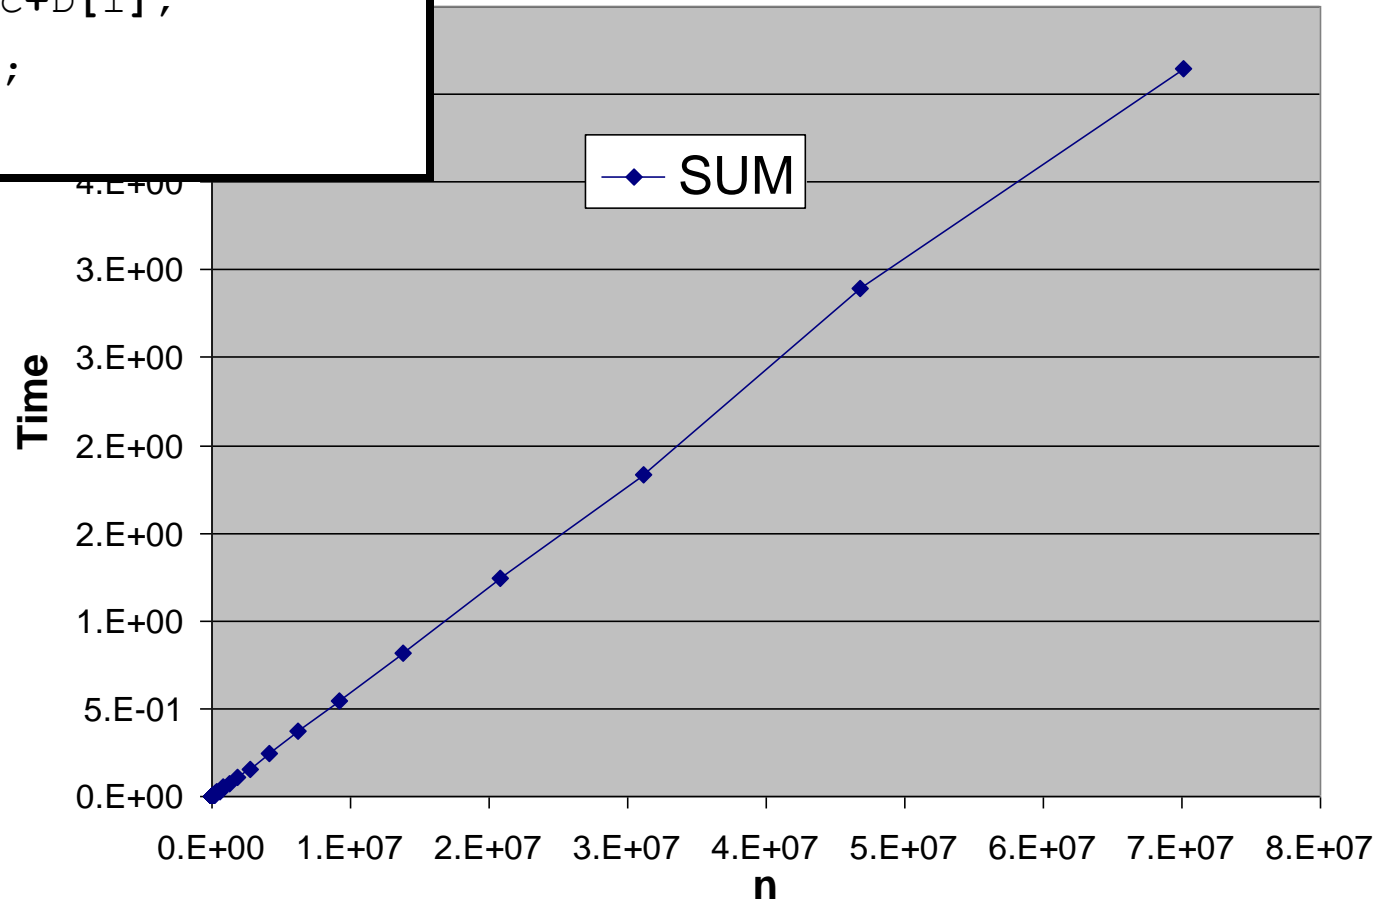
```
function [R]=randFFT(N)
    A=rand(N,1);
    B=rand(N,1);
    R=fft(A,B);
end
```

```
int Ack(int N)
{
    int A(int m, int n)
    {
        if(m==0) return n+1;
        if(n==0) return A(m-1,1);
        return A(m-1,A(m,n-1));
    }


    return A(N,N);
}
```



```
int SUM(int N)
{
    int acc=0;
    for(int i=0;i<N;i++)
        acc=acc+D[i];
    return acc;
}
```




Matrix-Matrix Multiply



```
function [R]=randMMM(N)
    A=rand(N,N);
    B=rand(N,N);
    R=A*B;
end
```

```
function [R]=Mul(A,B)
    for r=1:n
        for c=1:n
            R(r,c)=sum(A(r,:).*B(:,c));
        end
    end
end
```


Matrix-Matrix Multiply



```
function [R]=randMMM(N)
    A=rand(N,N);
    B=rand(N,N);
    R=A*B;
end
```

```
function [R]=Mul(A,B)
    for r=1:n
        for c=1:n
            R(r,c)=sum(A(r,:).*B(:,c));
        end
    end
end
```


Matrix-Matrix Multiply



```
function [R]=randMMM(N)
    A=rand(N,N);
    B=rand(N,N);
    R=A*B;
end
```

```
function [R]=Mul(A,B)
    for r=1:n
        for c=1:n
            R(r,c)=sum(A(r,:).*B(:,c));
        end
    end
end
```


Matrix-Matrix Multiply



```
function [R]=randMMM(N)
    A=rand(N,N);
    B=rand(N,N);
    R=A*B;
end
```

```
function [R]=Mul(A,B)
    for r=1:n
        for c=1:n
            R(r,c)=sum(A(r,:).*B(:,c));
        end
    end
end
```

Matrix-Matrix Multiply

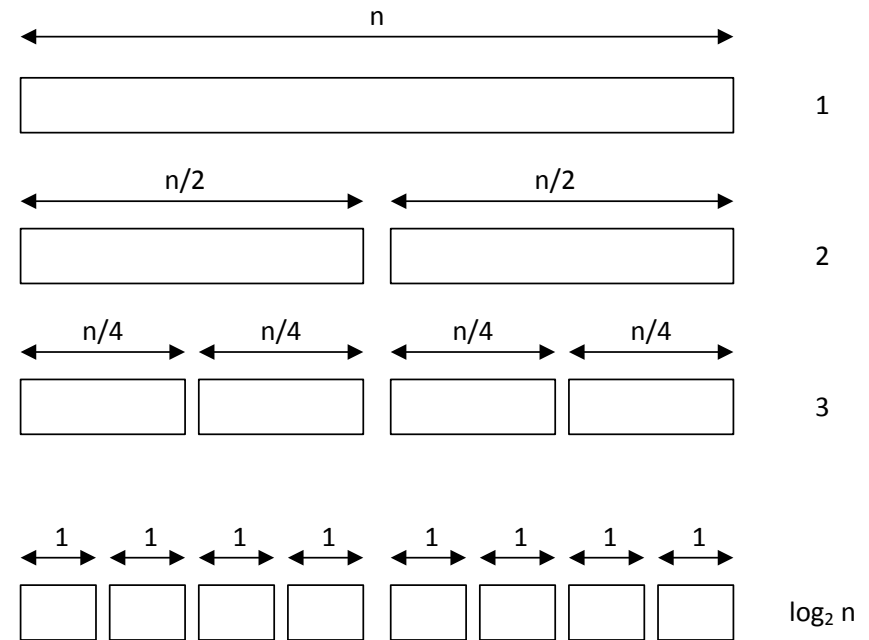


```
function [R]=randMMM(N)
    A=rand(N,N);
    B=rand(N,N);
    R=A*B;
end
```

```
function [R]=Mul(A,B)
    for r=1:n
        for c=1:n
            R(r,c)=sum(A(r,:).*B(:,c));
        end
    end
end
```

Quick-Sort

```
void Qsort(int N, int *data)
{
    partition(N, data); // O(n)
    if(N>1){
        Qsort(N/2, data);
        Qsort(N/2, data+N/2);
    }
}
```



How fast do functions grow?

$O(1)$

$O(n^4)$

$O(n)$

$O(1.01^n)$

$O(n!)$

$O(2^n)$

$O(n \log n)$

Polynomial Algorithms

- Polynomial-time algorithms are considered “easy”
 - That’s mathematically easy, not necessarily practical
- $O(1)$ – **Constant time**
 - *The best complexity class! Not much interesting in it though...*
 - Read an item from RAM
- $O(n)$ – **Linear time**
 - Vector addition
 - Search through an un-ordered list
- $O(n^2)$ – **Quadratic time**
 - Matrix-vector multiply
- $O(n^3)$ – **Cubic time**
 - Dense matrix-matrix multiply
 - Gaussian elimination

— In theory it’s lower, but in practise it often isn’t – see Strassen’s algorithm

Log and Log-Linear Algorithms

- Algorithms which recursively sub-divide some space
 - These are *actually* easy, not just mathematically
- $O(\log n)$ – **Logarithmic time**
 - Find an element in a **sorted** list
 - Root finding through bi-section
 - See if an element belongs to a set / add an element to a set
- $O(n \log n)$ – **Log-Linear time** (also called *linearithmic*)
 - Sorting a vector of items
 - Fast-Fourier-Transform (FFT)
- Both are huge improvements over closest polynomial
 - $O(\log n)$ preferred to $O(n)$
 - $O(n \log n)$ is *massively* better than $O(n^2)$

Exponential Algorithms

- Algorithms which explore some multi-dimensional space
- Class of algorithms with complexity $O(a^n)$ for $a > 1$
 - Brute-force search of all length- n binary patterns : $O(2^n)$
 - Brute-force search of all length- n decimal strings : $O(10^n)$
- Exponential time algorithms are generally very bad
 - Scale extremely poorly with n
- Occasionally useful as long as you are careful
 - $O(2^n)$ algorithm can be useful for $n < 32$
 - $O(a^n)$ algorithm with a close to 1 is sometimes feasible

Combinatorial Algorithms


- Looking at permutations and combinations of things
- Lots of specific sub-classes, but generally $O(n!)$
- Optimal mapping: bind abstract resources to physical ones
- Find the best sub-set from a larger set of resources
 - Many interesting engineering problems are combinatorial

Avoid anything more than polynomial




Thought Experiment

- We have four applications
 - All currently run in one second
 - All currently handle problems of “size” 16
- Each application has different complexity
 - Log-linear: $O(n \log n)$
 - Cubic: $O(n^3)$
 - Exponential: $O(2^n)$
 - Combinatorial: $O(n!)$
- The customer wants to handle problems of twice the size
 - How much do we need to accelerate the existing applications?




Speedup by about 2.5 times – probably use multi-core




Speedup by about 8 times – maybe use a GPU?





Speedup by 65000 times – err, cluster of FPGAs? Cloud?



Speedup by 10^{22} times – turn every atom in 1kg of iron into a Pentium?

The limits of computation

- New systems are not magic
 - Multi-core CPUs: ~16x speedup
 - GPUs: ~500x speedup
 - FPGAs: ~1000x speedup
 - Cloud: ~10,000x speedup
 - $O(n!)$, from $n=16$ to $n=32$ ~ 10^{22} required
- Lots of problems are completely intractable
 - Travelling Salesman: find shortest path to visit n cities
 - Bin packing: pack objects into the minimum number of bins
 - Boolean satisfiability: find values to make equation true

How to deal with intractable problems?

- Circuit place-and-route has ridiculously high complexity
 - But we regularly create designs with millions of logic gates...
- Must make decision about quality versus runtime
 - Wait 1 hour : design runs at 250MHz
 - Wait 10 hours : design runs at 310MHz
 - Wait ? hours : design runs at 317 MHz
- Some algorithms are progressive and approximate
 - Quality of solution improves as more compute time applied
 - Monte-Carlo, Genetic Algorithms, Simulated Annealing, ...
 - No guarantee of optimality – ***but at least you get an answer***

Why parallelism fails: Amdahl's Law

- Split a given compute task **X** into two portions
 - **A** : The parts that cannot be easily optimised or accelerated
 - **B** : The parts that can be sped up significantly
- Assume we achieve a speed-up of S_B times to part **B**
 - What is the speed-up S_X for the entire task?

$$T_X = T_A + T_B \quad \text{Original execution time}$$


$$T_{X'} = T_A + T_B / S_B \quad \text{New execution time}$$

$$S_X = T_X / T_{X'} \quad \text{Achieved speedup}$$

$$= \frac{T_A + T_B}{T_A + \frac{T_B}{S_B}}$$

$$= \frac{T_A + T_B}{T_A} \quad \text{as } S_B \rightarrow \infty$$

- Maximum speed-up is limited by the ***serial fraction*** : $T_A/(T_A+T_B)$
- Need a *tiny* serial fraction to achieve big speed-ups
- Are 1000x speed-ups realistic then?



Practical example

- Finite difference applications (fluid-mechanics, physics)
 - Discretise continuous space into cells
 - Discretise continuous time into distinct time-steps
- Goal of acceleration is to support finer resolution solutions
 - Usually increase resolution of space and time axis together
 - Let's take ***n*** as the resolution along each axis
- Tasks within finite-difference
 - Initialisation: initialise the *n* items in the first column
 - Processing: advance each column through *n* steps in time
 - Collection: retrieve answers from final column


```
g[1,1]=F1(0)
```

```
for s=2..n
```

```
  g[s,1]=F1(g[s-1,1])
```

```
end
```

```
for t=1..n-1
```

```
  g[1,t+1]=g[1,t]
```

```
  for s=2..n-1
```

```
    g[s,t+1]=F2(g[s-1,t],g[s,t],g[s+1,t])
```

```
  end
```


```
  g[n,t+1]=g[n,t]
```


```
end
```

```
return F3(g[1..n,t])
```


Space

Time







•Total work: $(n+1)(n+2) + C \in O(n^2)$




• Total work: $(n+1)(n+2) + C \in O(n^2)$

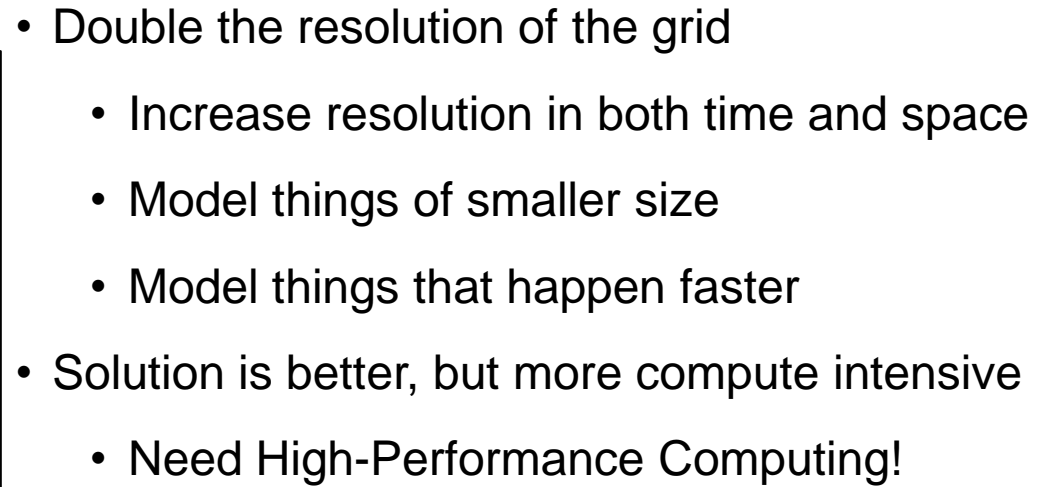





Critical path

- **Critical path:** longest path through dependency graph
 - Assume infinite processors, and zero communication overhead





- Double the resolution of the grid
 - Increase resolution in both time and space
 - Model things of smaller size
 - Model things that happen faster
- Solution is better, but more compute intensive
 - Need High-Performance Computing!





Serial execution: $(n+2) + n \times (n+2) + C$



Parallel execution: $(n+2) + n + C$

Speedup: $[n^2 + 3n + 2 + C] / [2n + 2 + C]$

Speedup is $O(n)$ - increases linearly with problem size

Why parallelism works: Gustafson's Law

- Split a task **X** into two portions **A** and **B**
 - **A** cannot be accelerated, while **B** can be parallelised
 - But now the execution time of **A** and **B** depends on problem size
 - $T_A(n)$: time to perform part **A** for problem of size n

$$T_X(n) = T_A(n) + T_B(n) \quad \text{Original execution time}$$

$$T_{X'}(n) = T_A(n) + T_B(n) / S_B \quad \text{New execution time}$$

$$S_X(n) = T_X(n) / T_{X'}(n) \quad \text{Achieved speedup}$$

$$= \frac{T_A(n) + T_B(n)}{T_A(n) + \frac{T_B(n)}{S_B}}$$

$$= S_B \quad \text{as } \xrightarrow{n \rightarrow \infty} \text{ if } O(T_A(n)) \prec O(T_B(n))$$