

# CW5 - the end is nigh

- Julia floating-point – sorry!
  - I forget that architecture 1 was a long time ago
  - Without the reference it is more difficult to get bit-accurate
    - though some did!
- Try to avoid last-minute performance hacks
  - Unlikely to be that much faster
  - Quite likely to have a failure mode
- Do your push, then pull+build+test in a *different* directory
  - Useful to find uncommitted files
  - Can highlight where stale files
- Decide it is working then... walk away!
  - Performance is engineered, not rushed

# Parallel Design Patterns

# What are Design Patterns?

- Design patterns are templates for “solving” a problem
  - Solving = {structuring, executing, writing, analysing, ...}
  - Range from application level down to scheduling instructions
  - Emphasis on **composability**: able to safely combine patterns
- Structured programming introduced the first design patterns
  - Ordering execution of statements: *sequences, branches, loops*
  - Grouping data into a unit: *structures*
  - Grouping statements into a unit: *procedures and functions*
- Object-oriented programming standardised common patterns
  - `struct X; F(X*,int); -> class X{ F(int); }`
  - Polymorphism (virtual functions) increase composability
  - Don't need to know what an object is, just what it does

# Describing Design Patterns

- Design patterns try to formally capture intuition or experience
  - Reduce the need for “rock-star” programmers: *engineering not art*
  - Increase productivity: *don’t re-invent the wheel*
  - Increase reliability: *record both the patterns **and** when they apply*
- Elements of a design pattern
  - **Name** : we need a common name so people can talk about them
  - **Problem** : simple description of what problem the pattern solves
  - **Context** : where does the problem occur, and any background info
  - **Forces** : any intrinsic tradeoffs that are being addressed
  - **Solution** : how the pattern is applied
  - **Examples** : application of the pattern to a real-world example

# Application of Design Patterns

- Concreteness of the design pattern varies hugely
  1. Both solution and conditions are described programmatically
    - Can be turned into compiler optimisations
    - Rare and usually local in scope: checking conditions is difficult
    - *Why does this course even exist – can't the compiler do it?*
  2. Solution can be captured in code, but not the conditions
    - Compiler/library will apply the pattern for us: `spawn/parallel_for`
    - We need to check whether the pattern can and should be applied
  3. Both solution and conditions cannot be formally captured
    - Programmer has to interpret the conditions and the solution
    - Requires human skill and (maybe) some thought

# Umm...

- This all sounds great in theory, but the practise is less good
  - People have different names for the same thing
  - People have the same name for different things
  - Currently a lack of standardisation in parallel pattern names
  - GPUs etc. are muddying the waters further
- I will broadly use the terms from these two sources
  - TBB patterns (part of TBB documentation):  
[http://software.intel.com/sites/products/documentation/doclib/tbb\\_sa/help/tbb\\_userguide/Design\\_Patterns/Design\\_Patterns.htm](http://software.intel.com/sites/products/documentation/doclib/tbb_sa/help/tbb_userguide/Design_Patterns/Design_Patterns.htm)
  - Berkely parlab: <http://parlab.eecs.berkeley.edu/wiki/patterns/patterns>

# Map-Reduce

- **Problem** : a function must be applied to many pieces of data, followed by an associative reduction
- **Context** : some container contains  $x_1..x_n$ , and we wish to calculate  $a( f(x_1), a(f(x_2), \dots f(x_n)))$ . Both  $f(.)$  and  $a(.)$  must be side-effect free, and  $a(.)$  must be associative
- **Solution** : apply  $f(.)$  to all data-items as independent parallel tasks (*map*), then use a recursive tree of parallel tasks to collect the results (*reduce*)

# Example : largest magnitude complex number

```
double max_magnitude(int n, const complex_t *data)
{
    double best = abs(data[0]);

    for(int i=1; i<n; i++) {
        double curr = abs(data[i]);
        if(curr > best)
            best=curr;
    }
    return best;
}
```

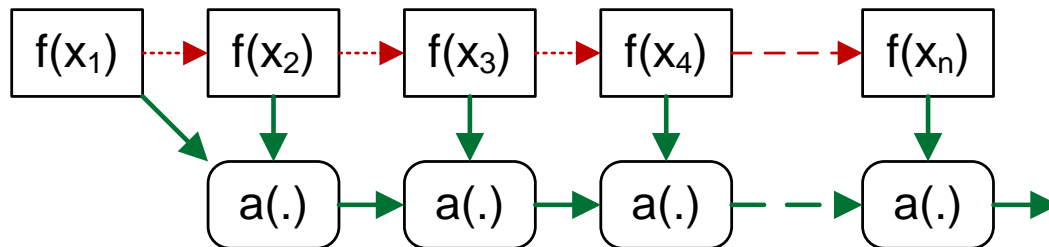
- $f(.) = \text{abs}(.)$  - Complex magnitude (modulus)
- $a(.) = \text{max}(.)$  – Maximum of two numbers
- Do we meet requirements of pattern?
  - $a(.)$  and  $f(.)$  are side-effect free
  - $a(.)$  is associative



```

double max_magnitude(int n, const complex_t *data)
{
    double best = 0.0;
    for(int i=0;i<n;i++){
        double curr = abs(data[i]);
        if(curr > best)
            best=curr;
    }
    return best;
}

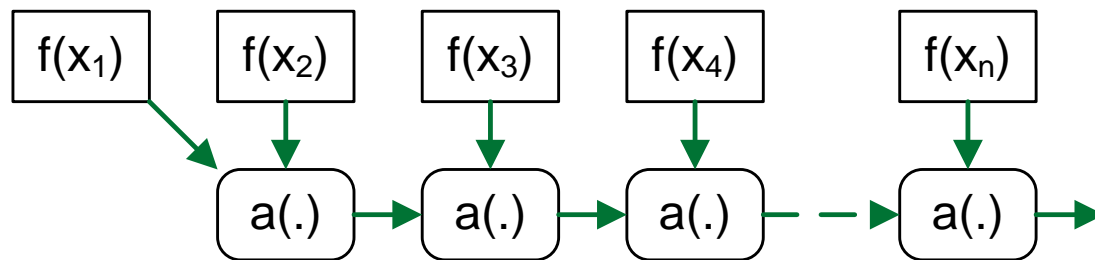
```



```

double max_magnitude(int n, const complex_t *data)
{
    std::vector<double> temp(n);
    tbb::parallel_for(0, n, [&](int i){
        temp[i] = abs(data[i]);
    });
    for(int i=0; i<n; i++){
        if(temp[i] > best)
            best=temp[i];
    }
    return best;
}

```

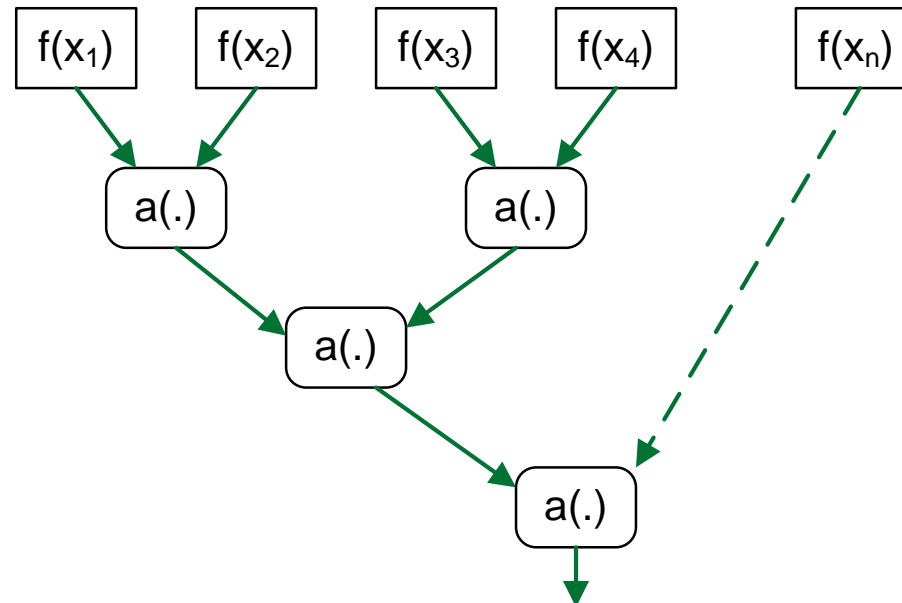


```

double max_magnitude(int n, const complex_t *data)
{
    if(n==1) return abs(data[0]);

    double left, right;
    tbb::parallel_invoke(
        [&]() { left = max_magnitude(n/2, data); },
        [&]() { right = max_magnitude(n-n/2, data+n/2); }
    );
    return max(left, right);
}

```

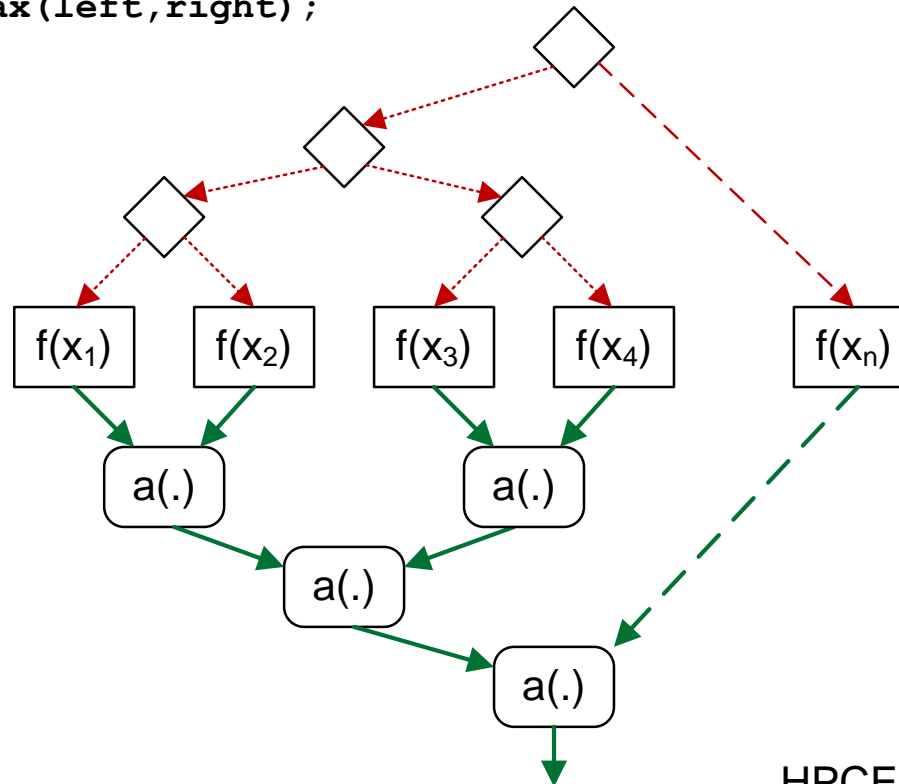


```

double max_magnitude(int n, const complex_t *data)
{
    if (n==1) return abs(data[0]);

    double left, right;
    tbb::parallel_invoke(
        [&]() { left = max_magnitude(n/2, data); },
        [&]() { right = max_magnitude(n-n/2, data+n/2); }
    );
    return max(left, right);
}

```



# Map-Reduce : Forces

- Balancing parallelism vs overhead
  - We want to maximise available tasks to increase parallelism
  - But the reduction graph contains lots of synchronisation
- Determinism vs speed
  - Want to have as much scheduling flexibility as possible
  - Prefer deterministic results with pseudo-associative operations
    - e.g. floating-point addition

*(I find the “Forces” section often ends up a bit vague)*

# Implementation of Map-Reduce

- Map-Reduce occurs everywhere
  - Sequential version is available in many languages and libraries
  - Often the map and reduce functions must be combined
    - $mr(x,y) = a(x, m(y))$
  - C++ : `std::accumulate`; Python: `reduce`; Haskell: `foldl/foldr`
- Parallel versions of Map-Reduce are very common
  - Very easy to understand for users
  - Applicable in large numbers of real-world scenarios
- Google use a tool called MapReduce with thousands of CPUs
  - Excellent paper on real-world scaling in distributed systems:
    - <http://research.google.com/archive/mapreduce.html>
    - *But they didn't invent the idea, been around for decades*
  - Many similar approaches: Hadoop etc.

# tbb::parallel\_reduce

```
class Func
```

```
{  
    Func(Func &src, tbb::split);
```

```
    template<class T>  
    void operator()( const T &r );
```

```
    void join(Func &rhs);
```

```
};
```

```
template<class TR, class TF>
```

```
void tbb::parallel_reduce(  
    const TR &range,
```

```
    TF &body
```

```
);
```

- Must pass object with specific members
- Splitting constructor
  - Used by TBB when the range is split
- operator() (r)
  - Similar to tbb::parallel\_for
  - Except – **non-const!**
  - Object must change during reduction
- join function
  - Merge results from two ranges
  - Accumulator of rhs is merged into lhs
- parallel\_reduce returns void
  - Results have to come out of body

```
struct Worker
```

```
{  
    double m_max;  
    const complex_t *m_data;  
}
```

```
Worker(const complex_t *data)  
{    m_data=data;    m_max=0.0; }
```

```
Worker(Worker &src, tbb::split)  
{    m_data=src.m_data;    m_max=0.0; }
```

```
template<class T>  
void operator()( const T &r ) {  
    double acc=0.0;  
    for(int i=r.begin();i<r.end();i++)  
        acc=std::max(acc,abs(m_data[i]));  
    m_max = std::max(m_max, acc);  
}
```

```
void join(Worker &rhs)  
{ m_max = std::max(m_max,rhs.m_max); }  
};
```

- Two member variables
  - m\_data : What we are working on
  - m\_max : The maximum so far
- Standard constructor
  - Creates the “root” object
- Splitting constructor
  - Called when TBB splits the range
  - Need to store where the data is
- Function operator
  - Do local maximum over range
  - Add to maximum seen so far
- join: merge two sub-ranges
  - rhs will then disappear



```

double max_magnitude(
    int n,
    const complex_t *data
){
    return tbb::parallel_reduce(
        // range
        tbb::blocked_range<int>(0,n) ,
        // identity
        DBL_MAX,
        // reduction over a range
        [](tbb::blocked_range<int> &r) -> double
        {
            double acc=0.0;
            for(int i=r.begin(); i<r.end(); i++)
                acc=std::max(acc,abs(data[i]));
            return acc;
        },
        // reduction of two values
        [](double left, double right) -> double
        {
            return std::max(left, right);
        }
    );
}

```

- There is also a lambda form
- Needs an identity element
  - $f(x, \text{identity}) = x$

# Element-wise / Data-Parallel

- **Problem** : apply a function to all items in a container
- **Context** : an identical function is to be applied to all items in a container, either transforming the items themselves, or performing some other side-effect
- **Forces** : standard (no. of tasks vs task overhead)
- **Solution** : create work by sub-dividing the container
- **Examples** : `tbb::parallel_for`, CUDA, OpenCL

# Agglomeration

- **Problem** : an application has too much fine-grain parallelism
- **Context** : many algorithms can be easily decomposed to the level where each task executes a single instruction, but this results in an extremely high cost of work
- **Forces** : we need to balance the need for average parallelism versus the cost of work
- **Solution** : halt production of new tasks when the remaining work drops below a certain level, and switch to sequential execution

# Agglomeration Examples

- Recursive FFT
  - Switching to sequential when the matrix gets too small
  - Potential problem: a fixed threshold may give sub-optimal results
    - What is optimal on one platform may not be on another
    - Compiler optimisations may increase the agglomeration required
- `tbb::parallel_for` and `tbb::parallel_reduce`
  - Function object is given a **range** of indices, not a single index
  - Size of range is optimised to try to balance execution time
  - Potential problem: auto-optimisation may guess wrong
    - Splits a large but extremely fast range into individual tasks
    - Doesn't split a small loop where each iteration is very slow
- Usually auto-tuning is better, unless you perform experiments
  - Humans are pretty bad at guessing how fast things run
  - 8-CPU's used to be rare, now 32-CPU's is fairly common

# Divide and Conquer / Recursive

- **Problem** : a sequential program with recursive functions or data-structures must be parallelised
- **Context** : the existing program repeatedly splits the data into independent sub-sections, which can then be further decomposed. The splitting may be over a 1D or 2D range, or follow some graph-like data-structure
- **Forces**: standard (creating parallelism vs task size)
- **Solution** : when the program splits the problem, spawn new tasks for each sub-problem, then sync for all child tasks

```

// Split the range in two, returning
// the pivot element. Takes
// O(end-begin) steps
double *partition(
    double *begin,
    double *end
);

void QSort(
    double *begin,
    double *end
){
    if(end-begin<1000){
        // Agglomeration - drop to serial
        std::sort(begin, end);
    }else{
        double *mid=partition(begin, end);
        tbb::parallel_invoke(
            [&]() { QSort(begin, mid-1); },
            [&]() { QSort(mid, end); }
        );
    }
}

```

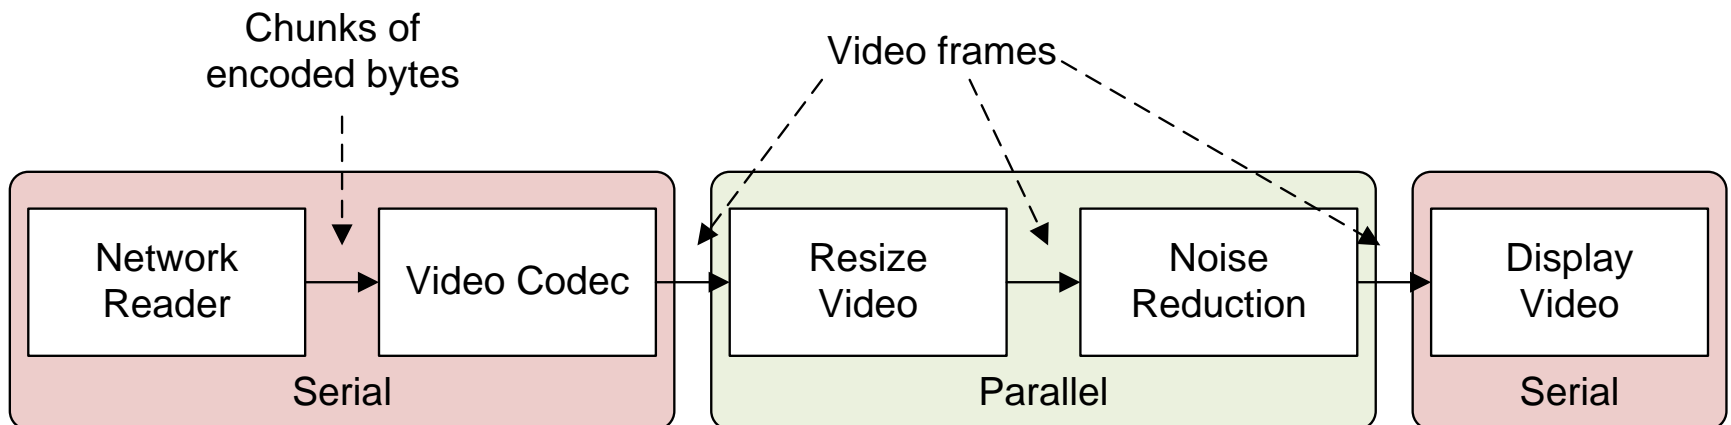
- Sorting: fundamental operation
  - Quick-Sort repeatedly splits problem into smaller problems
  - What is the big-O complexity?
- Sorting is so common it is supported directly in TBB
  - `tbb::parallel_sort`
  - Same interface as `std::sort`
  - Guaranteed **deterministic**
    - Always gives the same result
  - Not guaranteed **stable**
    - Different objects with the same key may be re-ordered

# Pipeline

- **Problem** : a stream of data needs to be processed in stages
- **Context** : some source function produces a stream of individual chunks of data, which must then be transformed using a set of operations, then sent to some sink function
- **Forces** : allowing many parallel data items increases available parallelism, but can greatly increase memory requirements
- **Solution** : create a pipeline of filter objects which can be applied as parallel tasks, then manage lifecycle of a limited number of tokens through the pipeline

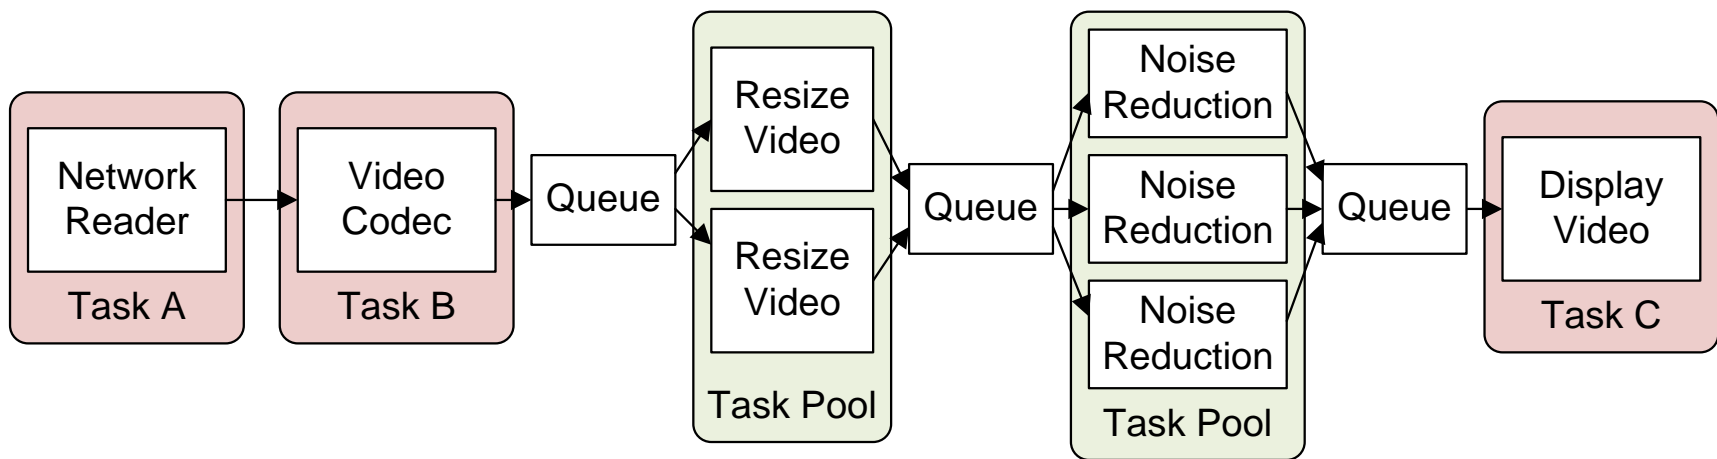
# Example : Video Processing

- A pre-recorded video stream is being decoded and displayed
  - Fetch and decode: get encoded bytes and turn into frames
  - Resize and de-noise: transform individual frames before display
  - Display : present the frames to the user
- Consider the metrics : optimise throughput, not execution time
- Video decoding and display have an intrinsic order
- Individual frames can be transformed in parallel in any order





- Use queues as a buffer between different tasks
- Ordered/serial processes are scheduled as a single task
  - Can only be executed on one CPU at a time
- Unordered/parallel processes can expand to multiple tasks
  - If there is lots of work in the queue create new task for each item
  - Task set expands to soak up CPUs
- Problem: process creates data faster than consuming process
  - Solution: limited size queues. e.g. only 16 frames in flight at once
- Implemented in TBB as `tbb::pipeline`



# Compare and Swap Loop

- **Problem** : a read-write variable is shared between tasks
- **Context** : multiple threads are communicating via some shared variable, such as the best or worst solution seen so far
- **Forces** : the update needs to be thread-safe, but also needs to be fast (so task dependencies cannot be used)
- **Solution** : use an atomic variable combined with a compare and swap loop
- **Examples** : see earlier lecture

# Further design patterns

- There are many more parallel design patterns
  - Some are very specialised for particular applications
  - Some seem extremely obvious
  - People who are into design patterns go a bit over the top
- Best design patterns are those that can be turned into a library
  - Require programmer's analysis to decide whether it applies
  - Provide a way to quickly apply solution ***if*** it is appropriate
- TBB has a number of design patterns as algorithms
  - Section 4 of ref. manual : `tbb::parallel_do`, `tbb::parallel_scan`
- Further useful patterns in the TBB Design Patterns Guide
  - *Fenced Data Transfer, Local Serializer, etc.*

CW6

# “Real” problems

- You’re at a research-intensive university
  - Theory is that good researchers make good teachers...
- We’re supposed to combine teaching and research
- Giving you problems I already solved is not real
- So....

# POETS

Andrew Brown (Southampton)  
Simon Moore (Cambridge)  
Andrey Mokhov (Newcastle)  
David Thomas (Imperial)



# POETS consortium



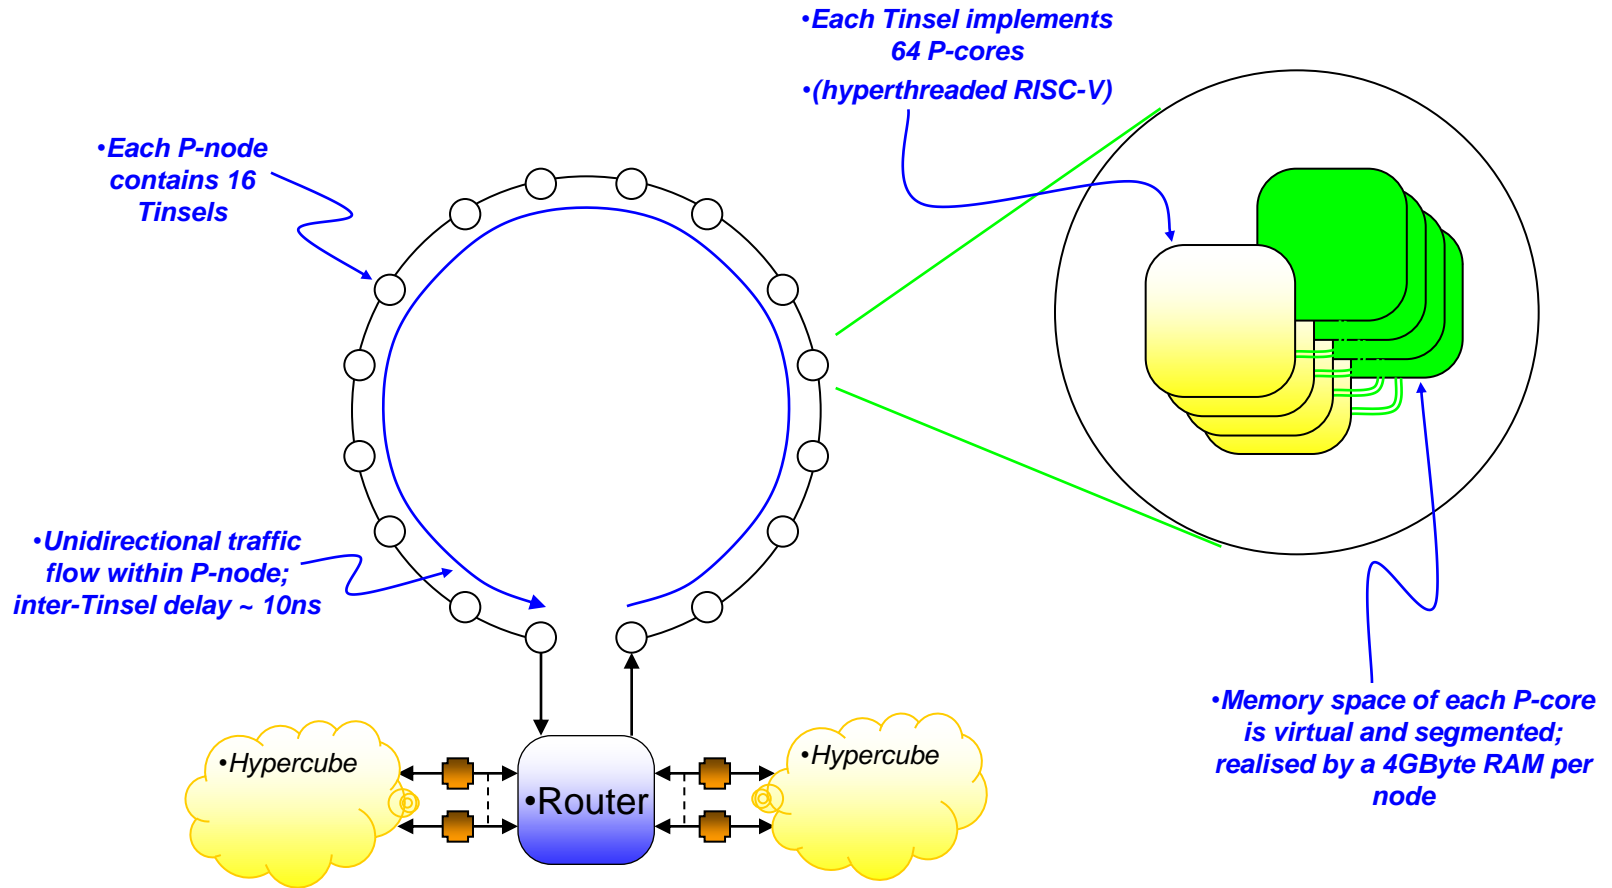
# POETS : Partially Ordered Event Triggered Systems

- A key barrier to high performance is synchronisation
  - Core-to-core : cache coherency
  - Workitem-to-workitem : OpenCL barrier
  - Thread-to-thread : mutexes
  - CPU to GPU : OpenCL events
  - Machine-to-Machine: MPI and networking
- What if we could avoid synchronisation?

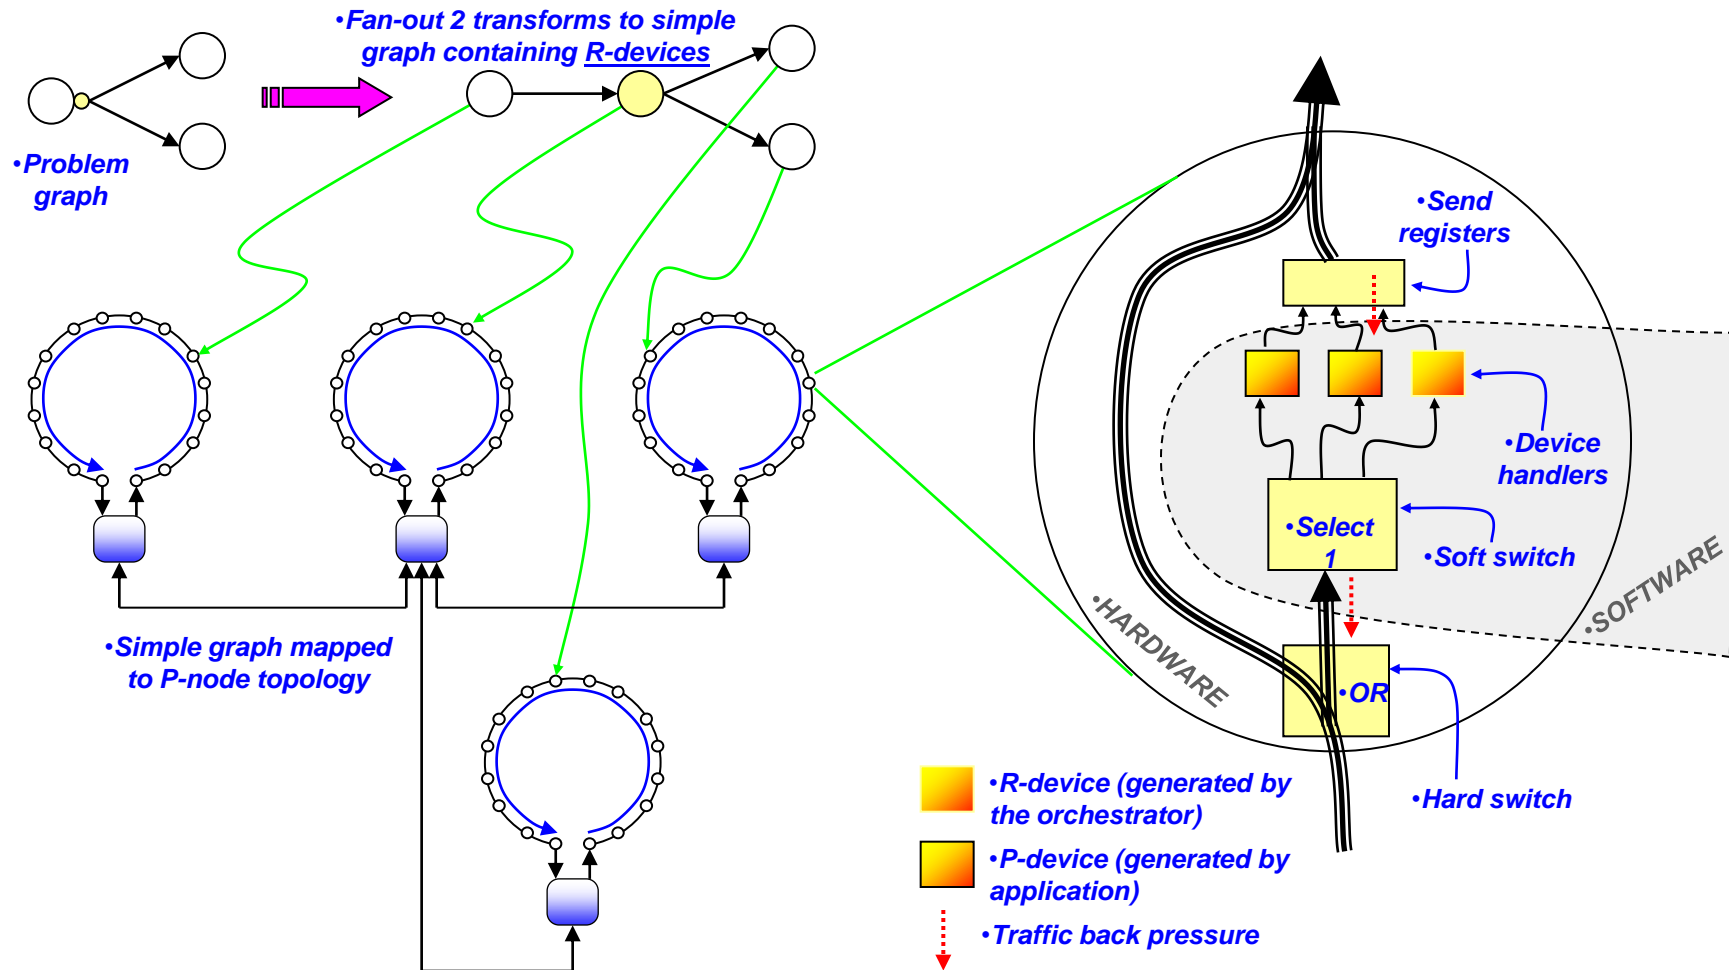
*Imagine... : 10,000 cores shouting at each other in parallel*



# The P-node

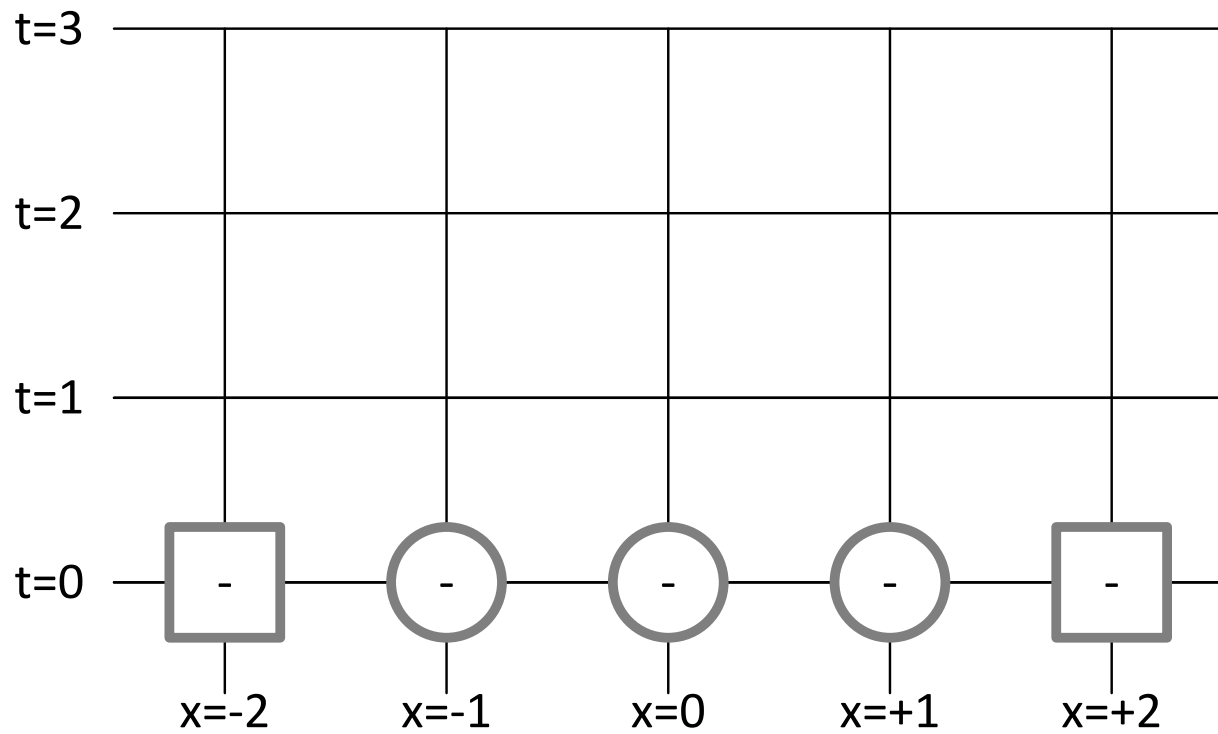


# The P-core

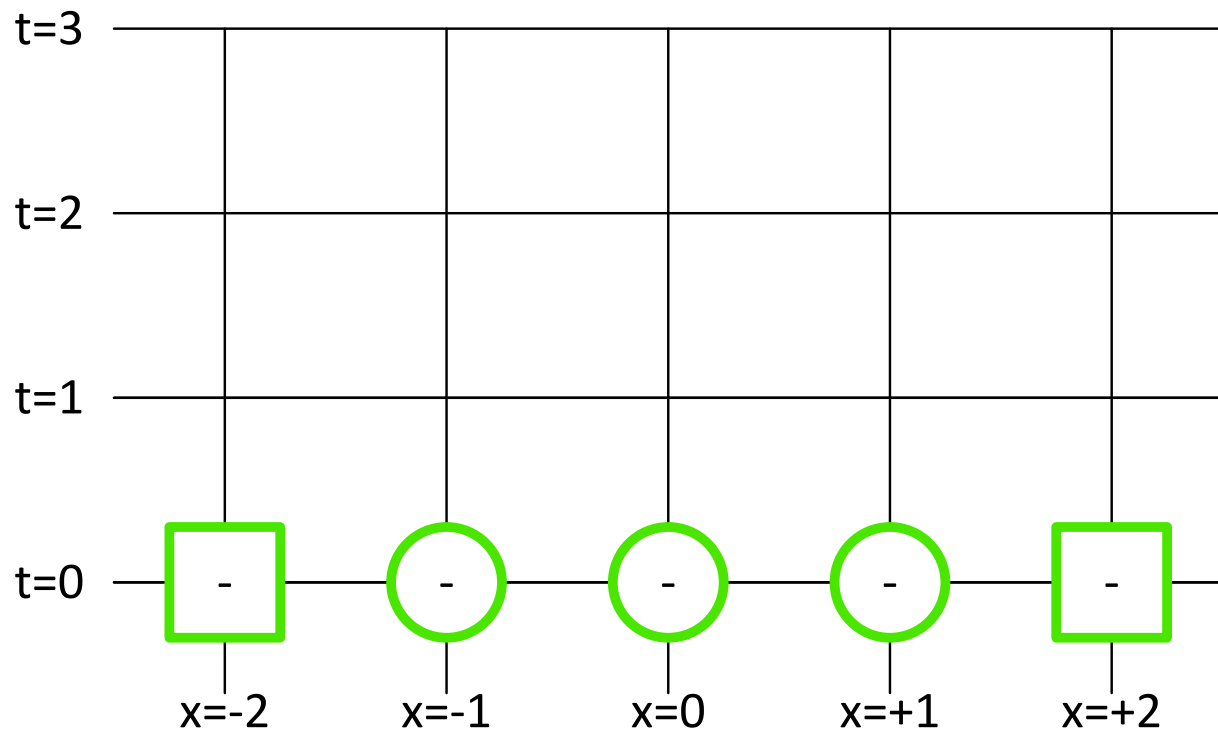


# POETS applications are asynchronous

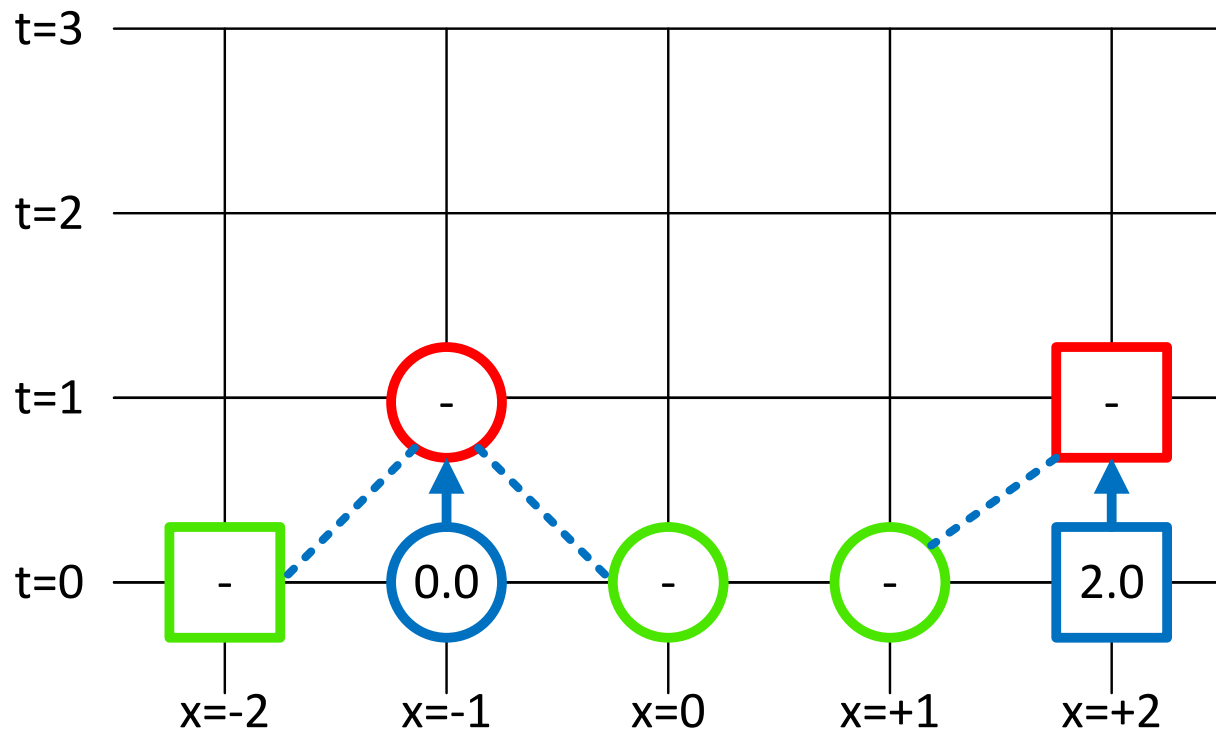
- Applications are split into thousands of “devices”
  - A device represents a point in space
  - Recall coursework 3 : world discretised into heat-cells
- Devices send messages to each other
  - For example in a heat equation:
    - Device A -> Device B : “My temperature is 0.6”
    - Device B : Ok, I’ll update to temperature 0.7
    - DeviceB -> DeviceA : “My temperature is now 0.7”
  - Messages are very efficient
    - No synchronisation
- But... messages are sent in any order
  - How do you maintain a notion of time?



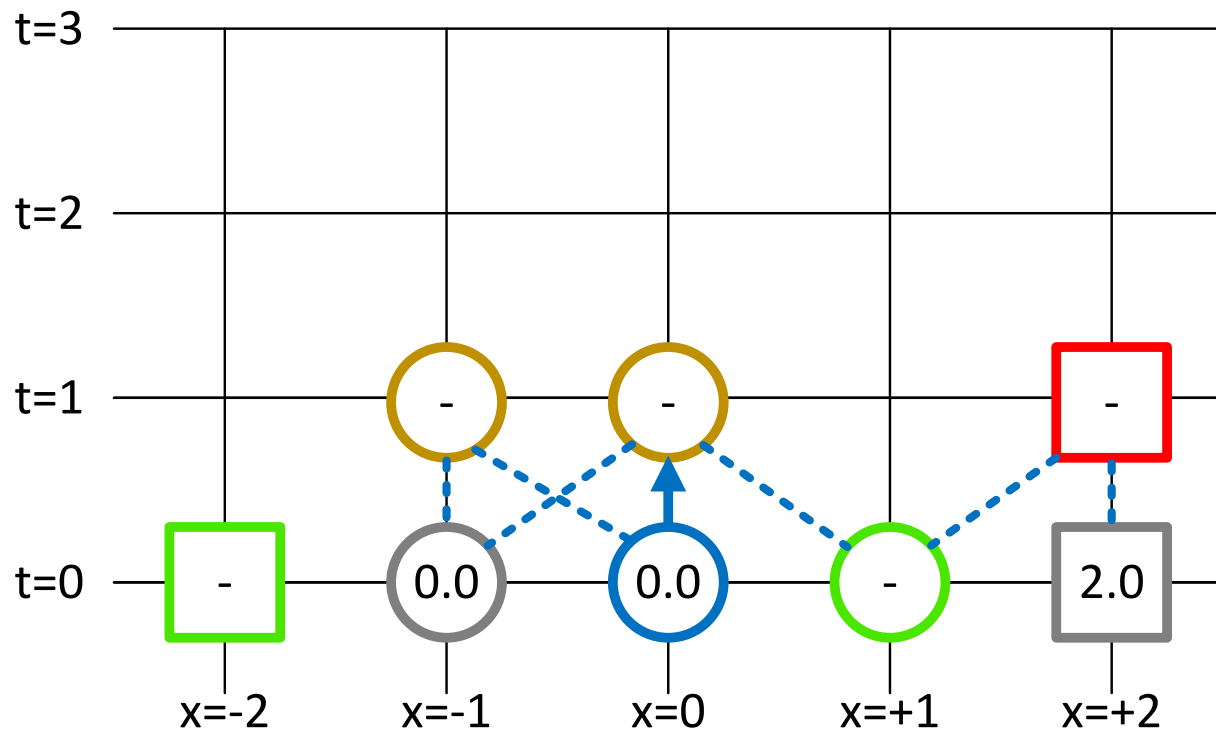
id:0



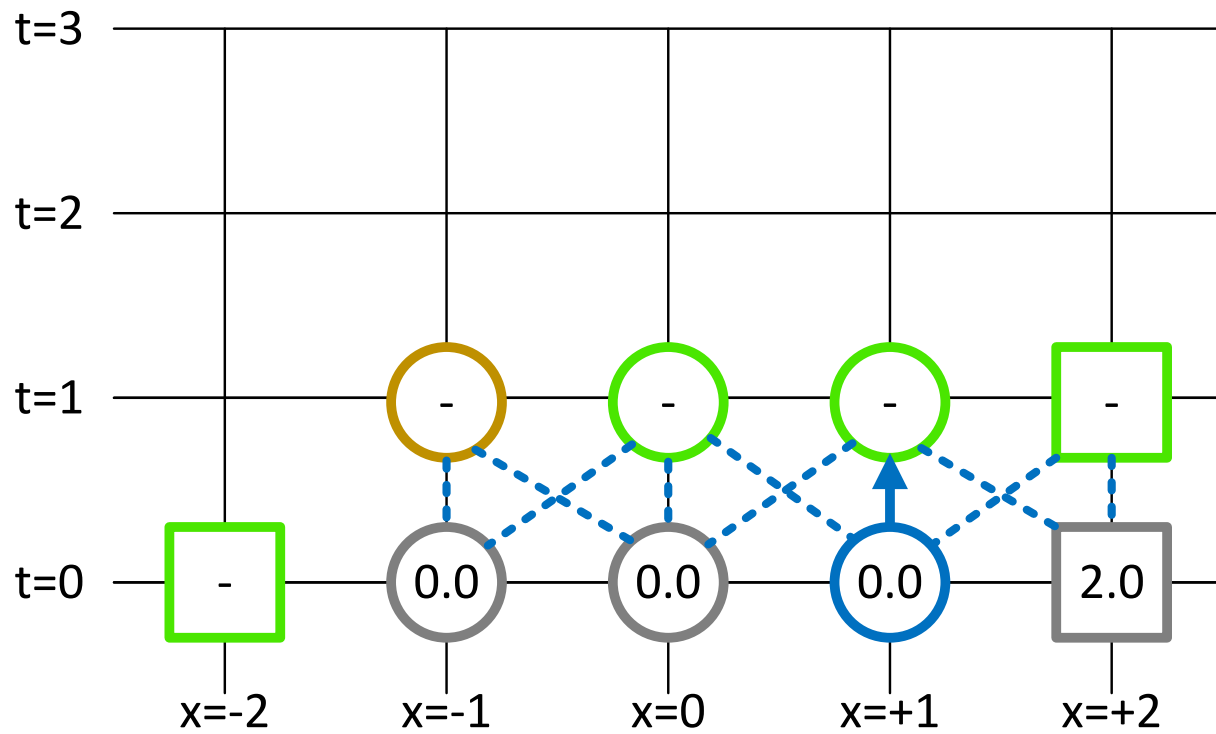
id:1



id:2

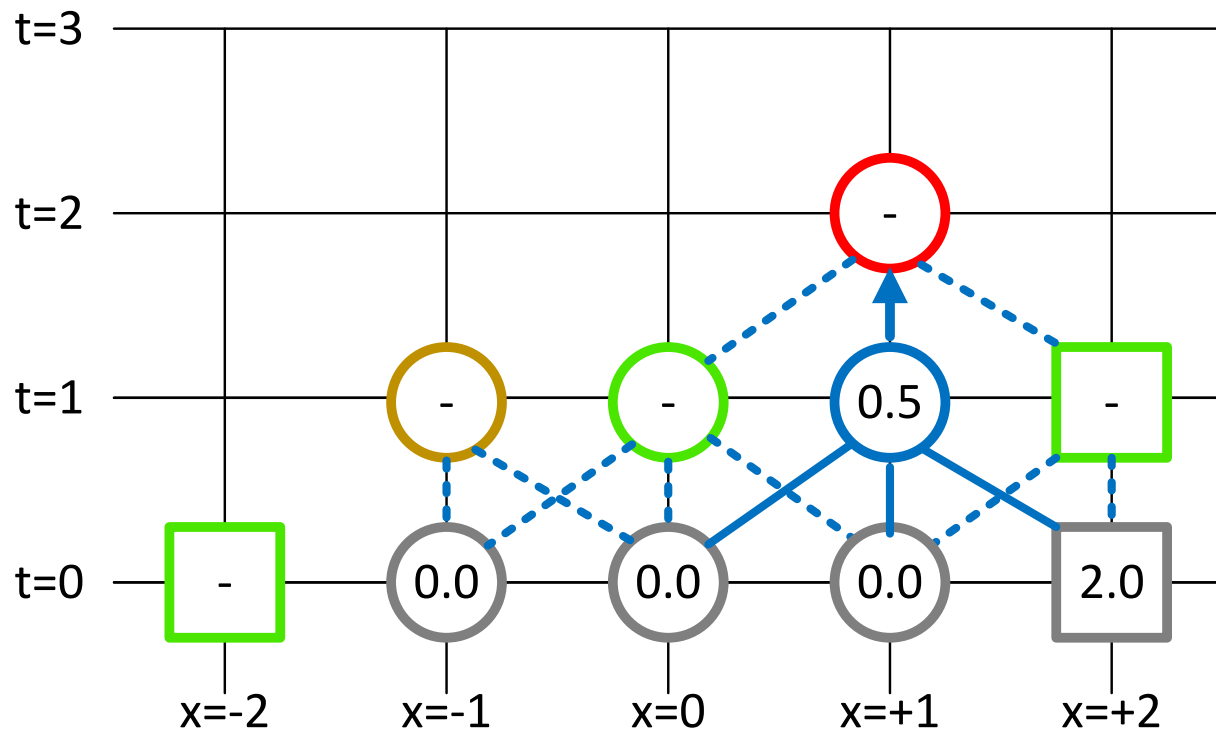


id:3

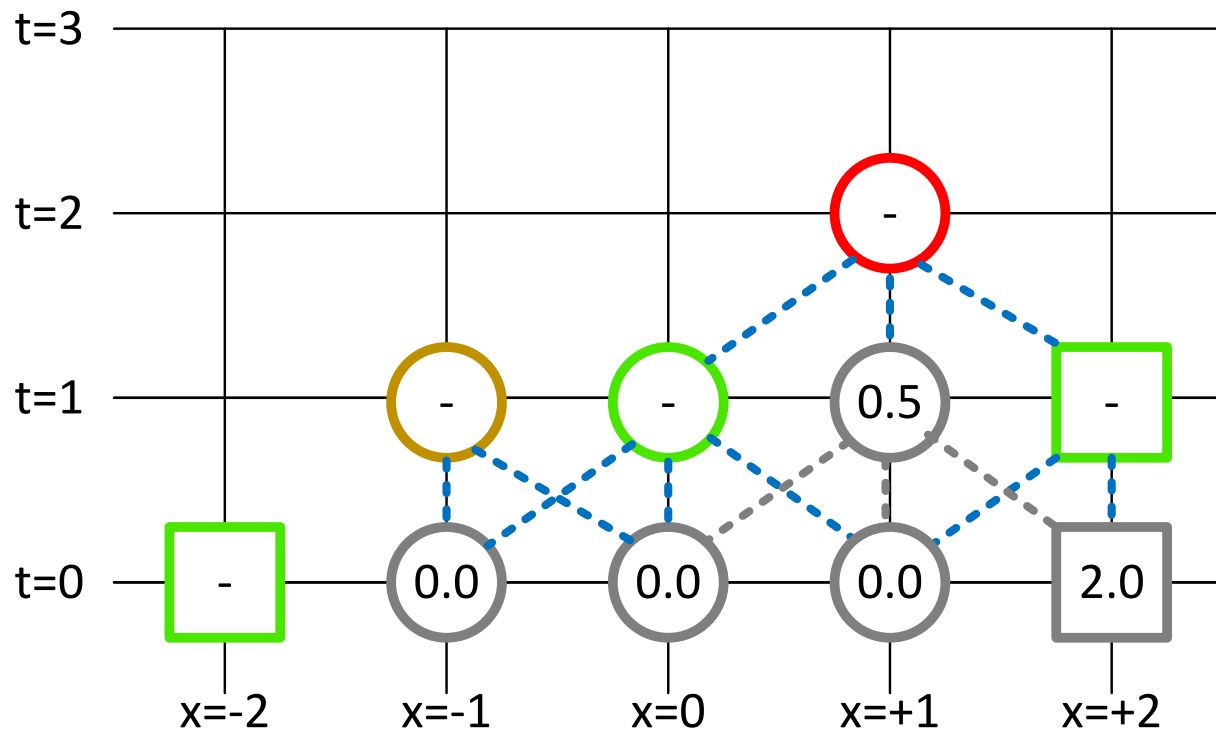


id:4

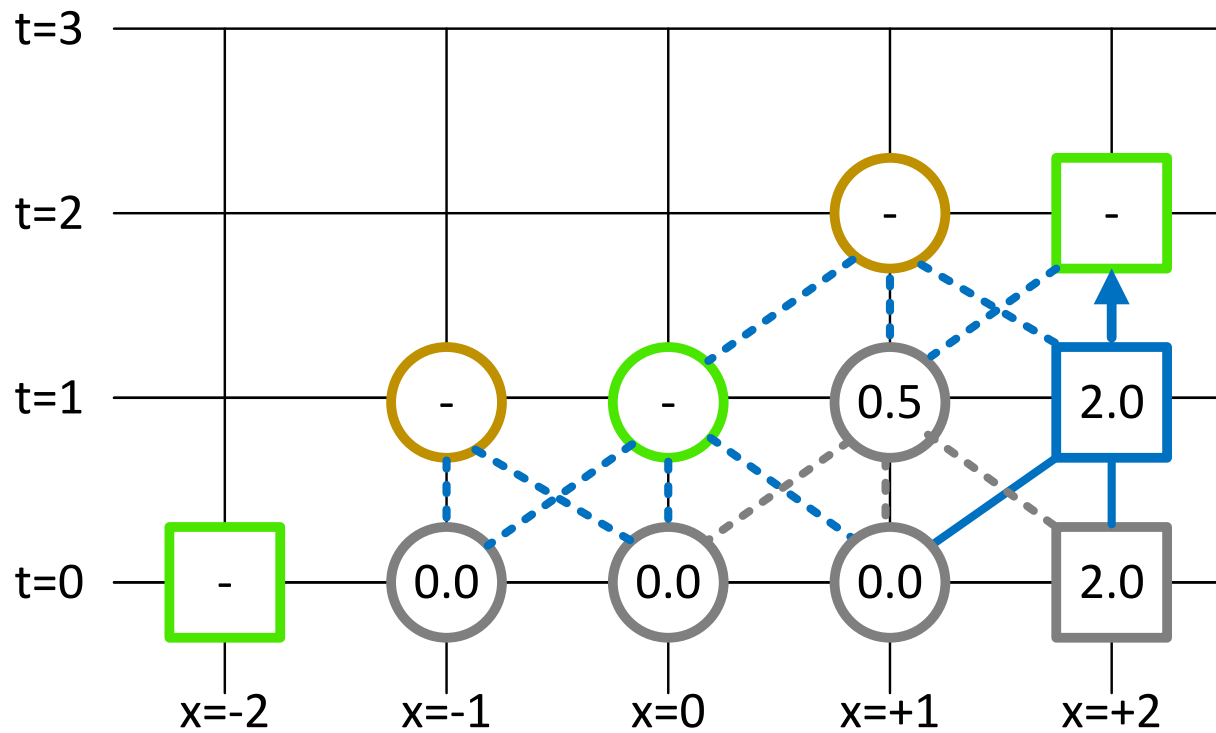




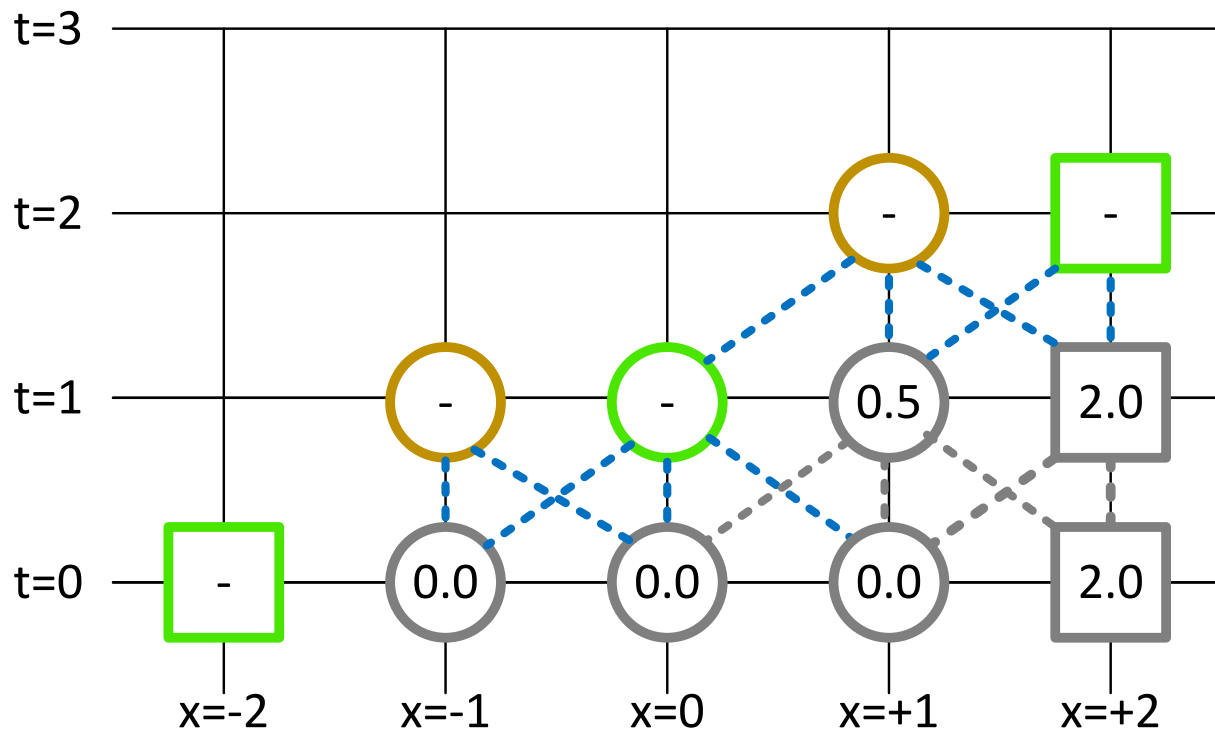
id:5



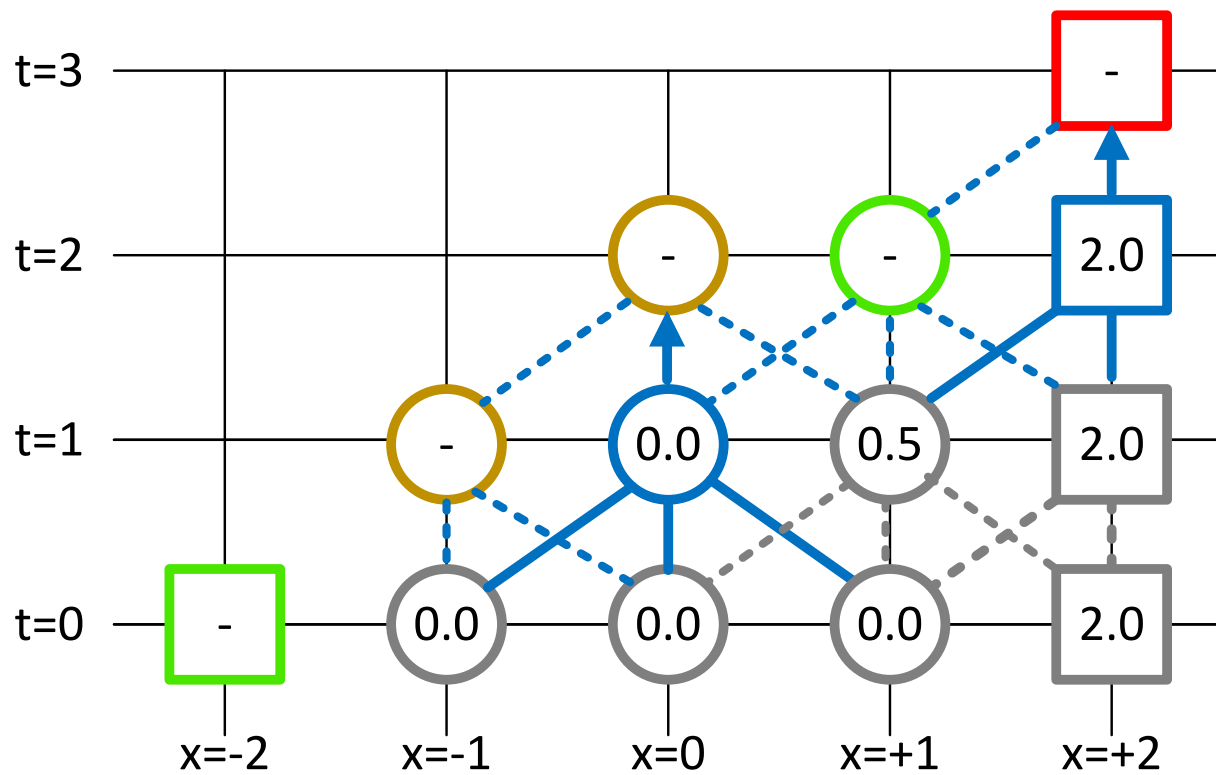
id:6



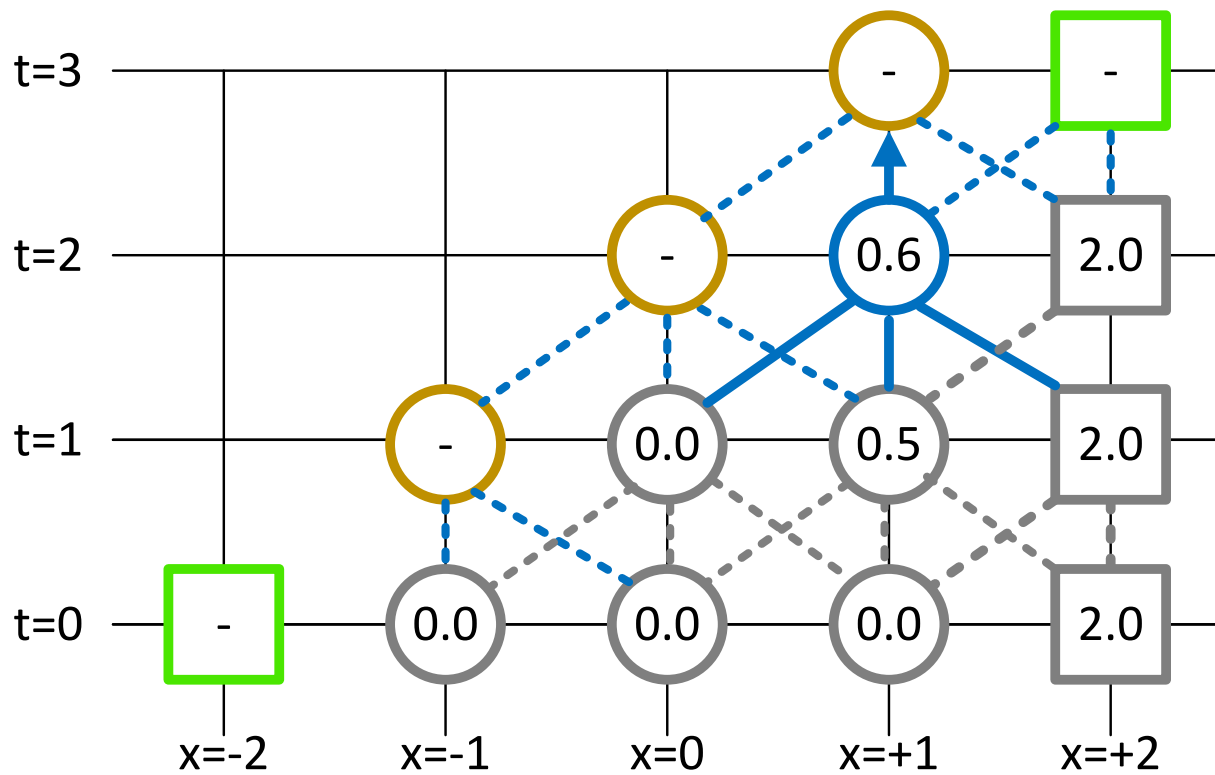
id:7



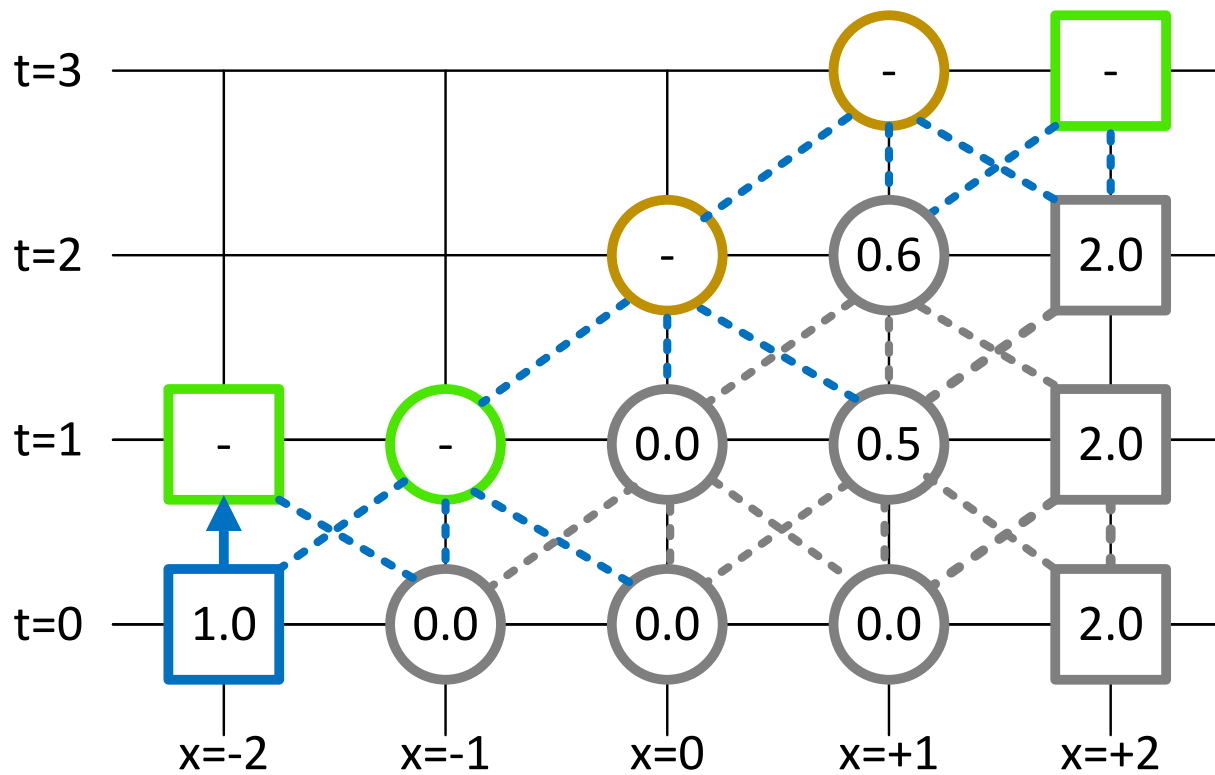
id:8



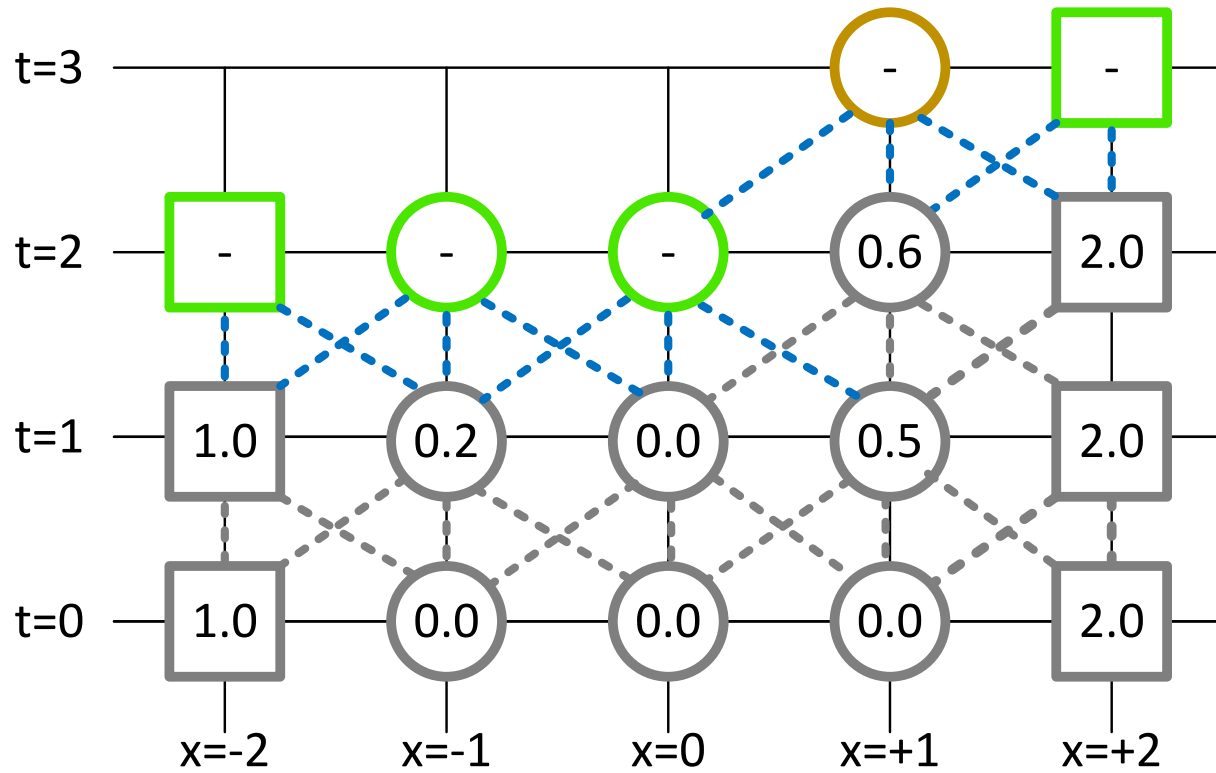
id:9



id:10

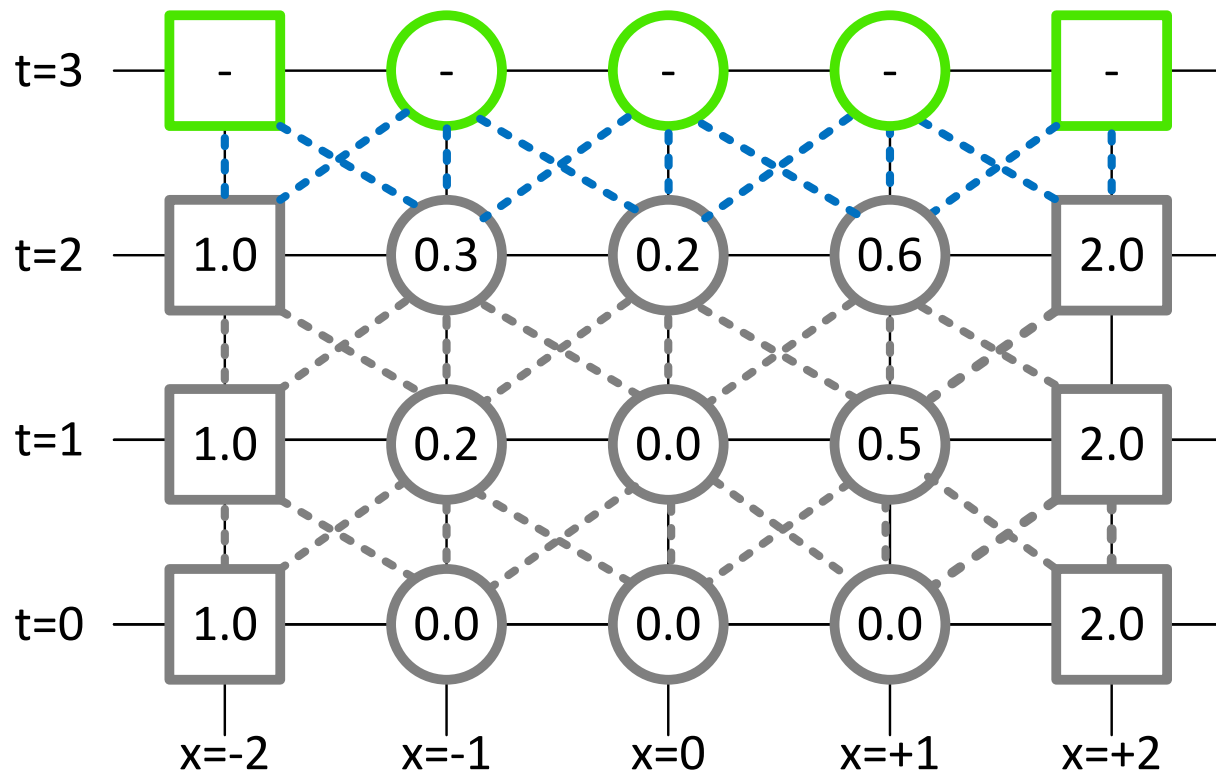


id:11



id:12



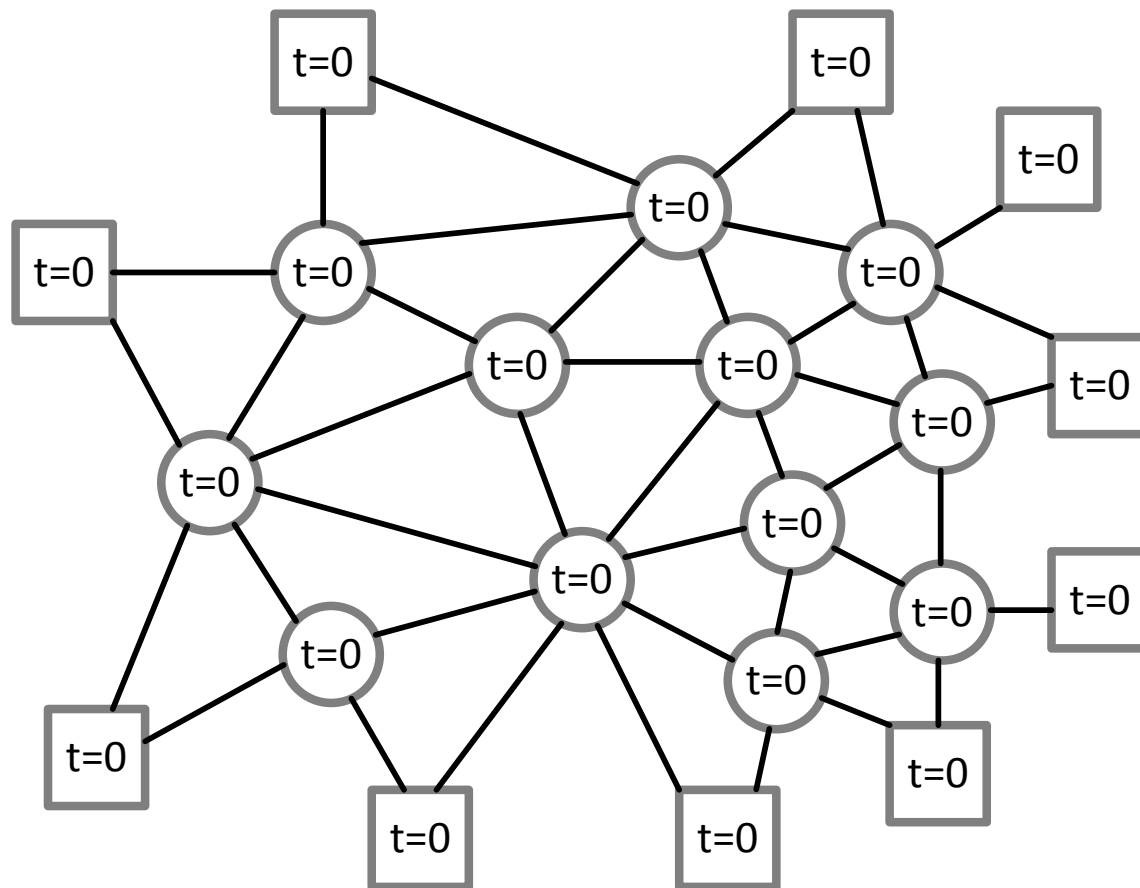


id:13

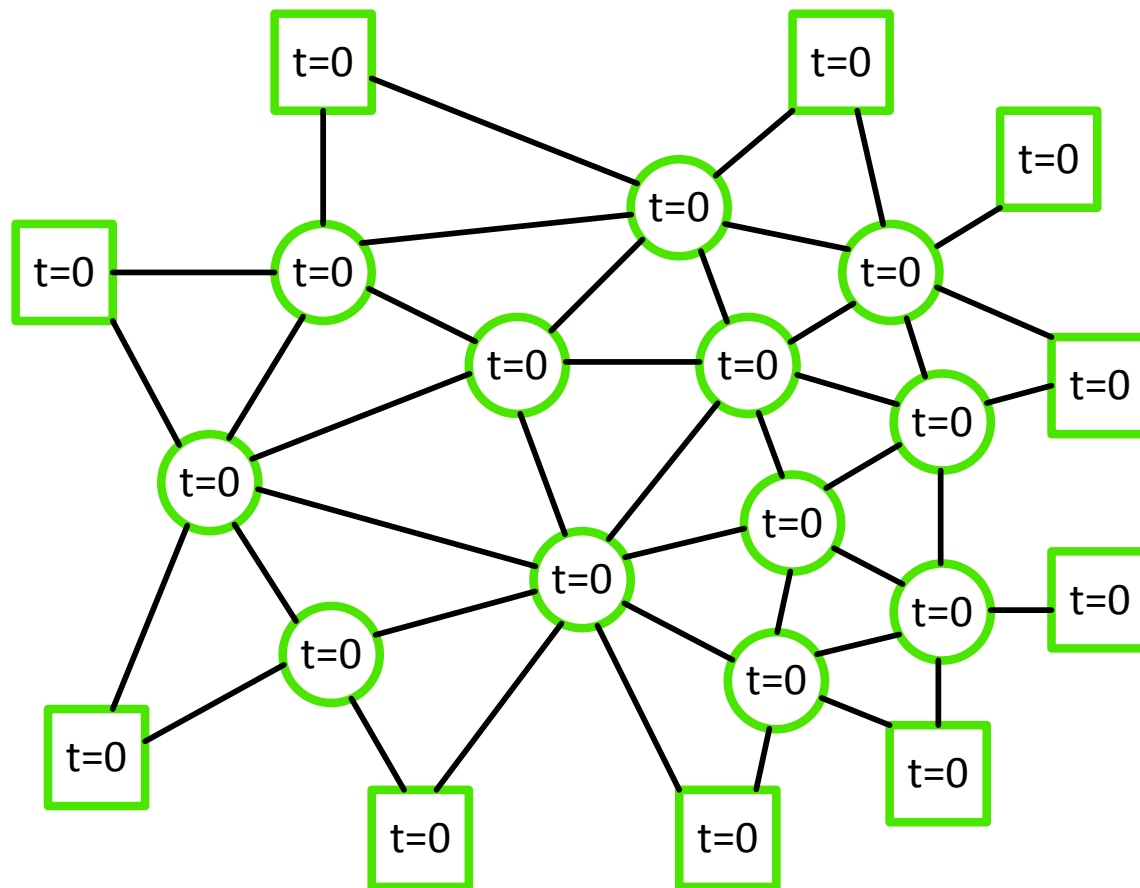


# Over to you

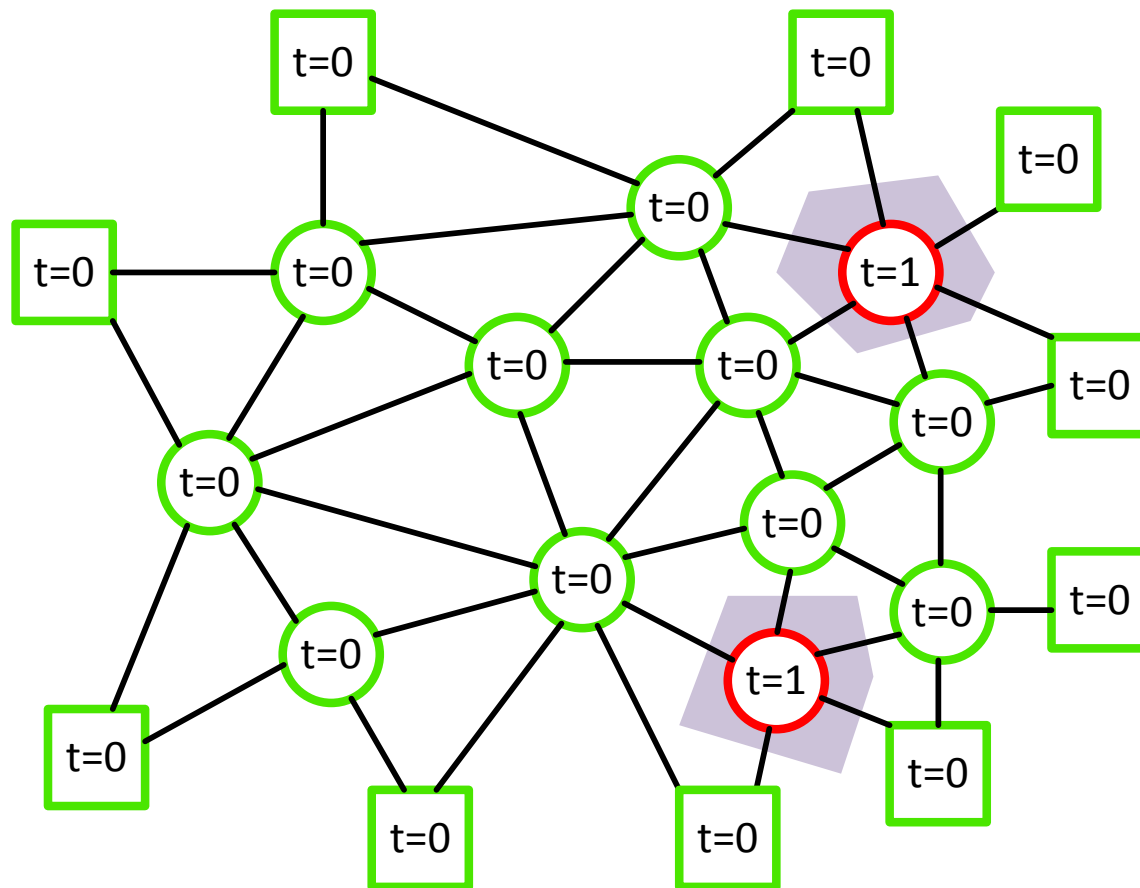
- There are diminishing returns of giving more and more detail
- Best approach is to think about how to solve:
  - Simple problems very well: CW5
  - Complex problems fairly well: CW6
- There will be the oral assessment in early summer term
  - Think about what you are doing as you do it
  - You will not have time to do everything
    - Remember what you could have done
    - Remember what you should have done (given time, hindsight, ...)



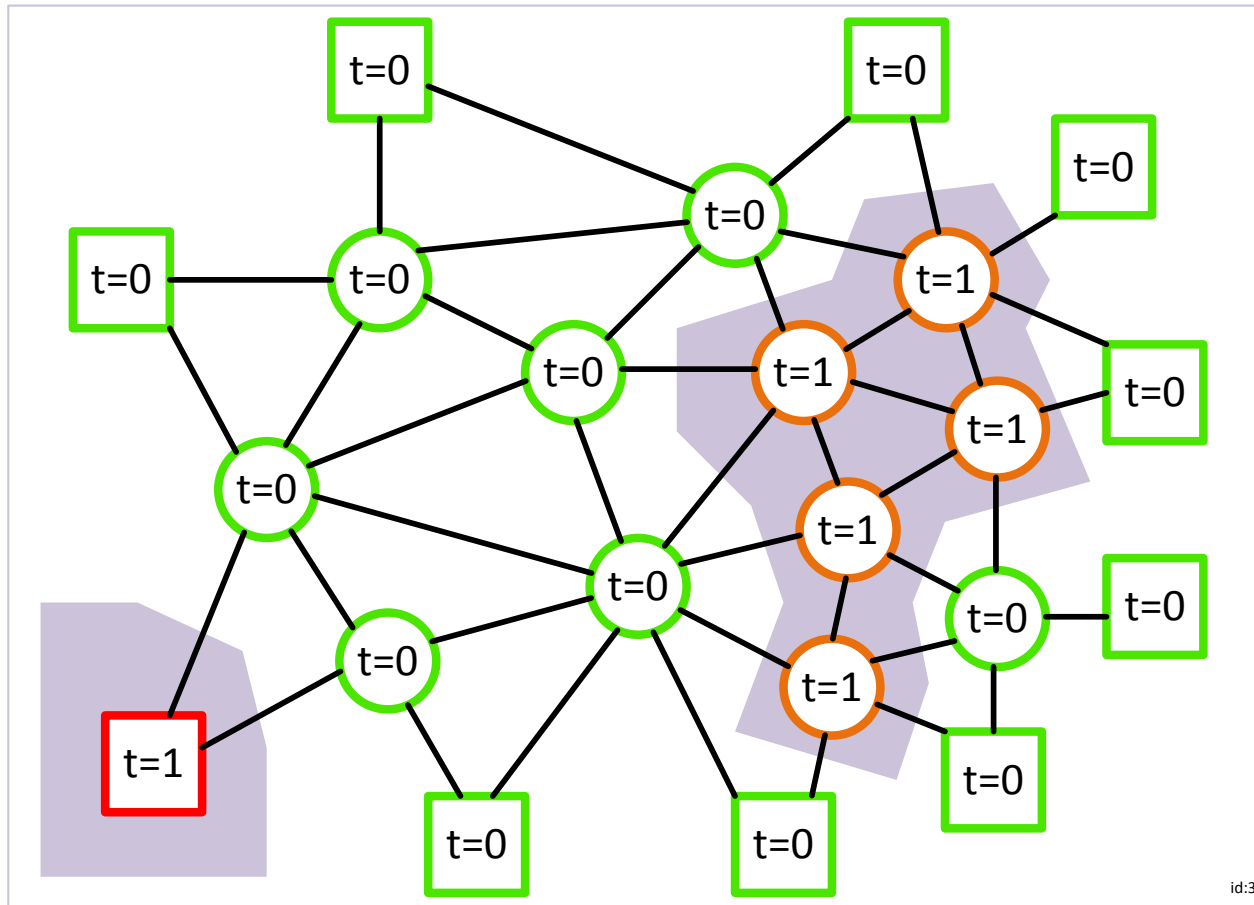
id:0

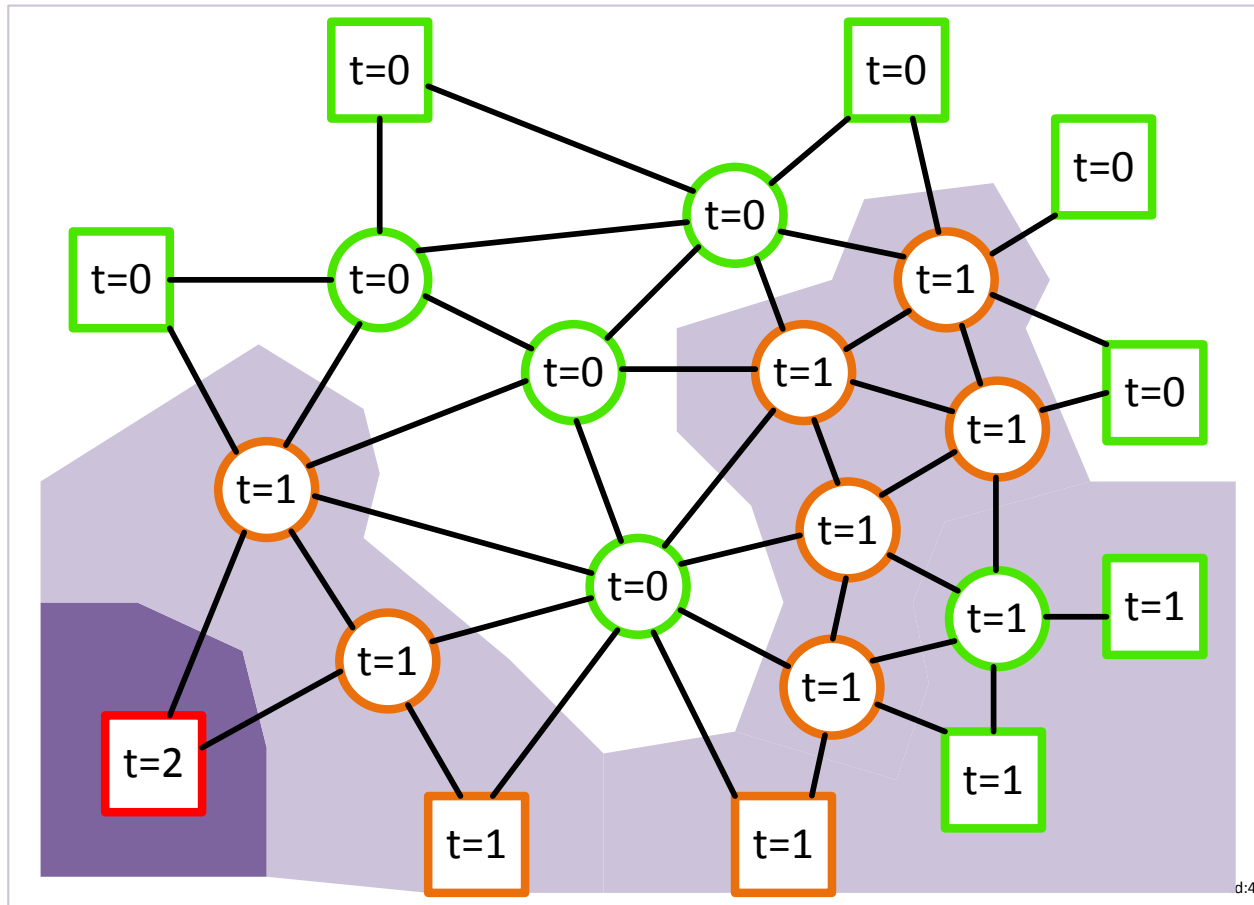


id:1



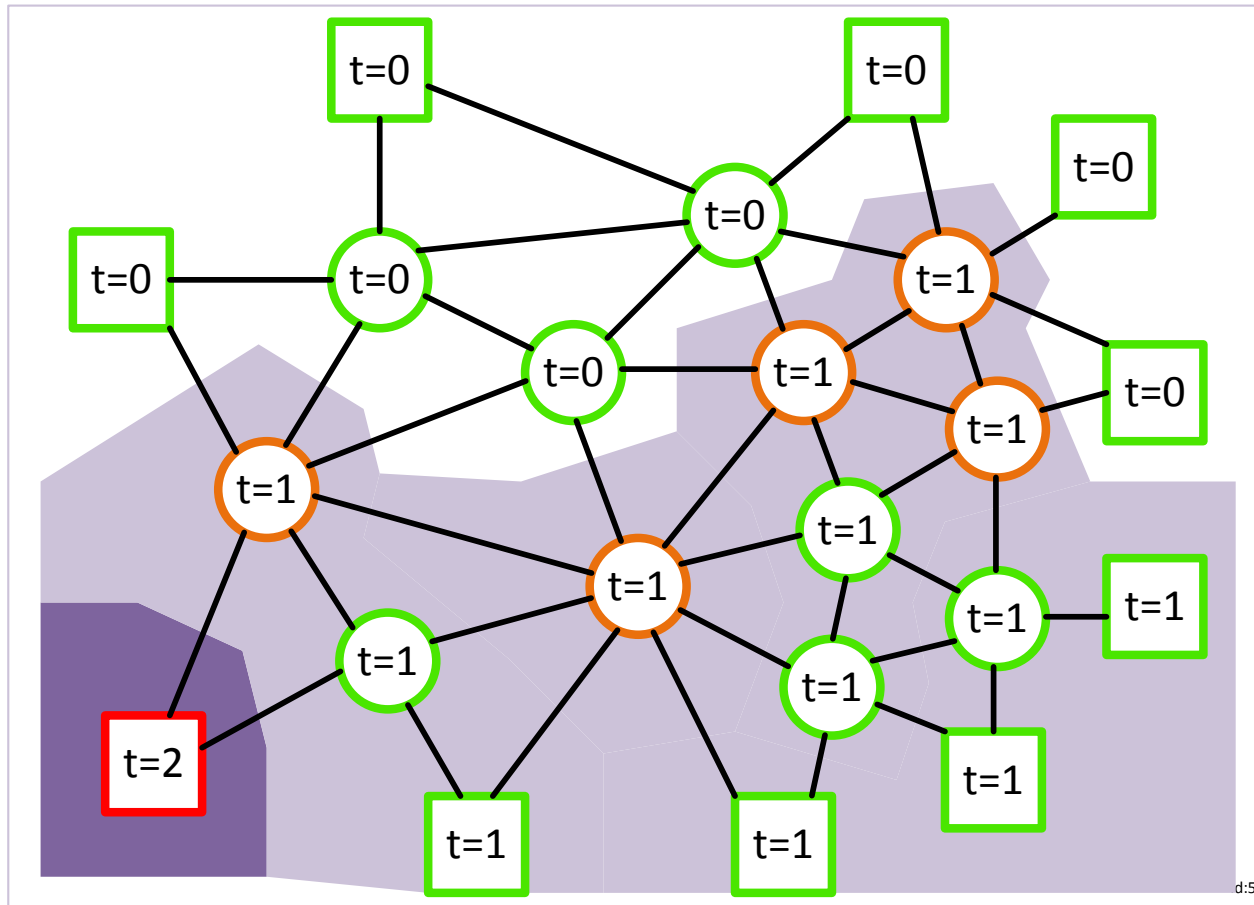
id:2

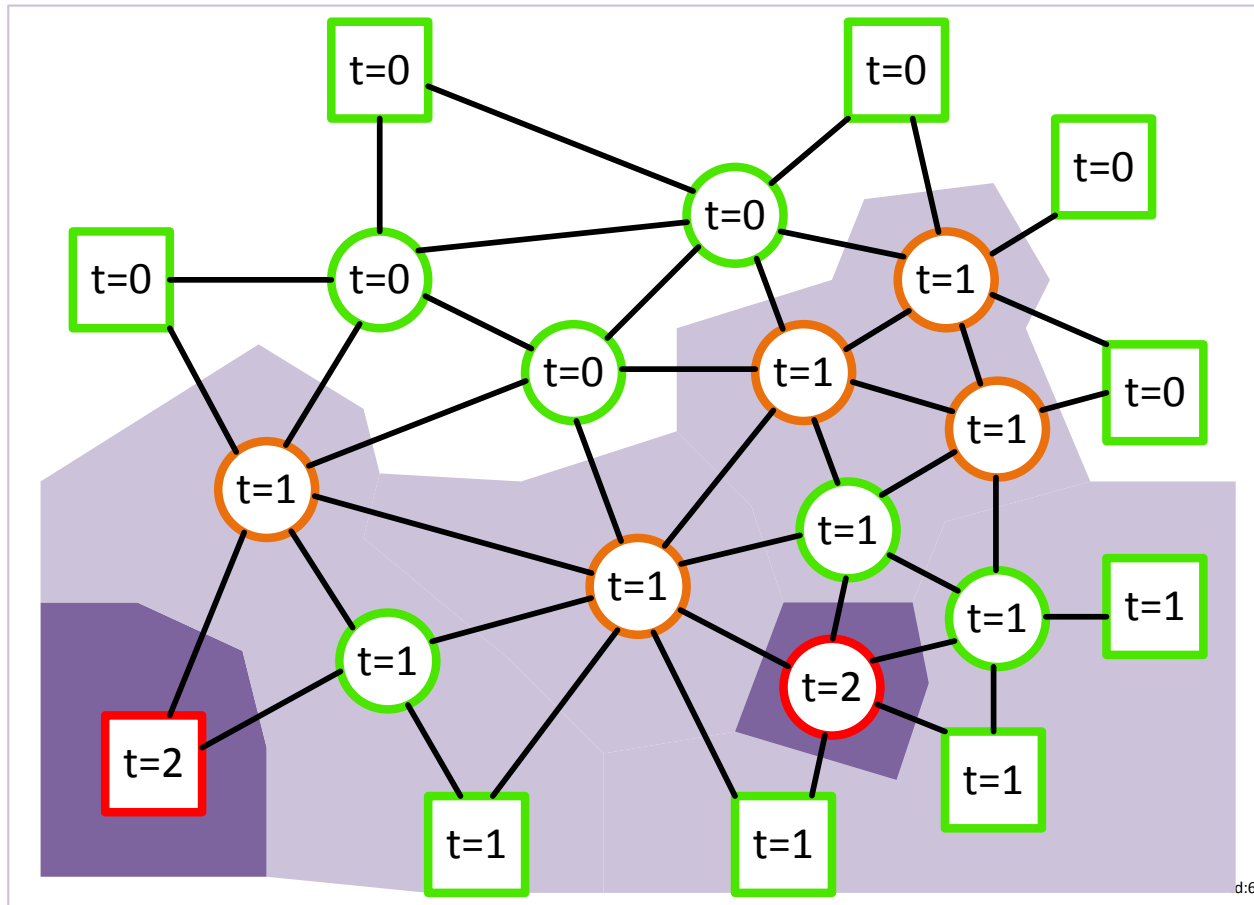




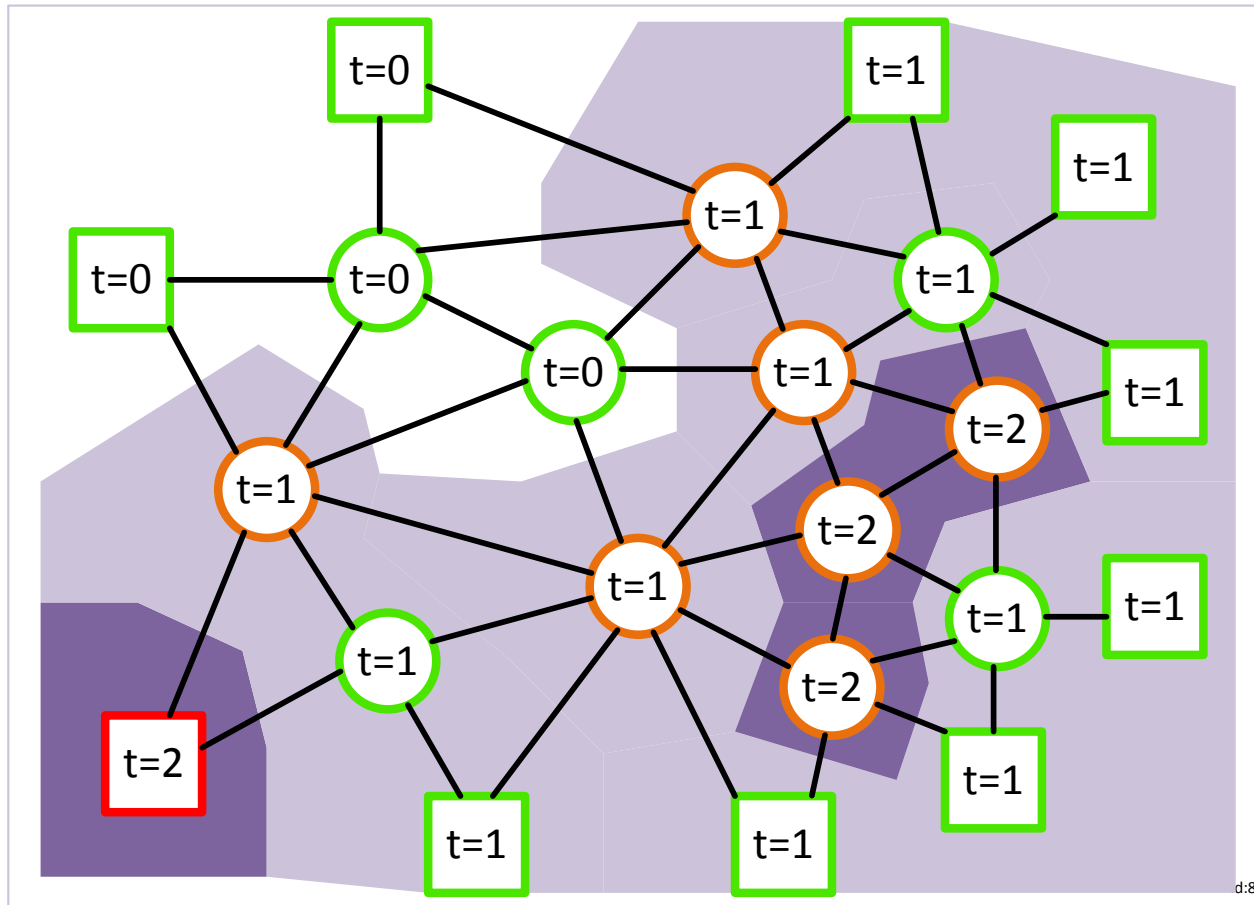
d:4

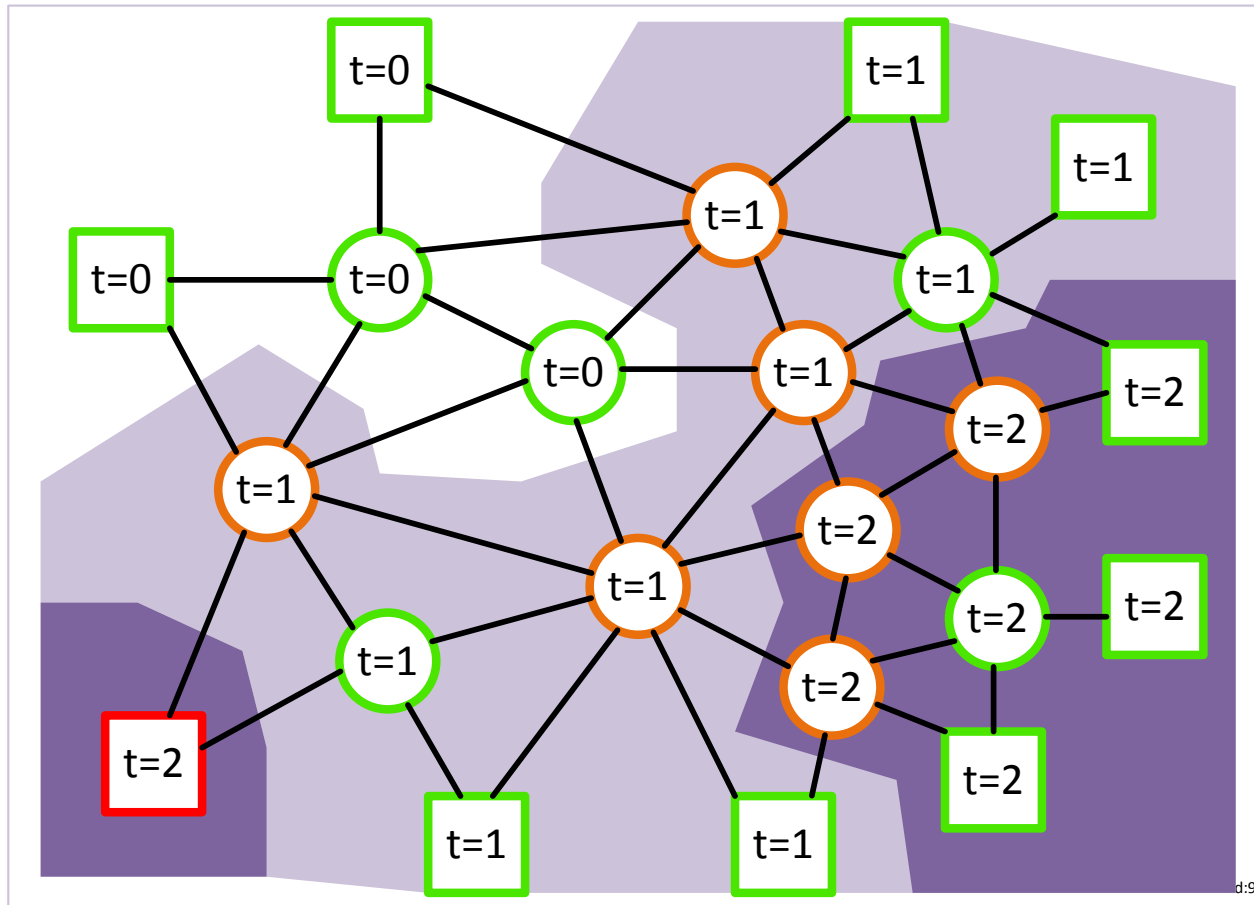


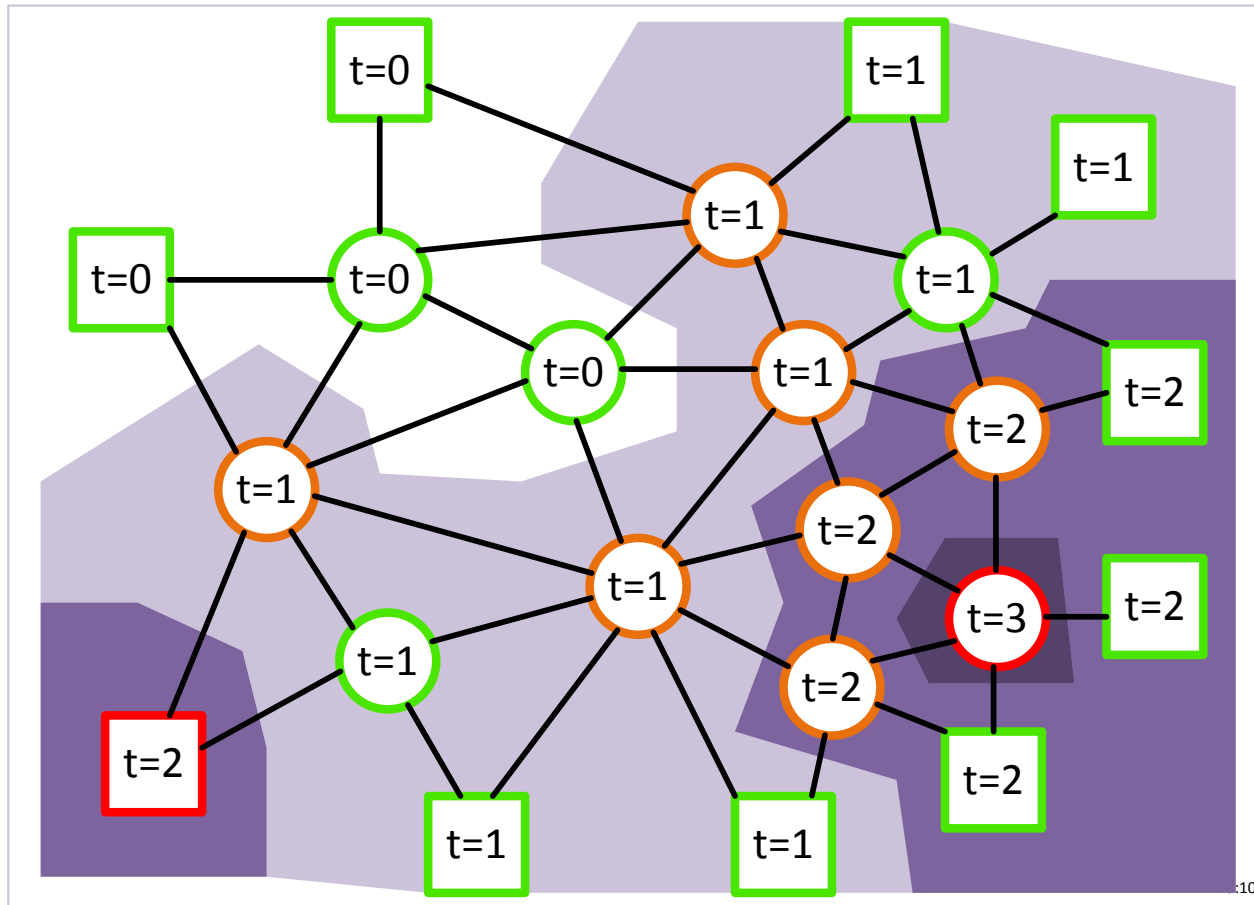




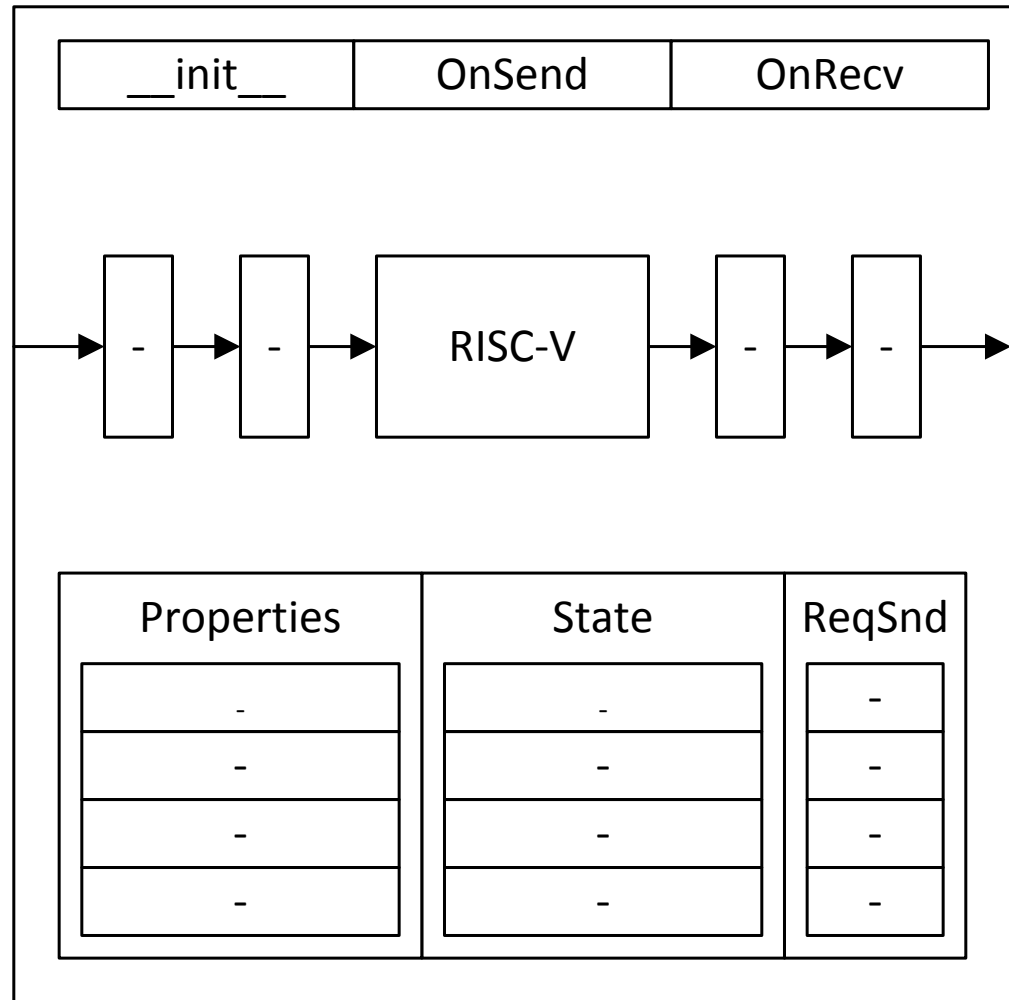
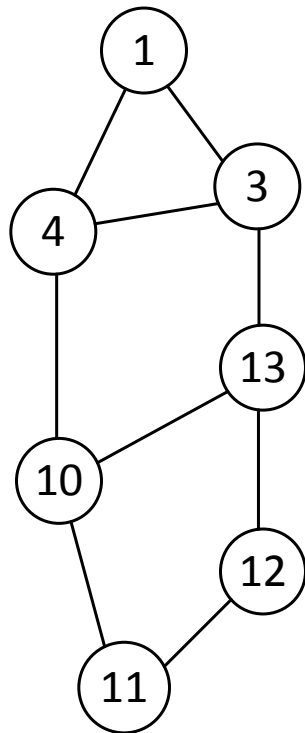




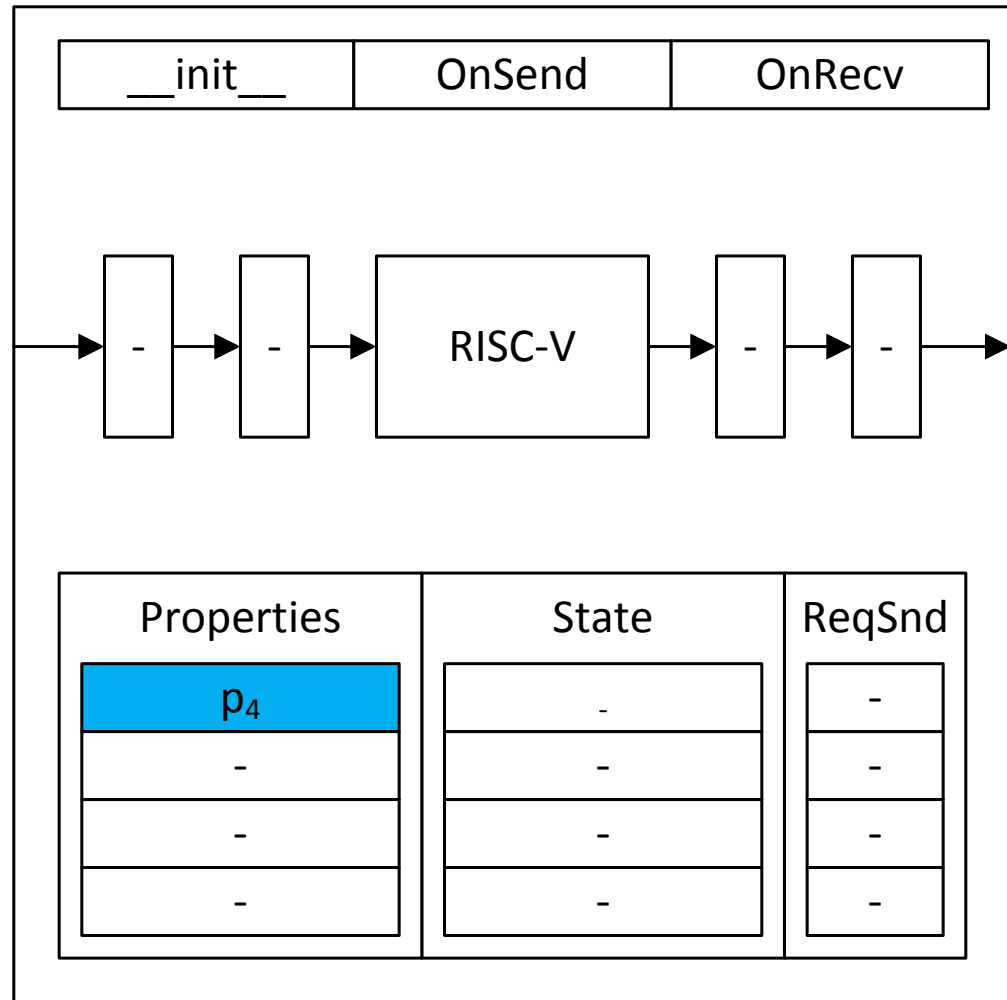
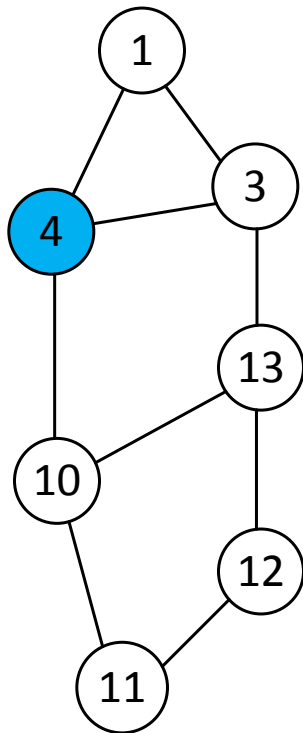


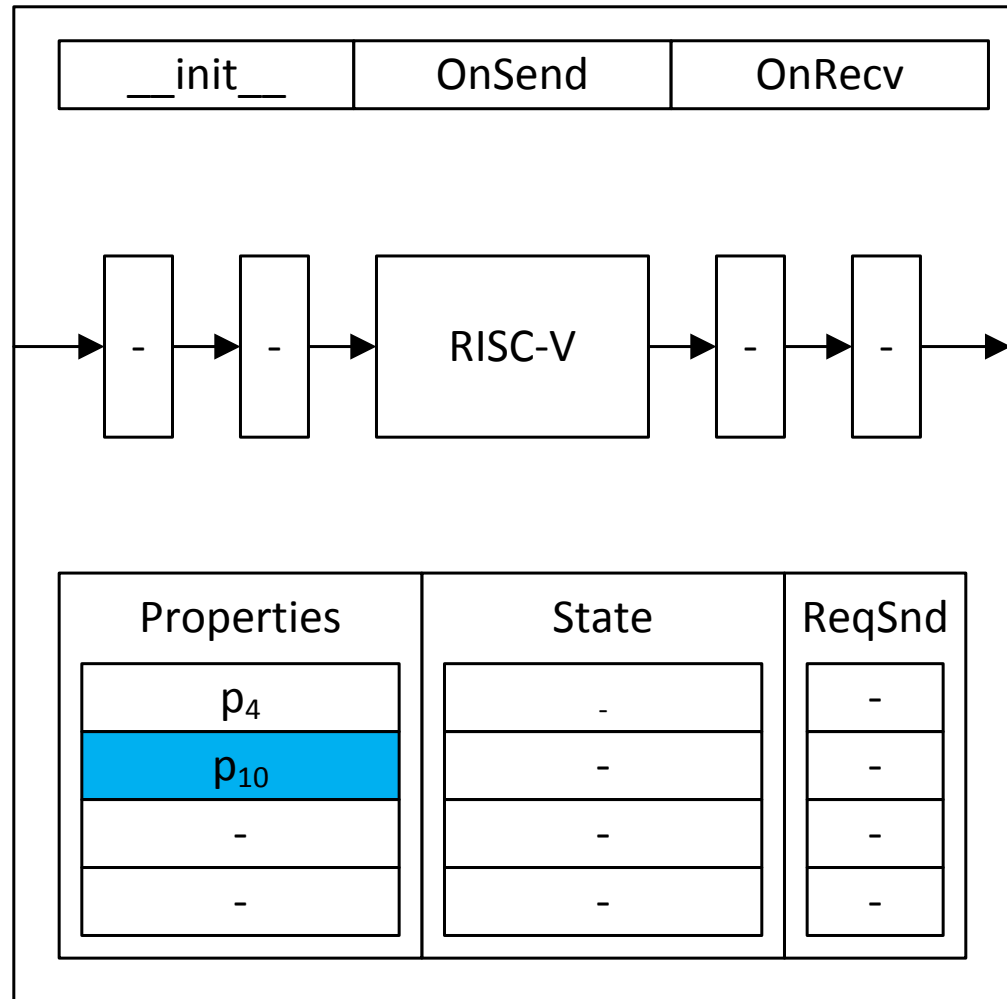
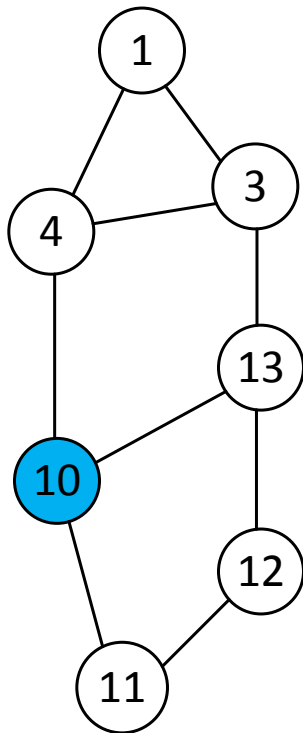


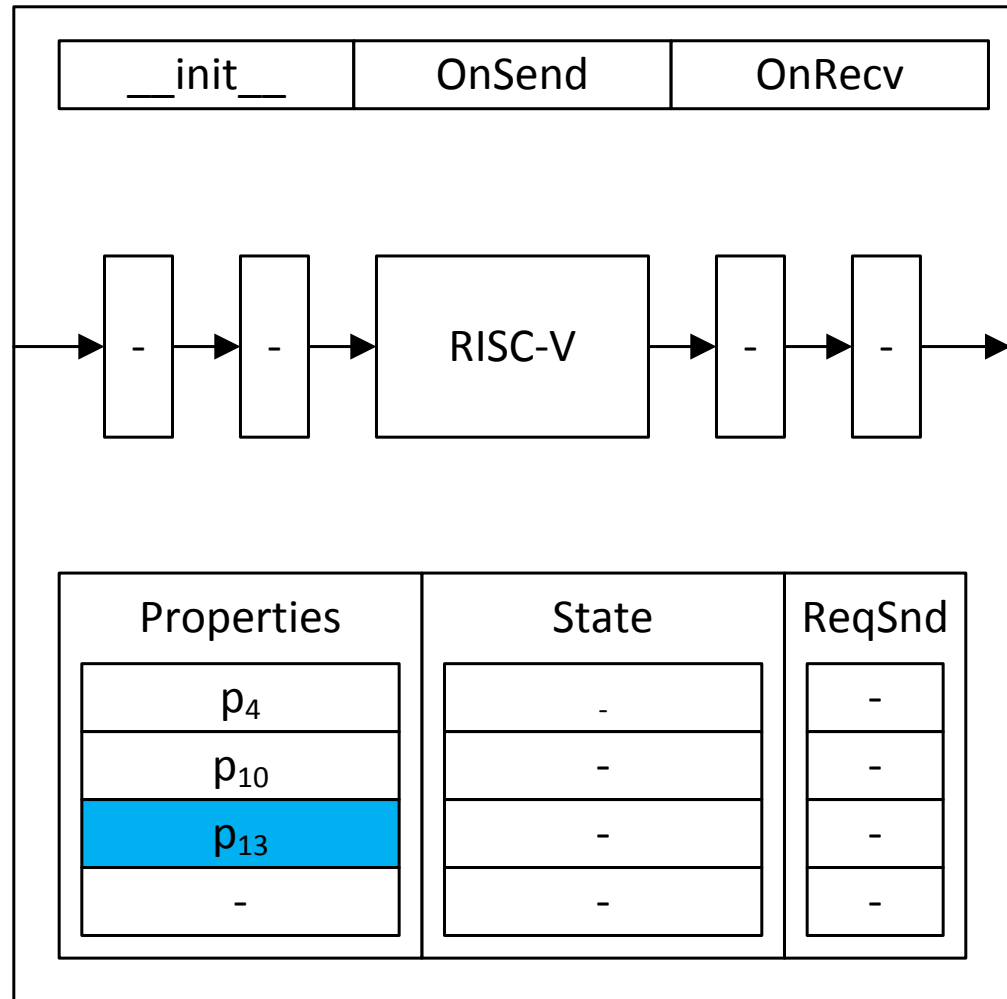
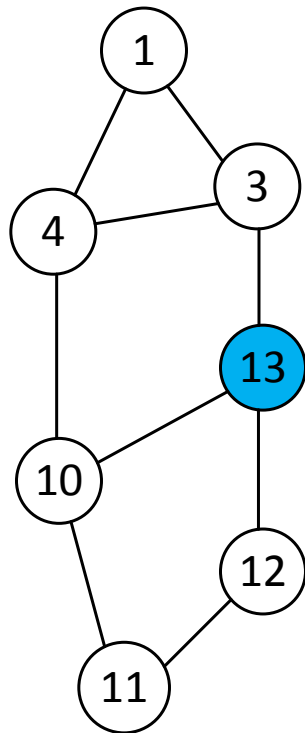
# device to core mapping

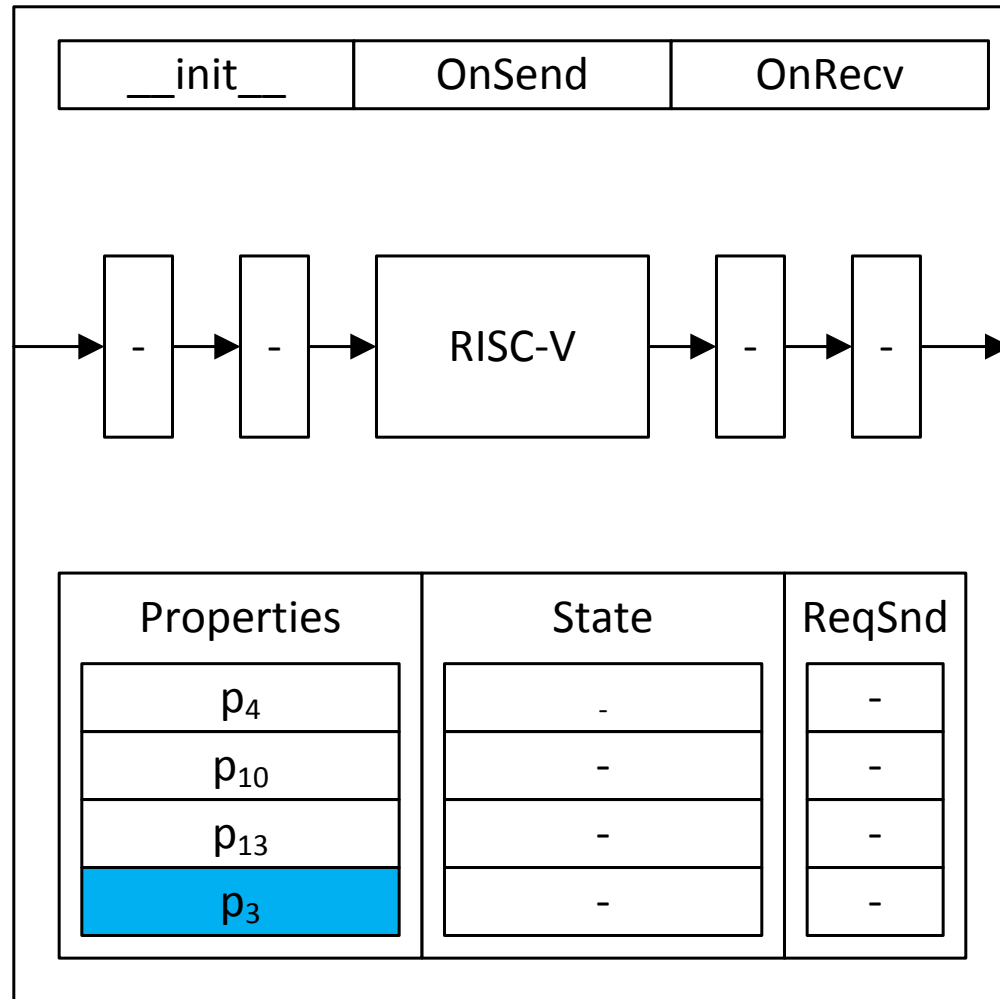
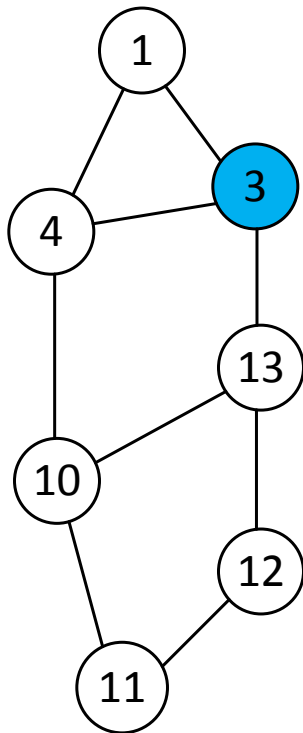


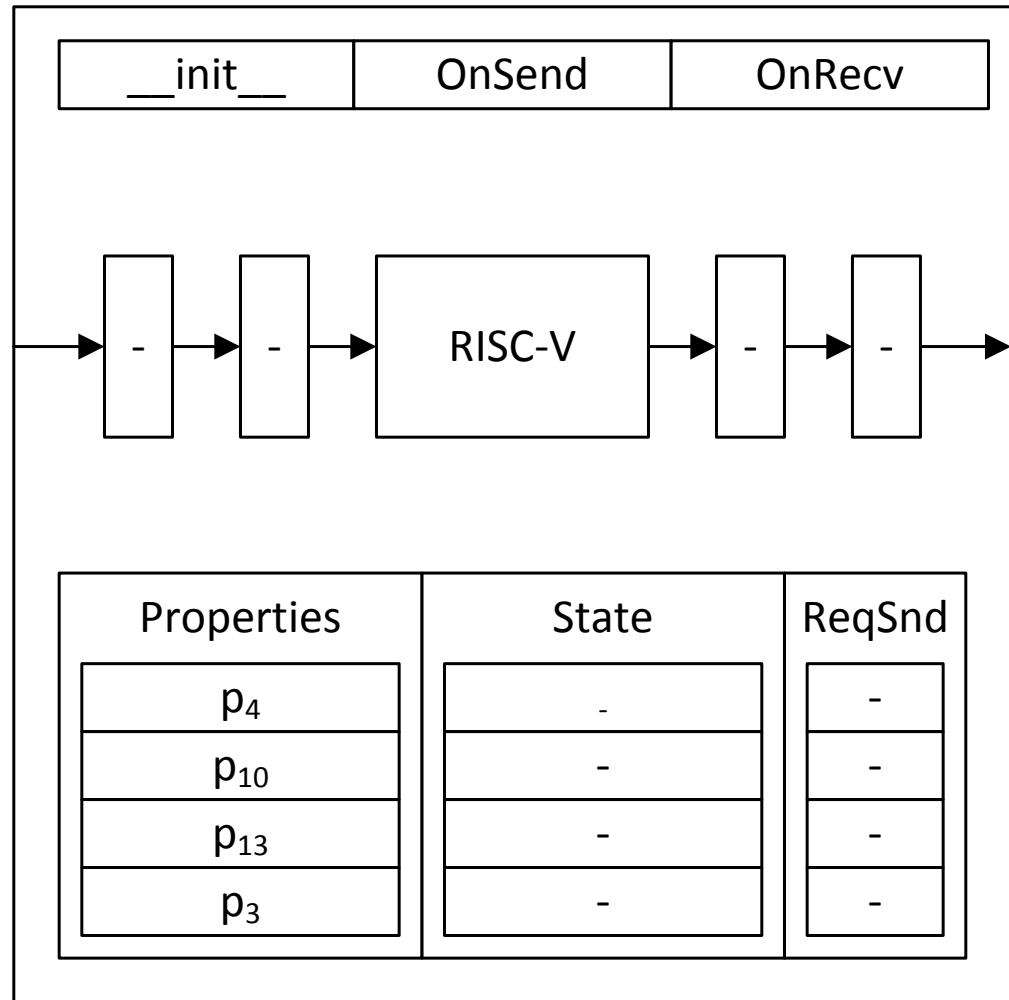
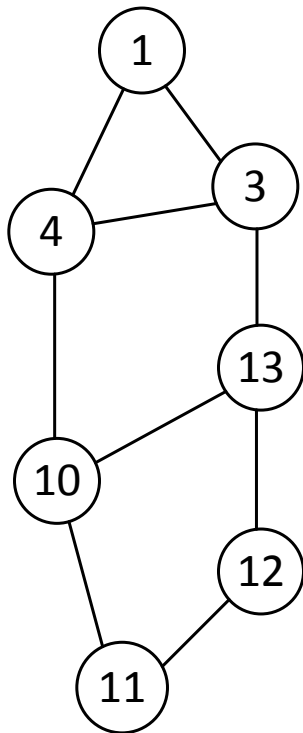


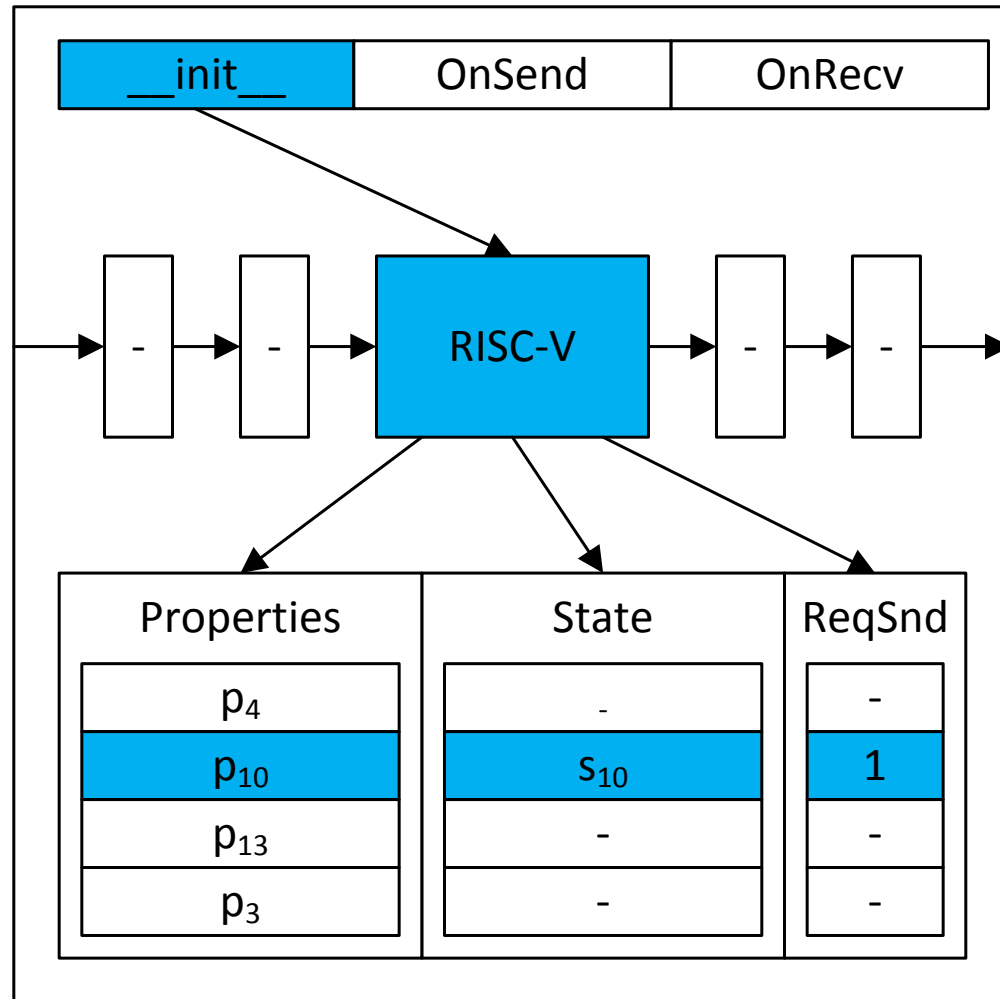
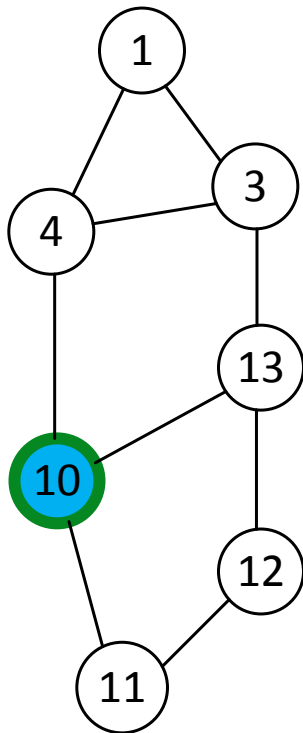


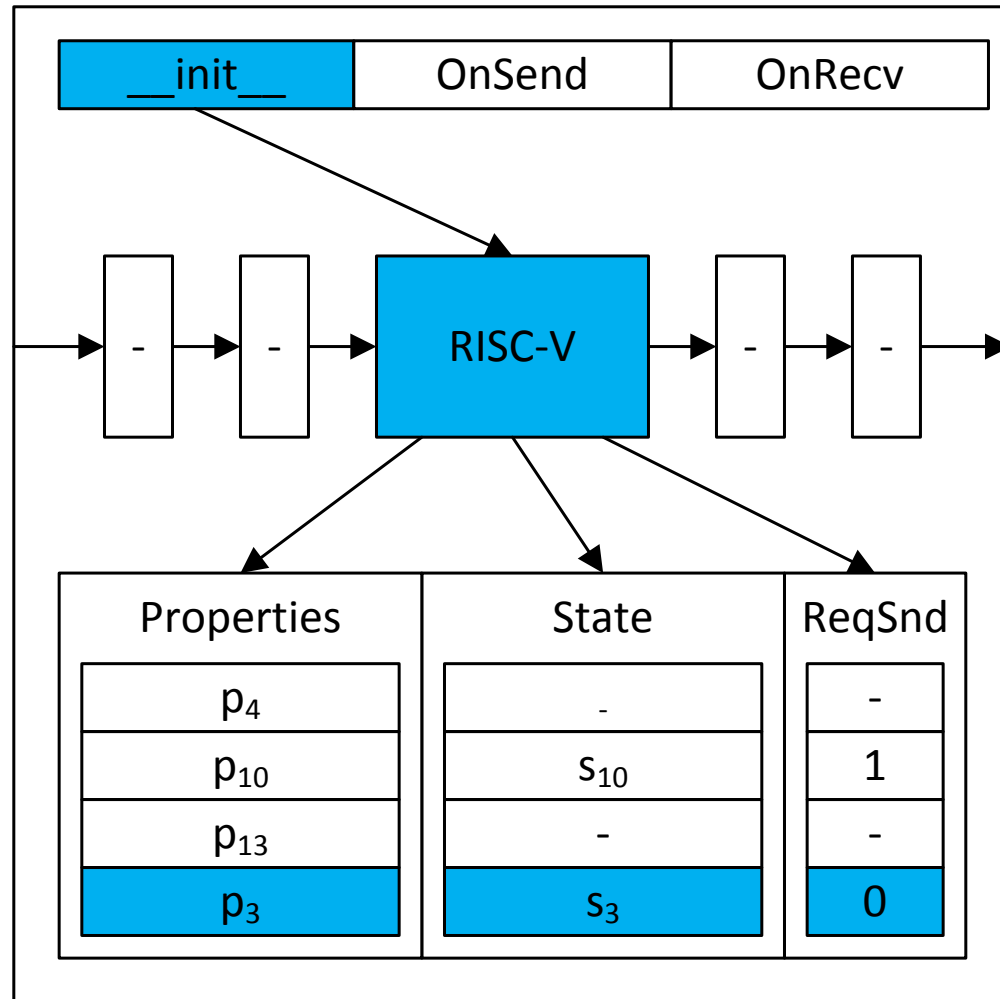
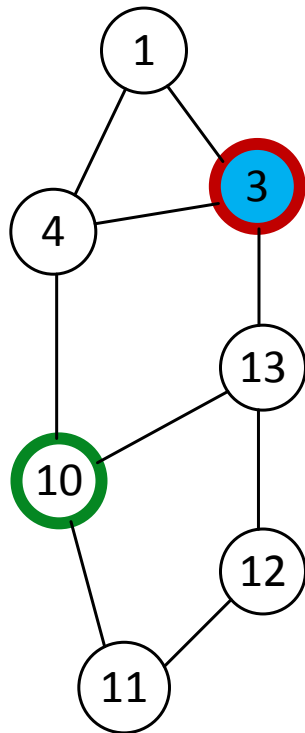


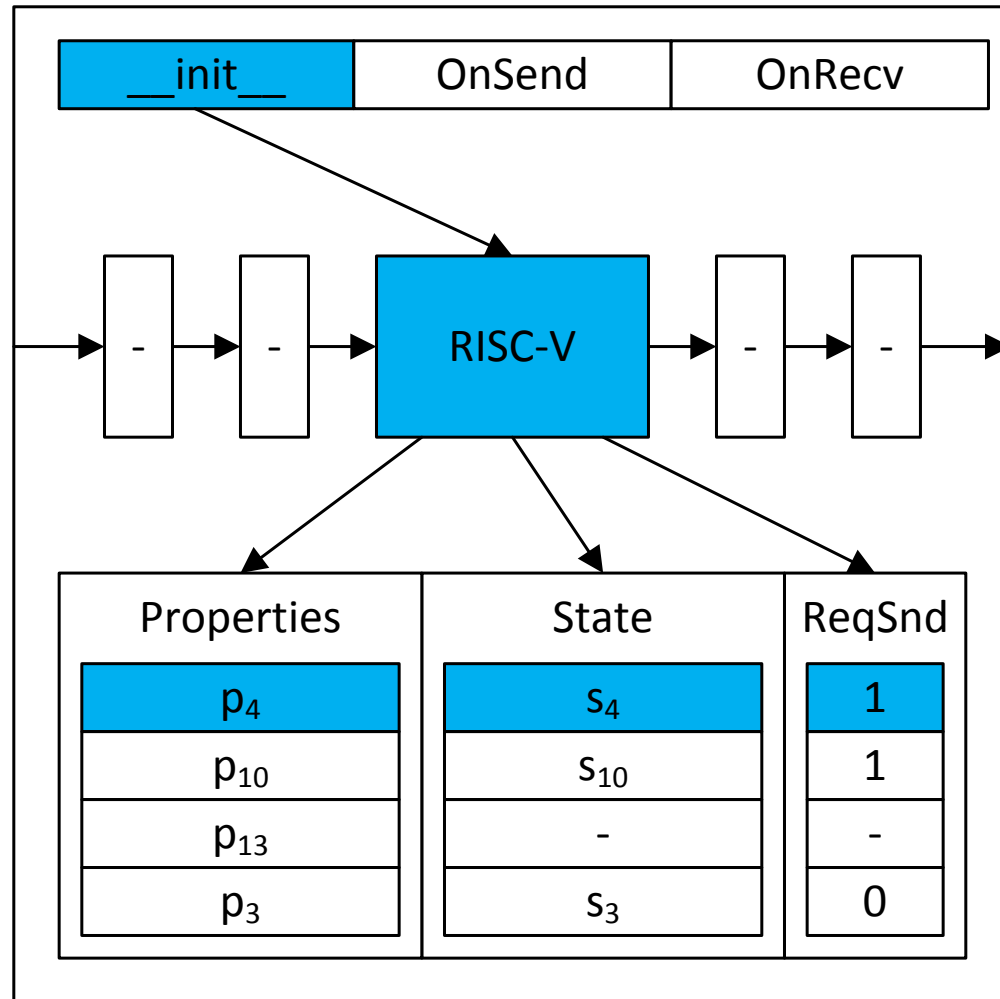
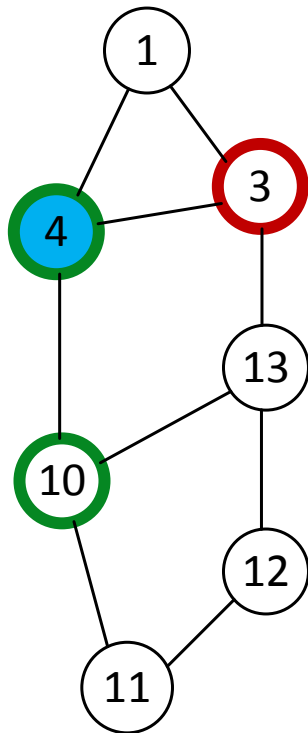




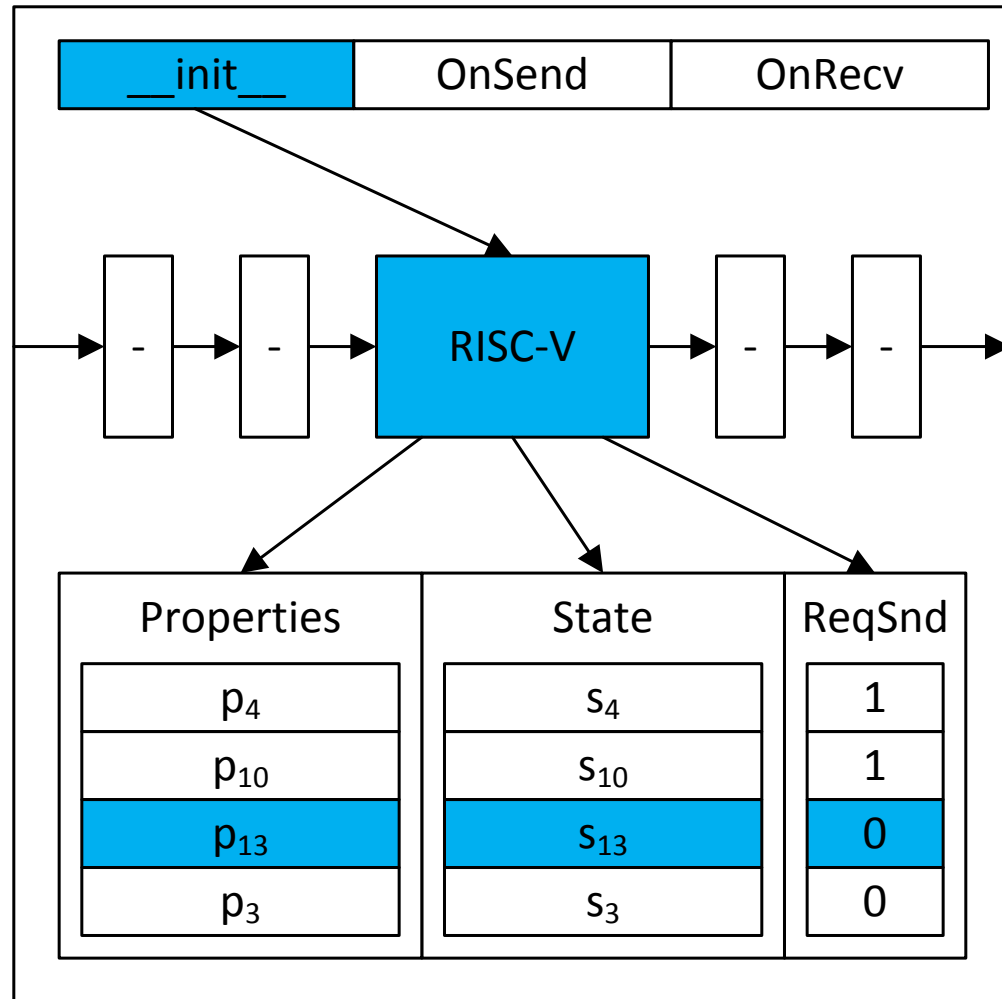
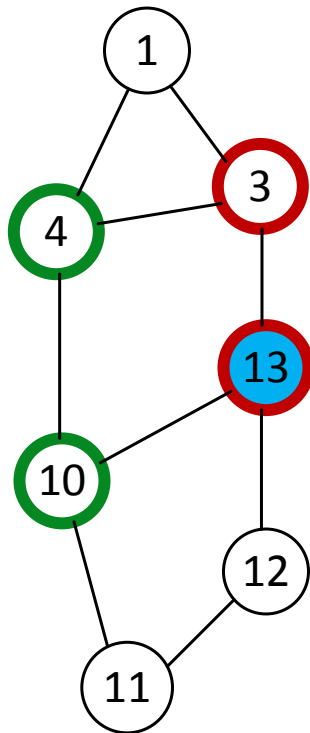


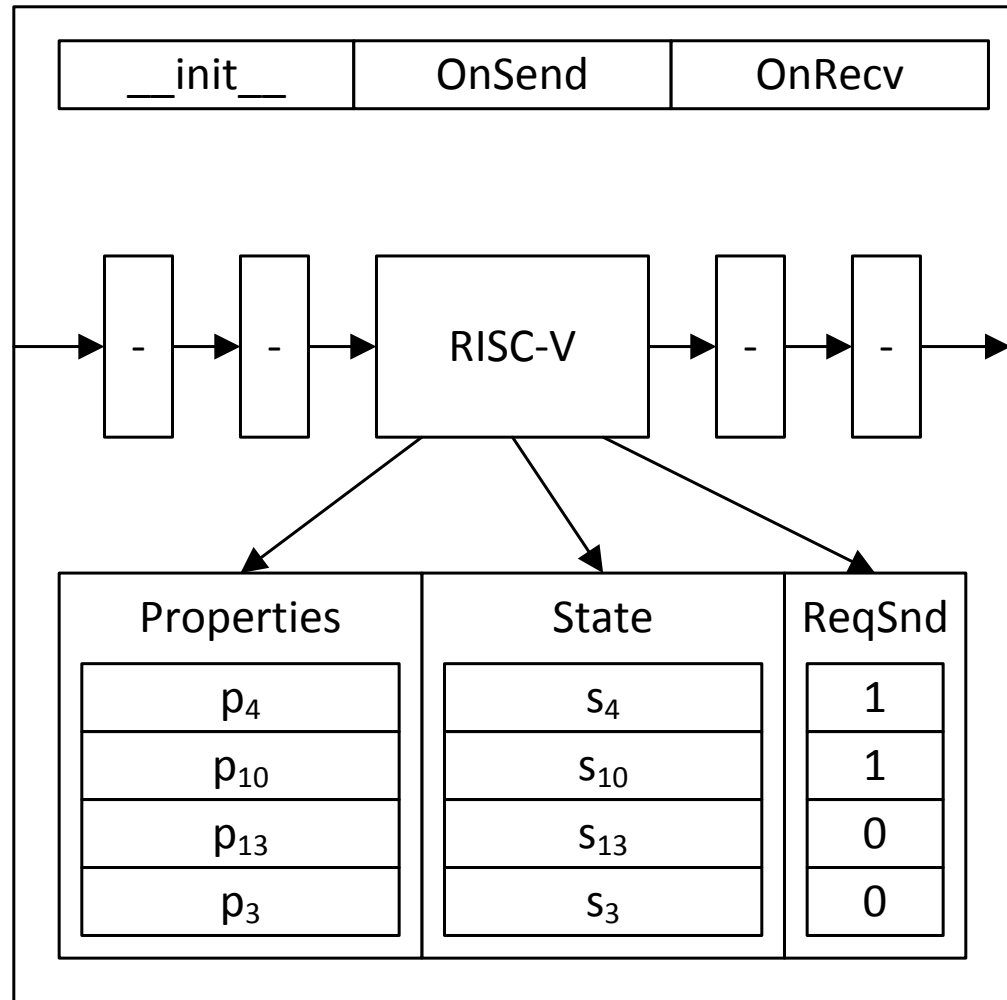
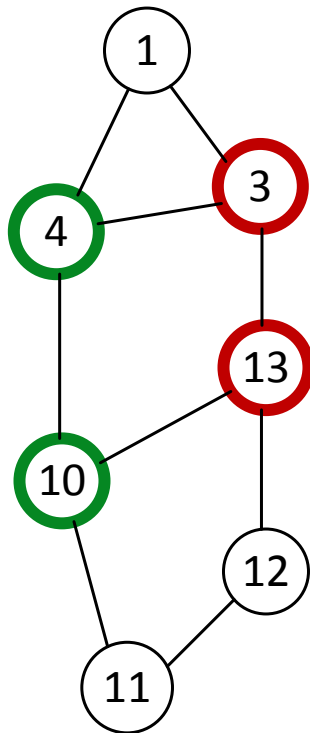


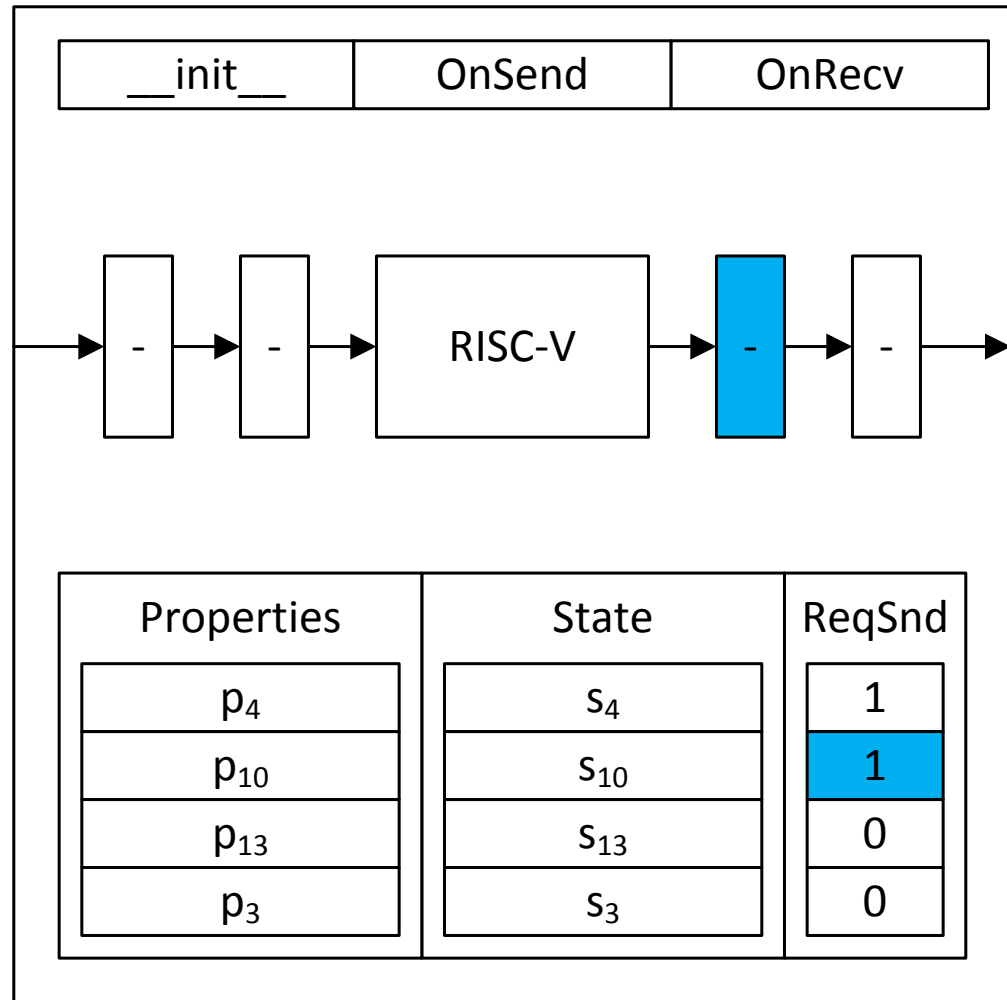
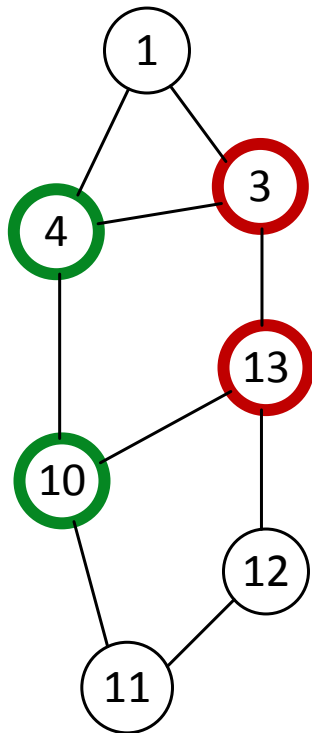


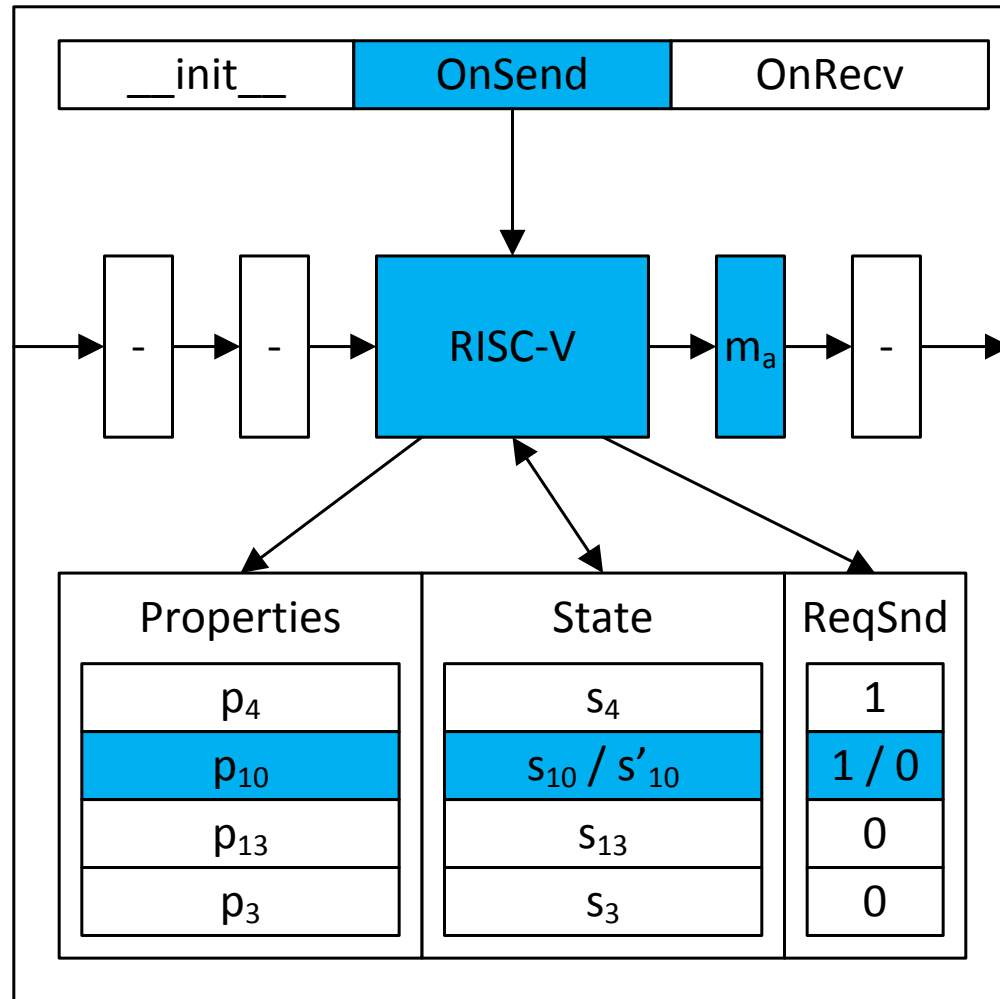
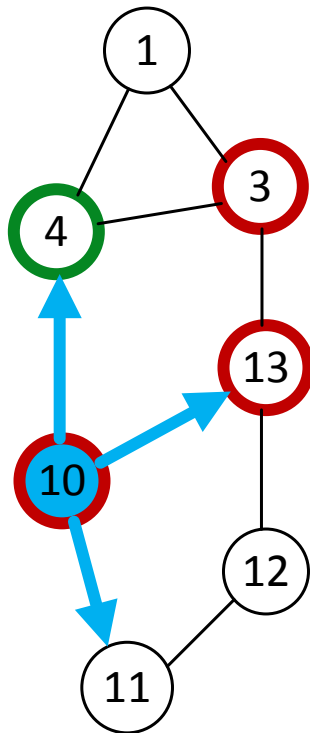


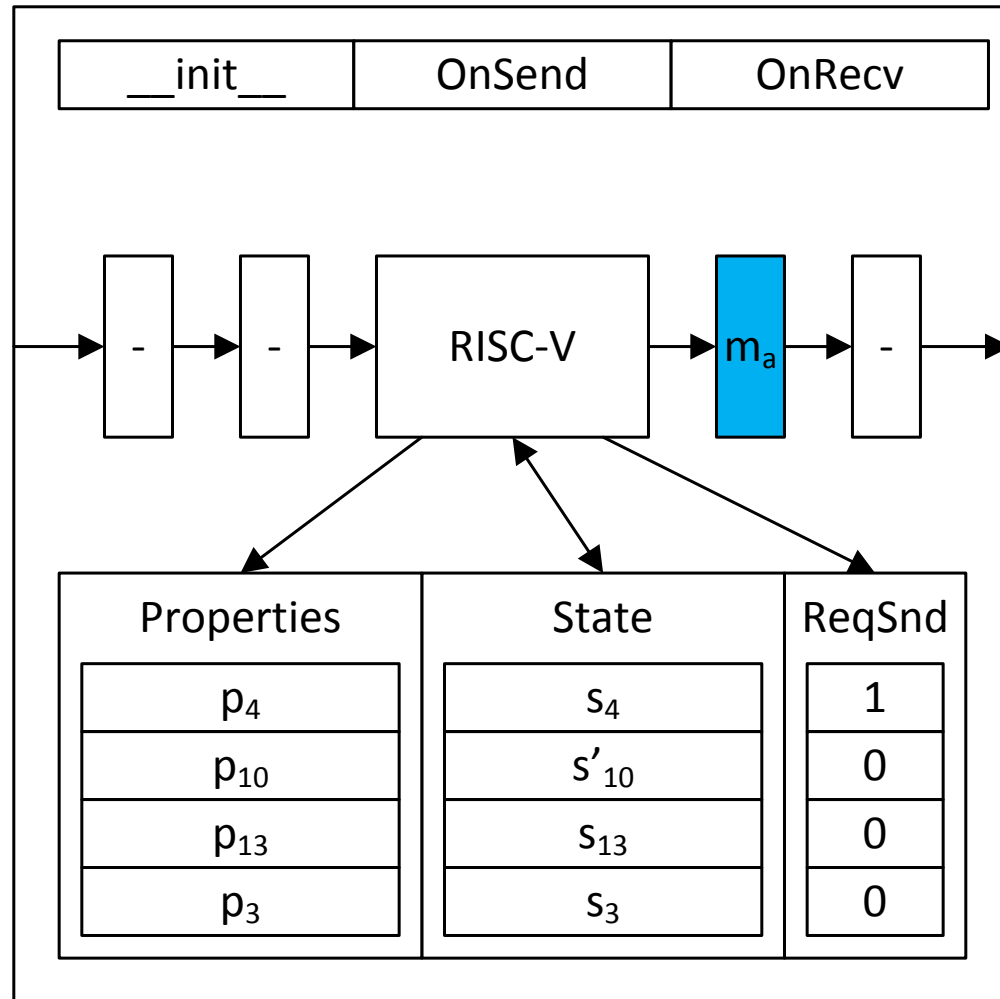
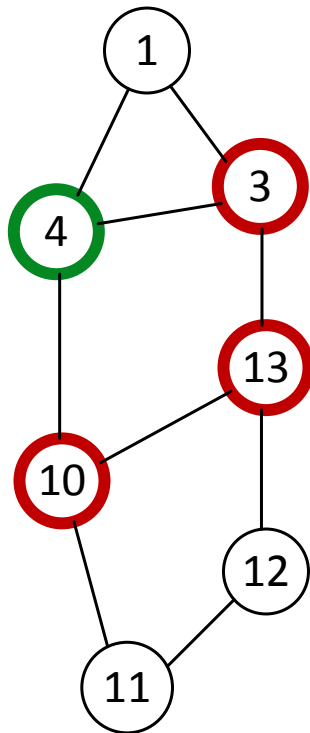


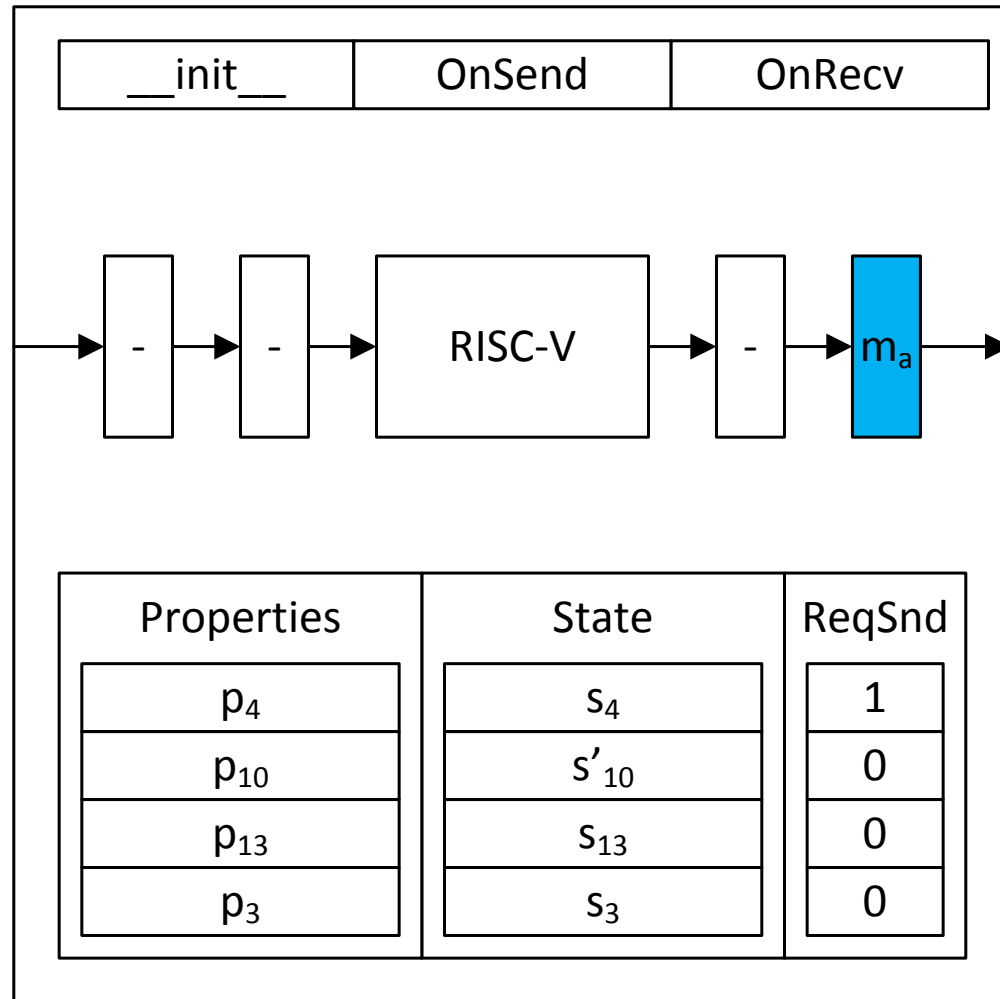
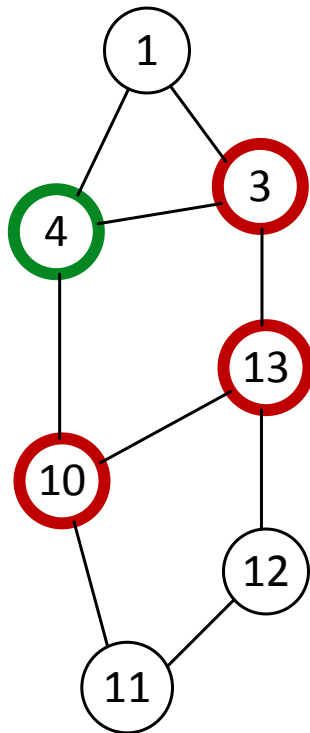


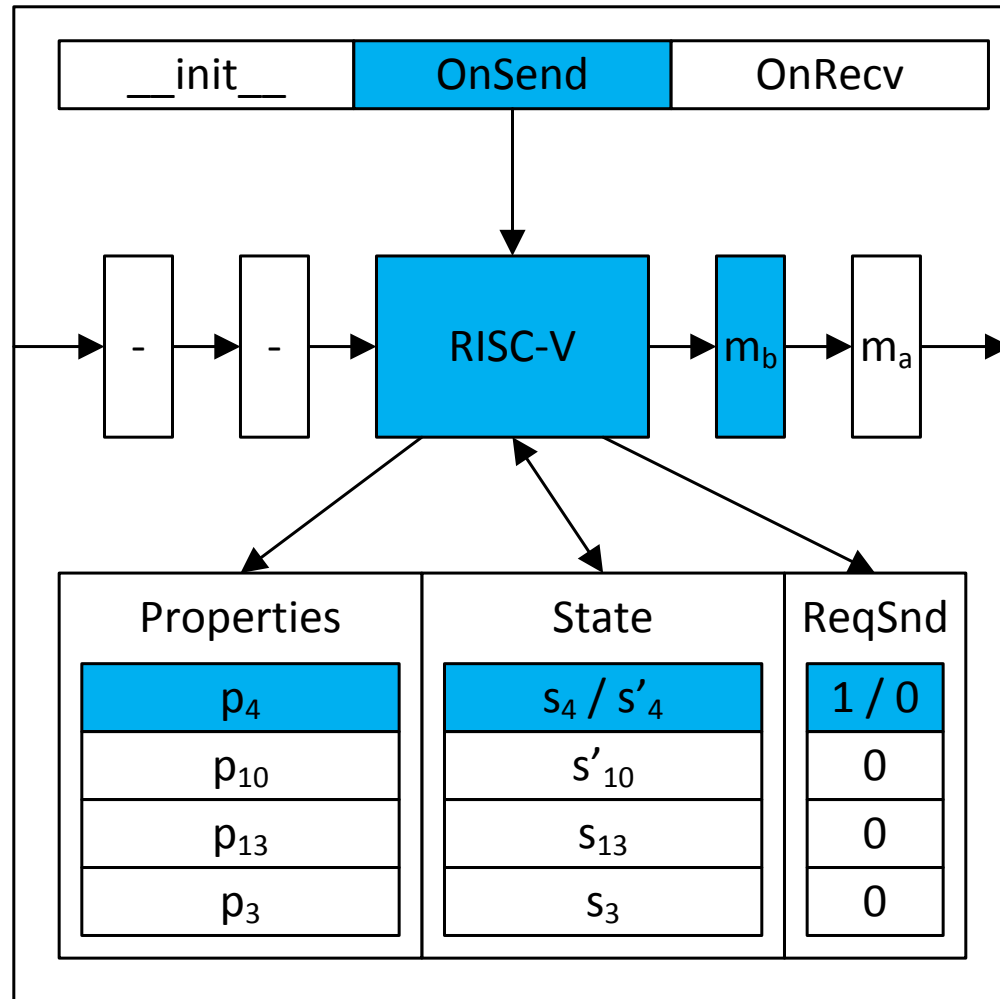
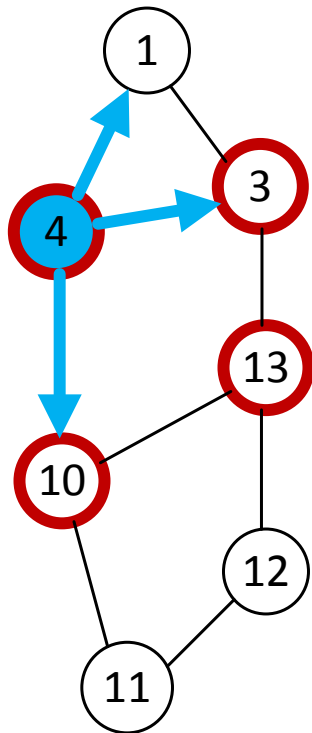


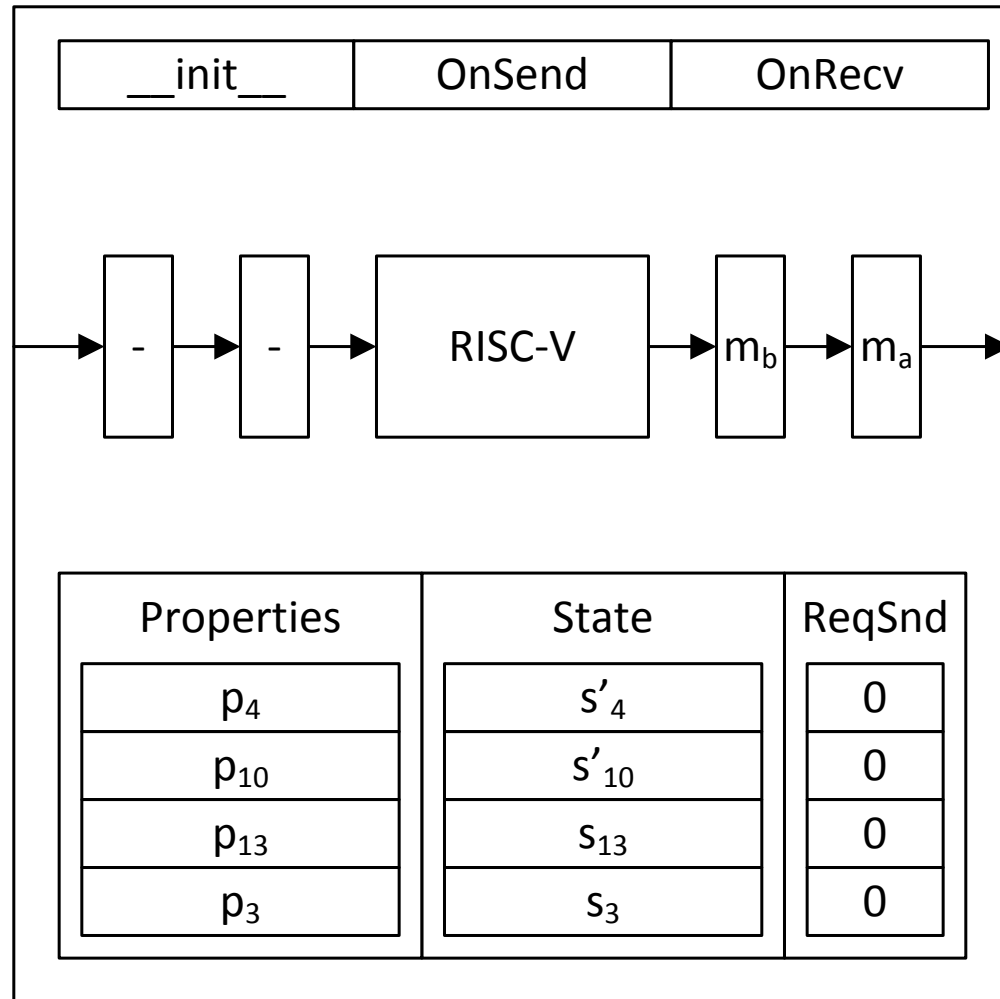
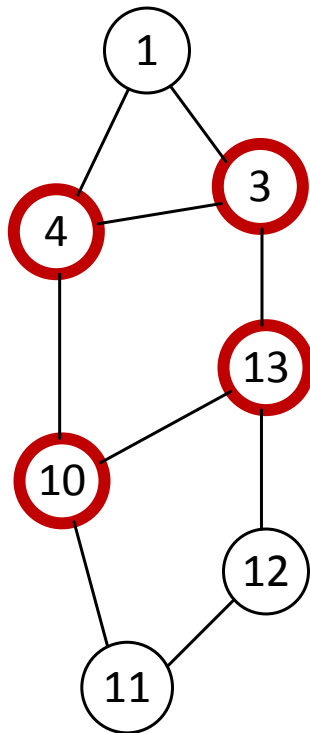




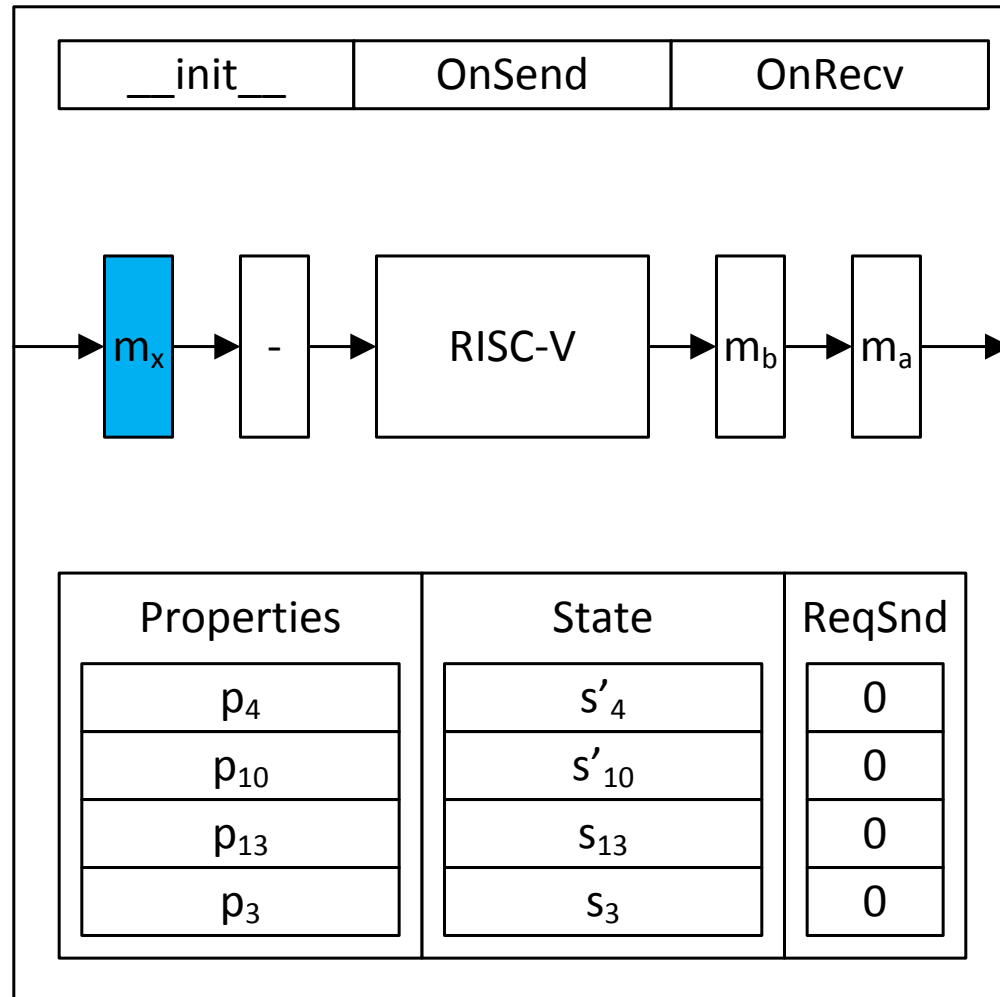
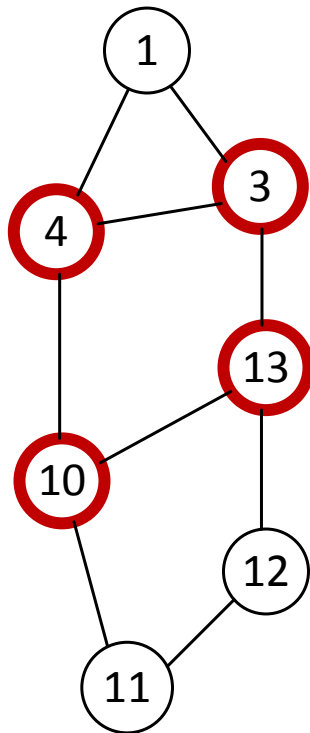


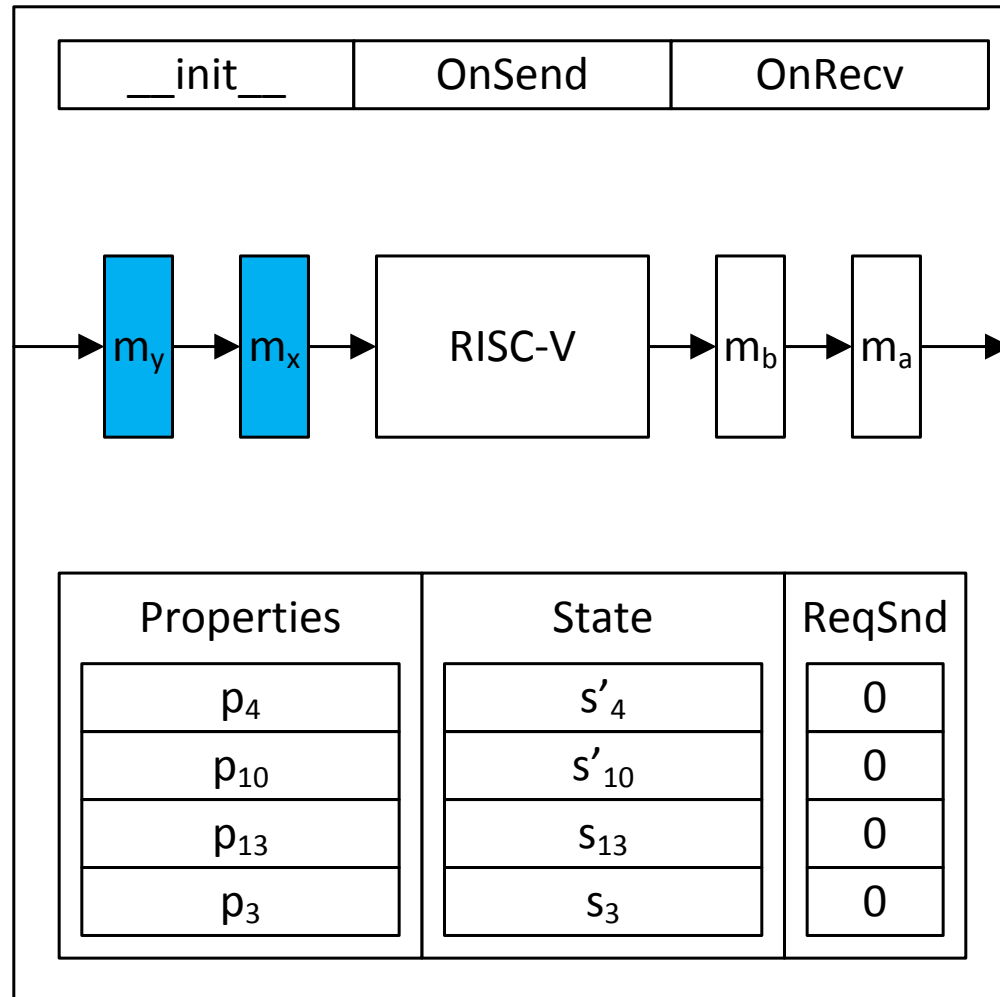
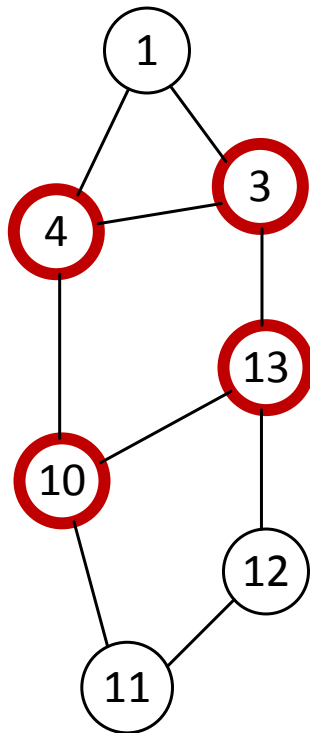


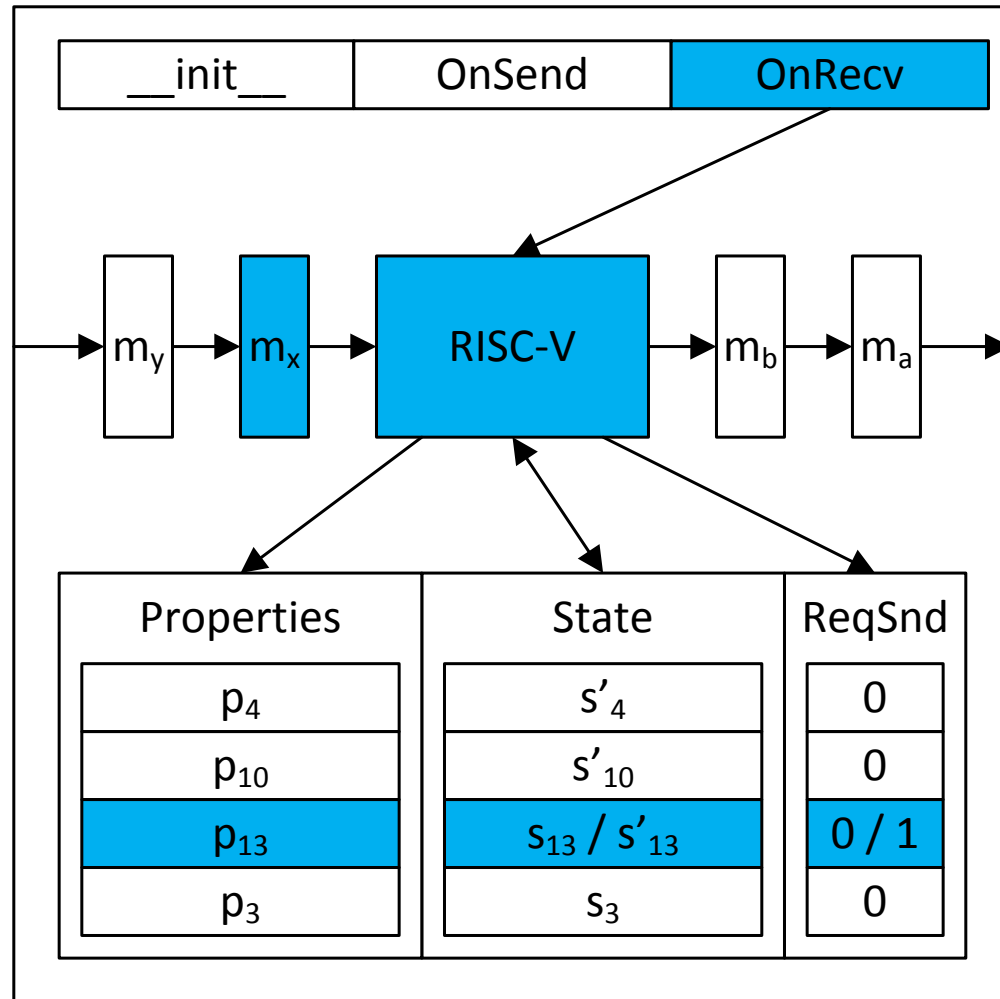
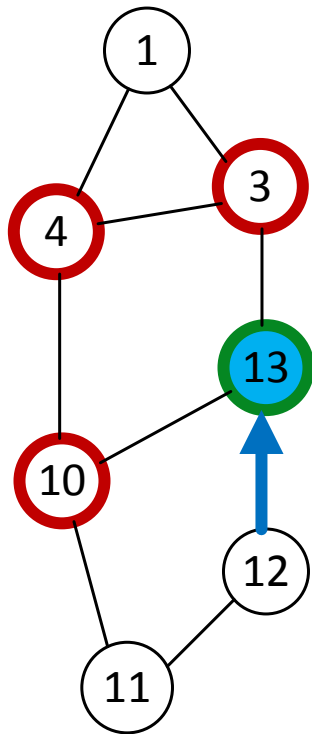


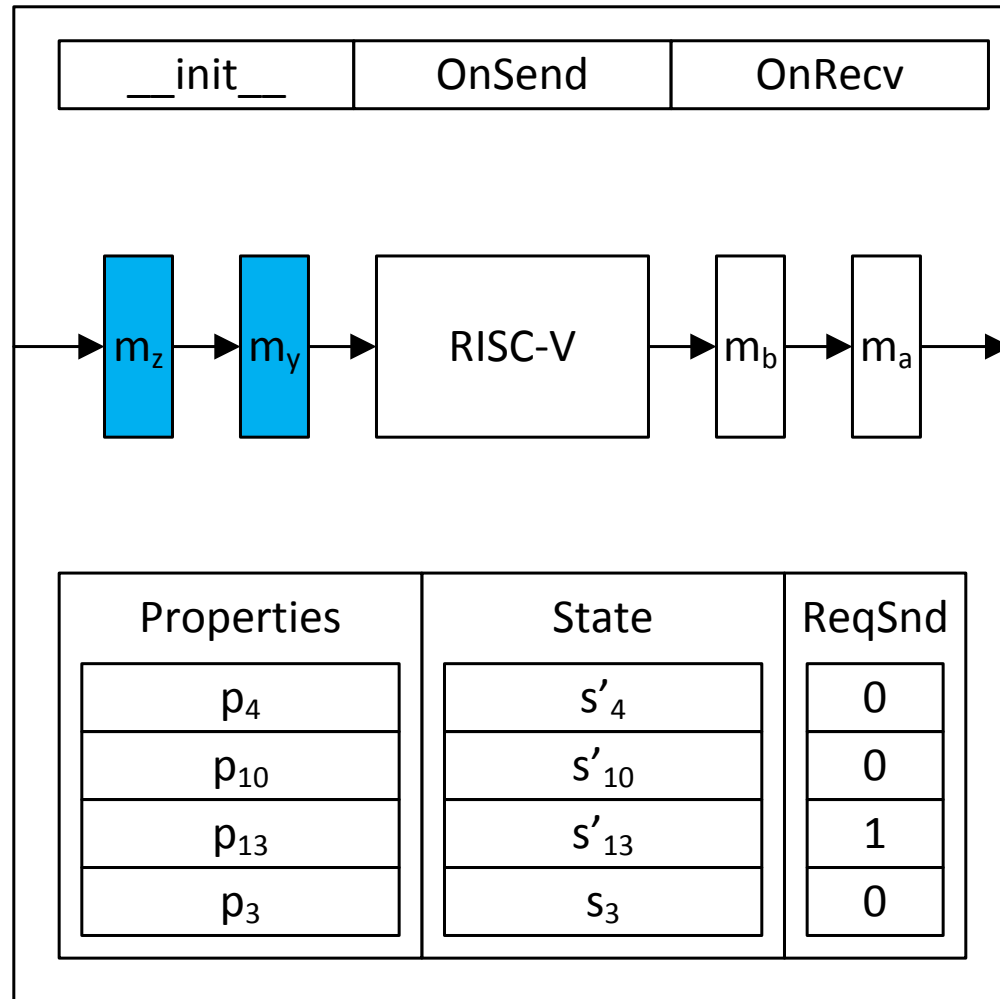
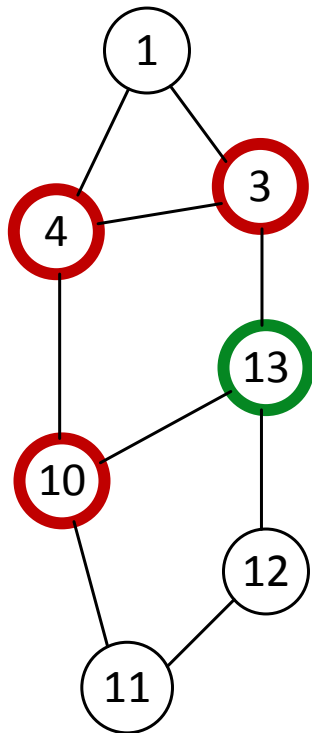


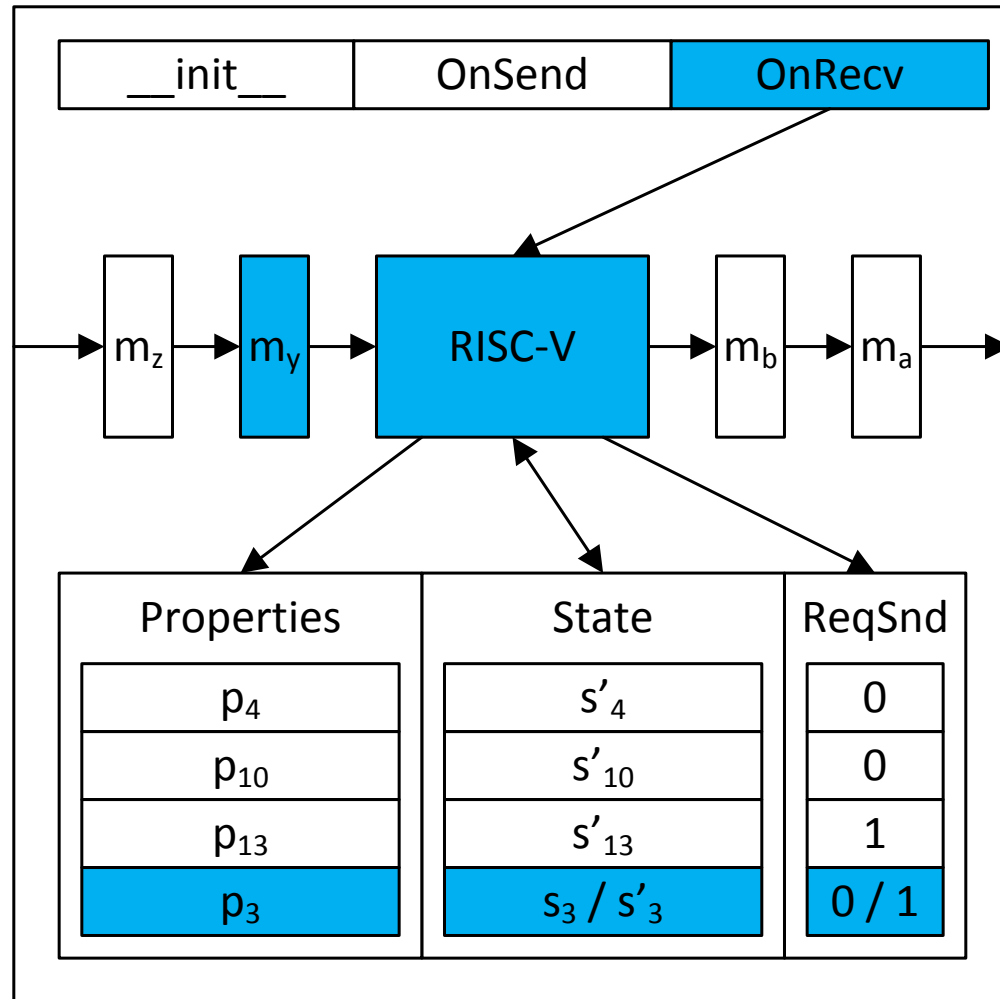
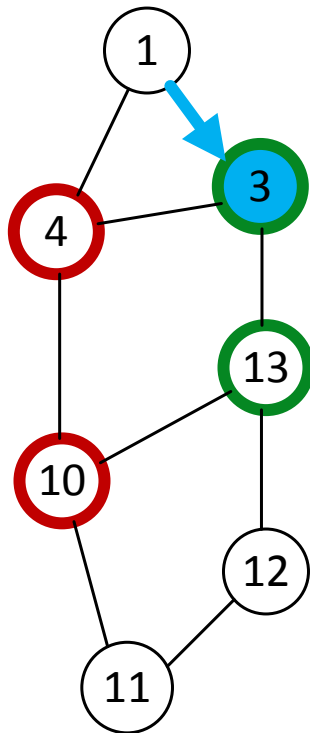


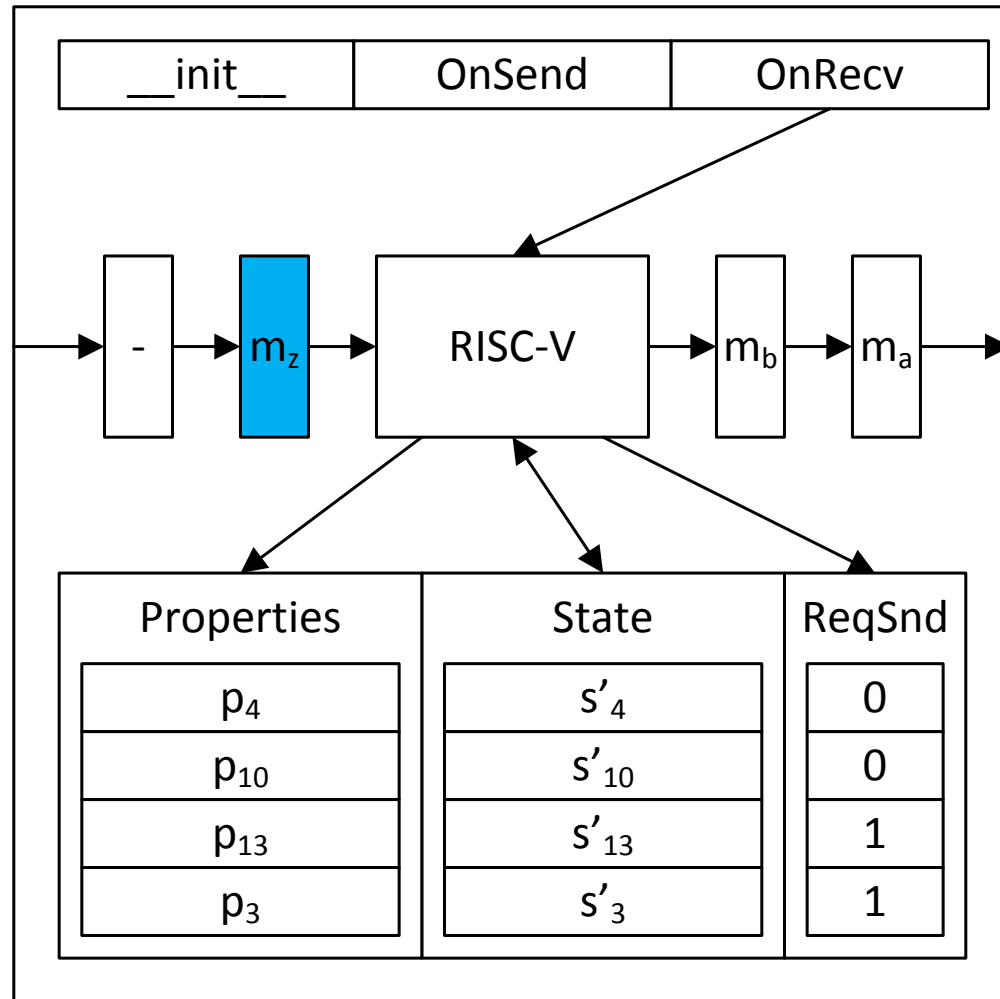
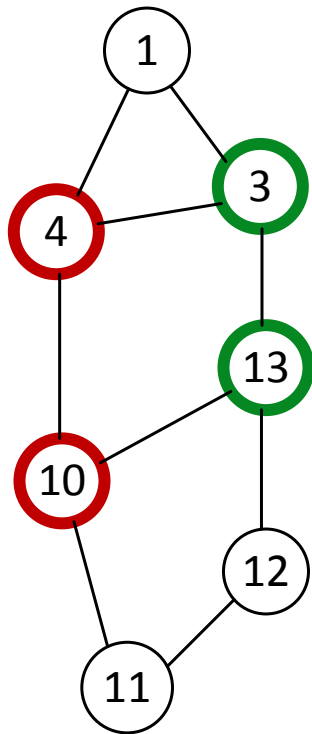


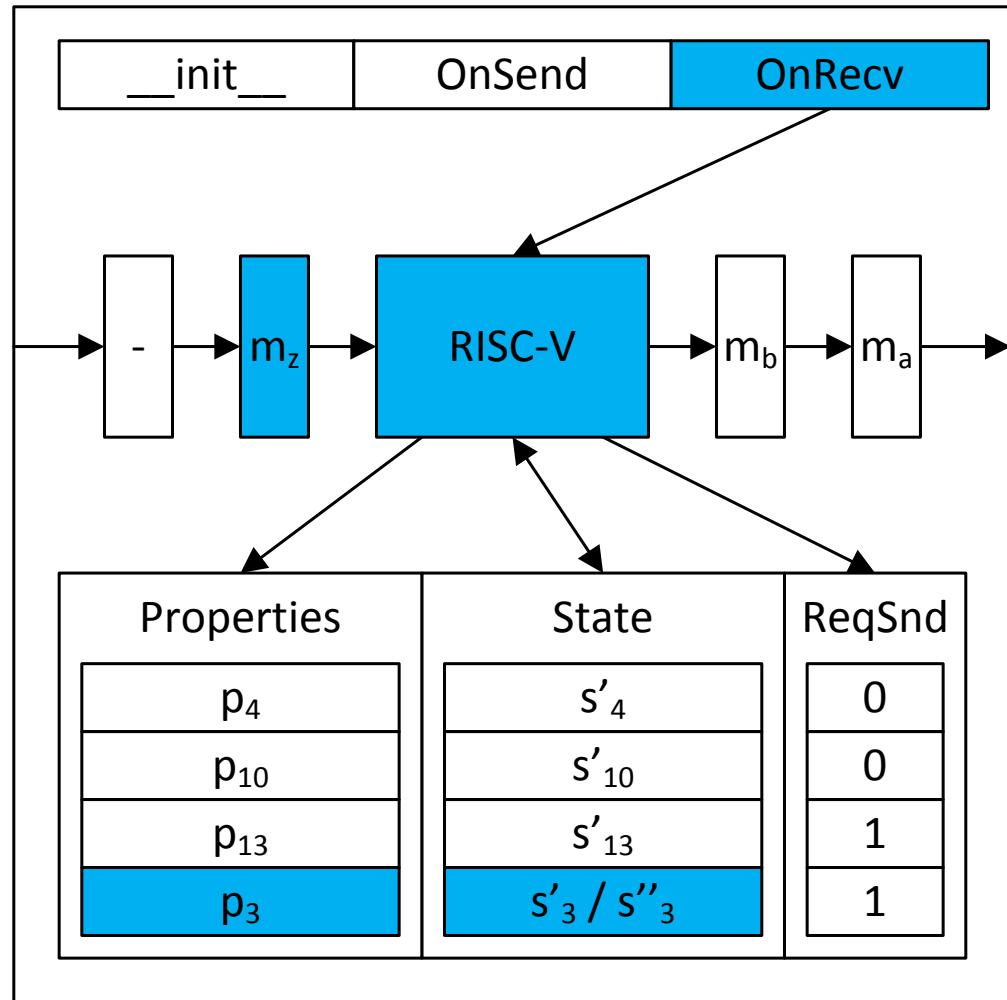
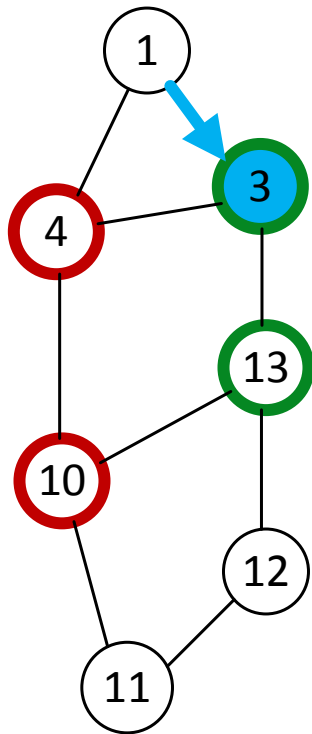


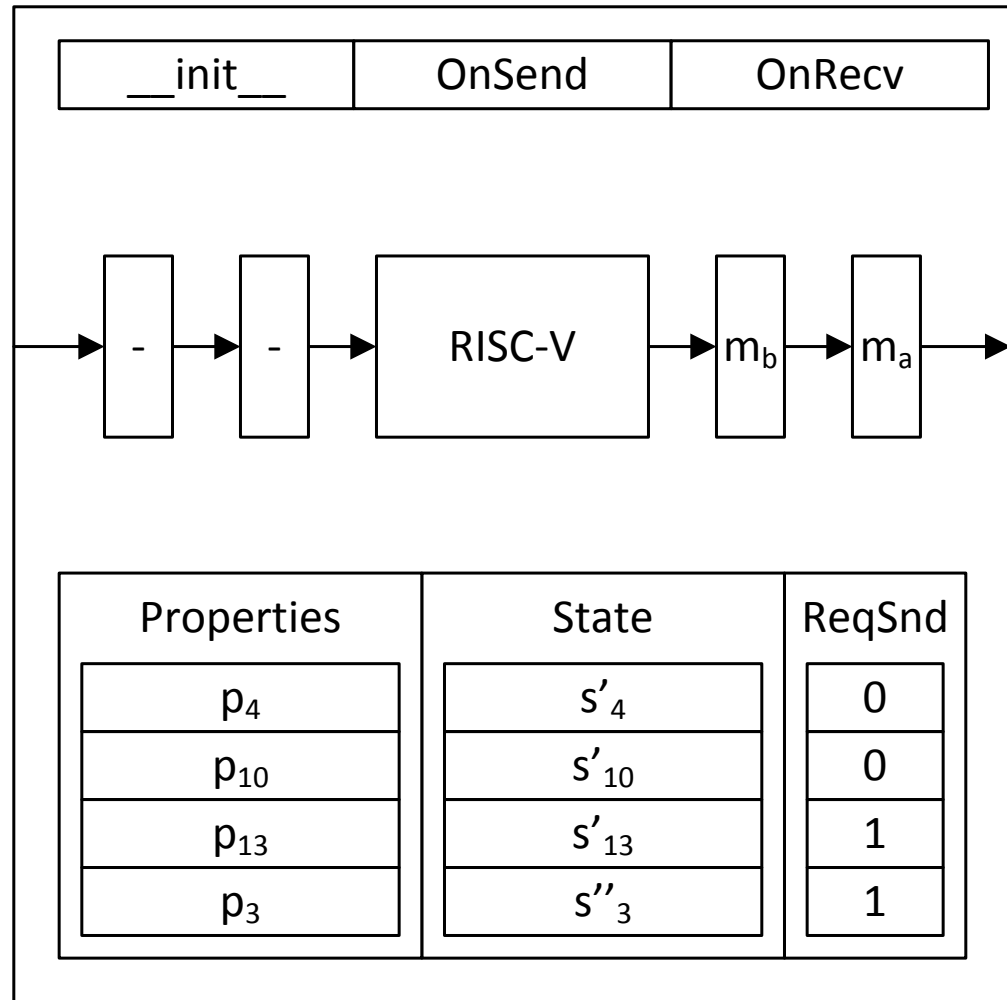
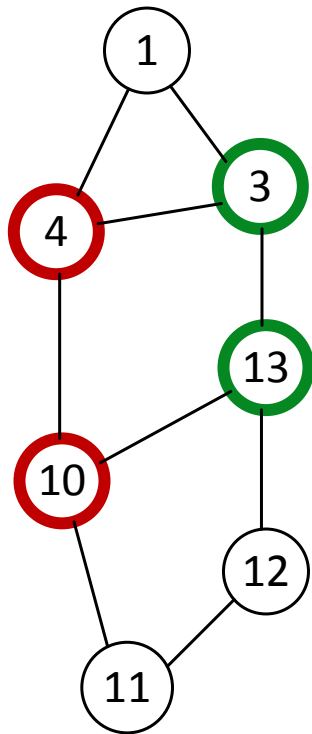




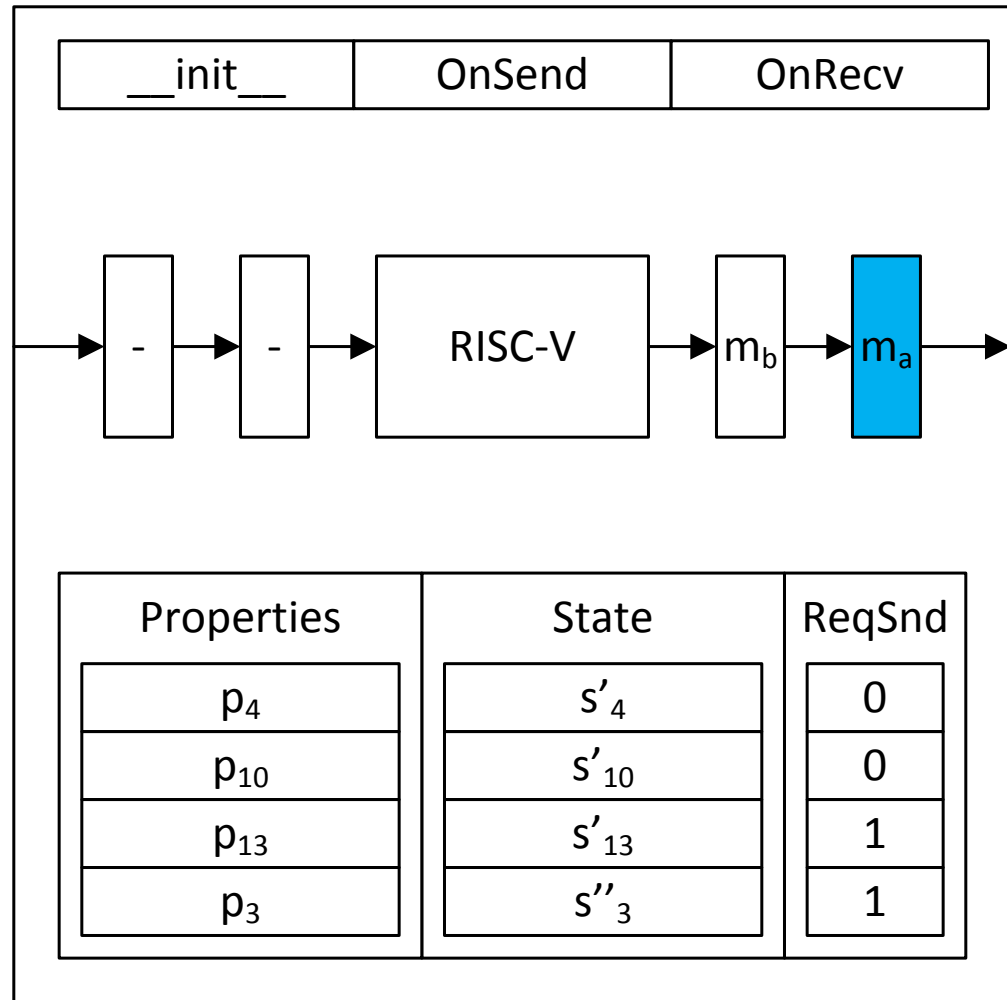
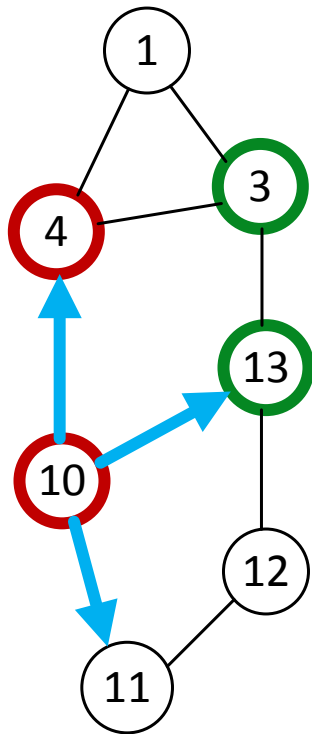


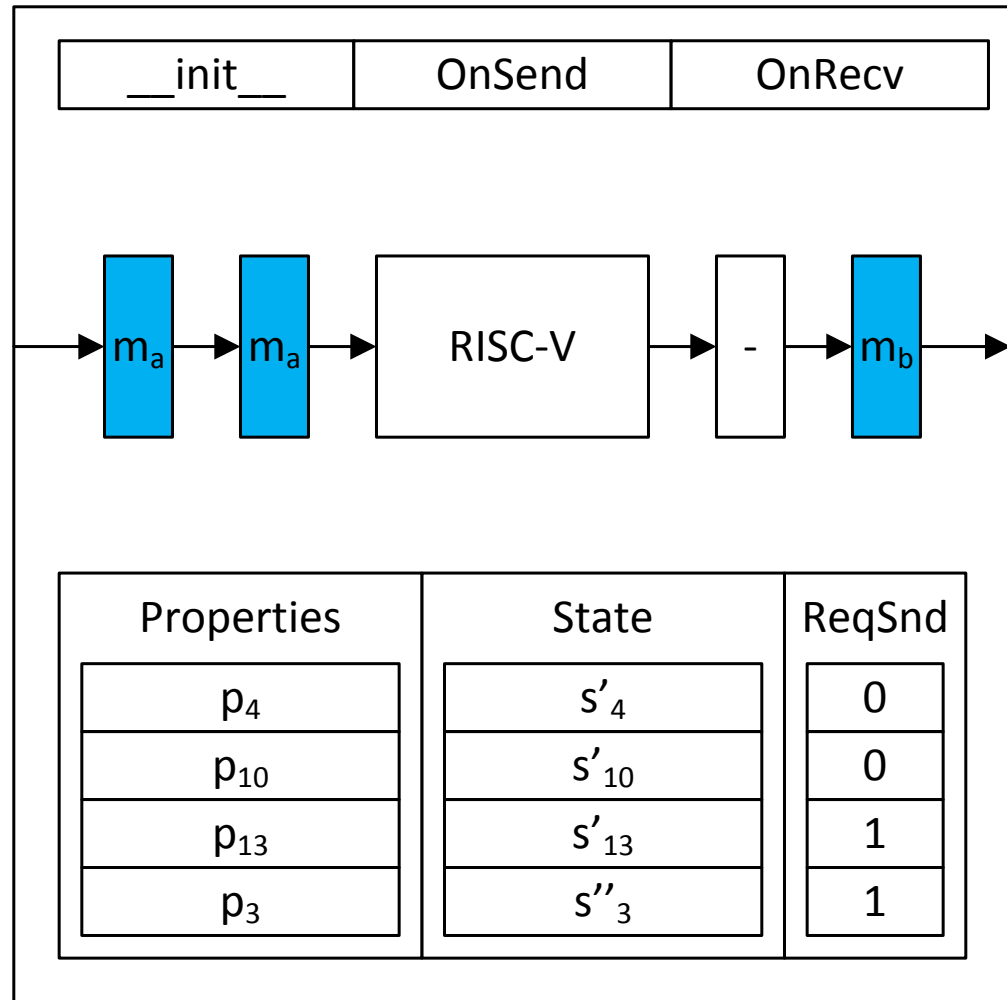
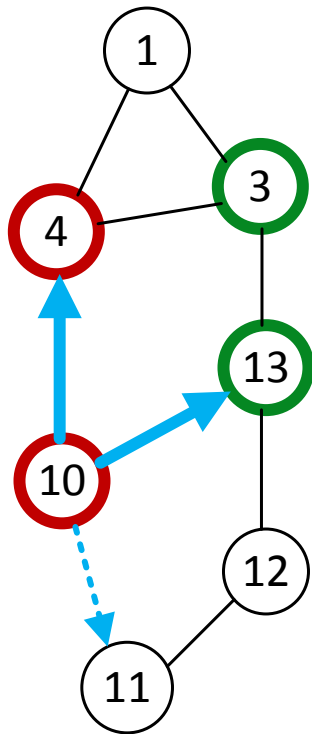


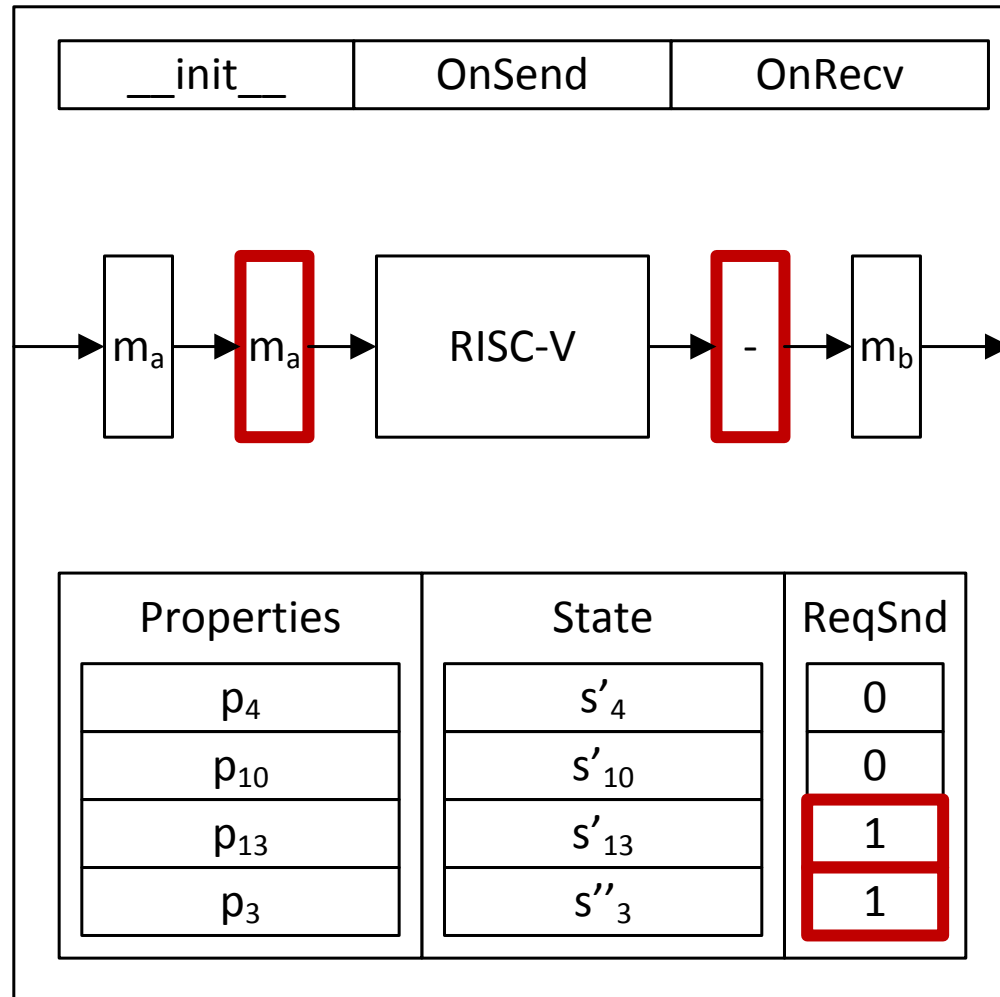
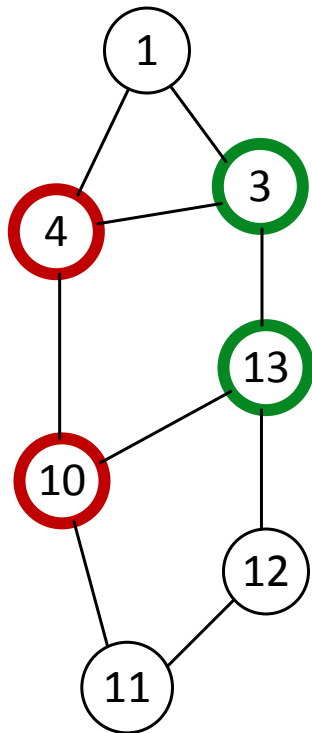


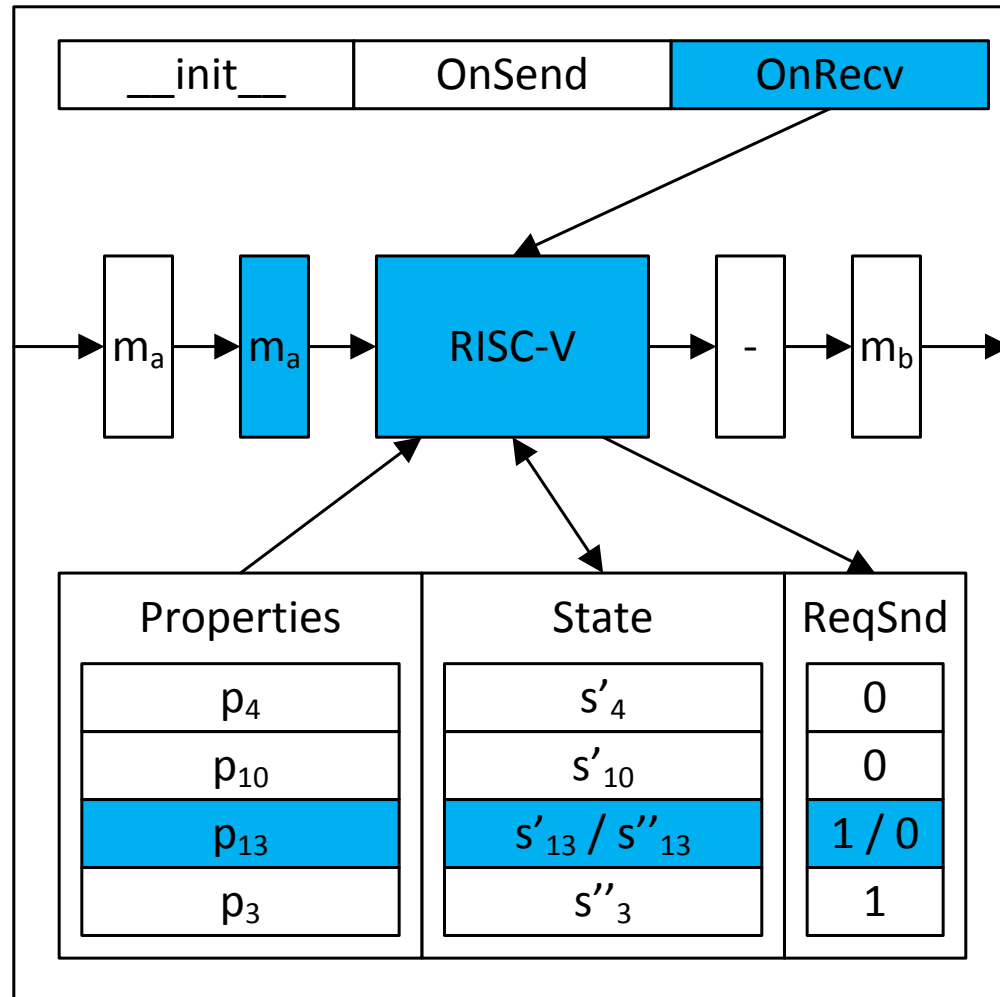
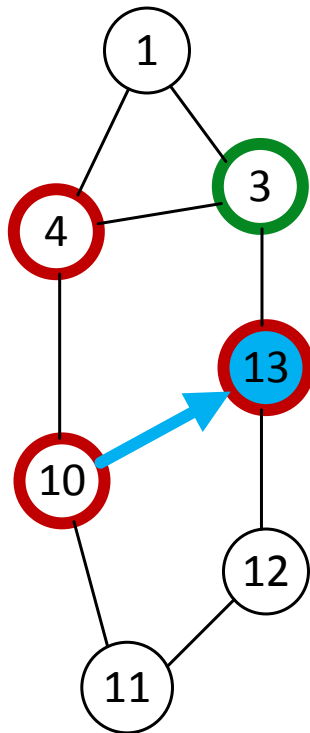


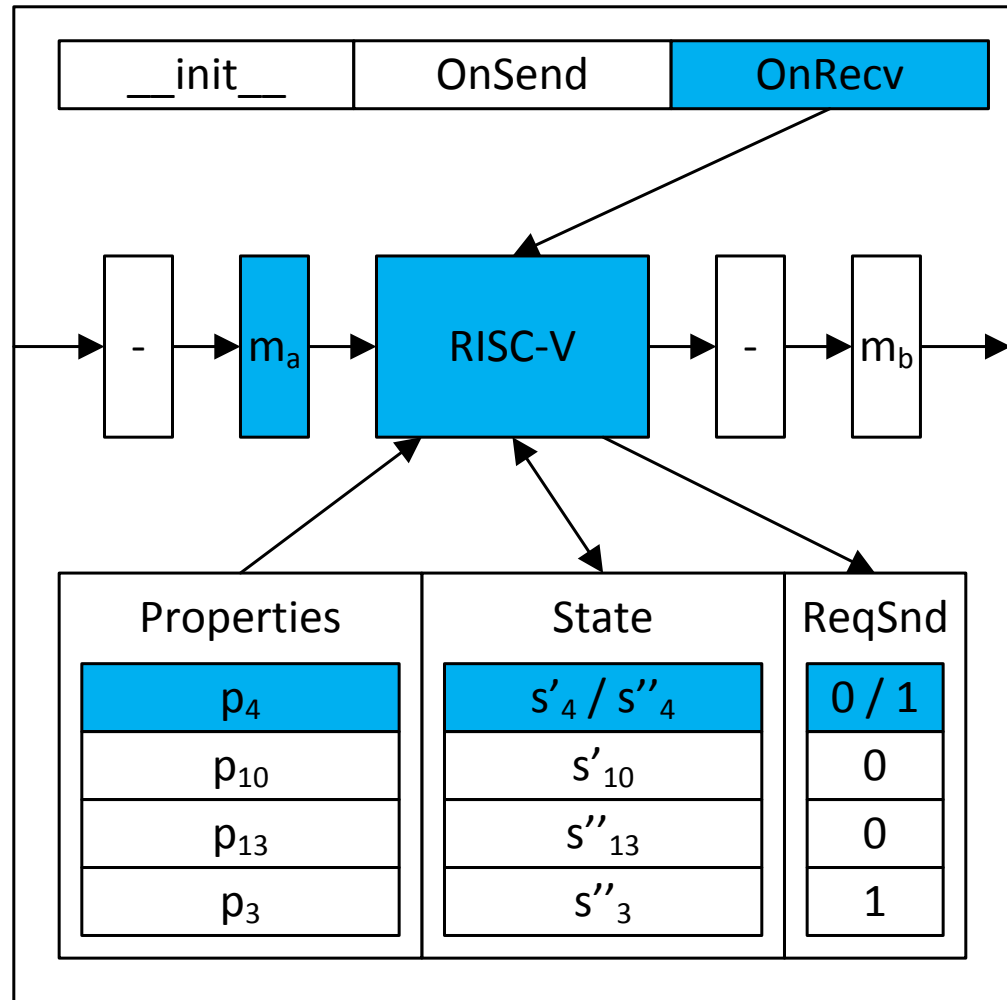
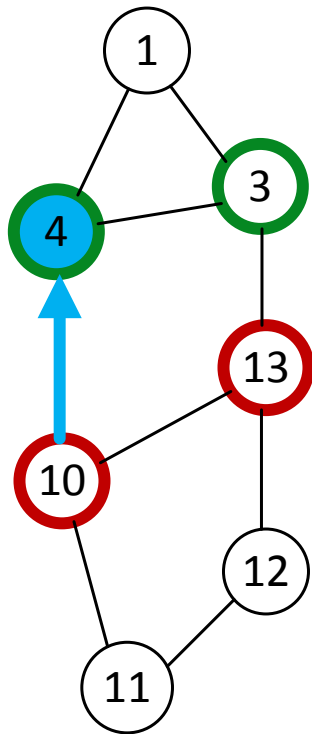


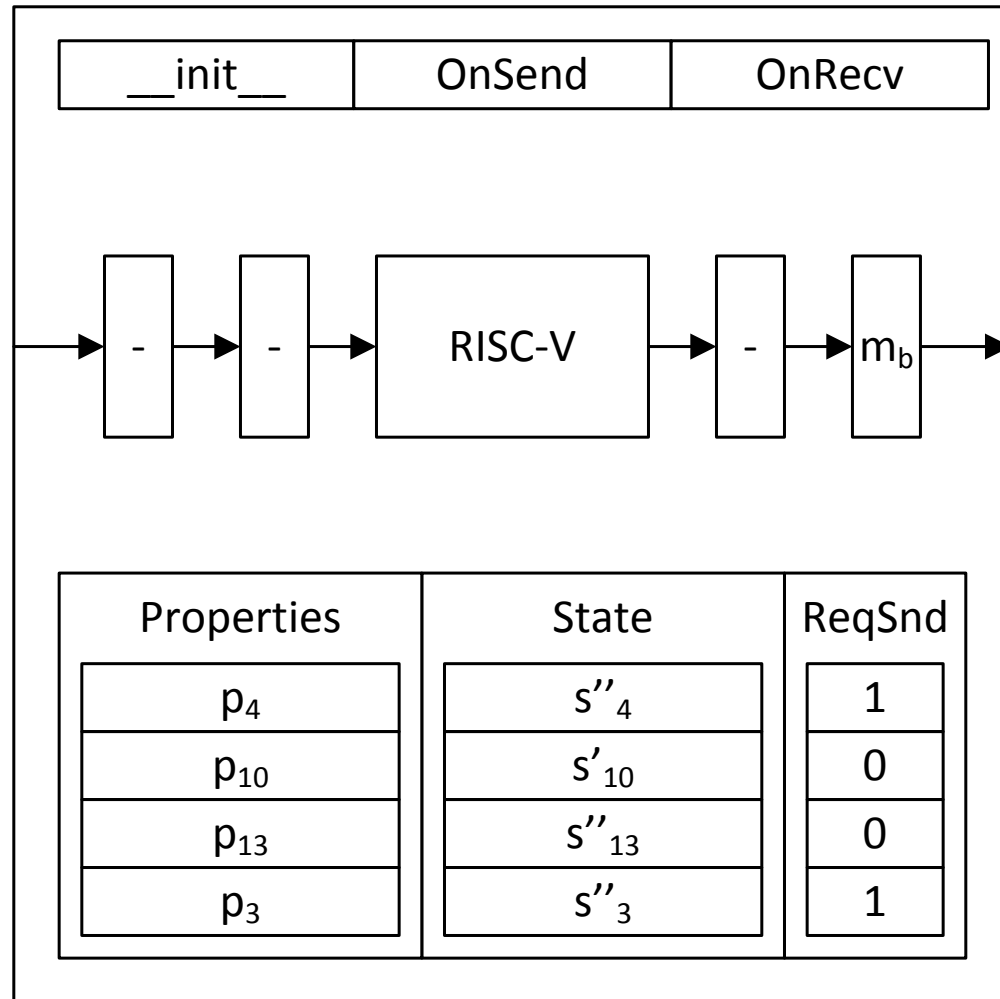
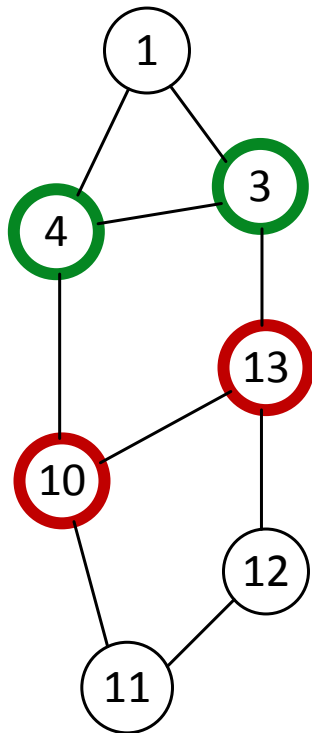


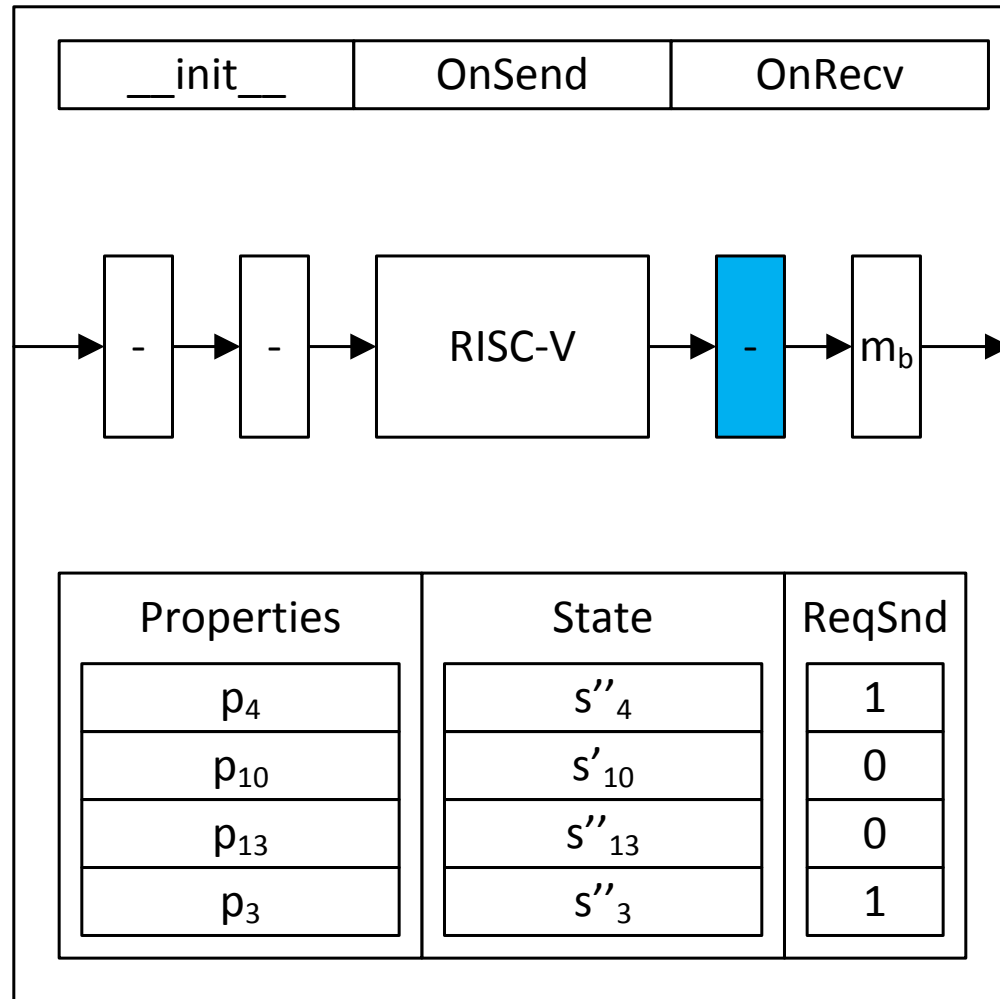
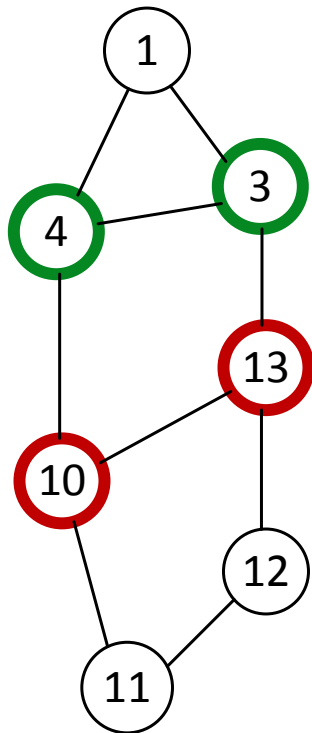


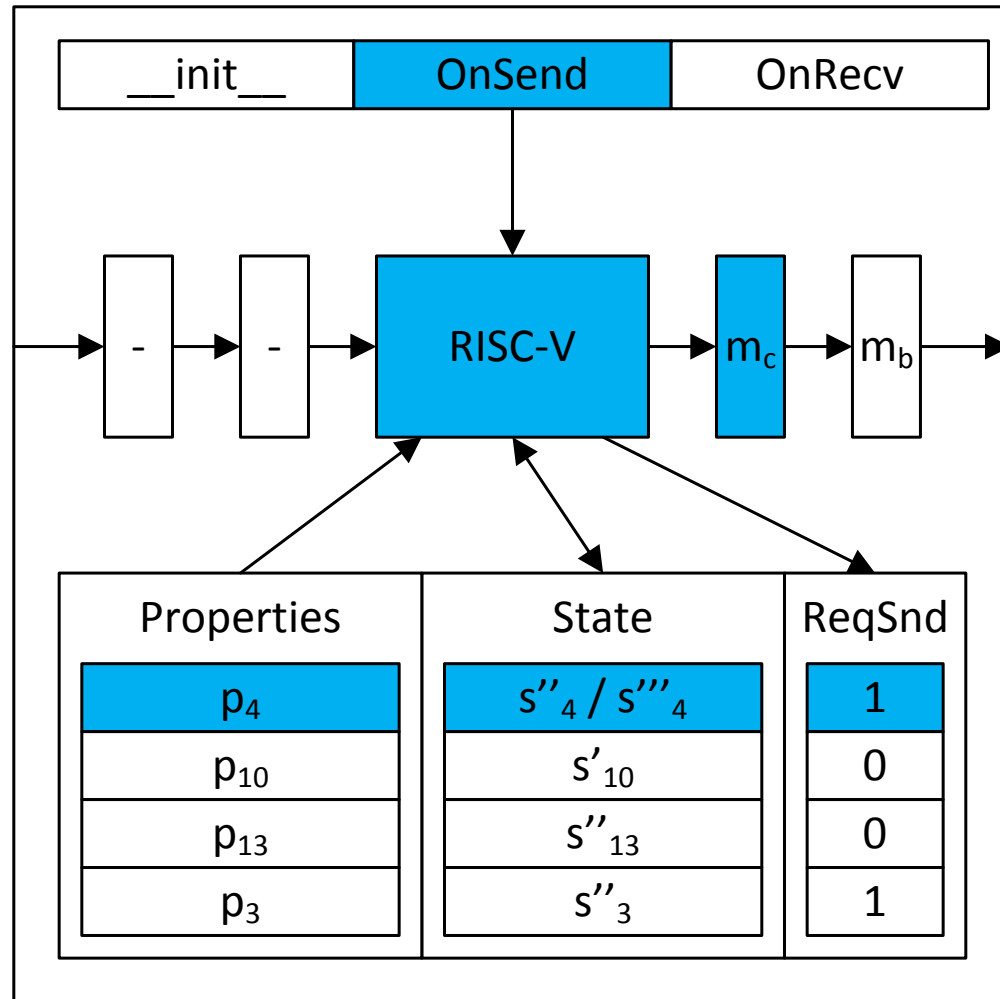
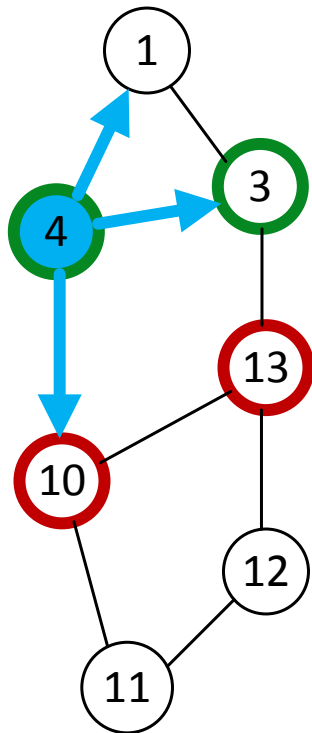




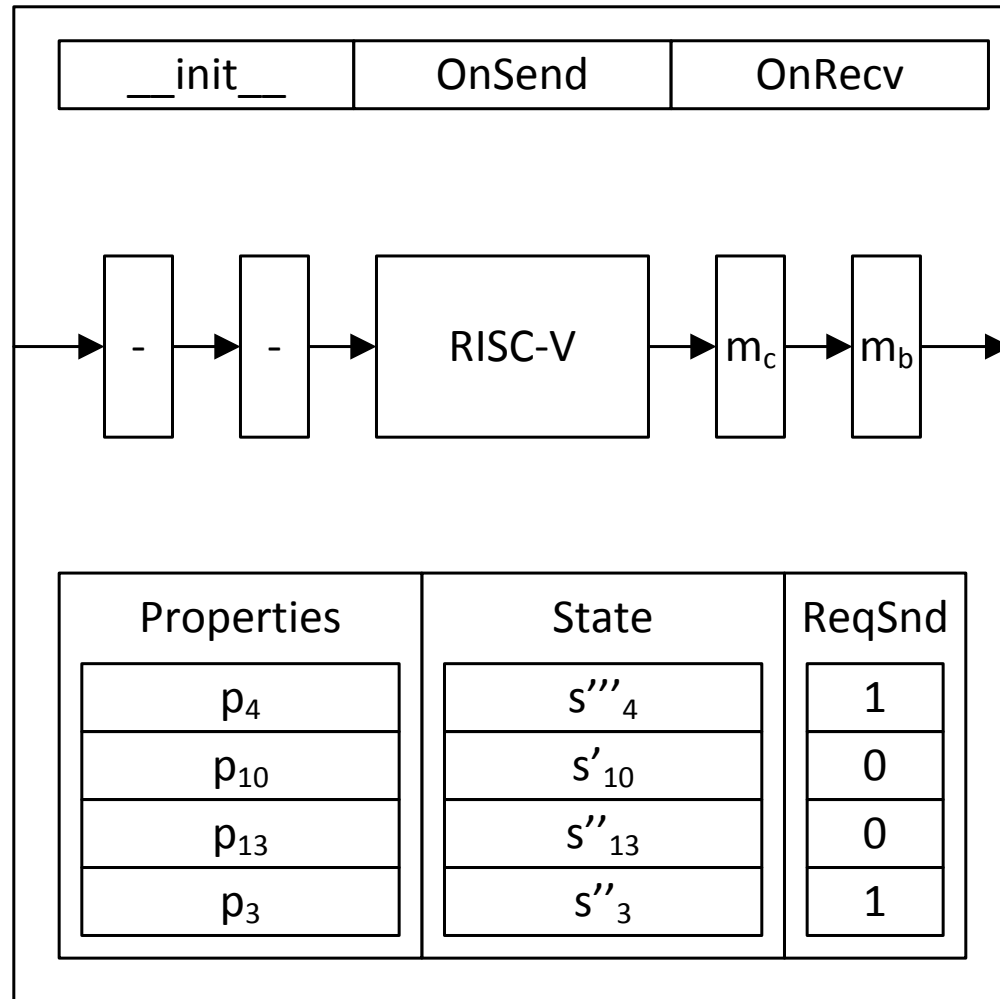
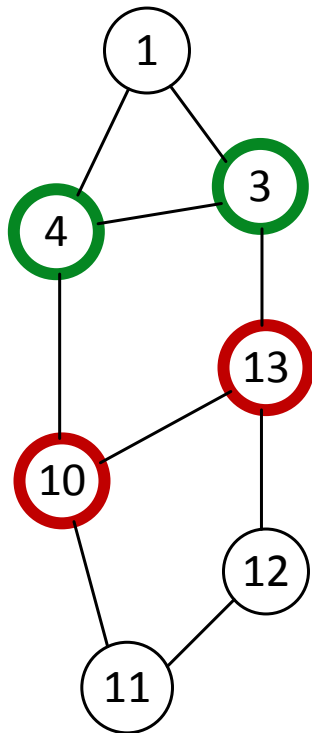


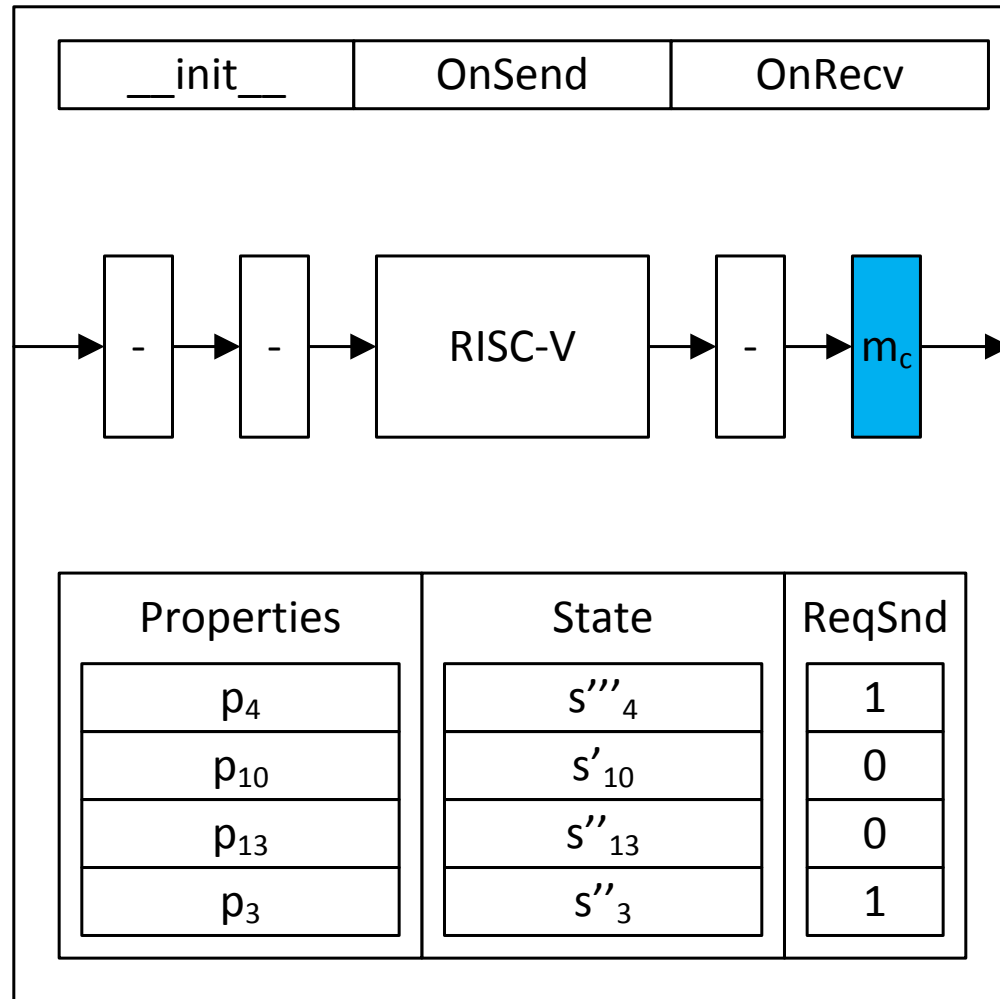
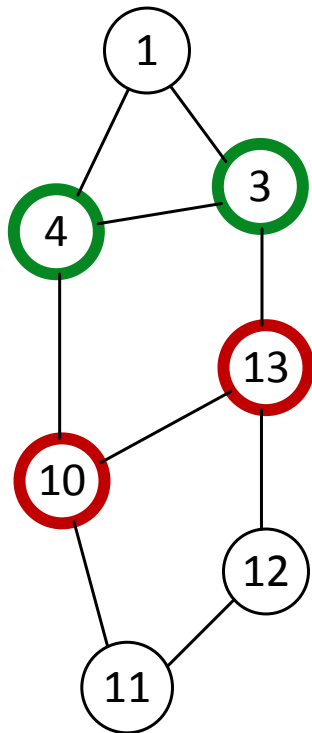


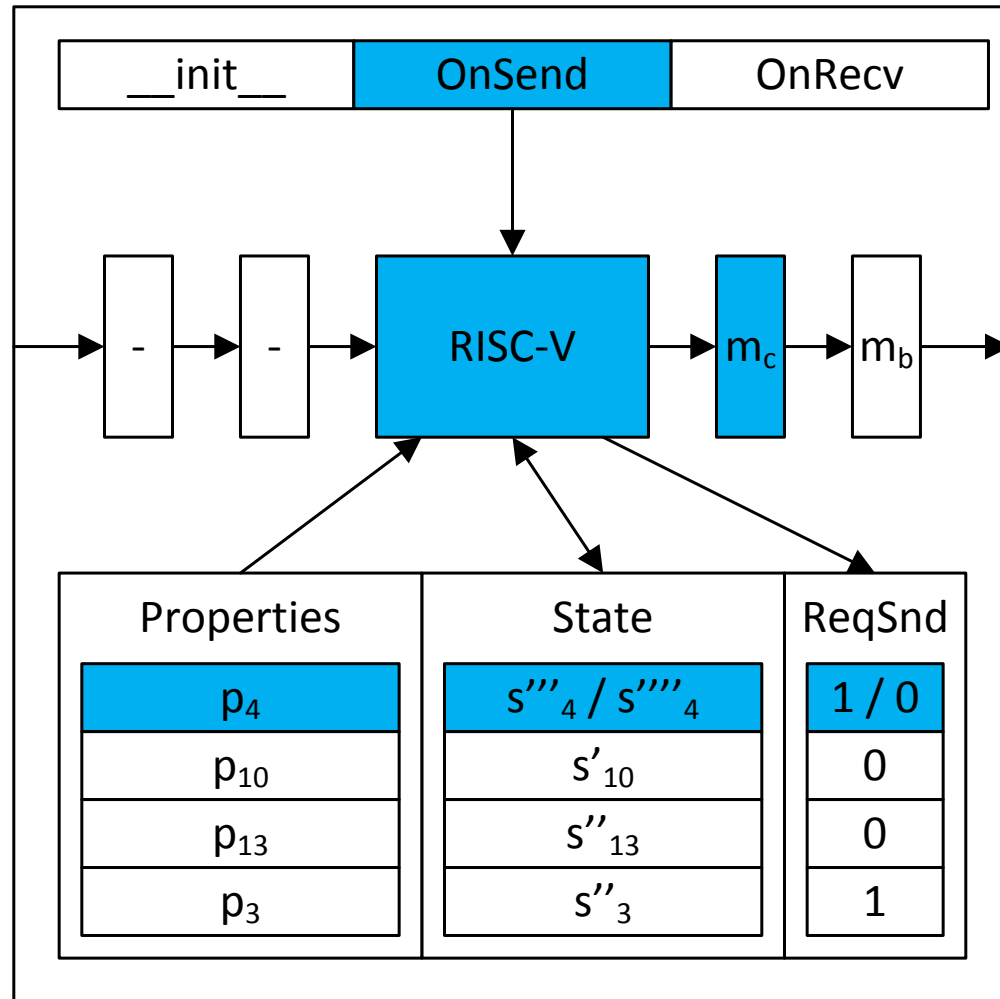
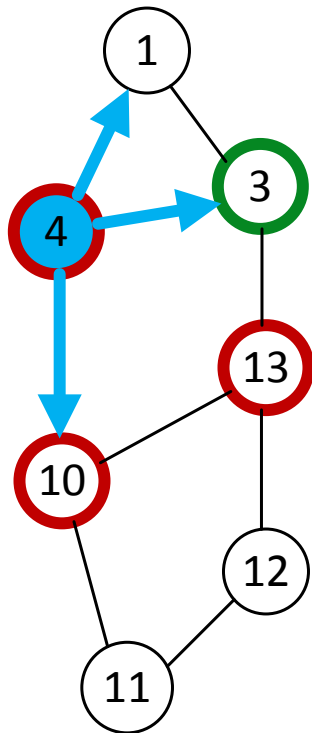


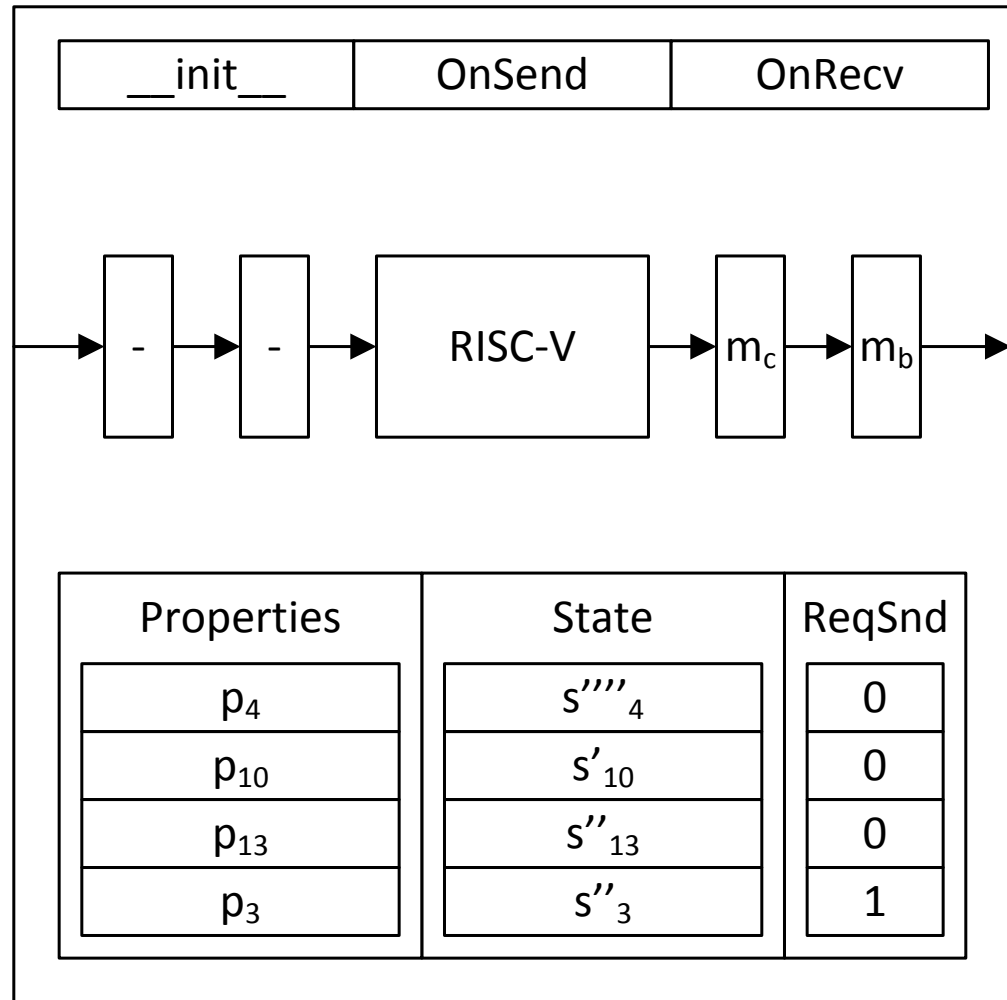
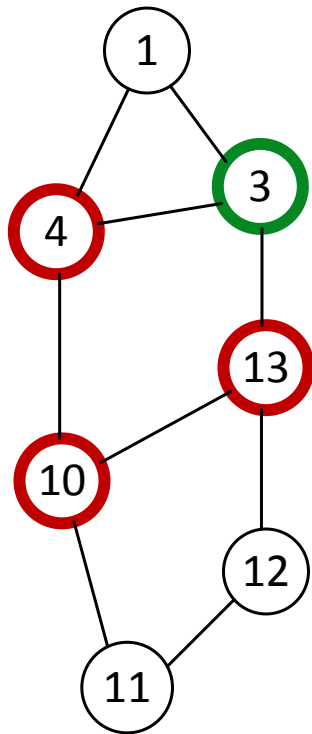


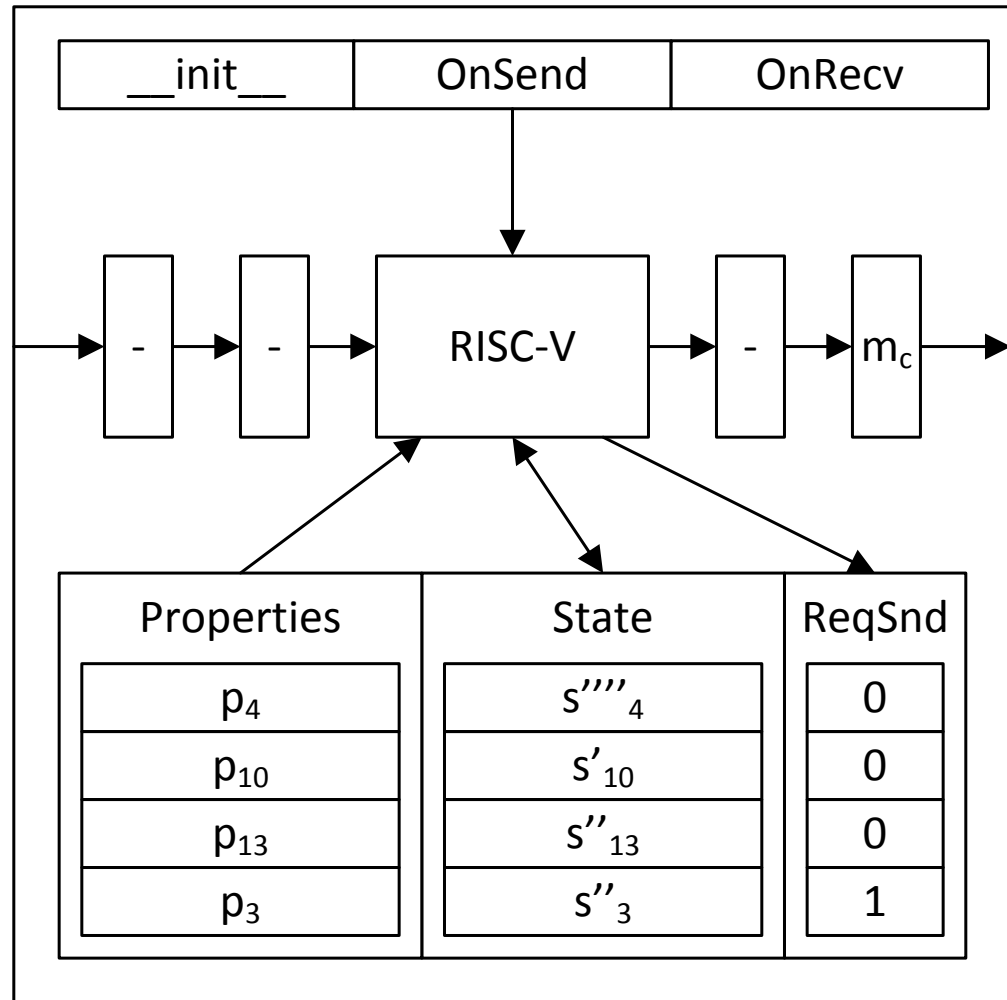
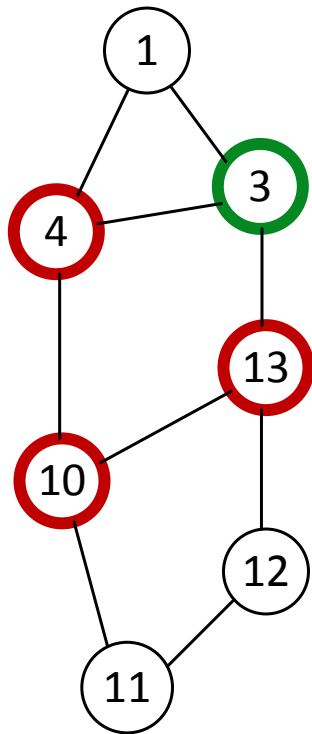


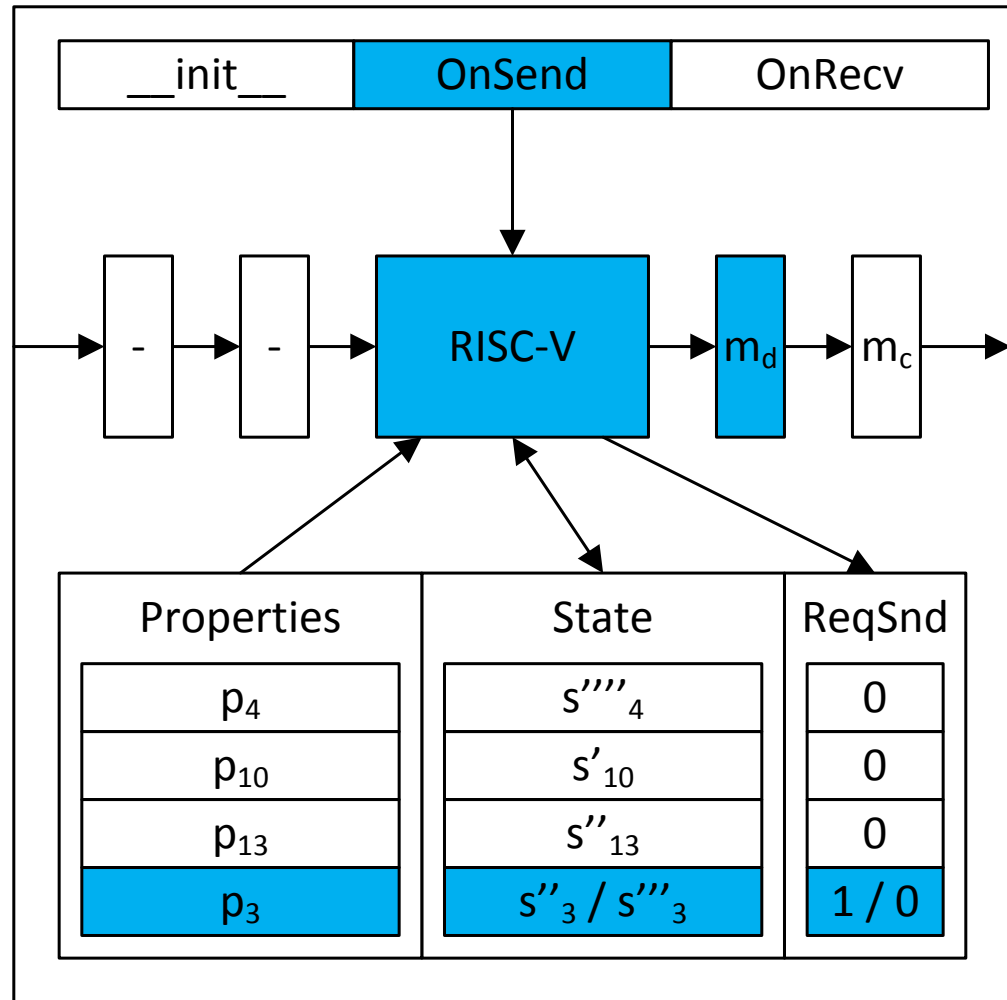
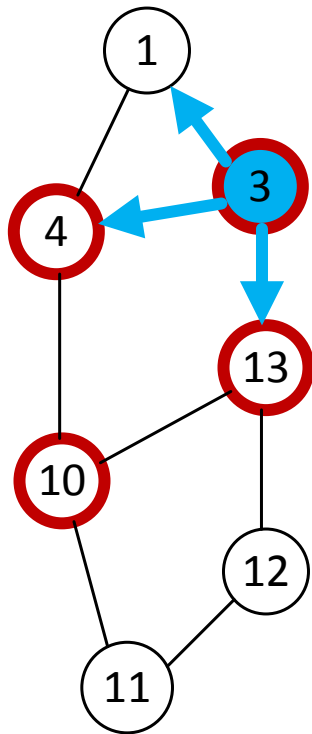


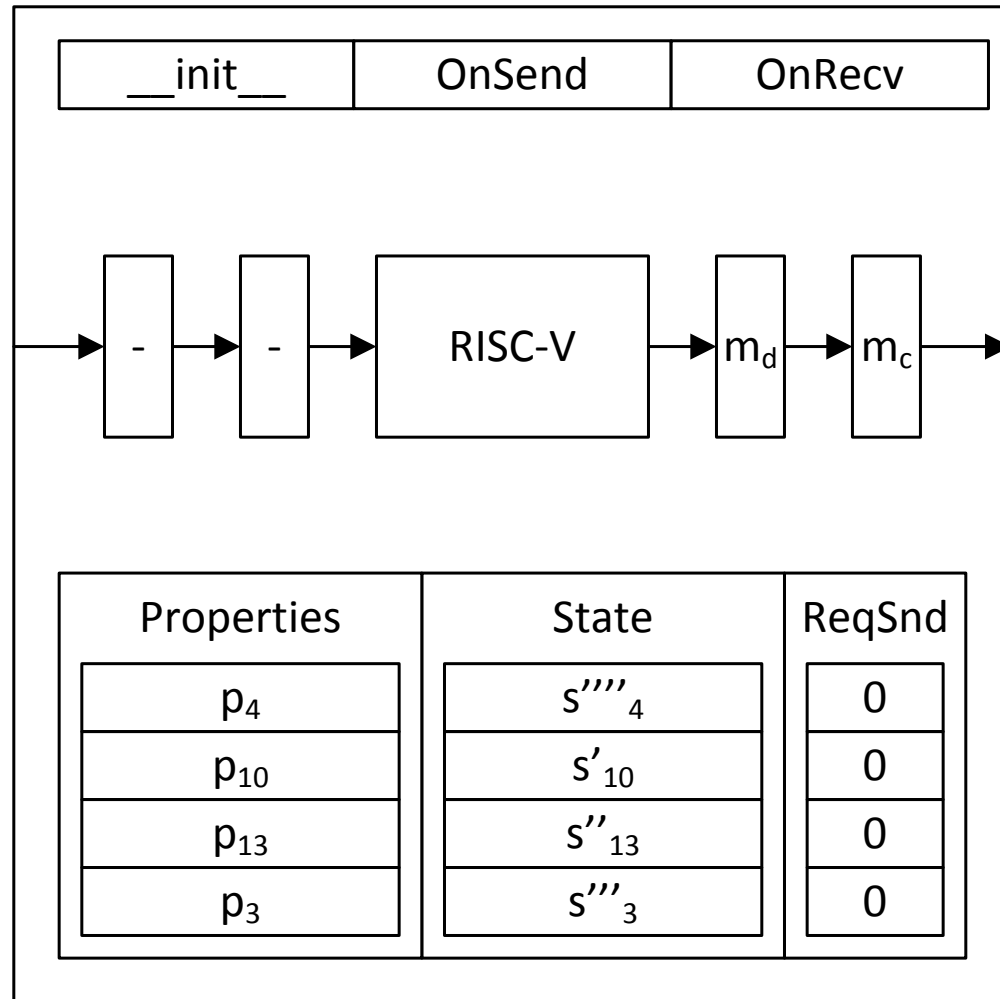
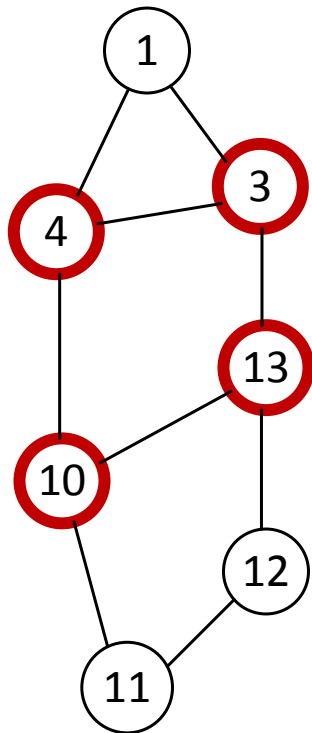






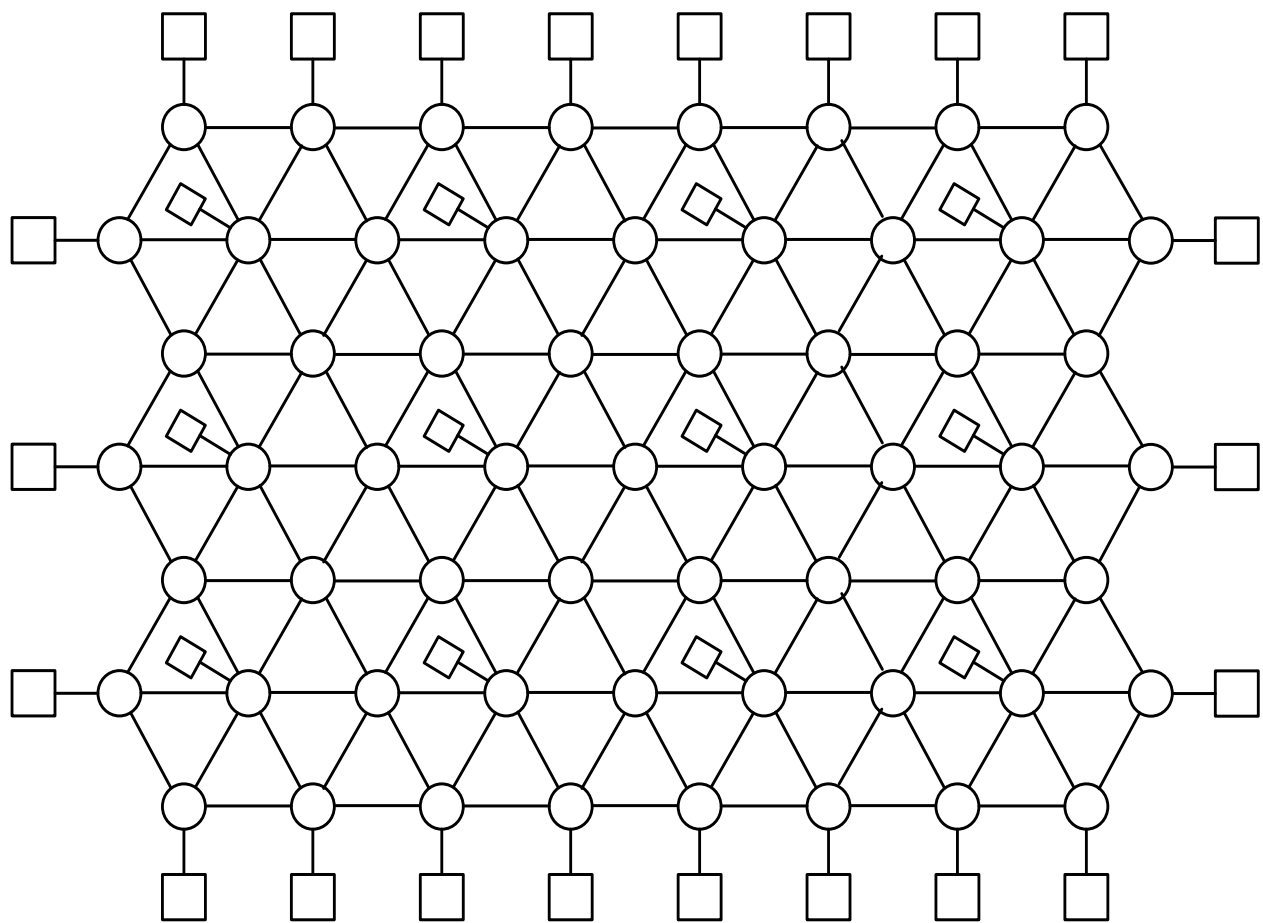


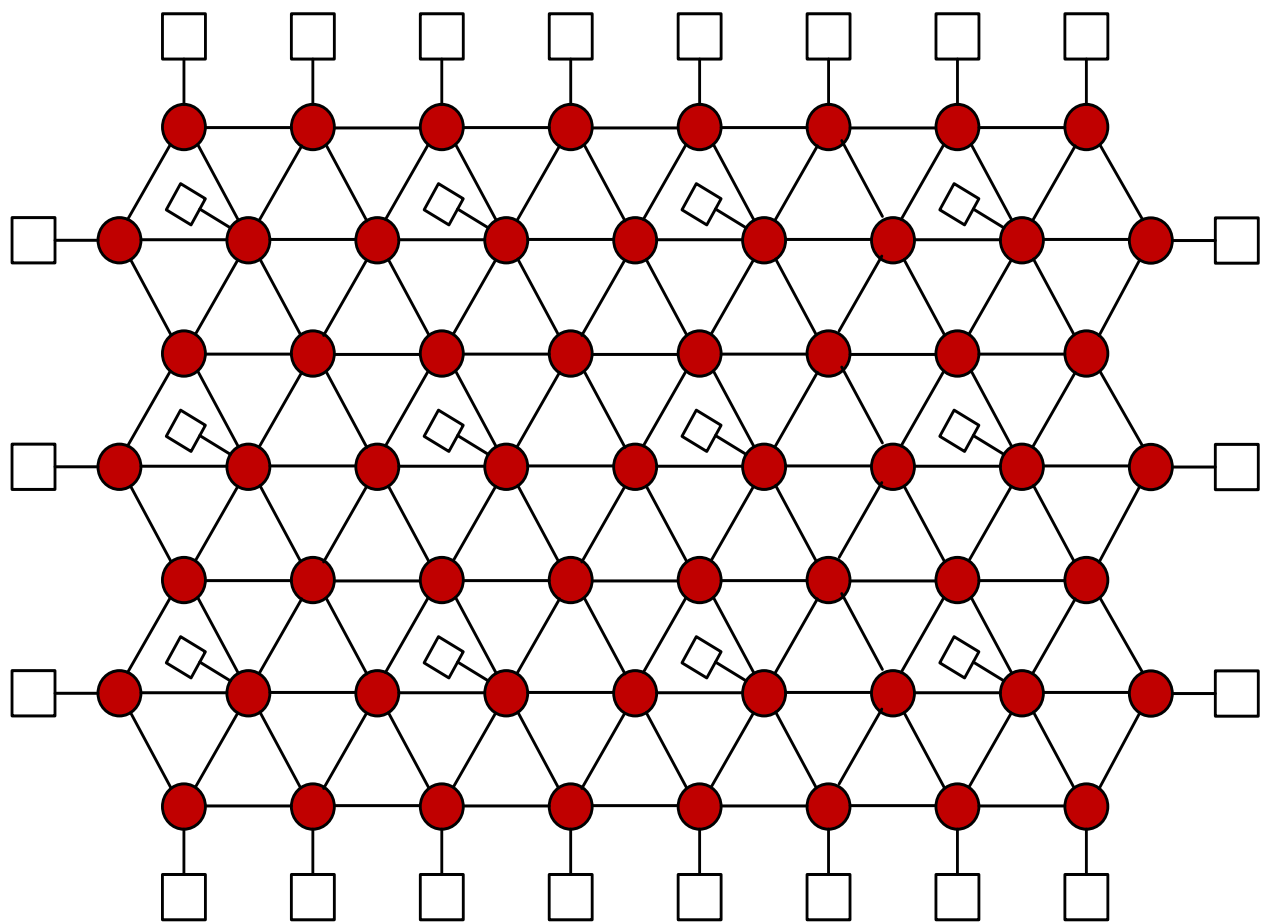


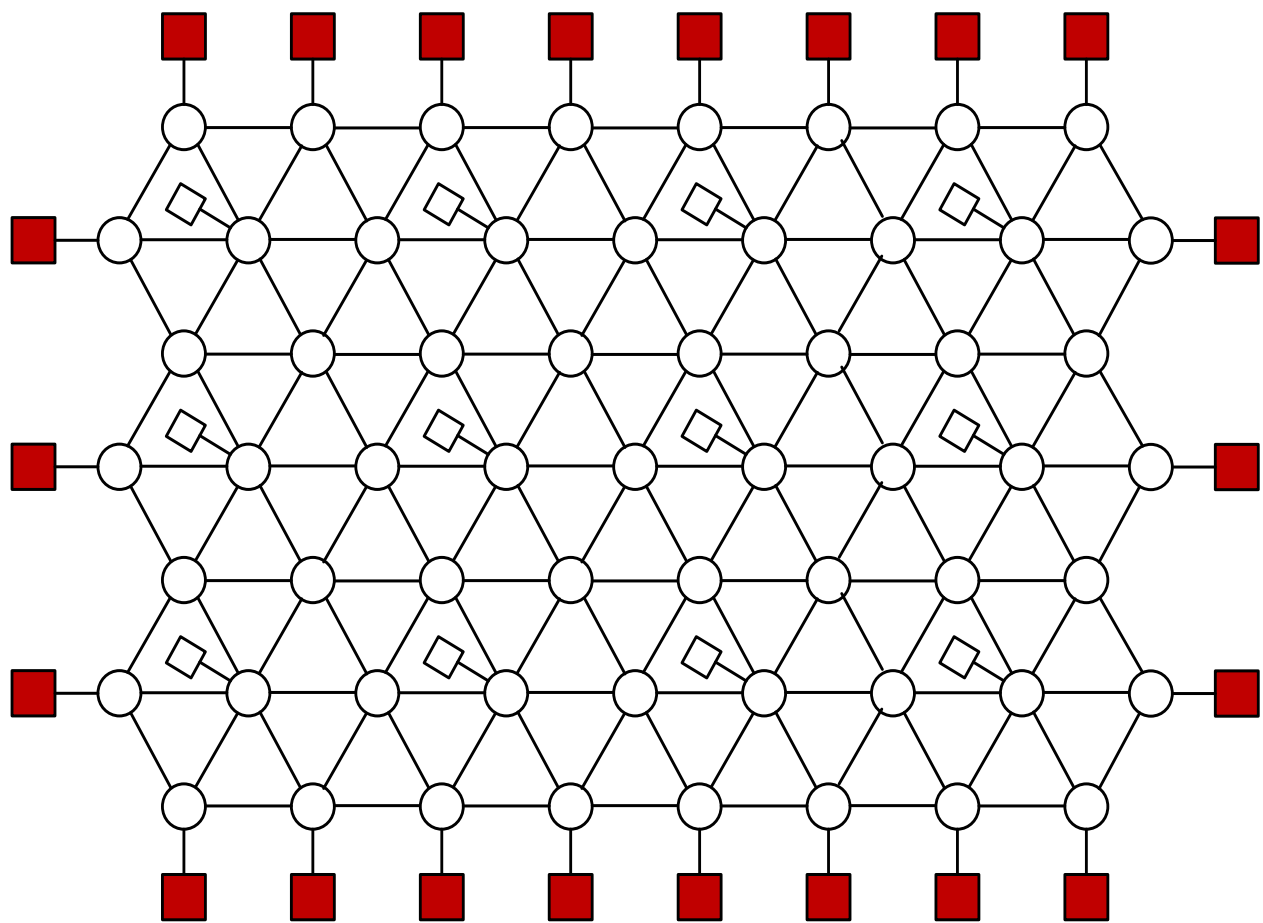


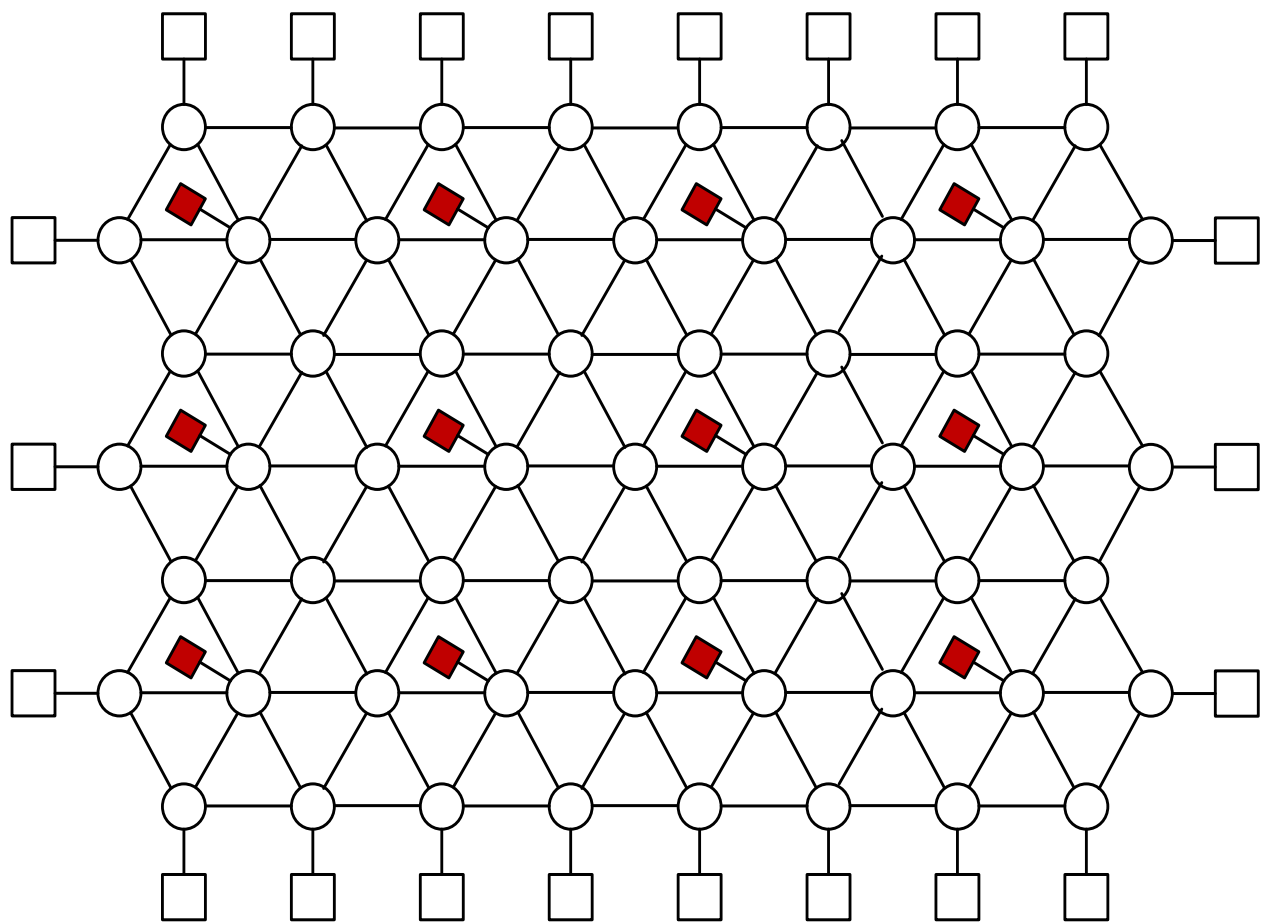
# Exfiltration (output)



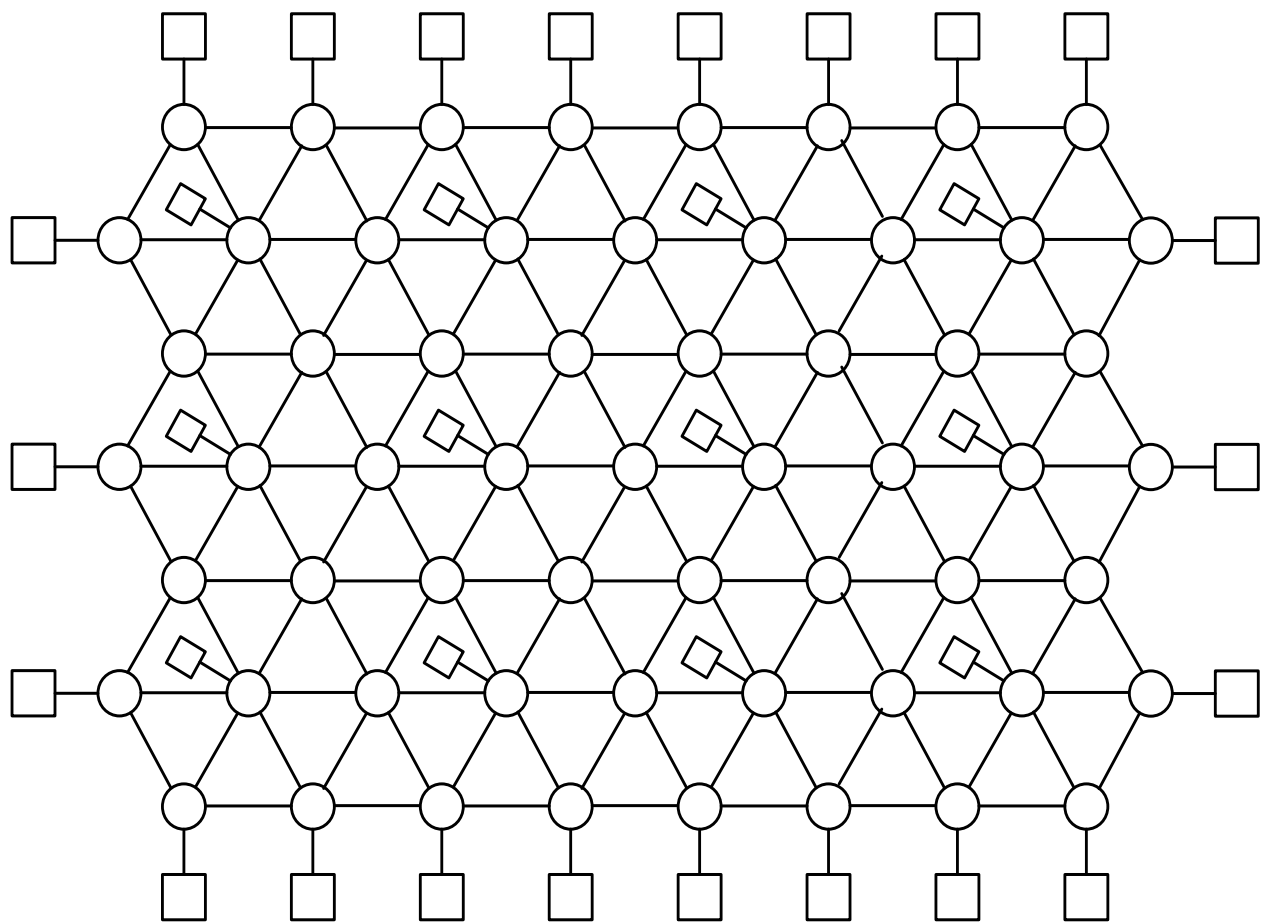






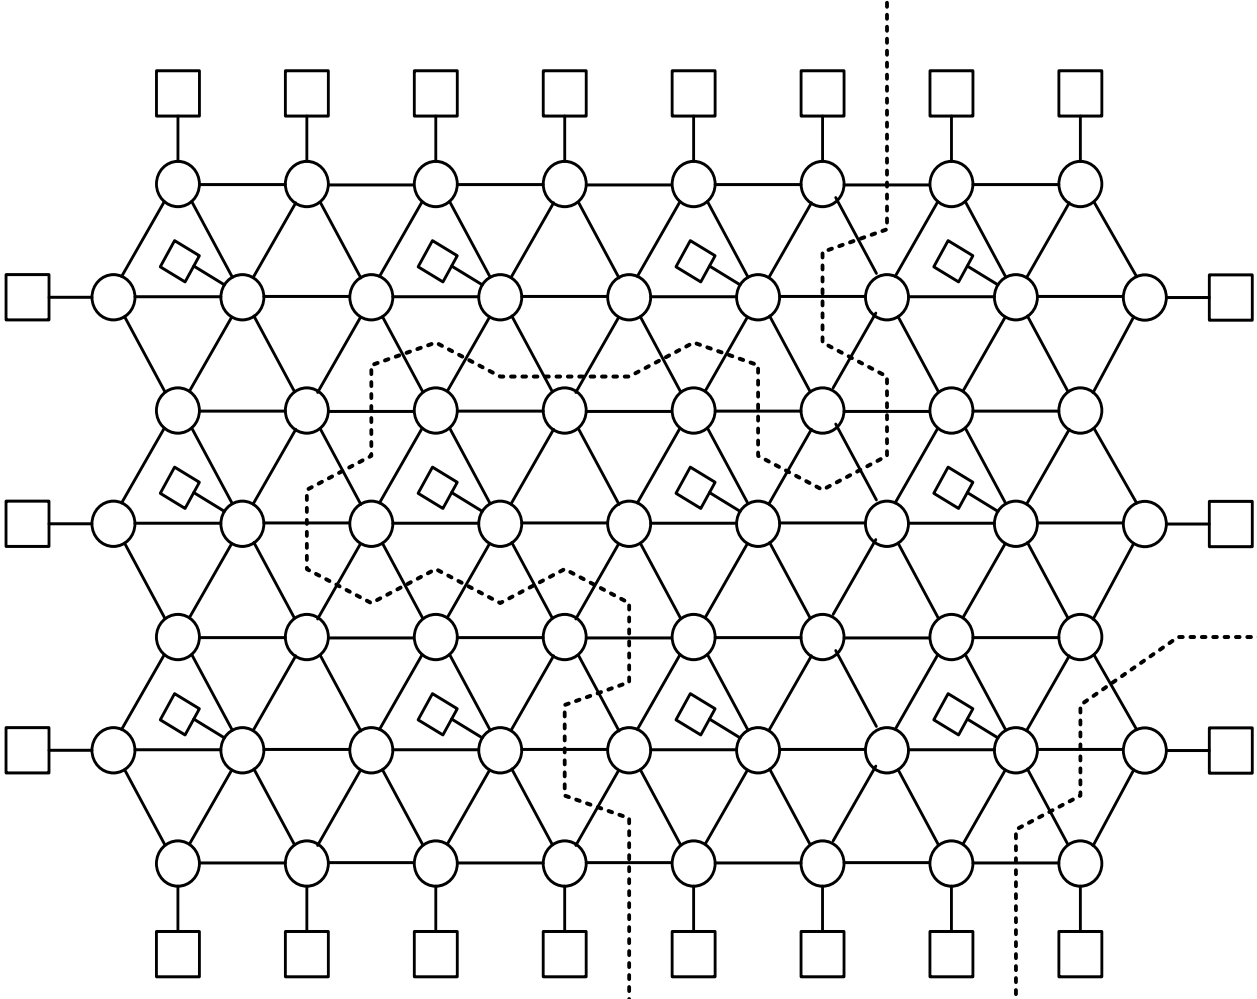


Supervisor 0



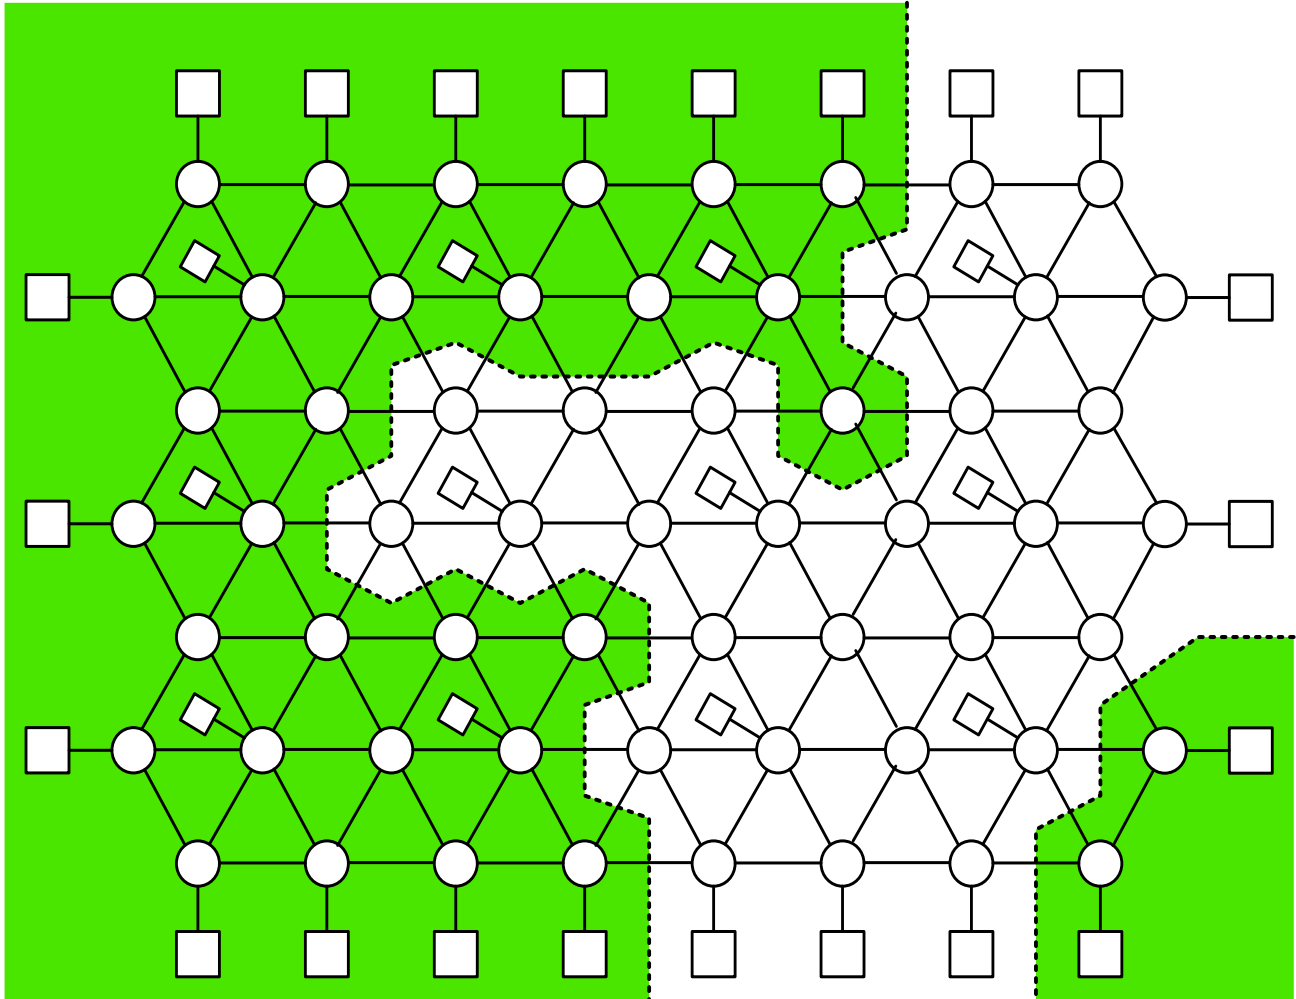
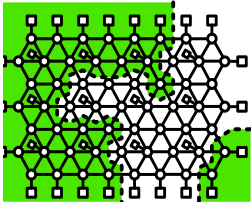
Supervisor 1

Supervisor 0



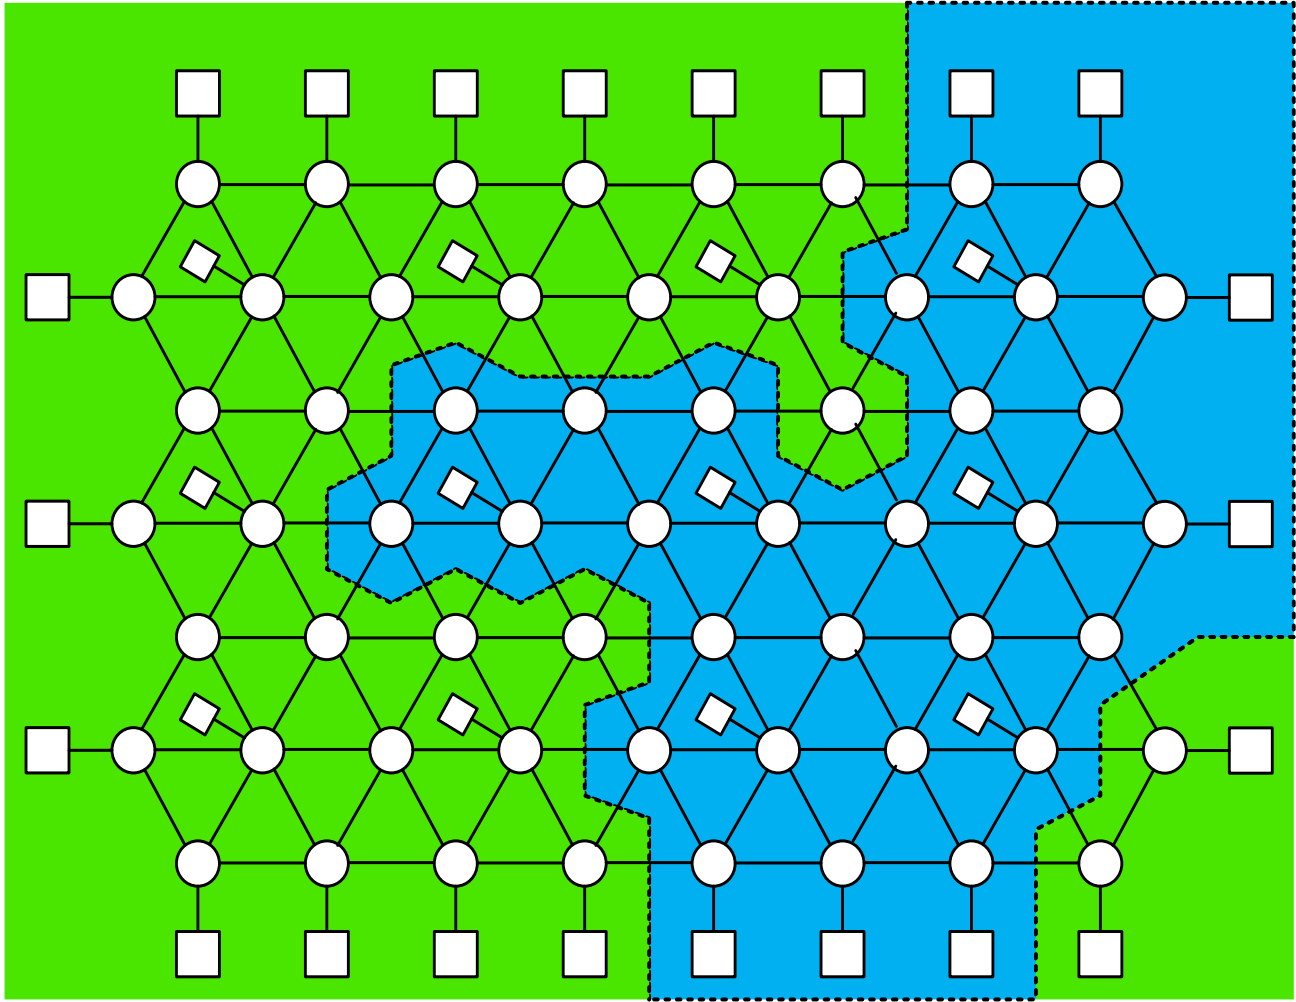
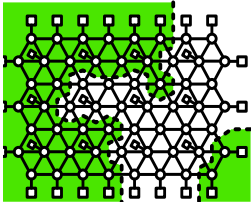
Supervisor 1

Supervisor 0

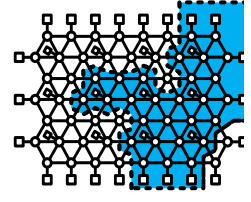


Supervisor 1

Supervisor 0

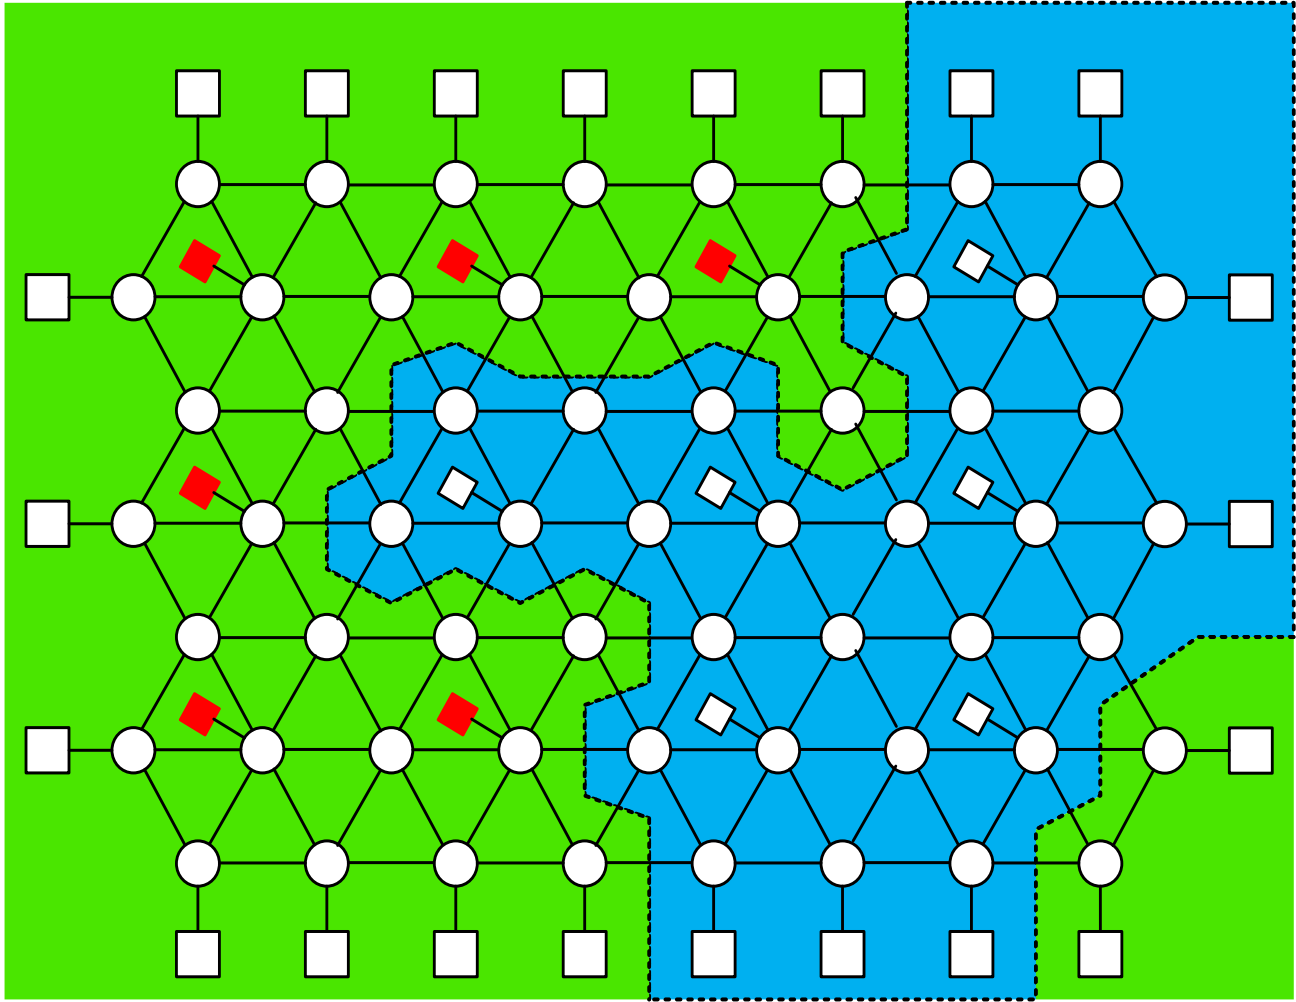
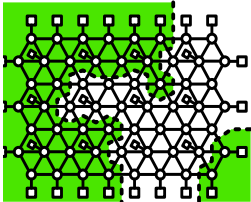


Supervisor 1

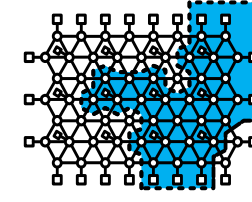




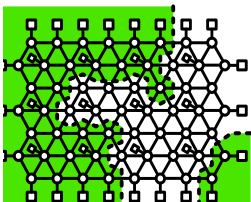
Supervisor 0



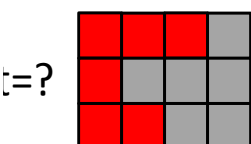
Supervisor 1



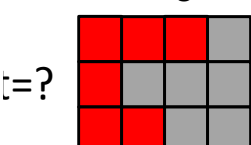
Supervisor 0



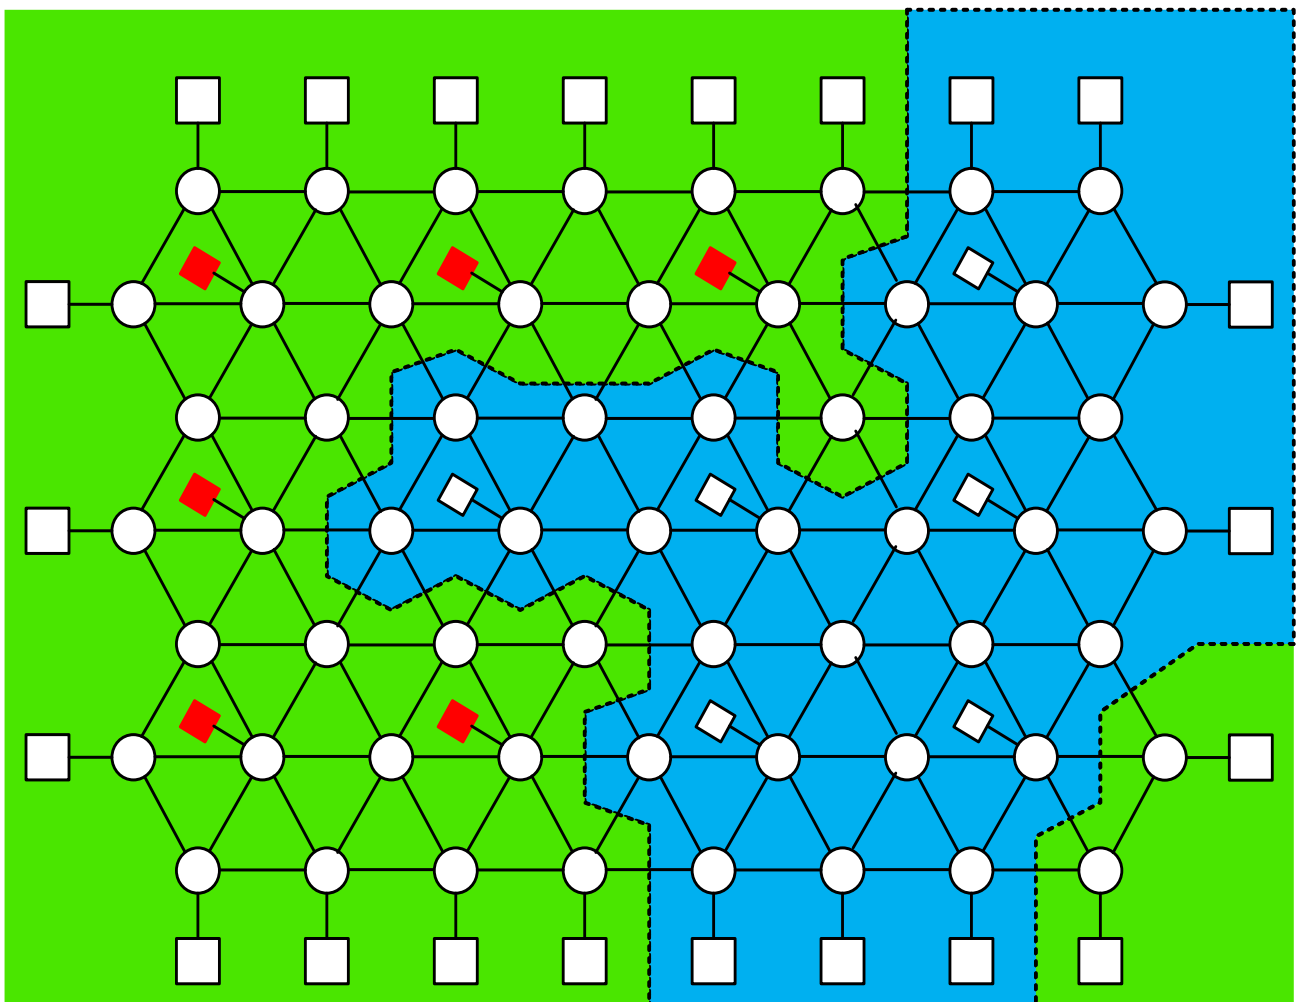
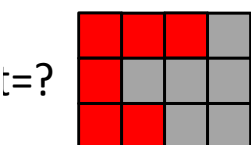
*Sending*



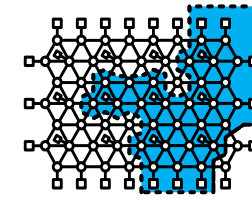
*Collecting*



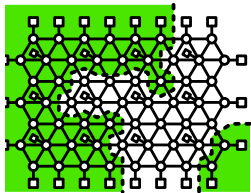
*Runahead*



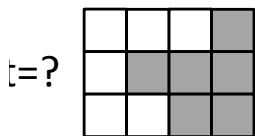
Supervisor 1



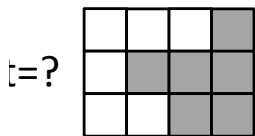
Supervisor 0



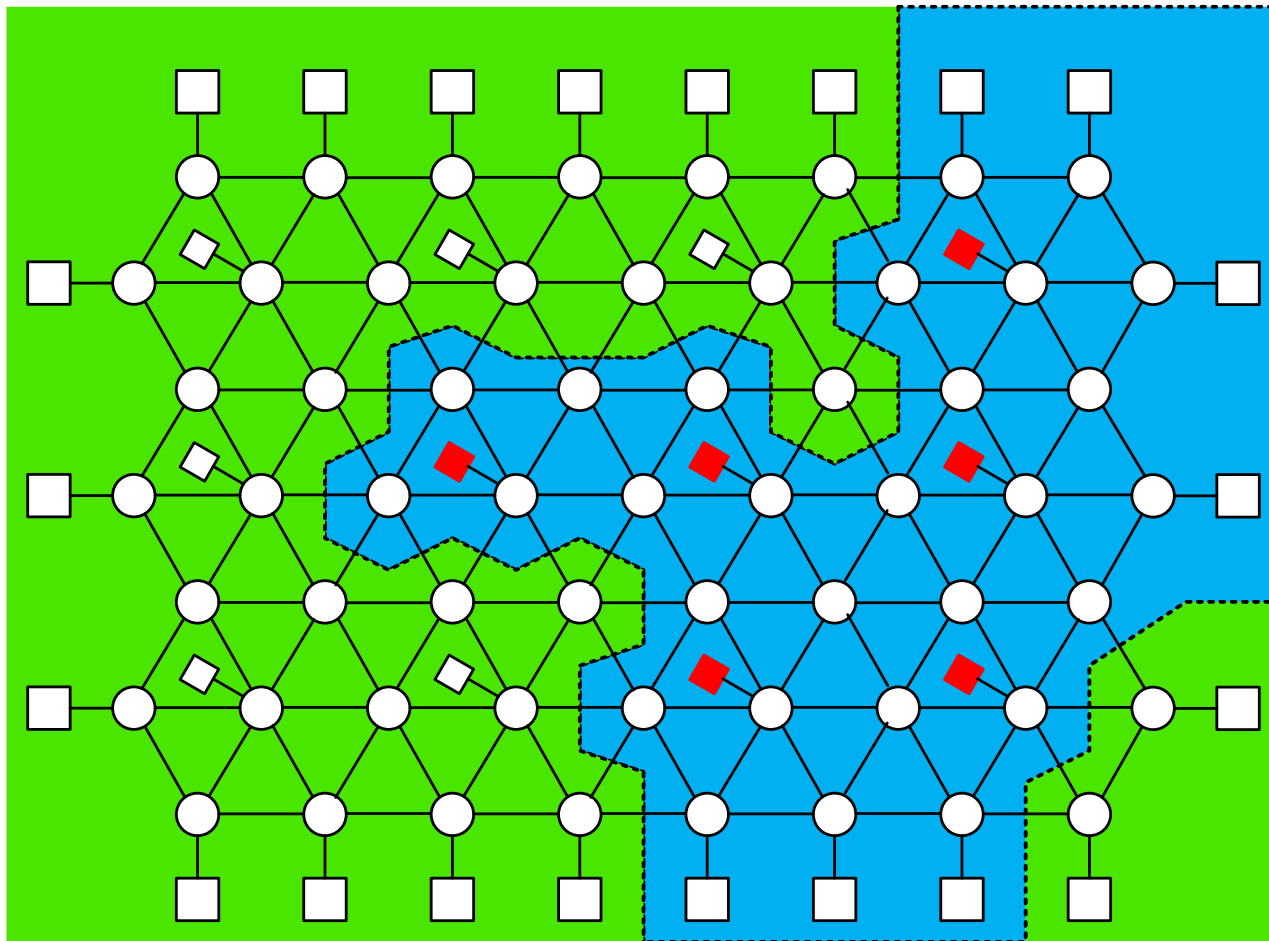
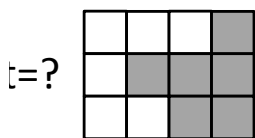
*Sending*



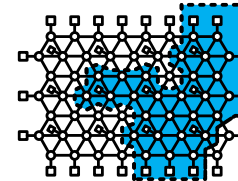
*Collecting*



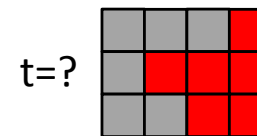
*Runahead*



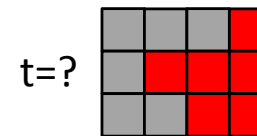
Supervisor 1



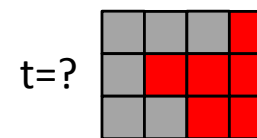
*Sending*



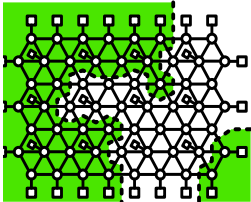
*Collecting*



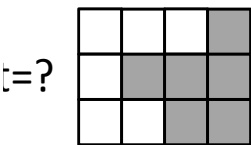
*Runahead*



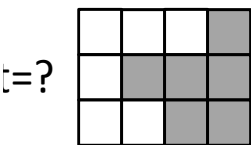
Supervisor 0



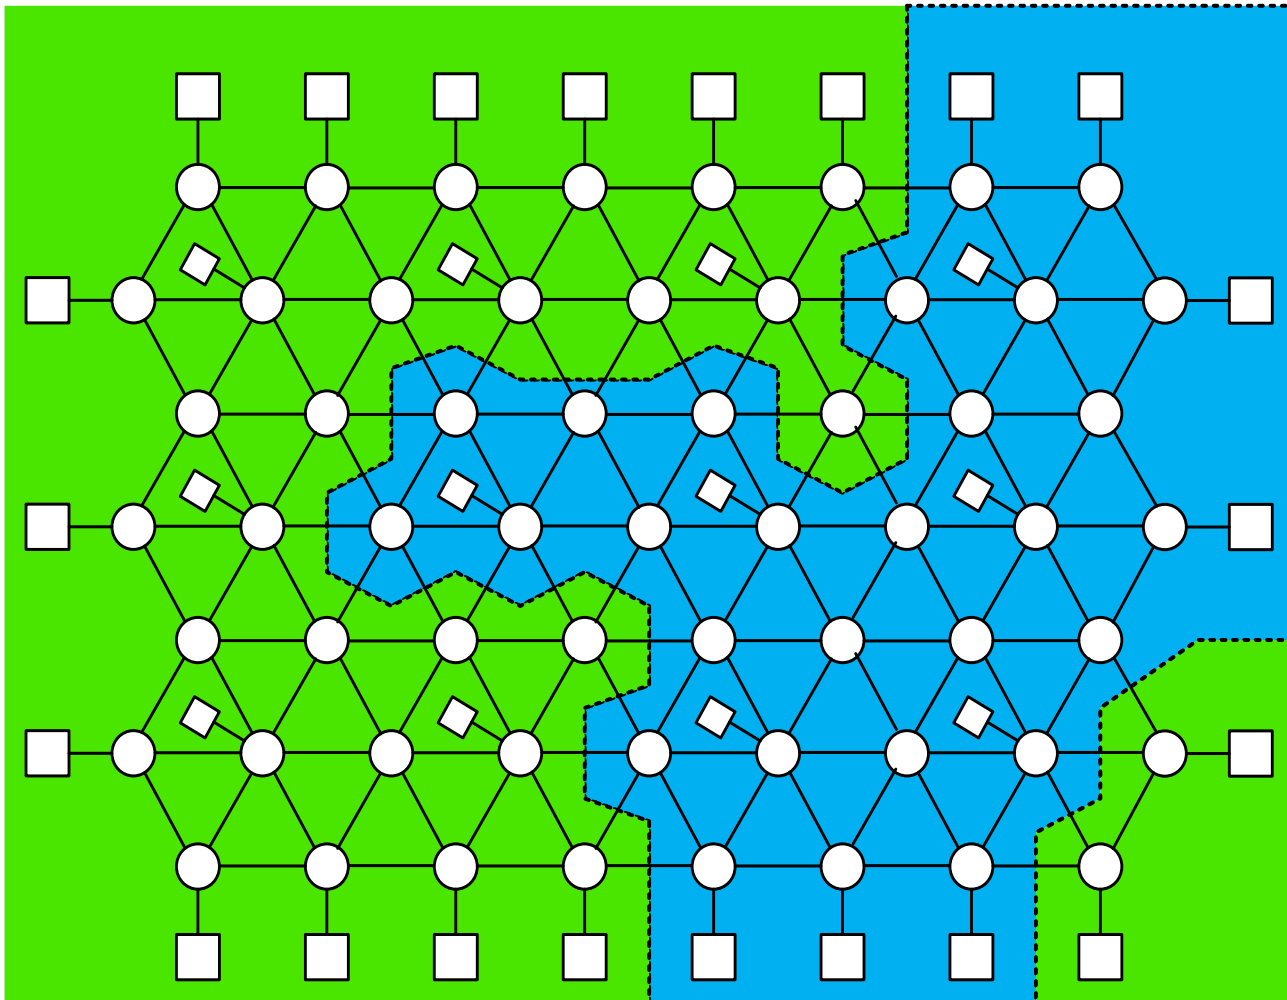
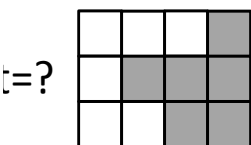
*Sending*



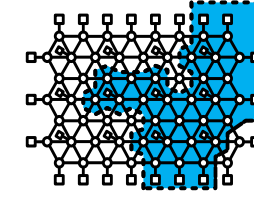
*Collecting*



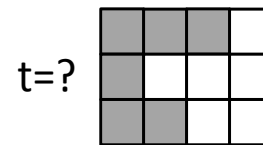
*Runahead*



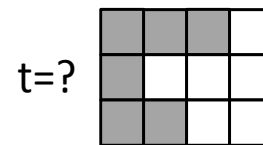
Supervisor 1



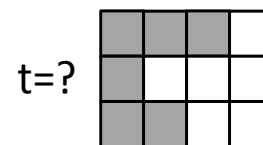
*Sending*



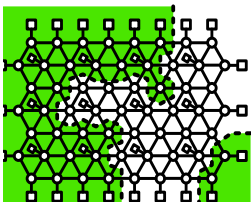
*Collecting*



*Runahead*

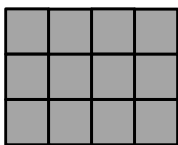


Supervisor 0



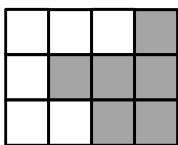
*Sending*

$t=?$



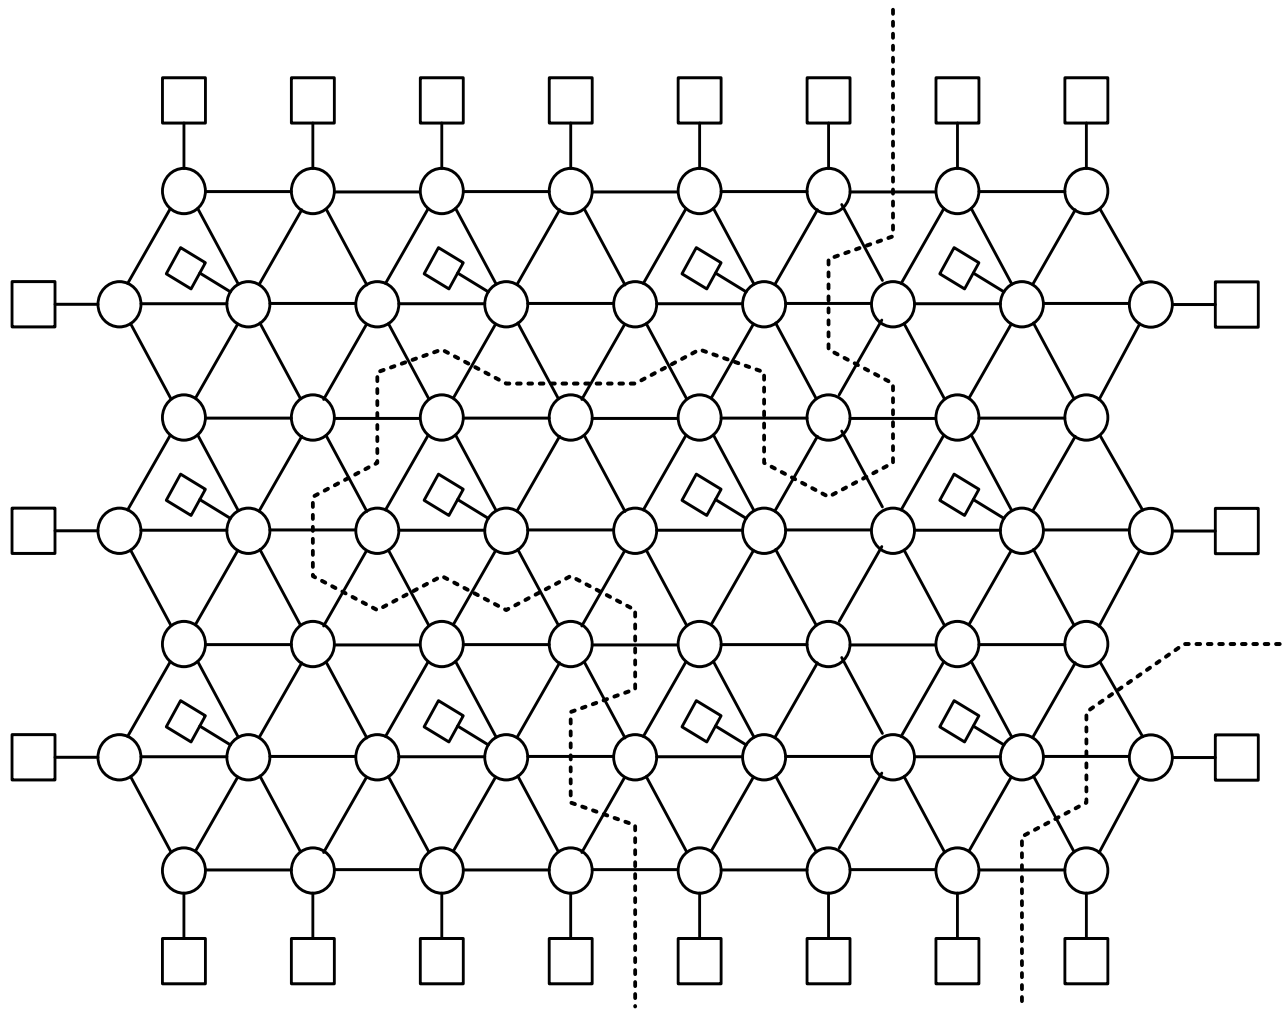
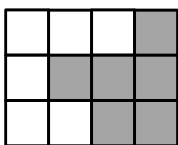
*Collecting*

$t=0$

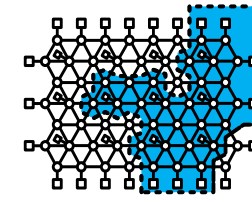


*Runahead*

$t=1$

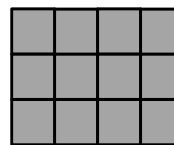


Supervisor 1



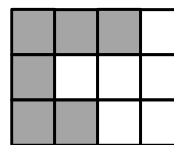
*Sending*

$t=?$



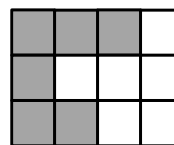
*Collecting*

$t=0$

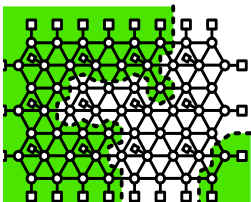


*Runahead*

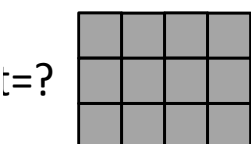
$t=1$



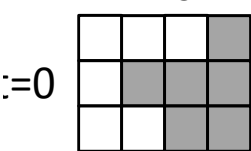
Supervisor 0



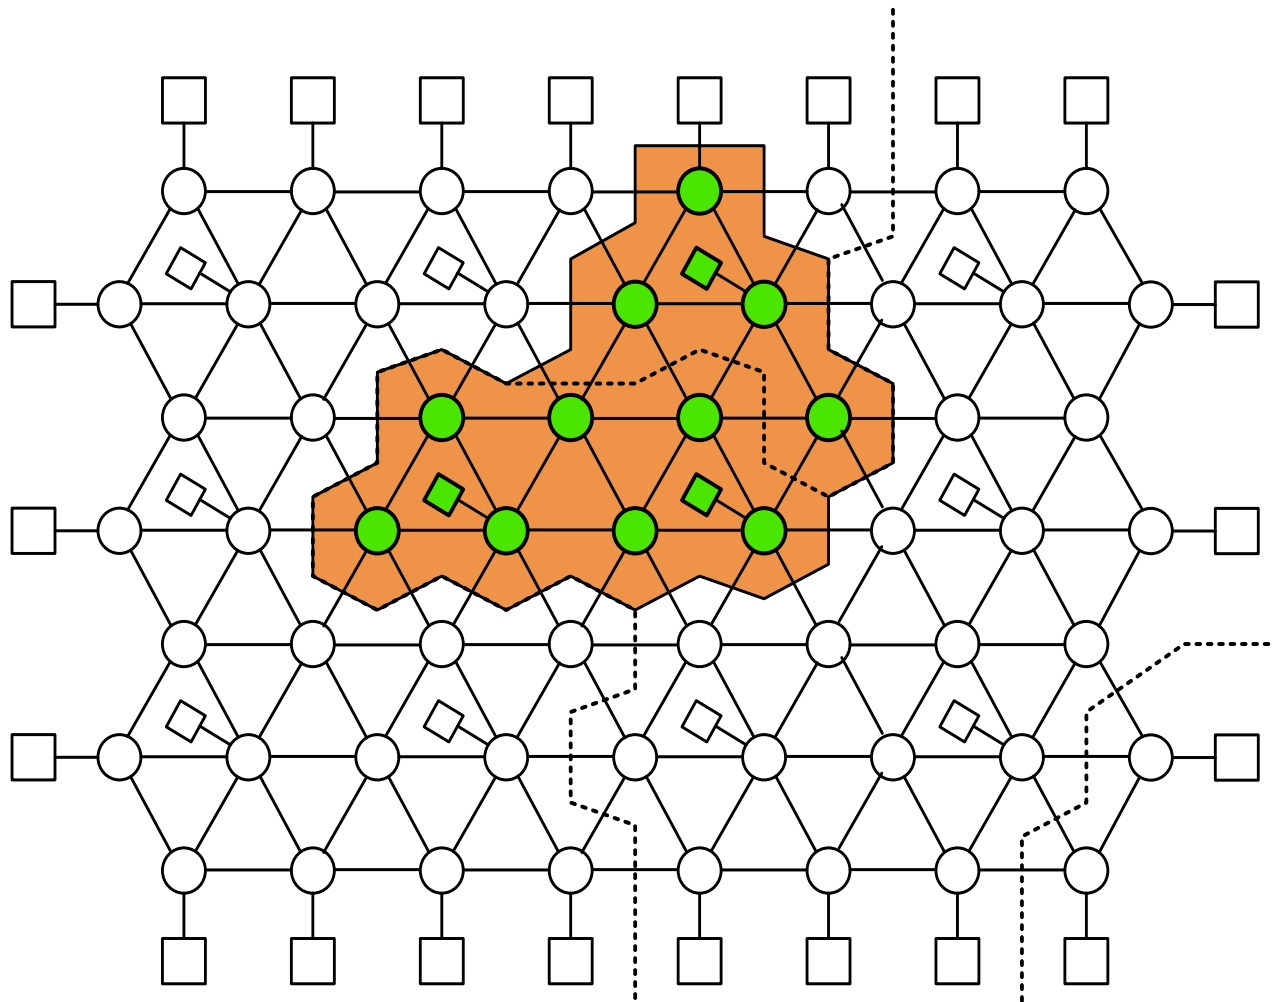
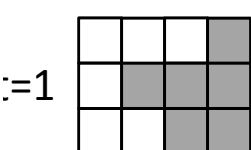
*Sending*



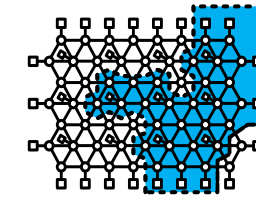
*Collecting*



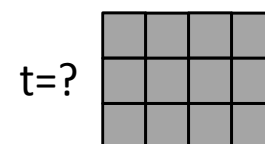
*Runahead*



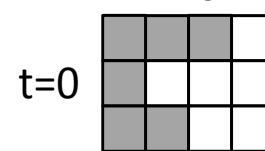
Supervisor 1



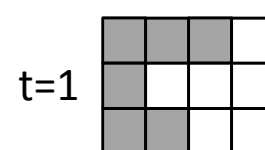
*Sending*



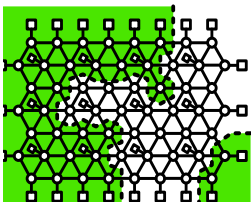
*Collecting*



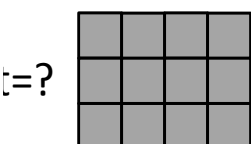
*Runahead*



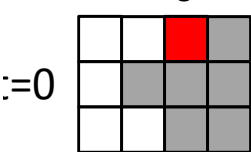
# Supervisor 0



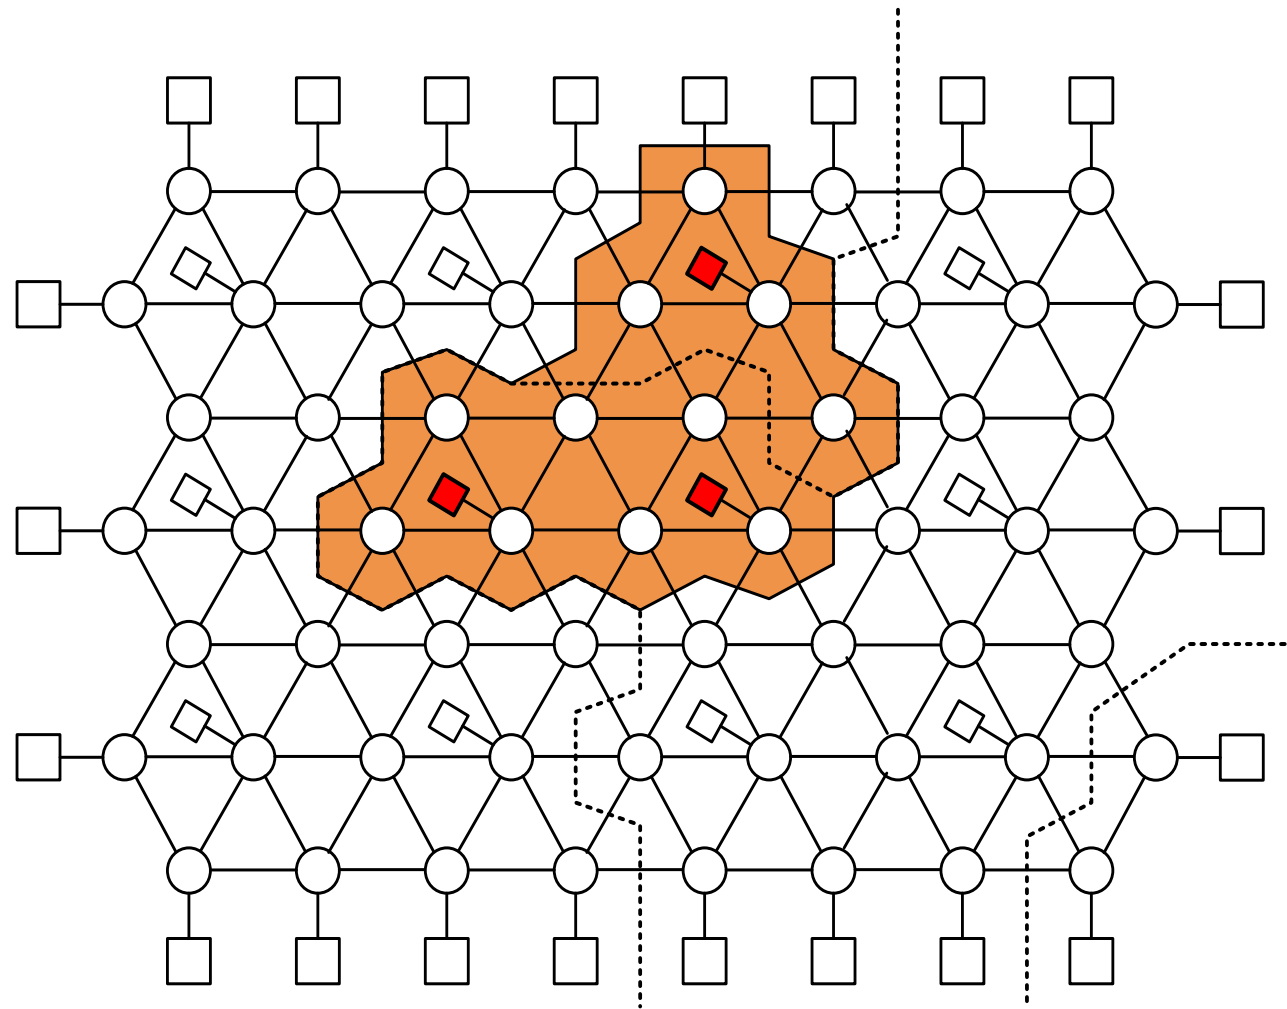
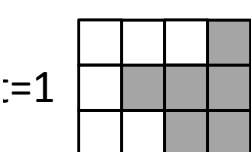
*Sending*



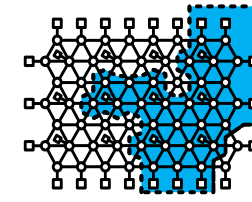
*Collecting*



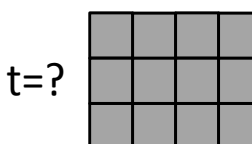
*Runahead*



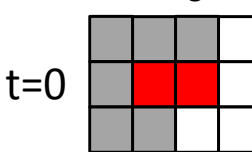
# Supervisor 1



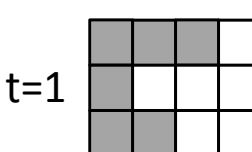
*Sending*



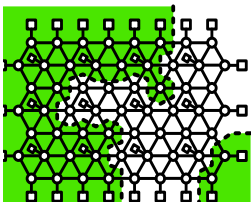
*Collecting*



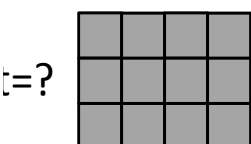
*Runahead*



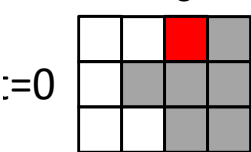
Supervisor 0



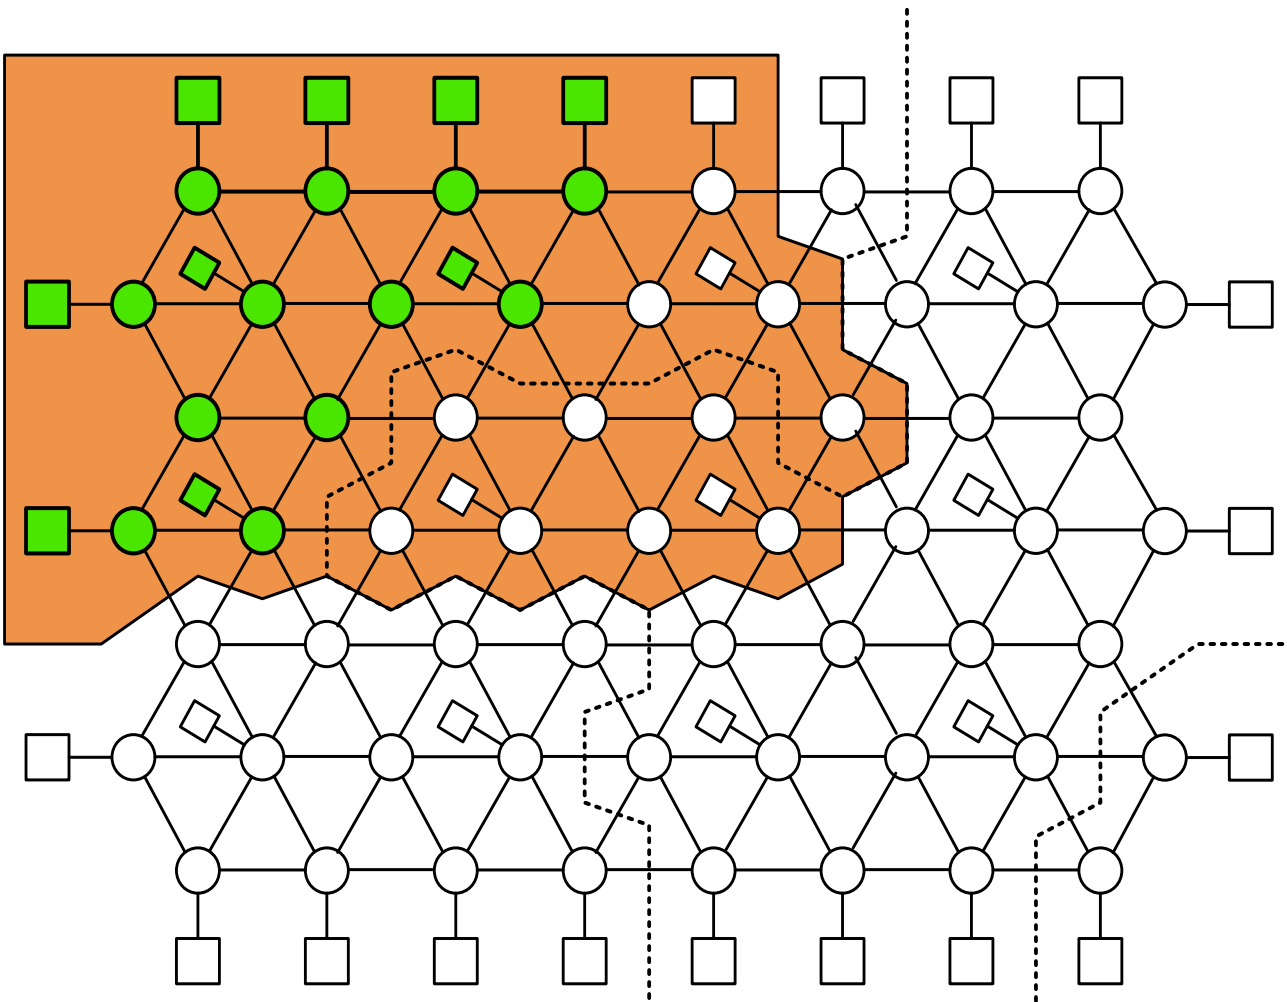
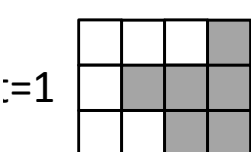
*Sending*



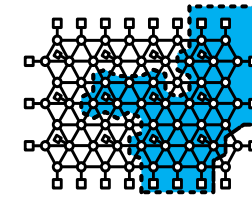
*Collecting*



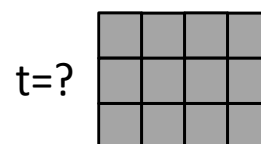
*Runahead*



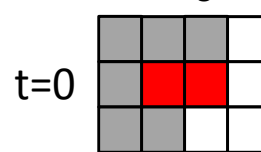
Supervisor 1



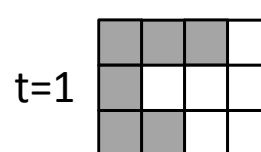
*Sending*



*Collecting*

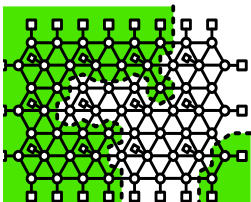


*Runahead*

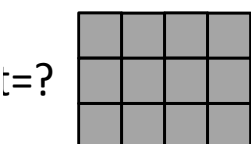




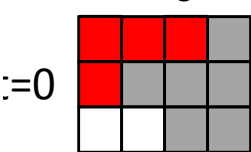
Supervisor 0



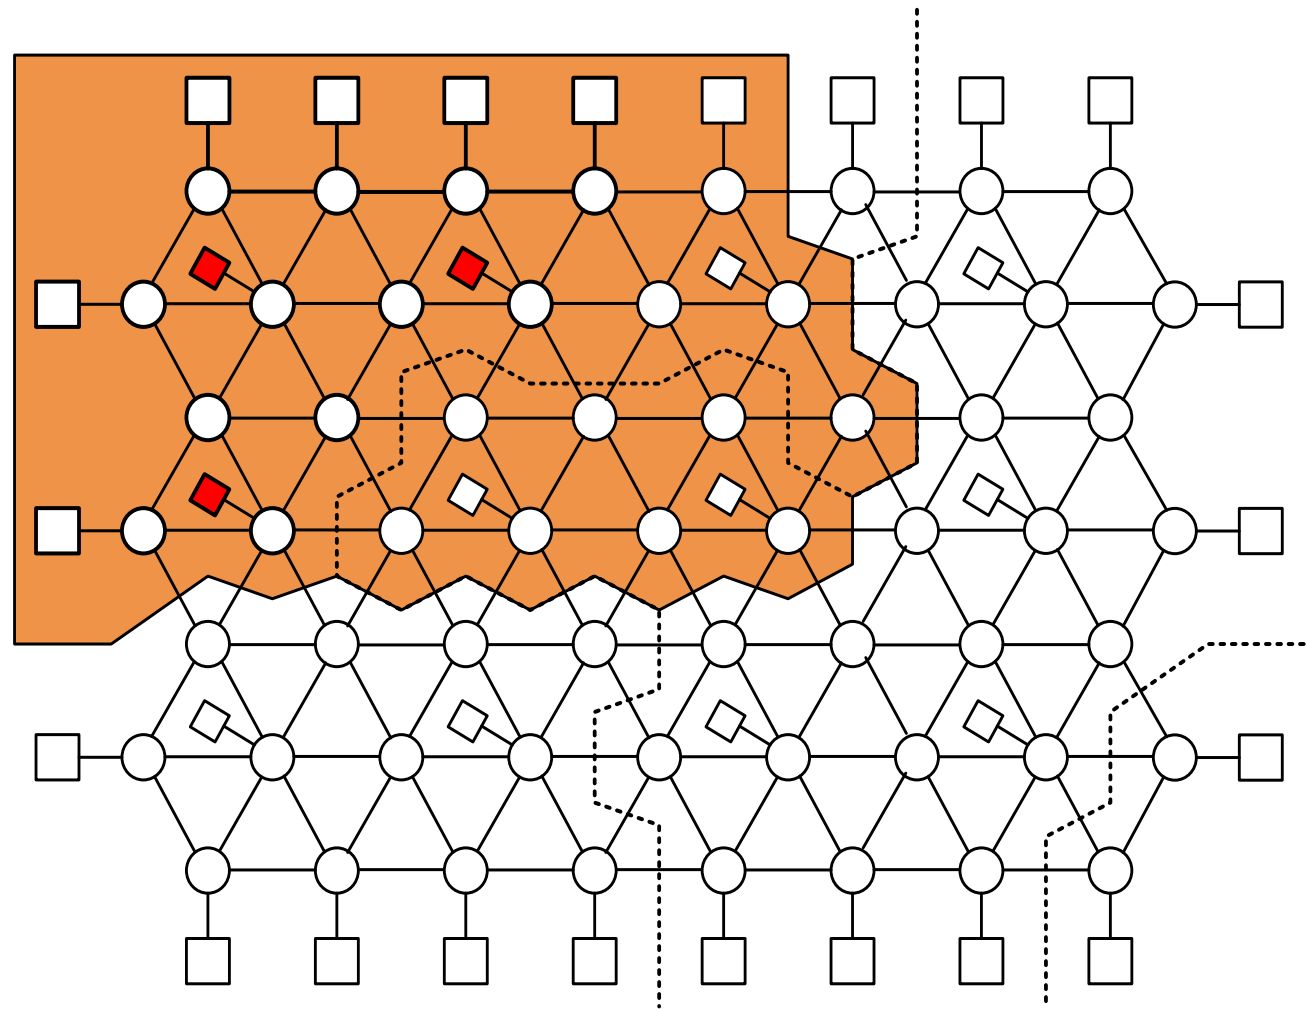
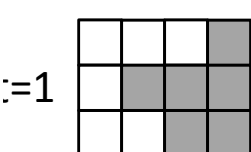
*Sending*



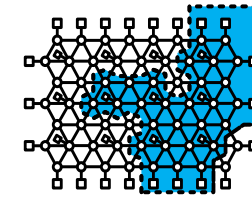
*Collecting*



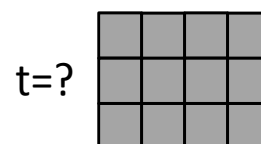
*Runahead*



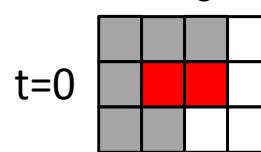
Supervisor 1



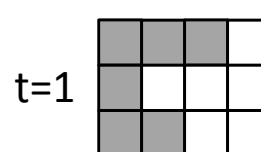
*Sending*



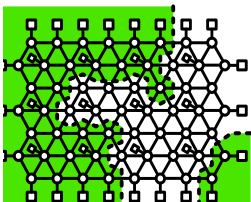
*Collecting*



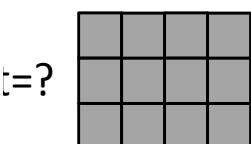
*Runahead*



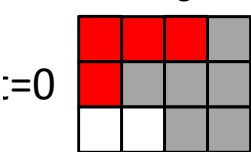
Supervisor 0



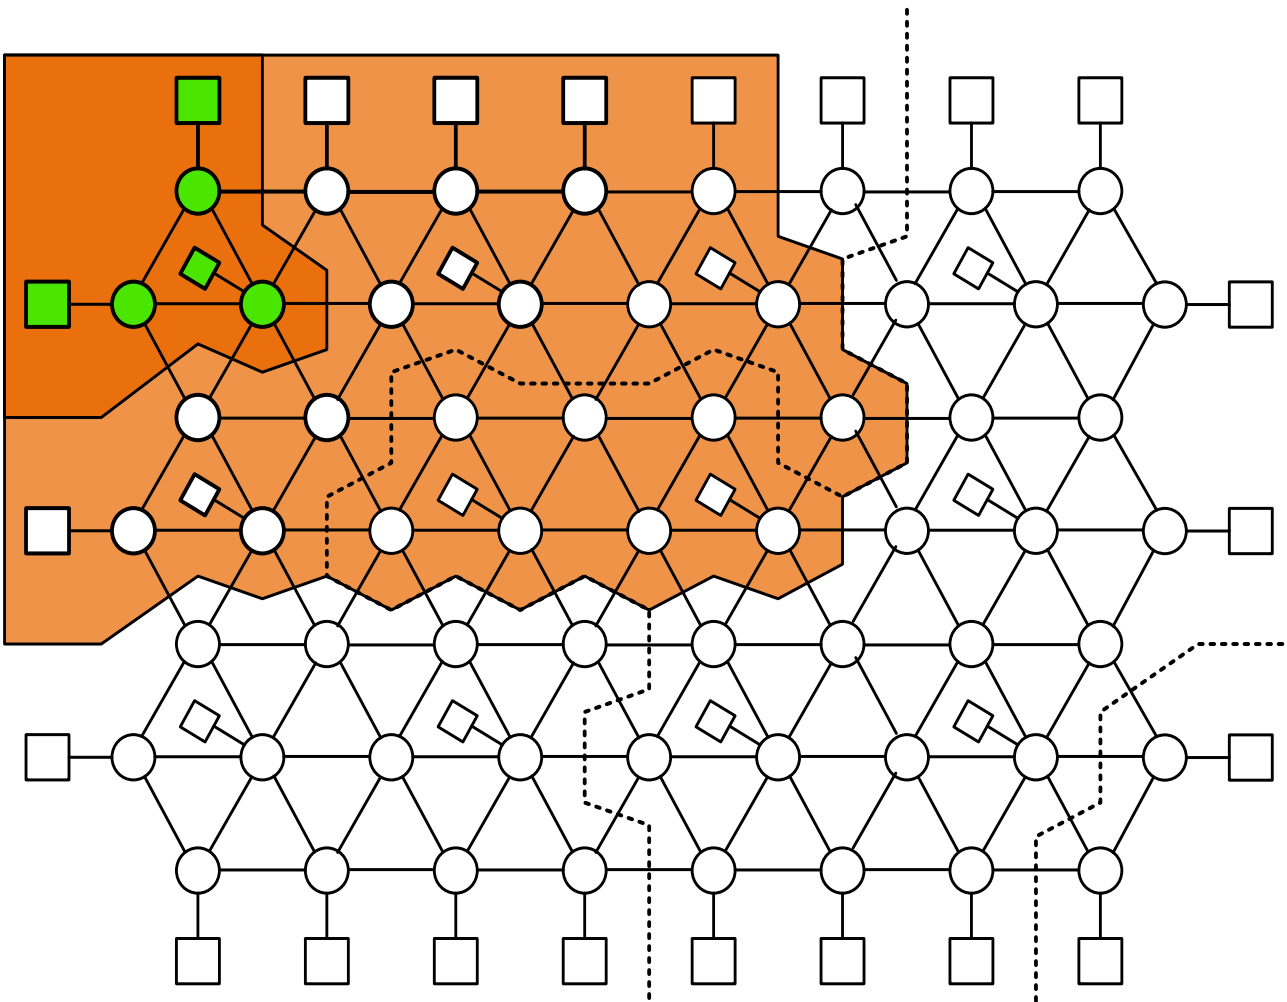
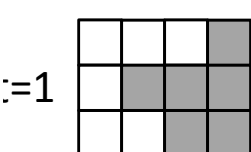
Sending



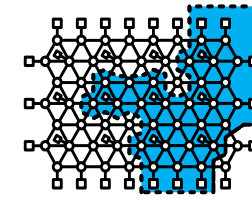
Collecting



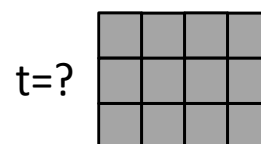
Runahead



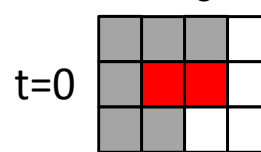
Supervisor 1



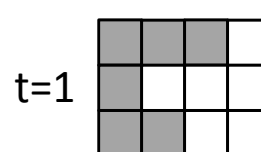
Sending



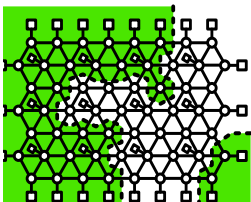
Collecting



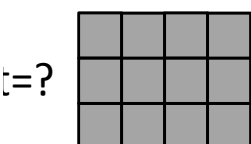
Runahead



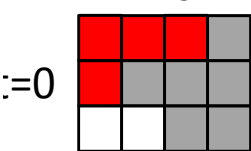
Supervisor 0



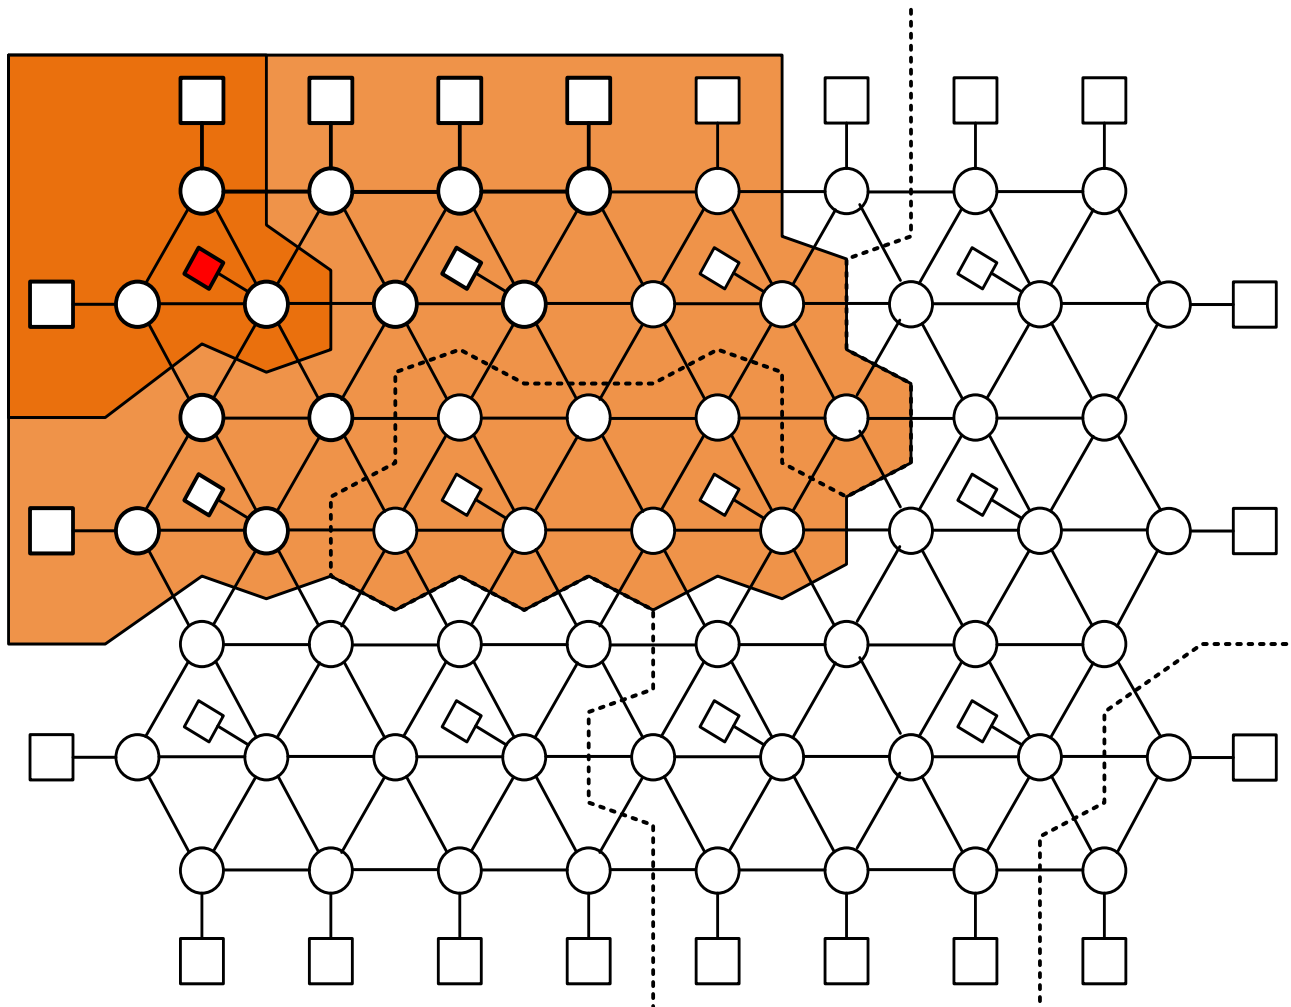
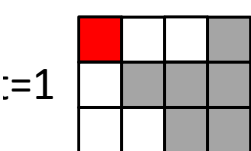
Sending



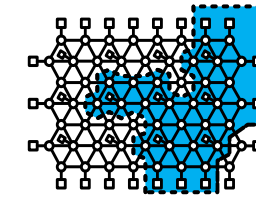
Collecting



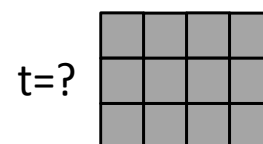
Runahead



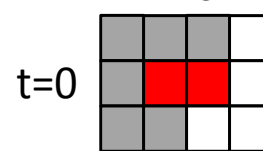
Supervisor 1



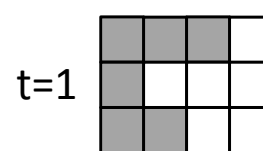
Sending



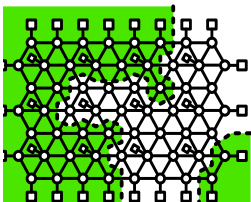
Collecting



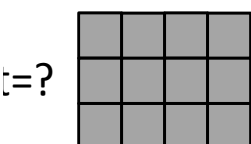
Runahead



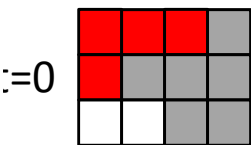
Supervisor 0



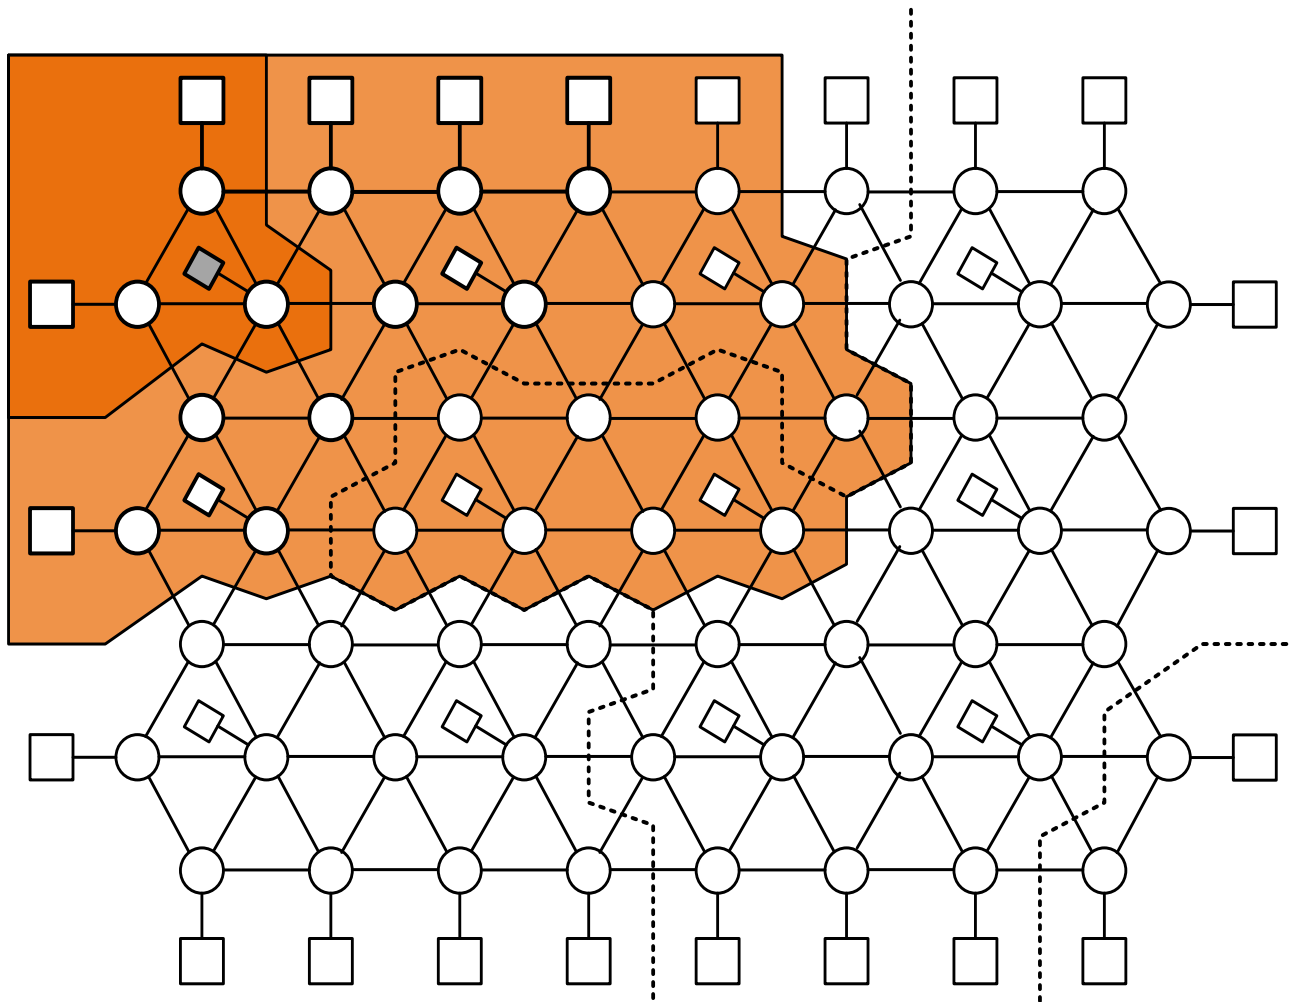
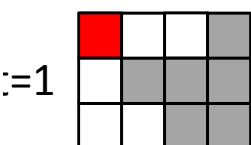
*Sending*



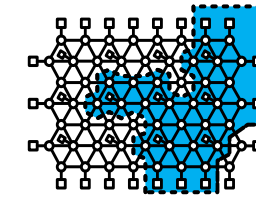
*Collecting*



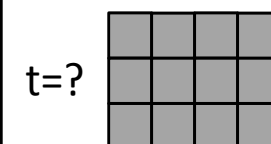
*Runahead*



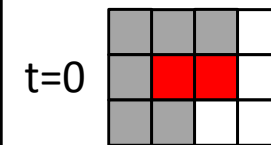
Supervisor 1



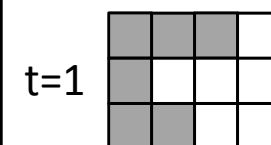
*Sending*



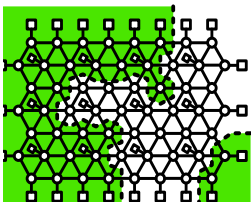
*Collecting*



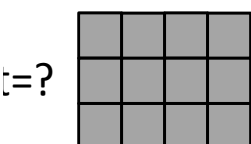
*Runahead*



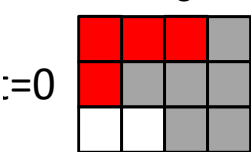
Supervisor 0



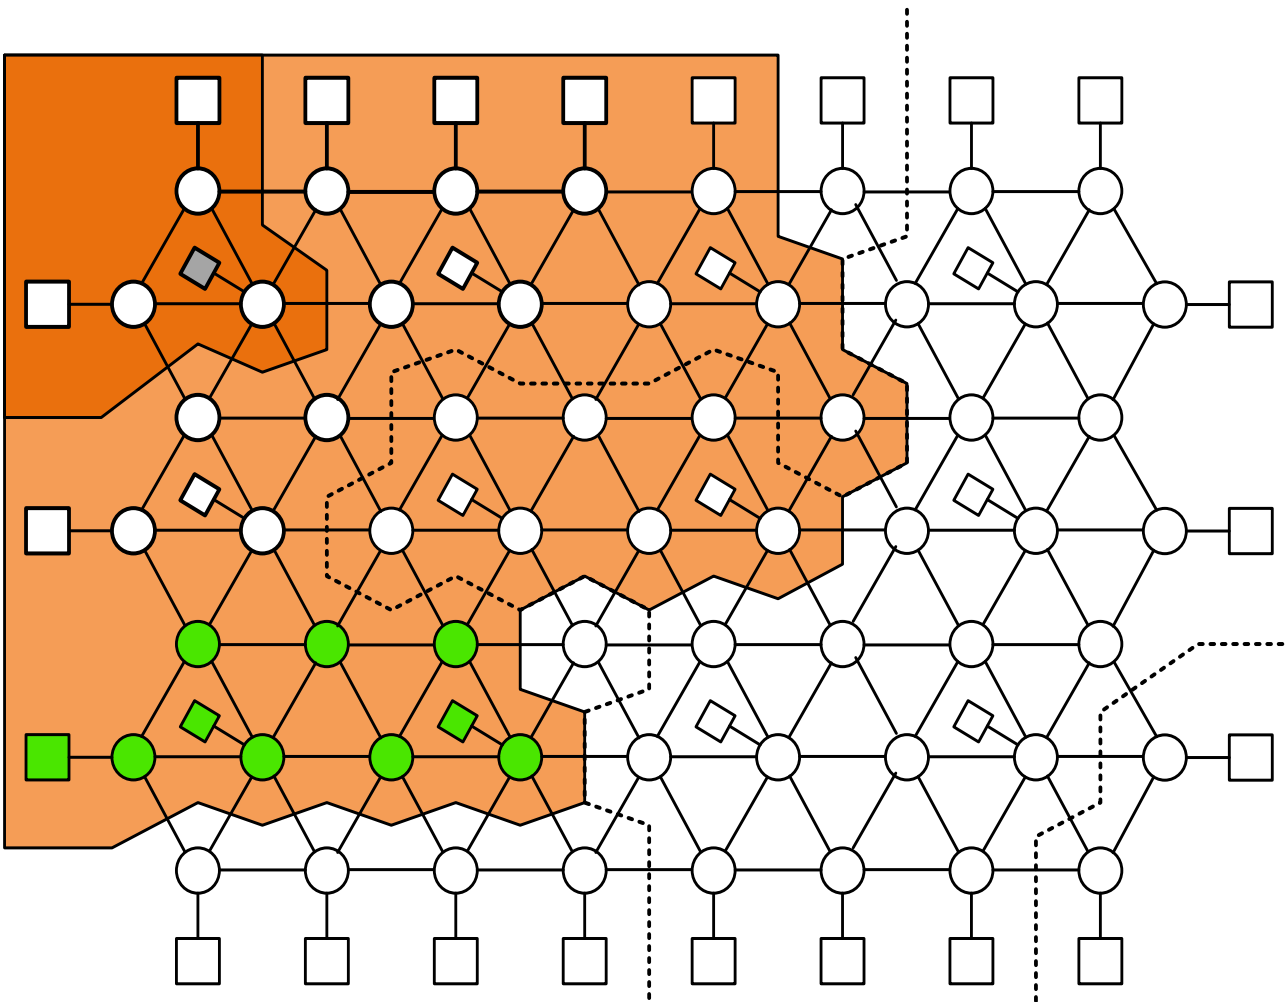
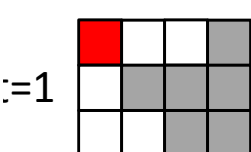
*Sending*



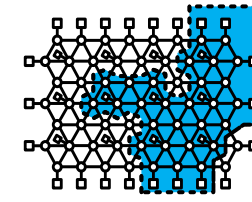
*Collecting*



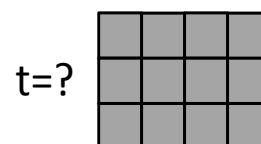
*Runahead*



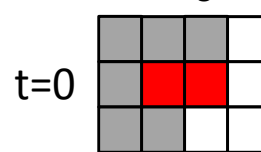
Supervisor 1



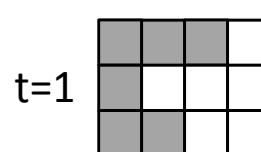
*Sending*



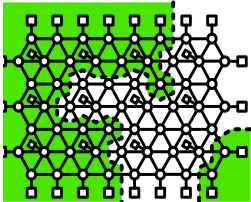
*Collecting*



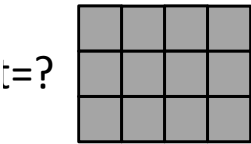
*Runahead*



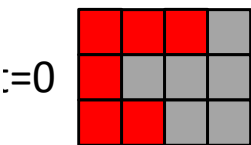
Supervisor 0



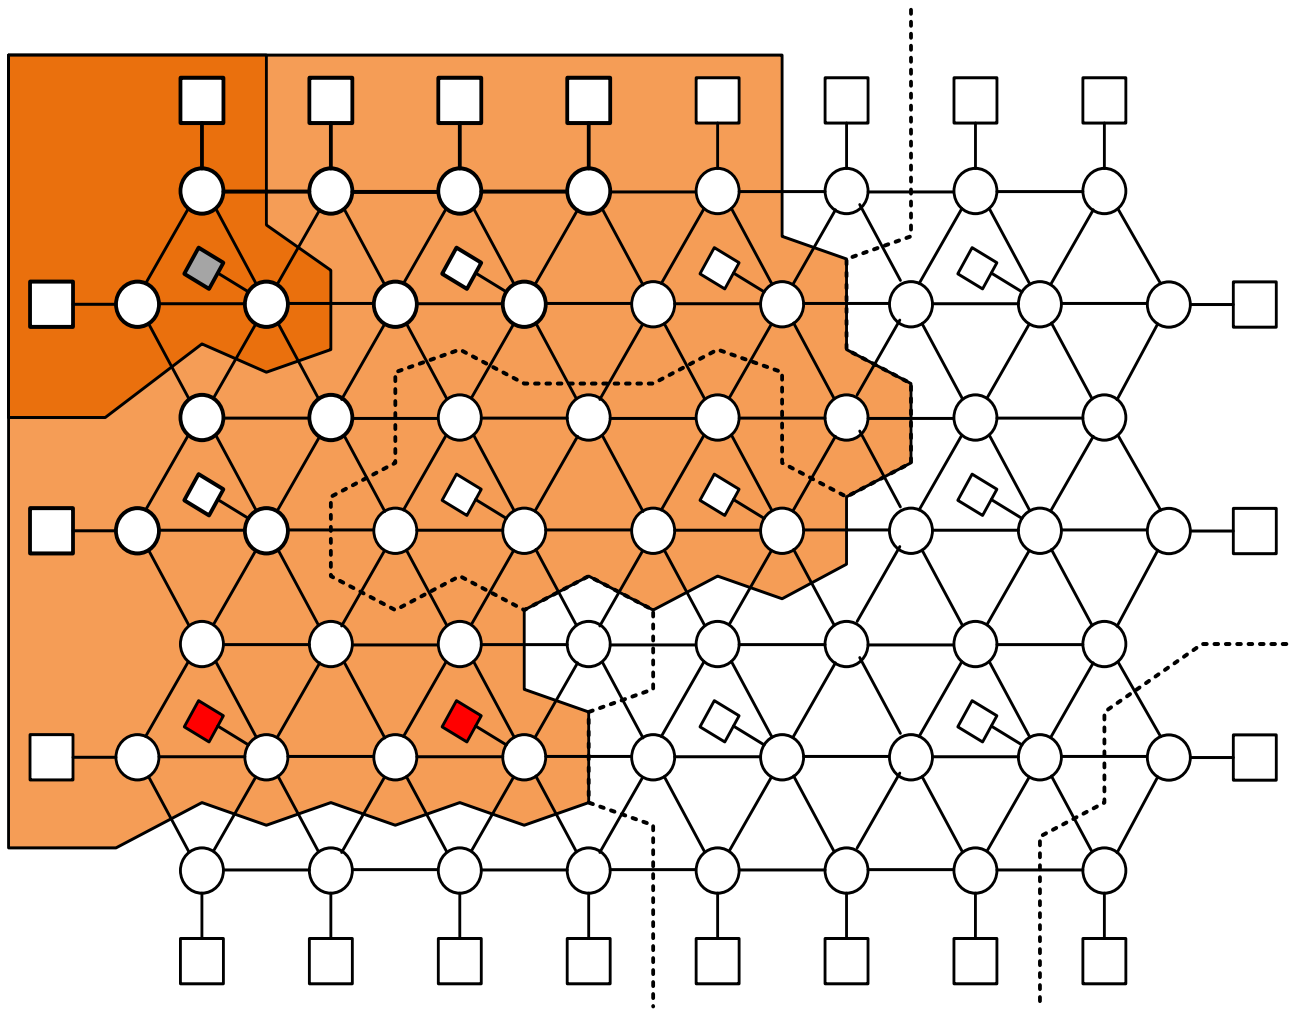
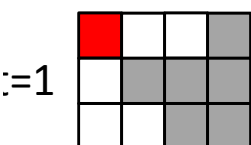
*Sending*



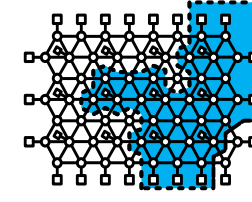
*Collecting*



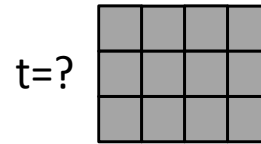
*Runahead*



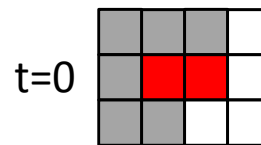
Supervisor 1



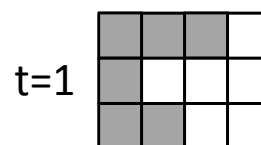
*Sending*



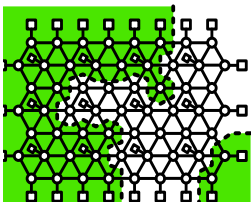
*Collecting*



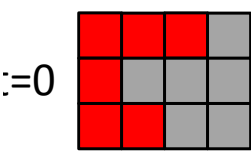
*Runahead*



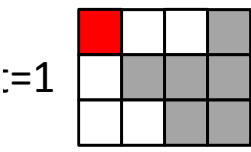
Supervisor 0



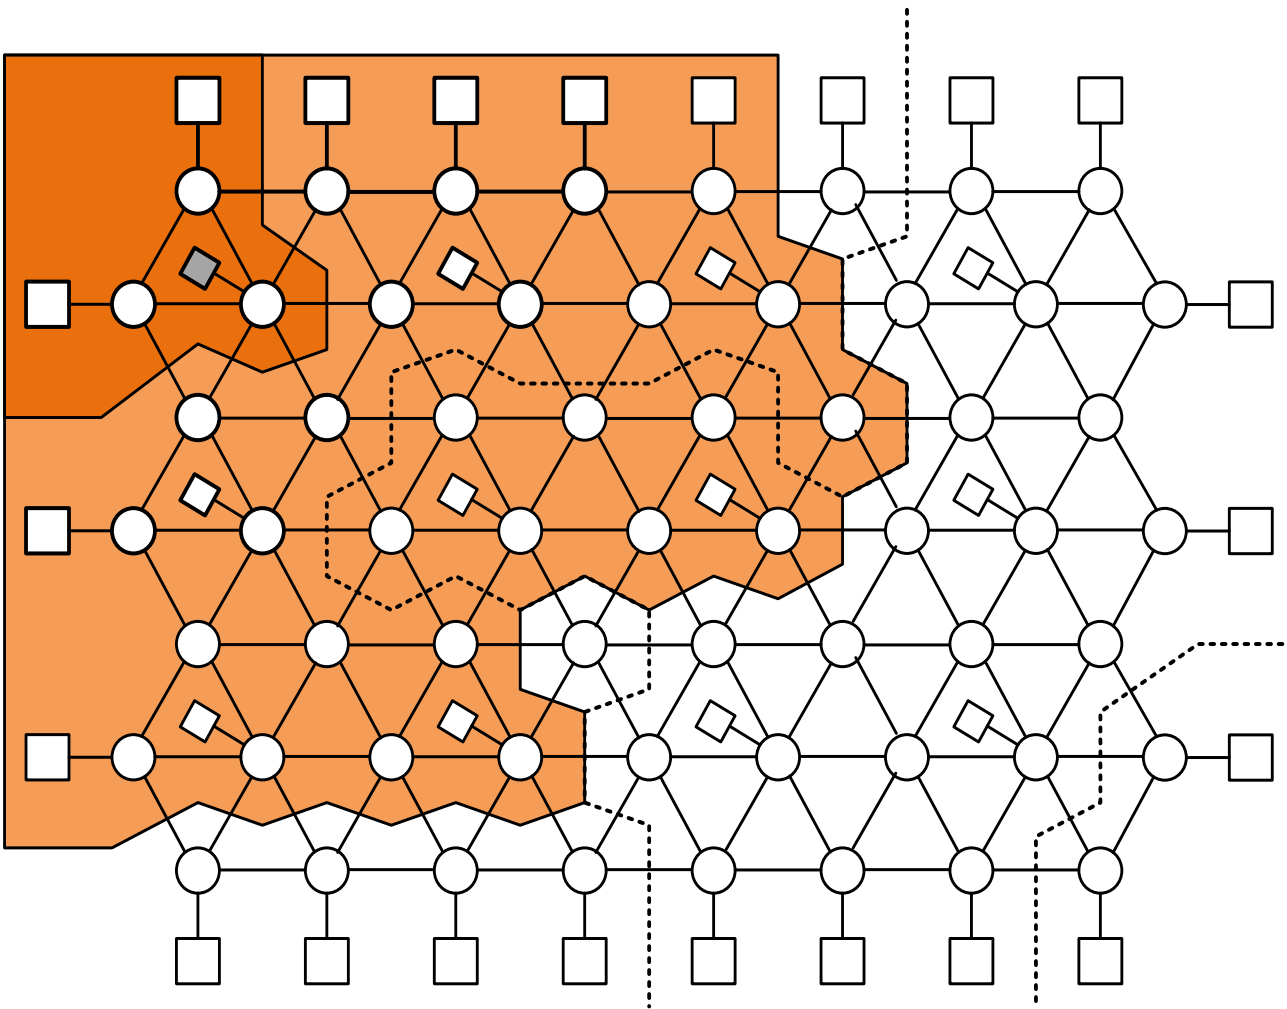
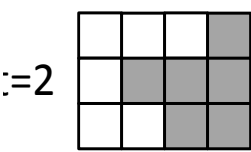
*Sending*



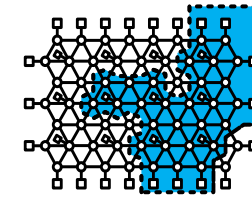
*Collecting*



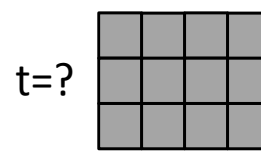
*Runahead*



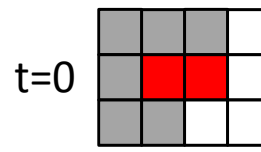
Supervisor 1



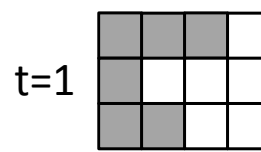
*Sending*



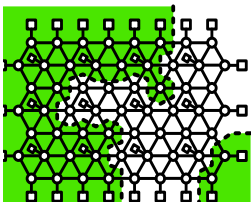
*Collecting*



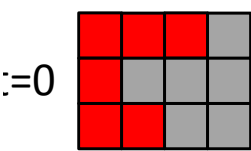
*Runahead*



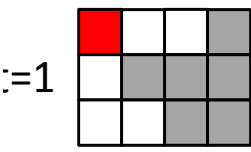
Supervisor 0



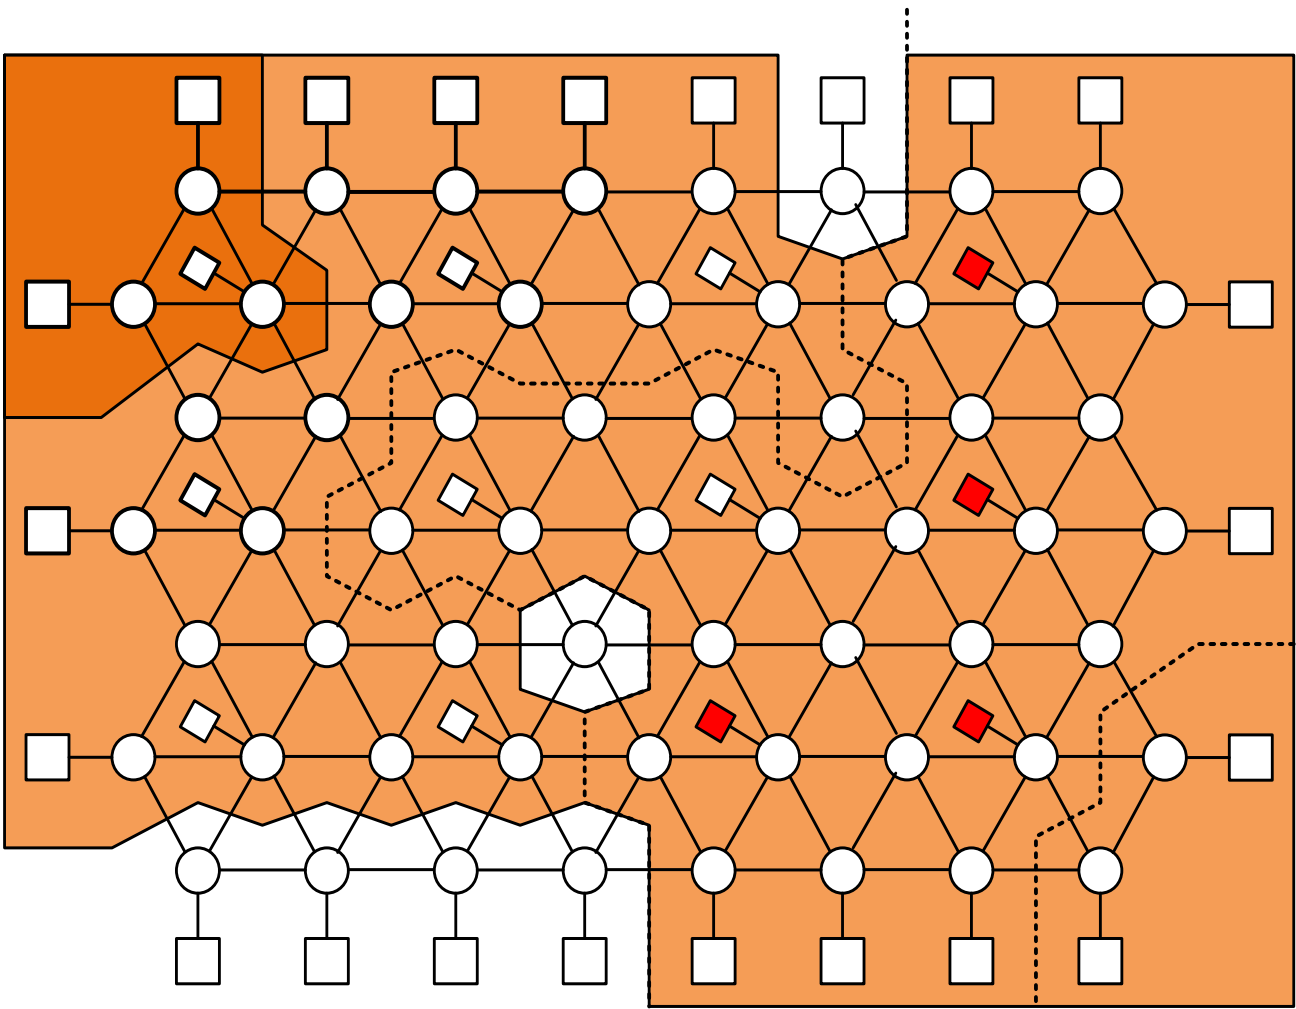
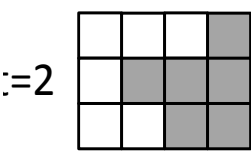
*Sending*



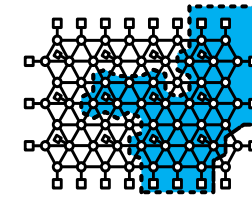
*Collecting*



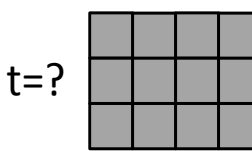
*Runahead*



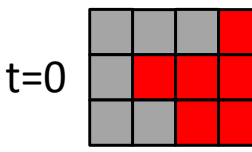
Supervisor 1



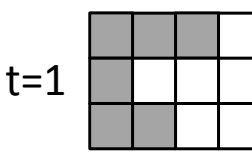
*Sending*



*Collecting*

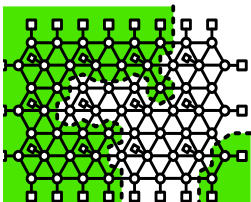


*Runahead*

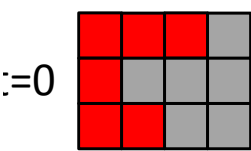




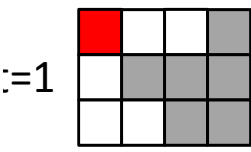
Supervisor 0



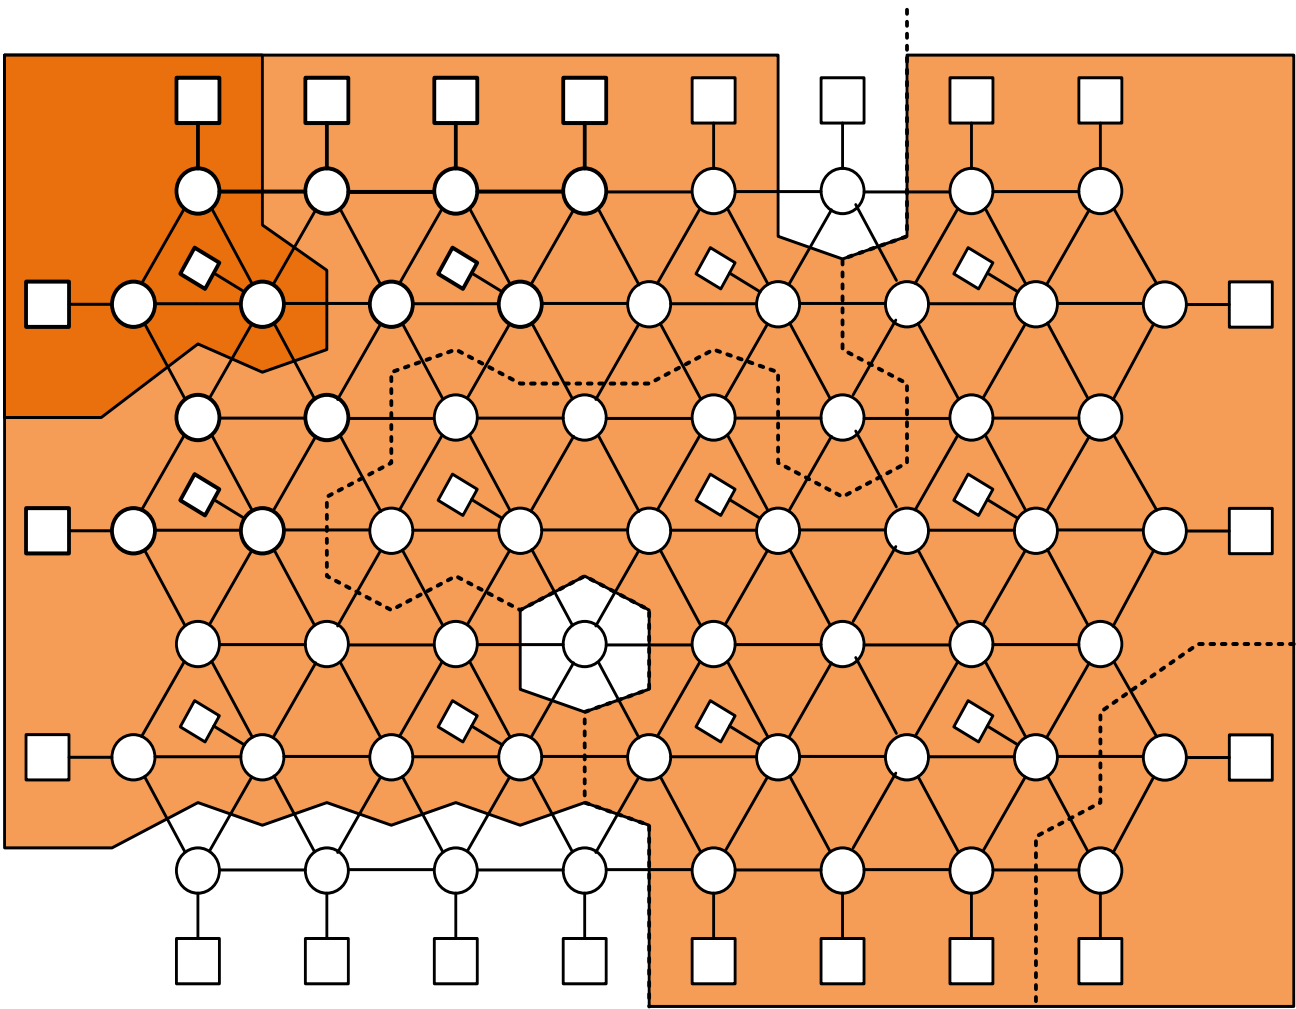
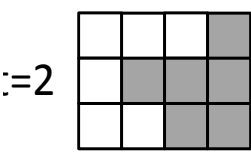
Sending



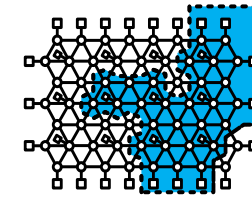
Collecting



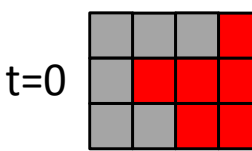
Runahead



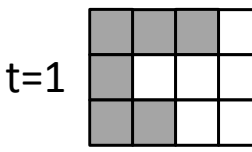
Supervisor 1



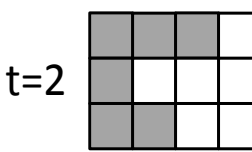
Sending



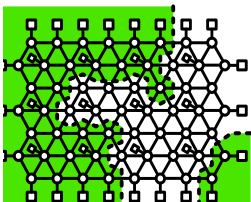
Collecting



Runahead

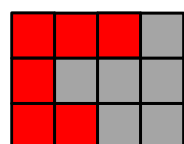


Supervisor 0



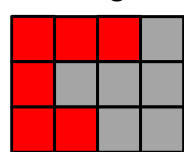
*Sending*

$t=0$



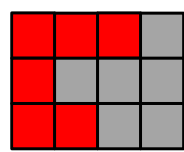
*Collecting*

$t=1$

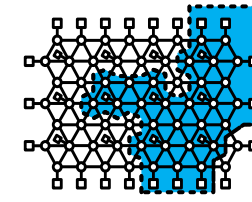


*Runahead*

$t=2$

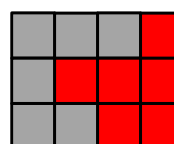


Supervisor 1



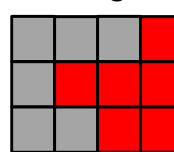
*Sending*

$t=0$



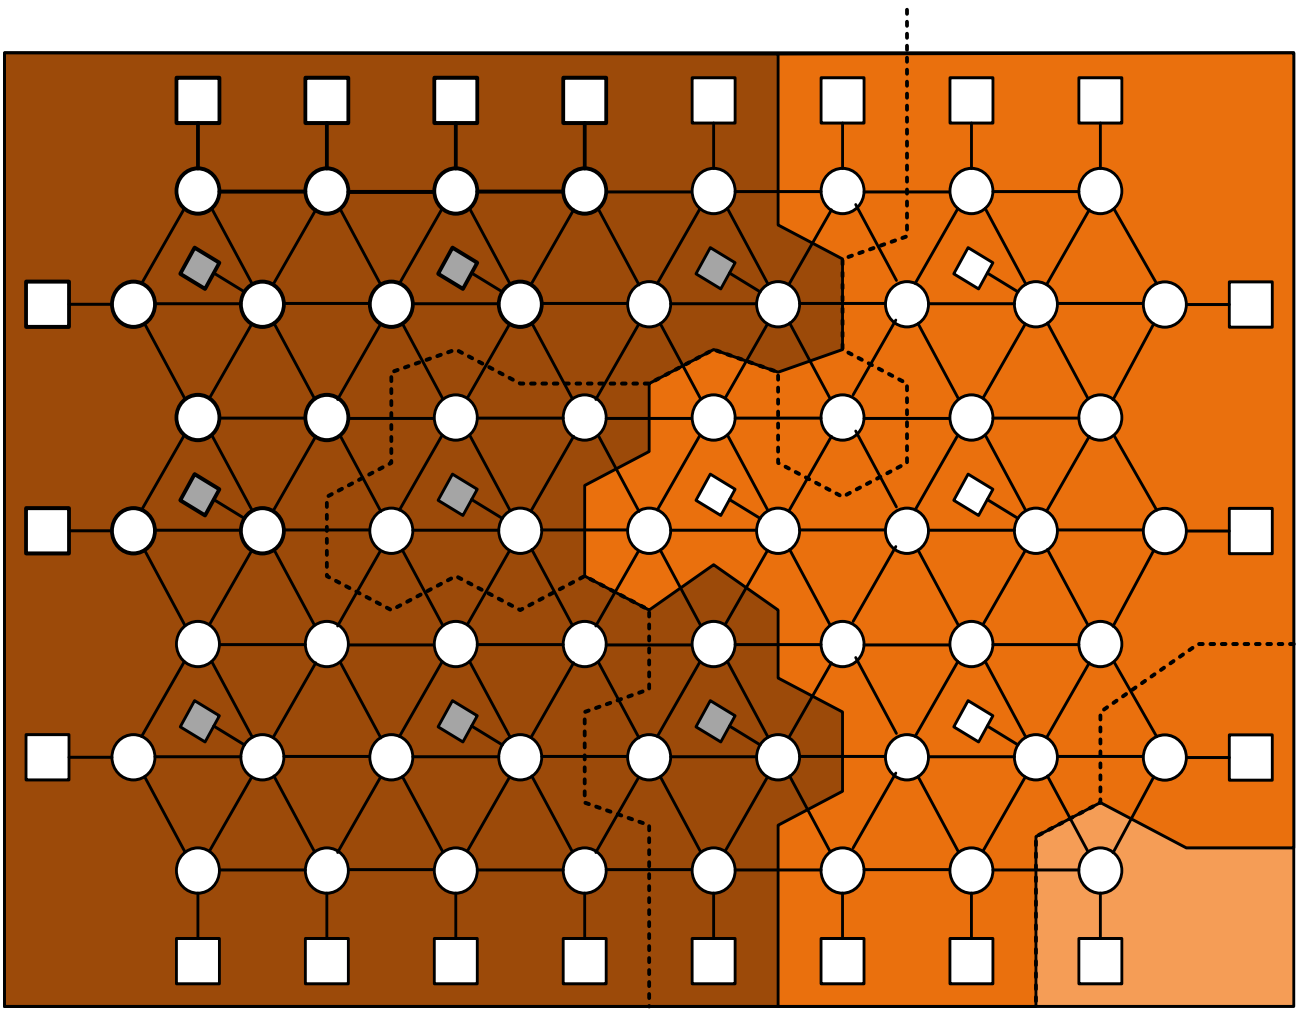
*Collecting*

$t=1$

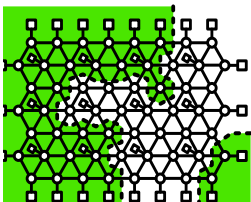


*Runahead*

$t=2$

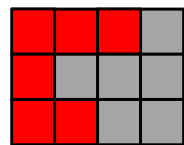


Supervisor 0



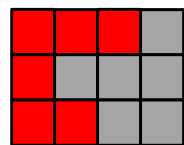
*Sending*

$t=0$



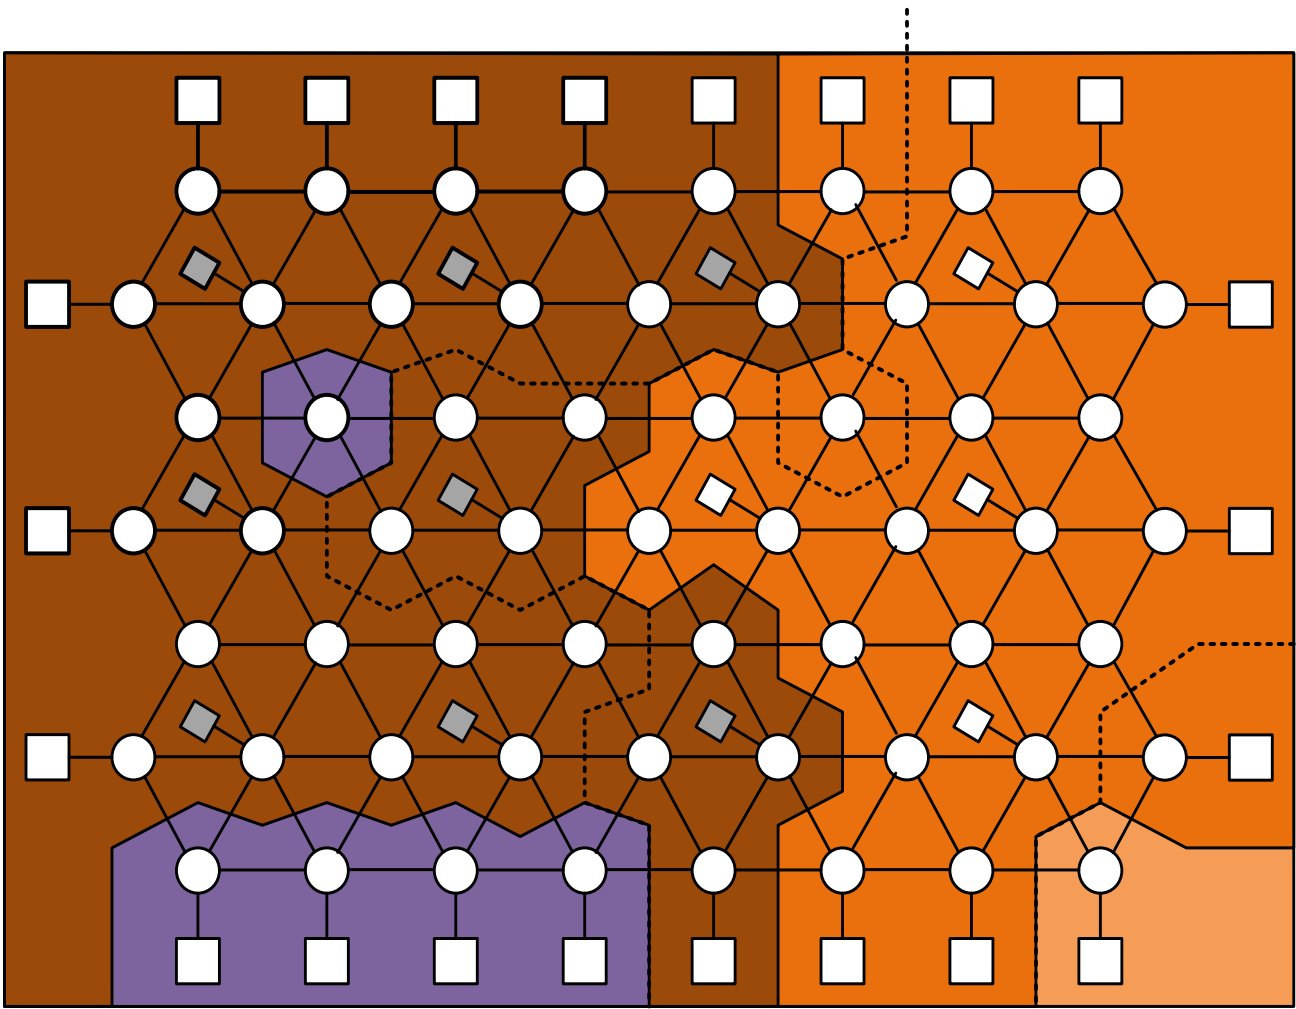
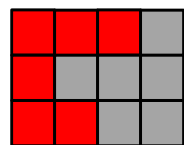
*Collecting*

$t=1$

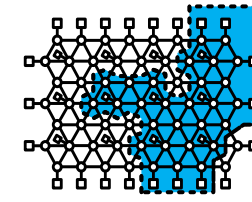


*Runahead*

$t=2$

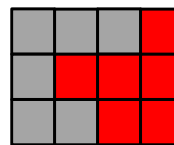


Supervisor 1



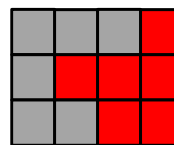
*Sending*

$t=0$



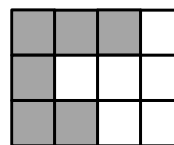
*Collecting*

$t=1$



*Runahead*

$t=2$





CW6

# CW6 – Speed up the simulator

- I have multiple applications written for POETS
  - Heat equation
  - Ising Spin
  - Helmholtz (EM field)
  - Navier Stokes (fluid mechanics)
- They all work, but...
  - Take a very long to simulate
  - [https://poetsii.github.io/graph\\_schema/heat/test.html](https://poetsii.github.io/graph_schema/heat/test.html)
  - It's difficult to estimate performance on real hardware
  - Bottlenecks are hard to find
- I'll give you a simulator for POETS
- You make it as fast as possible

# And that's it for lectures

- Available for consultation next Friday in the same slot (?)
- CW6 spec released tomorrow (Sat) morning
  - Have a good nights sleep...
- Feedback to keep coming
  - CW5 + CW6 feedback mostly given in orals
  - Orals scheduled on one-to-one basis in Spring term